



Micro-batching: High performance writes

Adam Zegelin, Instaclustr



Me

Adam Zegelin

- Co-founder of Instaclustr

Instaclustr

- Managed Apache Cassandra & DSE + Spark/Zepplin in the ☁
- Amazon Web Services, Azure, IBM Softlayer
- 24×7×365 Support

- Cassandra & Spark Consulting
- Enterprise support-only contracts

Problem Background

- Metrics — 2,000+ metrics (events) per-node, collected every 20 seconds
 - ~ 50k events *per-second*
- Streamed via RabbitMQ to Riemann
- Riemann by *Aphyr*
 - Event stream processor framework — Clojure (JVM)
 - Analyses/combines/filters events
 - Rules — events matching, over threshold, etc.
 - Actions — PagerDuty, email, Slack, etc.
 - Forwards *everything* to Cassandra

Problem Background cont'd

- Riemann:
 - Initially, single instance
 - Fairly complicated to scale (+HA)
See my blog post *Post 500 Nodes: High Availability & Scalability with Riemann*
 - Profiling: Writing to Cassandra = expensive (CPU intensive) *on the client*
 - Less CPU time to process events = backlog of events = 🔥
- End goal: reduce CPU-time spent writing to Cassandra
- Batching an applicable solution for our use-case

Micro-batching

- Insert data using small batches of Statements
- Improves throughput
- Reduces network overhead
- Less is more

Partition-aware Batching

- Batches of Statements, where each batch contains statements for the same partition
- `LoadBalancingPolicy.newQueryPlan(...)`
 - Returns `Iterable<Host>`, in order of preference
 - For `TokenAwarePolicy`, returns the replicas responsible for the Statement's partition key
- Group by the `first()` (most-preferable) host of each statement's query plan
 - ```
Multimap<Host, Statement> groups = Multimaps.index(statements,
 s -> lbp.newQueryPlan(s.getKeyspace(), s).next()
);
```

  
equiv. to `Map<Host, List<Statement>>`





# CASSANDRA SUMMIT 2016

## Benchmark

Insert *1 million* rows





# Benchmark Overview

- Write 1 million rows via a variety of strategies:
  - Individual statements, batches & host-aware batches
  - Batch sizes: 10, 50, 100, 250, 500 individual **INSERTs**
  - Consistency levels: **ALL**, **LOCAL\_QUORUM**, **LOCAL\_ONE**
- **UNLOGGED** batches
- Each strategy executed 10 times — average + std. dev.
- Determine:
  - Fastest
  - Lowest client CPU usage
  - Lowest cluster CPU usage



# Write Strategies

## Individual Writes

Each `PreparedStatement` is submitted to Cassandra via `executeAsync(...)`

## Batch Writes

Groups of  $n$  statements are combined into `BatchStatements`, one per group.

Each `BatchStatement` is submitted via `executeAsync(...)`

## Partition Aware Batch Writes

Groups of  $n$  statements, where every statement in the group shares the same partition key, are combined into `BatchStatements`, one per group. Submitted via `executeAsync(...)`

# Timing

The stopwatch is stopped once `ResultSetFuture.get()` unblocks for all futures returned by `executeAsync(...)`

## Runtime

Wall-clock time — `System.nanoTime()` via Guava's Stopwatch

## CPU Time

Total CPU time — `OperatingSystemMXBean.getProcessCpuTime()`  
Sum of the CPU time used by all threads in the JVM. Does not include I/O wait time.

## Cluster Average CPU Time

Average of the Total CPU time across all nodes in the cluster — collected via JMX



# Benchmark Setup

## Cluster

- Apache Cassandra 3.7
- 9 node cluster (3 racks, 3 nodes per rack)
  - m4.large — 250GB of EBS SSD, 2 vCPUs, 8 GB RAM
- NetworkTopologyStrategy, RF = 3
- `cassandra.yaml`:
  - Increased `batch_size_error_threshold_in_kb`  
(and `batch_size_warn_threshold_in_kb` to reduce log noise)
  - Increased `write_request_timeout_in_ms`
- Disabled compactions (`nodetool disableautocompaction`) & auto-snapshots  
Reduce variance between benchmark runs due to background tasks.

# Benchmark Setup cont'd

## Client

- *Single* c4.xlarge — 7.5GB RAM, 4 vCPUs
- OpenJDK 1.8, DataStax Cassandra Java Driver 3.1
- `TokenAwarePolicy + DCAwareRoundRobinPolicy`
- Tweaked connection pool & socket options:
  - Max requests per connection: 32,000 for `LOCAL` & `REMOTE`
  - Max connections per host: 20 for `LOCAL` & `REMOTE`
  - Pool timeout: 50 sec
  - Socket read timeout: 50 sec
- Upfront `prepare(...)` & `bind(...)` all 1 million statements




# Benchmark Setup cont'd

## Table/CF & Data

- Resembles our Riemann schema
- Keyspace DROPEd and re-CREATEd every benchmark run
- Generate:
  - 1000 random hosts (UUIDs) × 1000 random services (UUIDs)
  - Static values for bucket\_time, time, metric & state

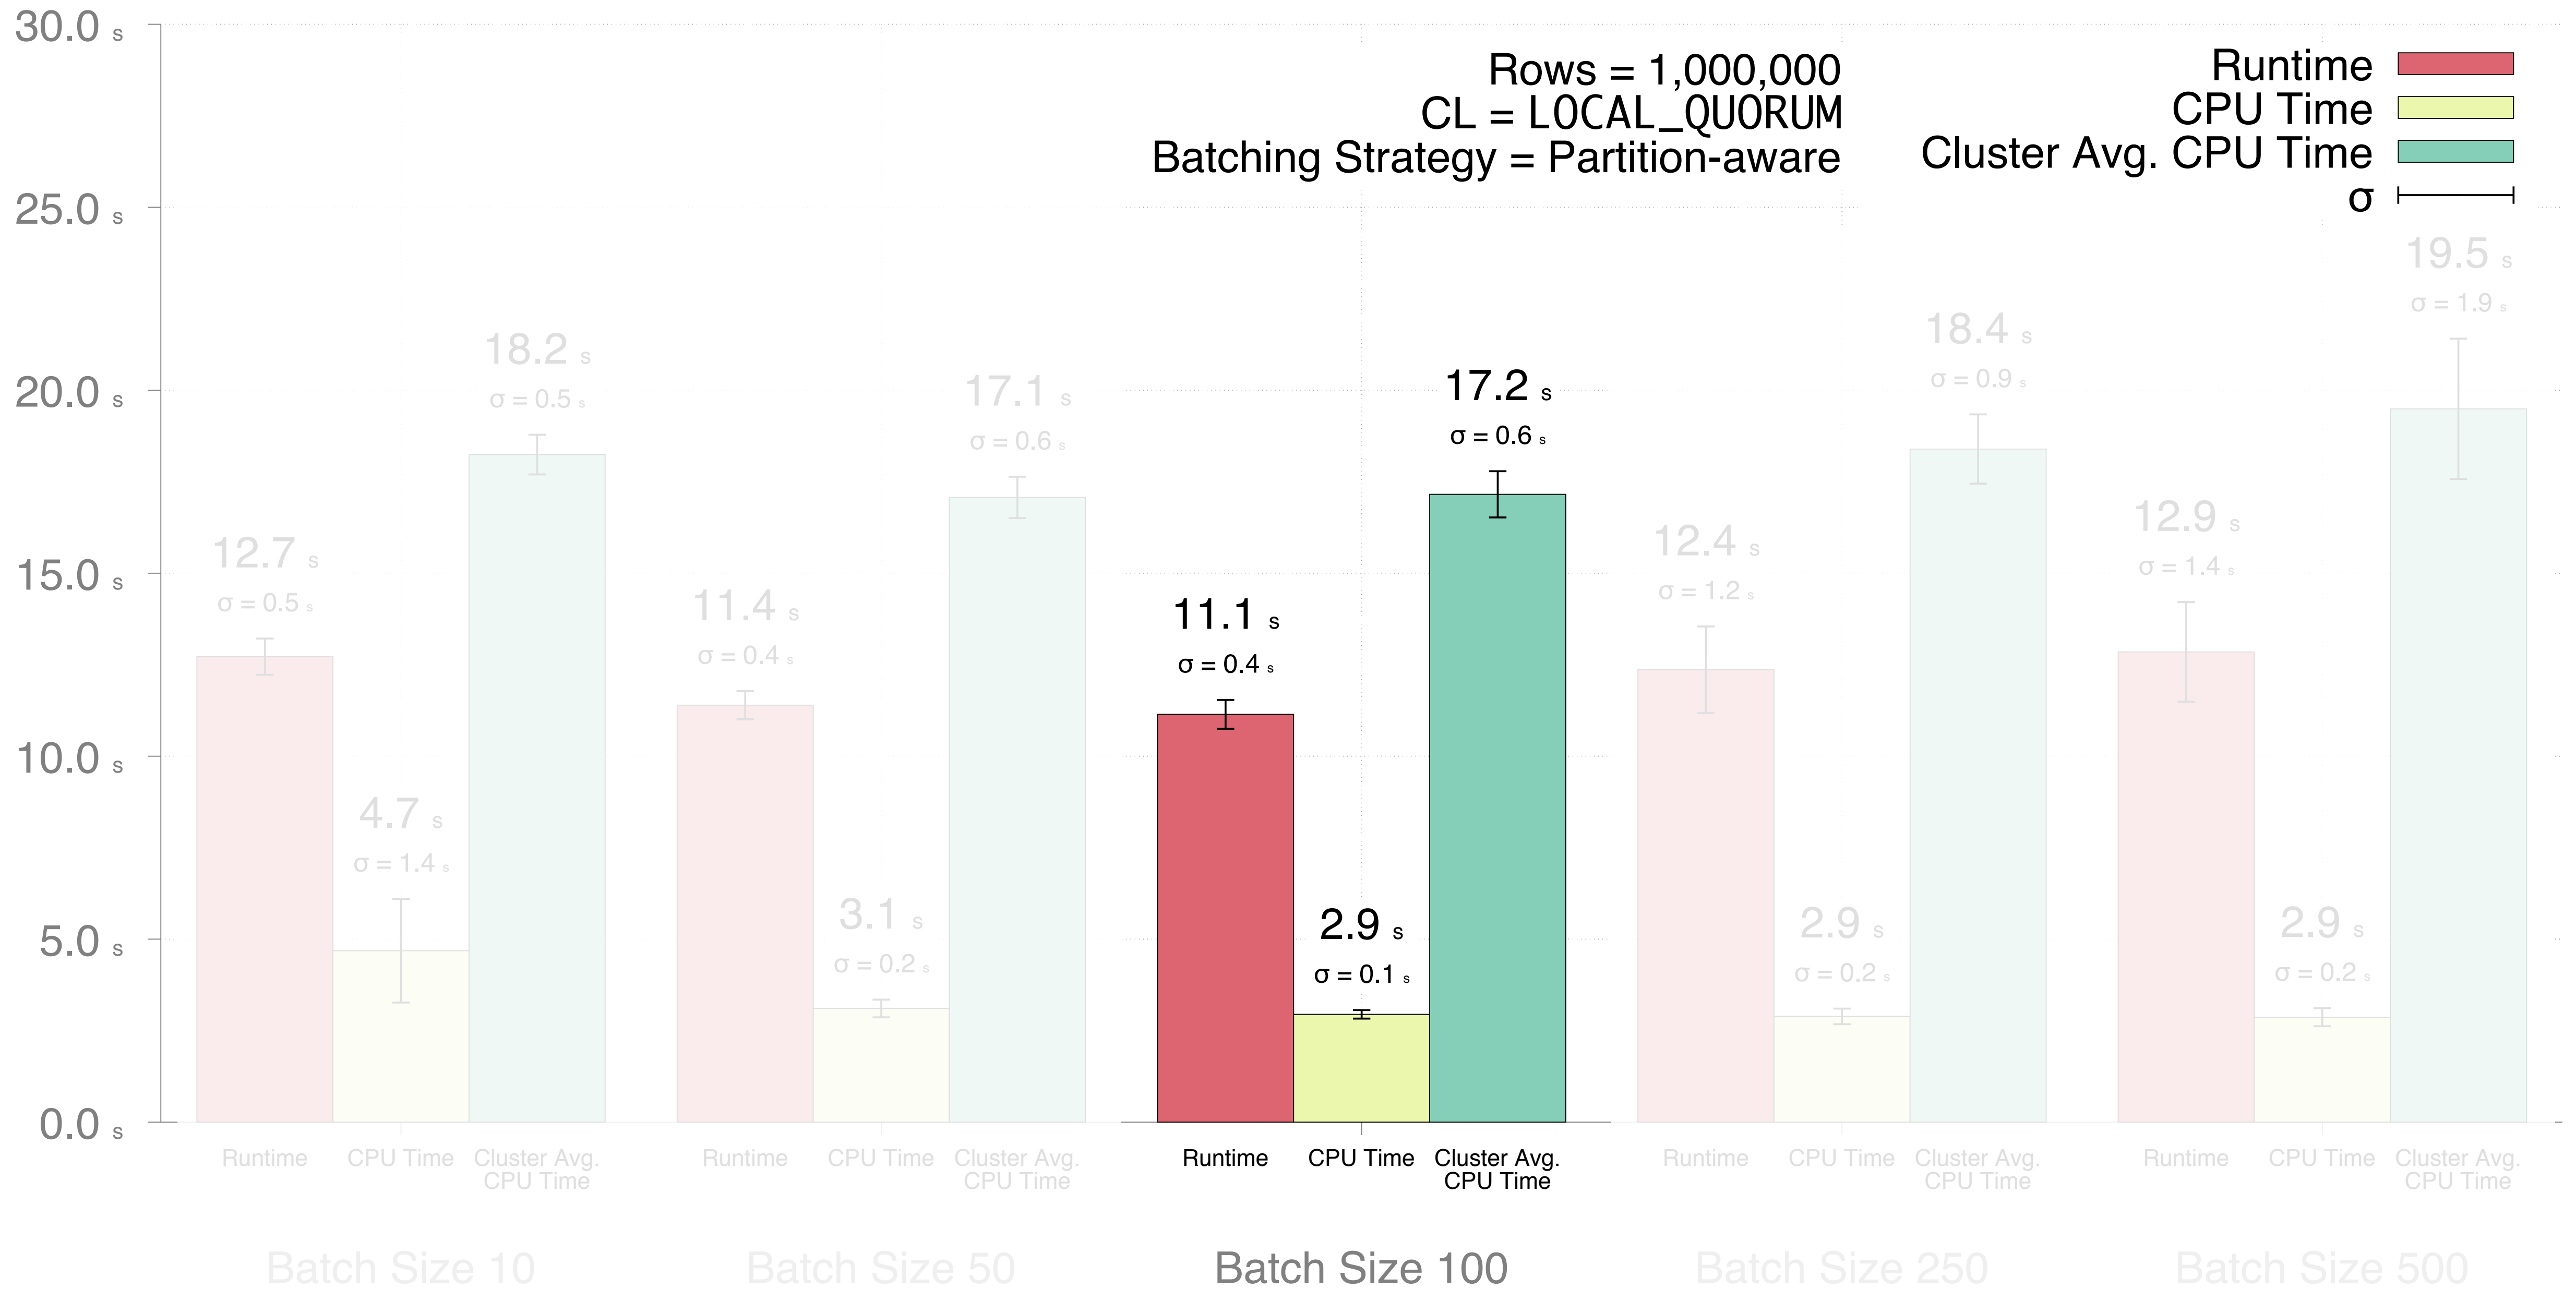
```
CREATE TABLE microbatching.example (
 host text,
 bucket_time timestamp,
 service text,
 time timestamp,
 metric double,
 state text,
 PRIMARY KEY ((host, bucket_time, service), time)
) WITH CLUSTERING ORDER BY (time DESC);
```

# Comparison of Write Strategies

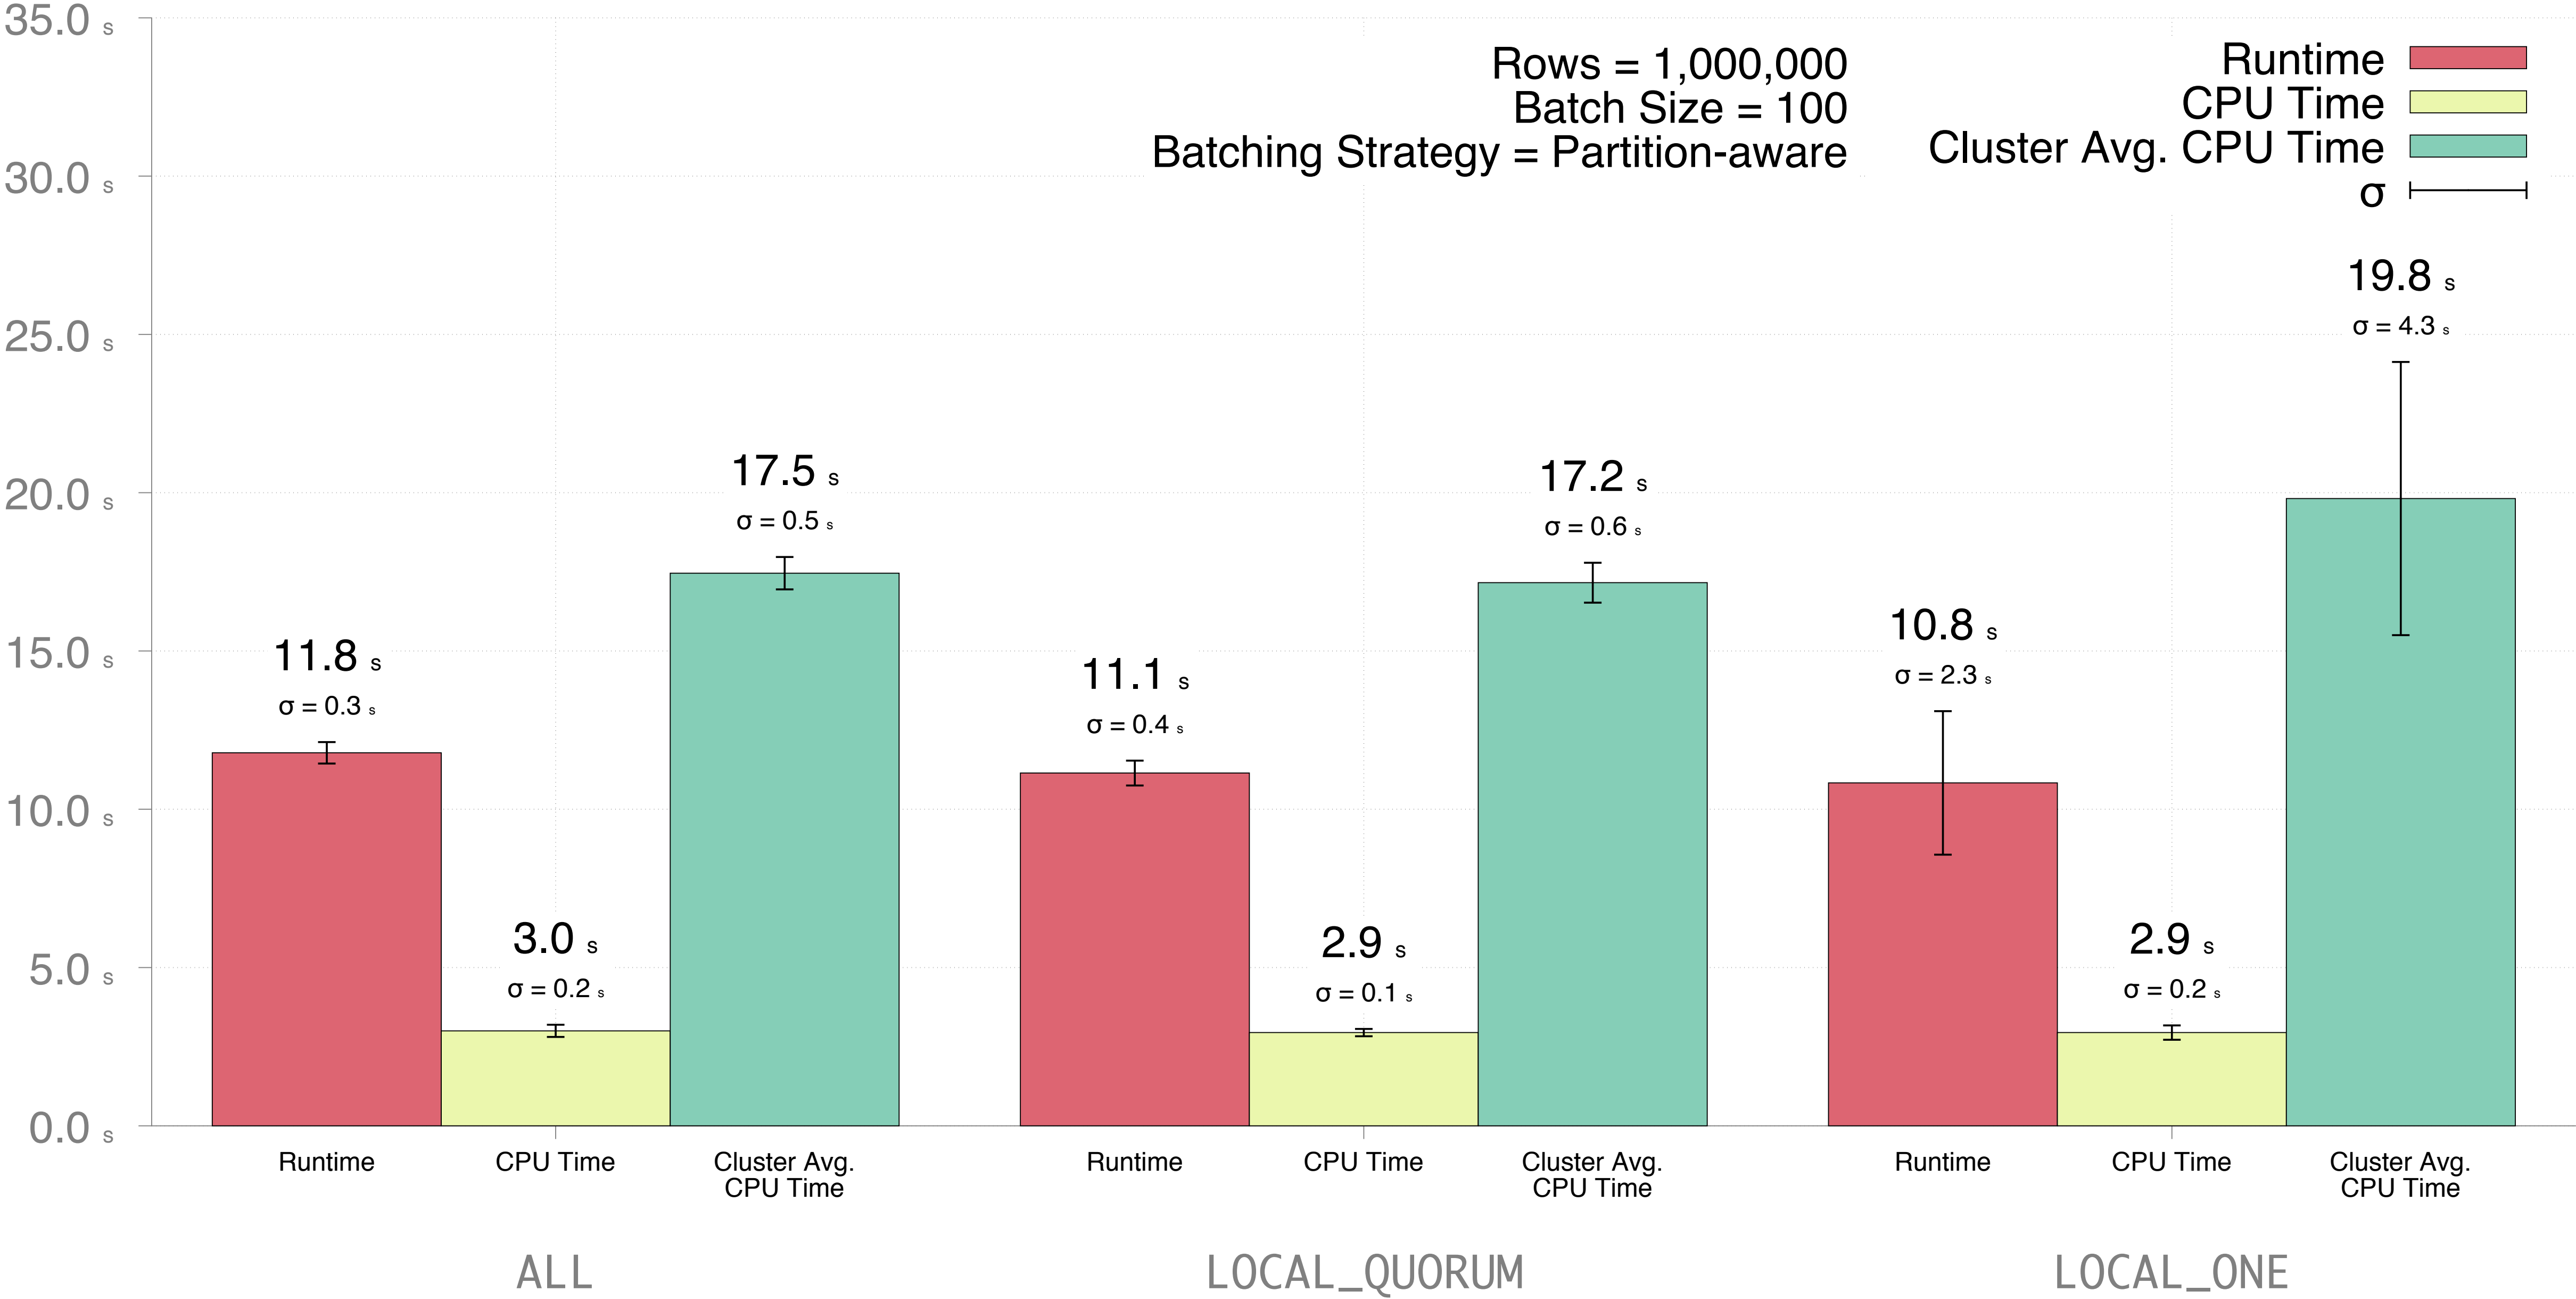




# Comparison of Batch Sizes



# Comparison of Consistency Levels





# Outcome

- More-performant
  - Shorter runtime
  - Lower client & cluster CPU load
- Possibly higher latency — more useful for bulk data processing
- Not a silver-bullet
- Standard **INSERTs** work well for most use-cases
- Cassandra benchmarking is hard

# Source

The Java source code for this benchmark is available at:

[bitbucket.org/adamzegelin/microbatching](https://bitbucket.org/adamzegelin/microbatching)





# CASSANDRA SUMMIT 2016

## Thank You!

Questions?

