



Alex Petrov

Scalable time-series applications with Cassandra

Scalable Time Series with Cassandra

@ifesdjeeen



Cyanite

past, present and future



**CASSANDRA
SUMMIT 2016**

Requirements



Throughput

Incoming data

Aggregates

Read queries



Scalability

Paths / metrics

Historical data

Readers / writers



CASSANDRA
SUMMIT 2016

Aggregated

vs raw



CASSANDRA
SUMMIT 2016

Predefined list of reports to aggregate
Reducing amount of data points
Faster queries
Precision loss due to aggregation

“servers.*.workers.busyWorkers”: “10s:7d,1m:21d,15m:5y”

Graphite



CASSANDRA
SUMMIT 2016

Store numeric time-series data
Render graphs of this data on demand

<https://graphiteapp.org/>

Ecosystem



carbon: listens for time-series data
whisper: DB for storing TS data
graphite-web: UI & API for graphs



Website Overview



Zoom Out

Last 3 hours



Logins

190



Sign ups

269



Sign outs

273



Memory / CPU




logins





Memory / CPU

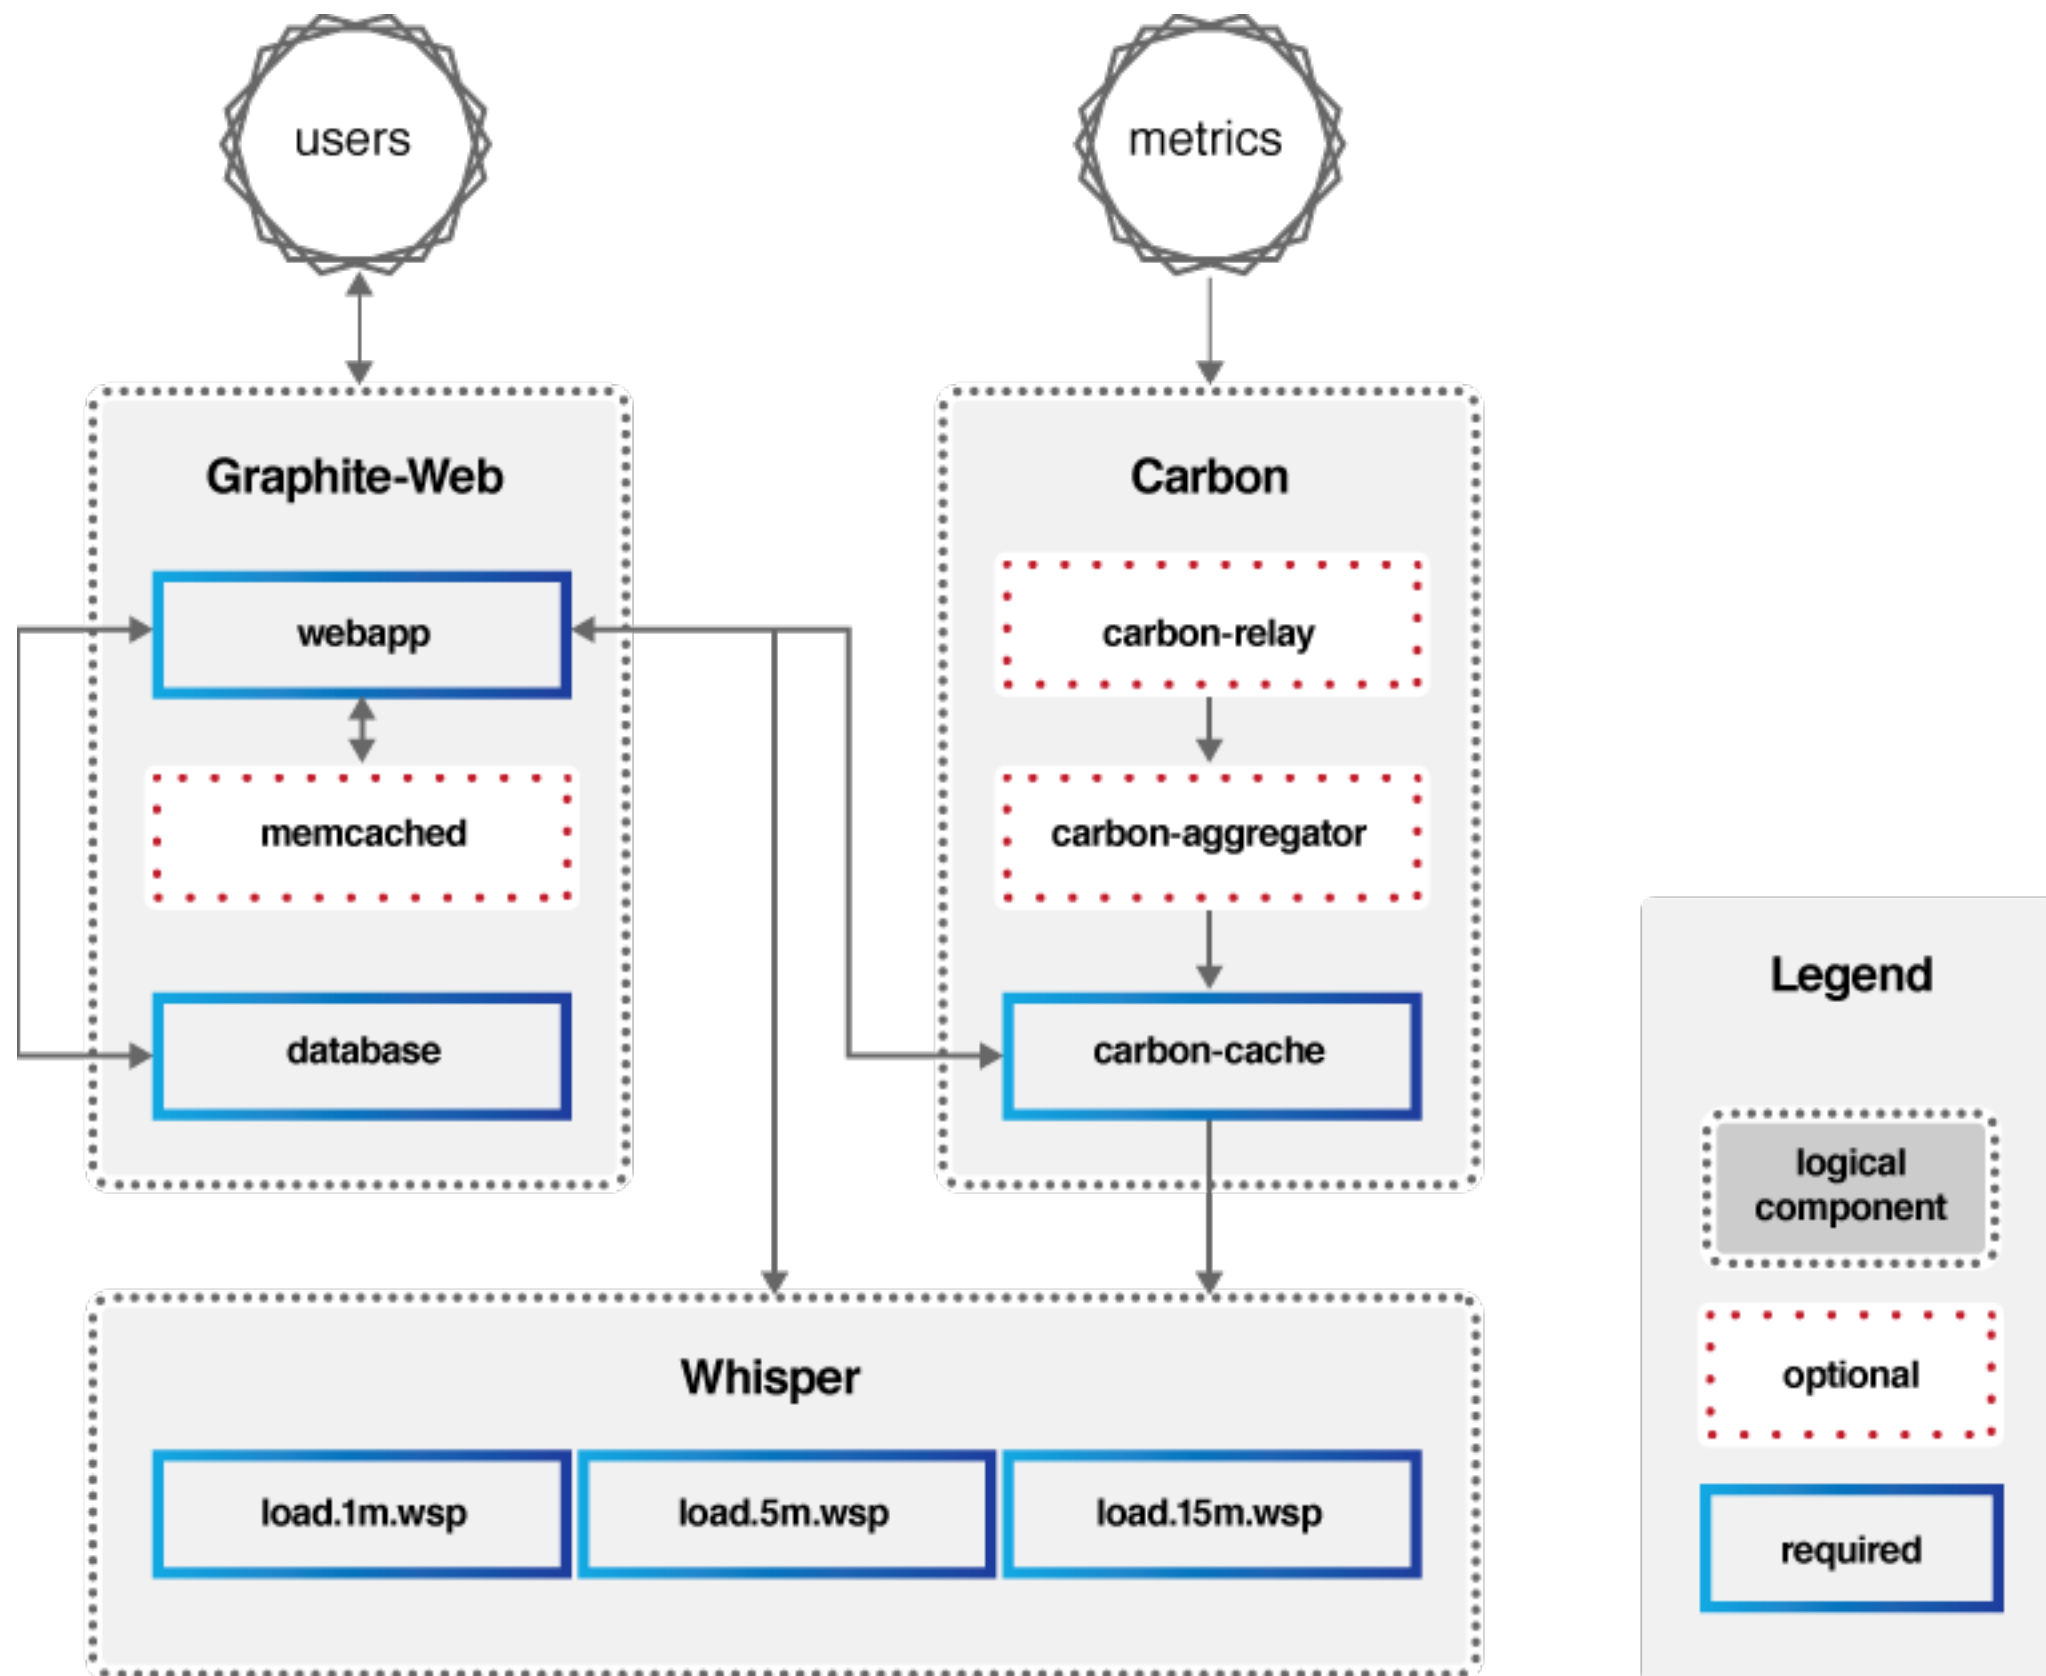


server requests



client side full page load





Downsides

Single-host solution

Plenty of disk seeks

Optimised for space

Sharding / replication is “manual”



Scaling

Cluster topology stored on every node

Manual replication

Changing cluster topology non-trivial



Scaling

Stateless service

Automatic shard assignment

Replication

Easy management

Easy topology changes



reminds you of anything?



**CASSANDRA
SUMMIT 2016**

Cassandra

perfect match

Cyanite

Stateless

Async I/O

Custom scheduler

No Whisper files

Distributed

Horizontally scalable



**CASSANDRA
SUMMIT 2016**

Cyanite responsibilities

Carbon-compatible listener

Aggregate data in-memory

Flush aggregates to Cassandra

Path storage

Retrieve paths for query

Aggregate the query results

Cassandra features

Metric expiry, TTL

User Defined Types

SASI Indexes, LIKE queries

Aggregate Functions

Offline data loading

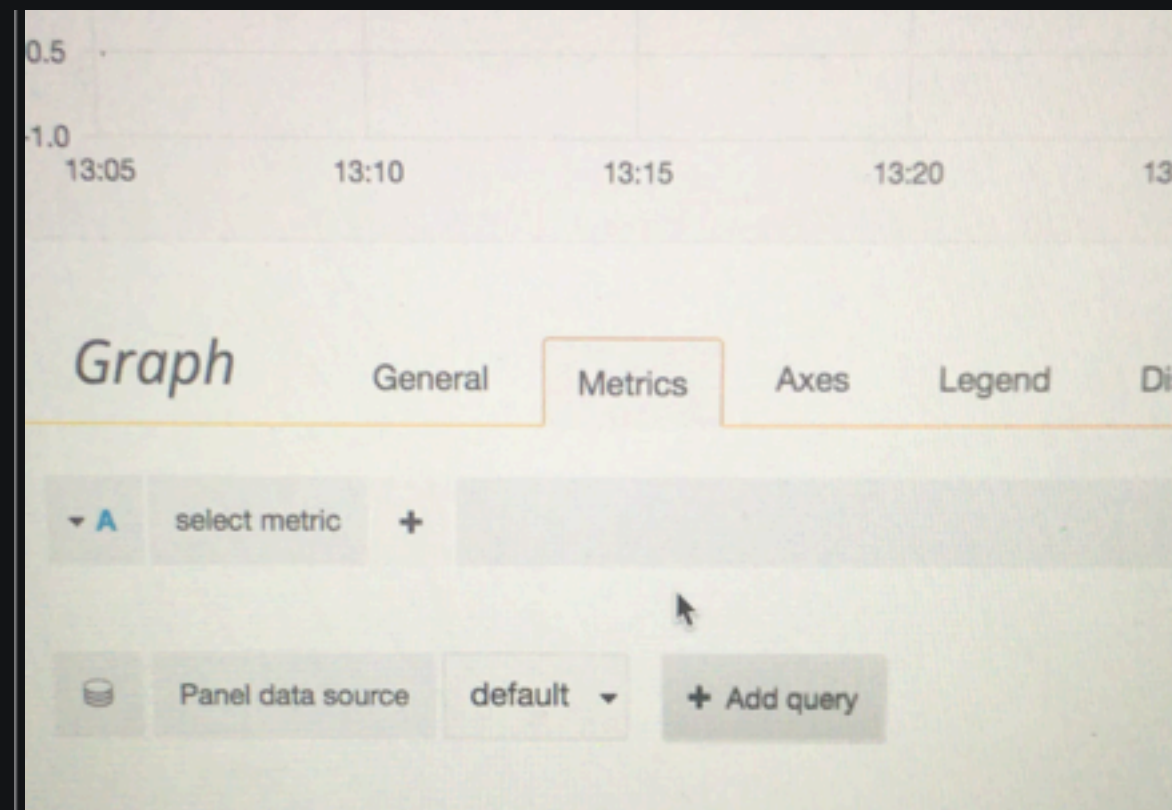
Clustering range queries

IN partition key locking



**CASSANDRA
SUMMIT 2016**

Globbing index



app.cluster.server.subsystem.metric

Built with SASI indexes
Fast, scalable queries
Optimised for glob

```
CREATE TABLE segment (  
    parent text,  
    segment text,  
    pos int,  
    length int,  
    leaf boolean,  
    PRIMARY KEY (parent, segment))
```



app.cluster.server.subsystem.metric

Wildcard: *

```
SELECT * FROM segments WHERE  
parent = 'root' AND  
pos = 1
```

Postfix: *..*.metric**

**SELECT * FROM segments WHERE
pos = 4**

Prefix: app.*

**SELECT * FROM segments WHERE
parent = 'app' AND
pos = 2 ALLOW FILTERING**



**CASSANDRA
SUMMIT 2016**

Suffix: abc.*.metric

```
SELECT * FROM SEGMENTS  
WHERE  
pos = 3 AND  
segment LIKE 'abc.%'  
ALLOW FILTERING
```


Query engine

Query engine

2-Step transformations:

Inner

Outer

Cross metric



**CASSANDRA
SUMMIT 2016**

scale(a.b.*, 10.0)



a.b.c1 a.b.c2



**{"a.b.c1" [1 2 3]
"a.b.c2" [5 6 7]}**



**{"a.b.c1" [10 20 30]
"a.b.c2" [50 60 70]}**



derivative(a.b.c)



a.b.c



{"a.b.c" [1 3 6]}



{"derivative(a.b.c)" [nil 2 3]}



**CASSANDRA
SUMMIT 2016**

sumSeries(a.b.c1,a.b.c2)



a.b.c1 a.b.c2



**{"a.b.c1" [1 1 1]
"a.b.c2" [2 2 2]}**



{"sumSeries(a.b.c,a.b.d)" [3 3 3]}

Data model

```
CREATE TYPE IF NOT EXISTS metric_resolution (  
    precision    int,  
    period       int  
);
```

```
CREATE TYPE IF NOT EXISTS metric_id (  
    path         text,  
    resolution    frozen<metric_resolution>  
);
```

```
CREATE TYPE IF NOT EXISTS metric_point (  
    max double,  
    mean double,  
    min double,  
    sum double  
);
```


Cassandra

Aggregates

Cassandra aggregates

Locked partition

Flexible query definition

Local aggregation

Reduced raw data chatter



CASSANDRA
SUMMIT 2016

CREATE OR REPLACE FUNCTION sumState
(state map<bigint, double>, ts bigint, val metric_point)
CALLED ON NULL INPUT RETURNS map<bigint, double>
LANGUAGE java AS \$\$

```
if (state.containsKey(ts)) {  
    state.put(ts, state.get(ts) + val.getDouble("mean"));  
} else {  
    state.put(ts, val.getDouble("mean"));  
}
```

```
return state;  
$$;
```

```
CREATE OR REPLACE FUNCTION sumFinal  
    (state map<bigint, double>)  
CALLED ON NULL INPUT RETURNS map<bigint, double>  
LANGUAGE java AS  
'return state;';
```

```
SELECT sum(time, point) FROM metric WHERE id  
IN ({path: 'a', resolution: {precision: 1, period: 3600}},  
    {path: 'b', resolution: {precision: 1, period: 3600}})  
AND time > 0  
AND time < 5;
```

Scaling



Scaling cyanite

Use DTSC

Cyanite is stateless

For high loads, split readers and writers

Load-balance with HAProxy

Colocate Cyanite & Cassandra

Coming

soon



CASSANDRA
SUMMIT 2016

Coming soon

P-Square histograms

T-Digest quantiles

Gorilla Compression

Cassandra aggregates

Custom glob indexes

Scheduler

Coming soon

Kafka ingester
Statsd protocol support
Standalone cyanite
Dynamic thresholds

That's all Folks!