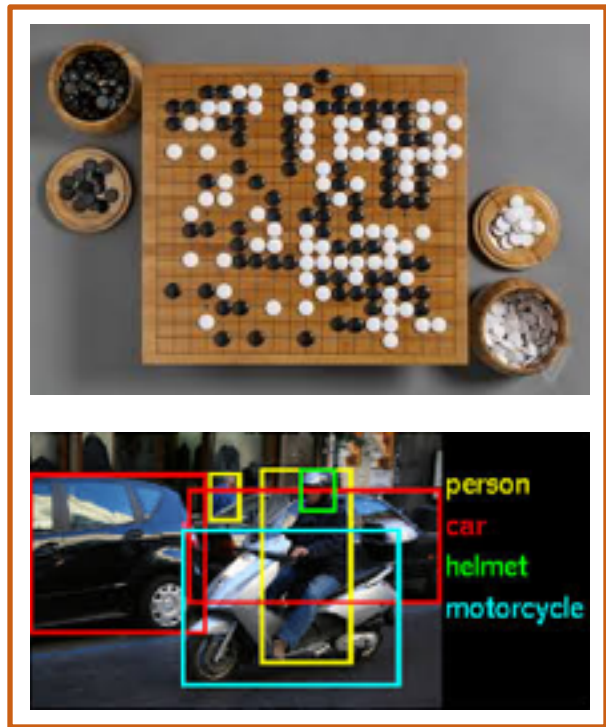




C\* for Deep Learning

# Tractable

- Artificial Intelligence products for businesses
- Combining Deep Learning / Neural Networks and traditional Machine Learning
- AI using Deep Learning has surpassed human intelligence:
  - Go
  - Image Recognition



# Deep Learning 101

## Deep Learning

A branch of **machine learning** based on a set of **algorithms** that attempt to model high-level abstractions in data by using a deep graph with multiple processing layers, composed of multiple linear and **non-linear transformations**.

---

1.

Requires large amounts of data to train networks

2.

Computations made feasible by use of GPUs for dramatic speedup

# Semantic Image Search

Search in the meaning of the image

**NOT** search on the image itself



Similar object



Similar Image



# Semantic Image Search



Similar object



Similar Image




# Semantic Search 2

- Intent can vary!
- With training AI can do both of these tasks
- Semantic search is more than just classification:
  - Ranking within a classification
  - Search for things that are not classification categories



# Feature Extraction




1 Image  $\longrightarrow$  4000 dimension vector

1 TB Images  $\longrightarrow$  300 GB of features

# Feature Extraction



GPU Servers



C\*



# Feature Extraction

- Processing images at 5 GigaBytes per second
- Features generated at 1.6 GigaBytes per second

```
CREATE TABLE features
  listing_id uuid,
  image_id uuid,
  feature_vector blob,
  PRIMARY KEY ((listing_id), image_id)
```

# Search

“I want a truck like this”

- 300 GB of features
- Millions of rows in C\*




# Spark

Distributed, in-memory computation

- Map-reduce
- Graph analysis
- SQL abstraction

DataStax Spark Connector

- Understands Cassandra partitioning
- Push-down queries to C\* keys where possible
- Join on C\* tables



# Spark Similarity Search

1. Cache our 300GB features in Spark
2. Score every feature vector with our search image
3. Aggregate score for images in each listing
4. (optional) Join on 'live' C\* data
5. Write result back to C\*

```
CREATE TABLE features
  listing_id uuid,
  image_id uuid,
  feature_vector blob,
PRIMARY KEY ((listing_id), image_id)
```

```
CREATE TABLE stock_level
  listing_id uuid,
  inventory int,
  etag uuid,
PRIMARY KEY (listing_id)
```

```
CREATE TABLE query_results
  query_id uuid,
  listing_id uuid,
  score float,
PRIMARY KEY ((query_id), score)
```

# Search

Search returns results within 5 seconds

---

**100x** Speedup from holding data in-memory in Spark

**10x** Speedup from co-locating Spark and C\*

**20x** Speedup from partitioning on the group key (listing id)

**5x** Speedup from “Cassandra sort”

# Search Results



# Search Results



# Sophisticated Analytics

- Visualisation
- Clustering
- Outlier Detection

Real time, interactive  
exploration of large visual  
data sets

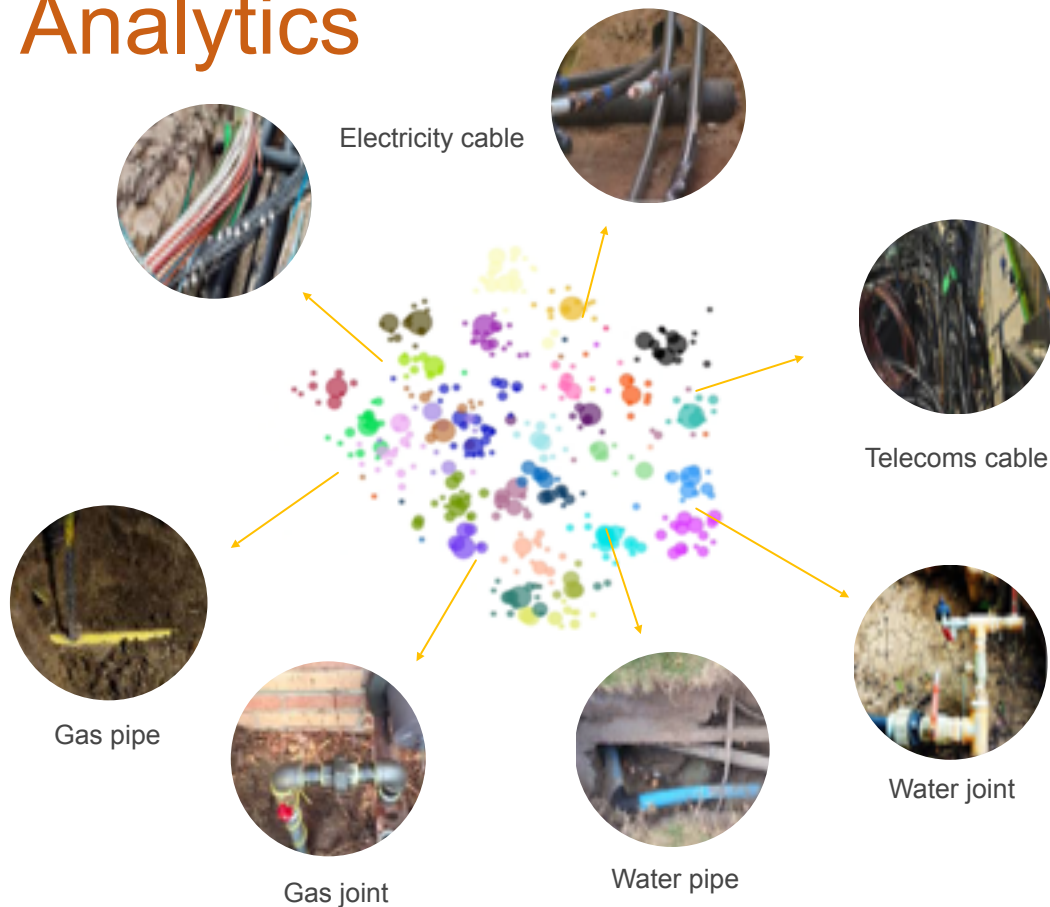


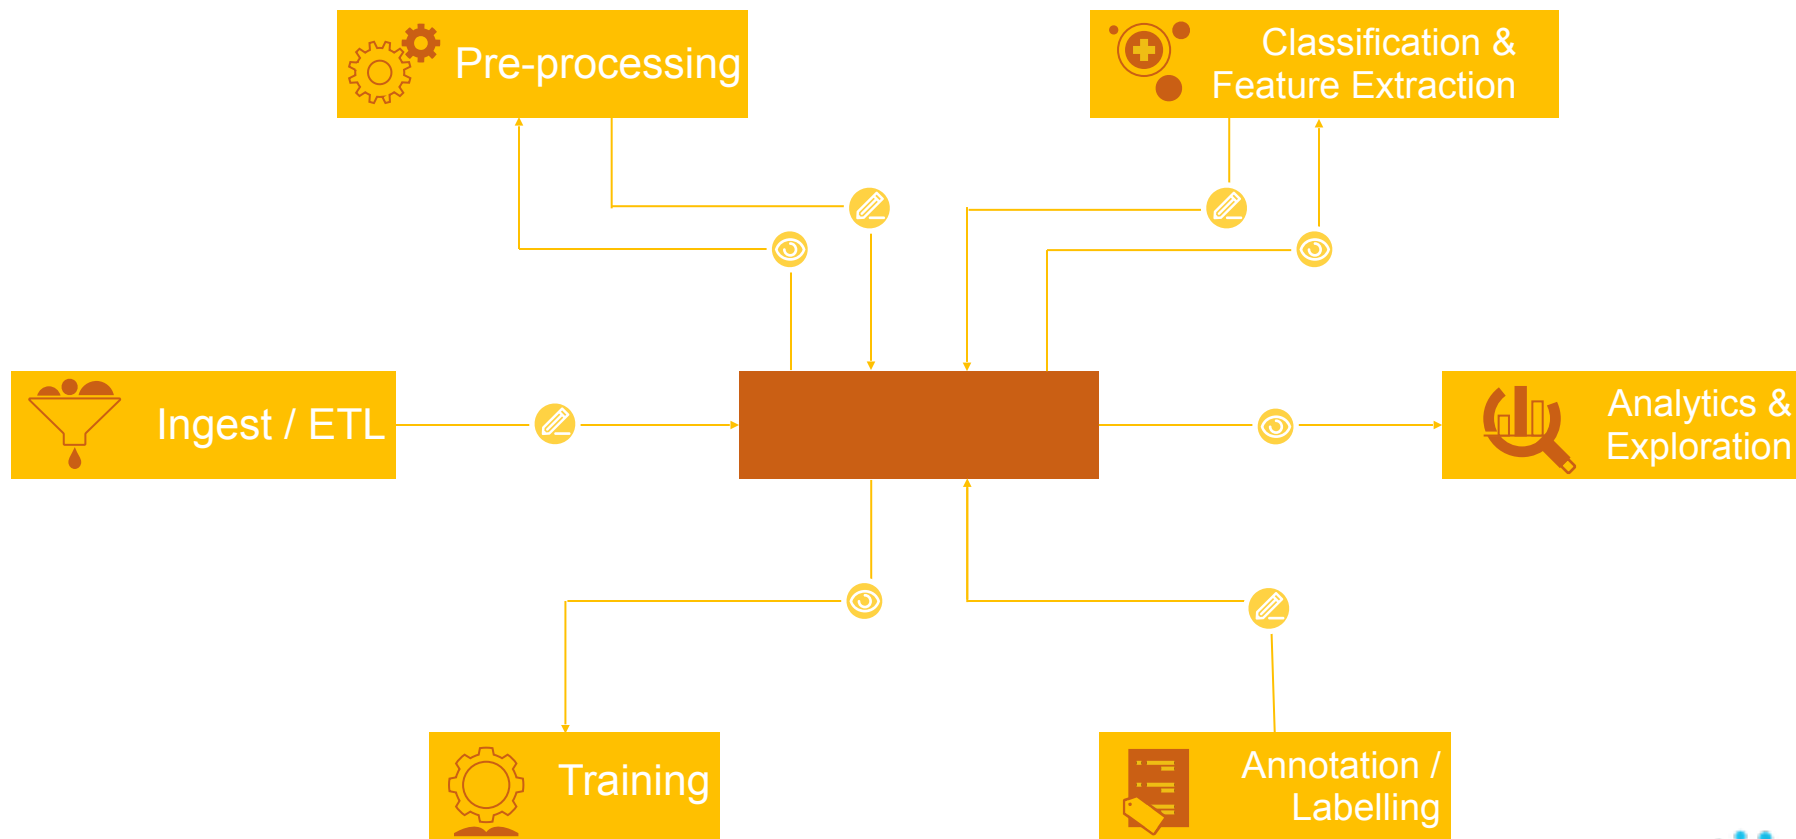


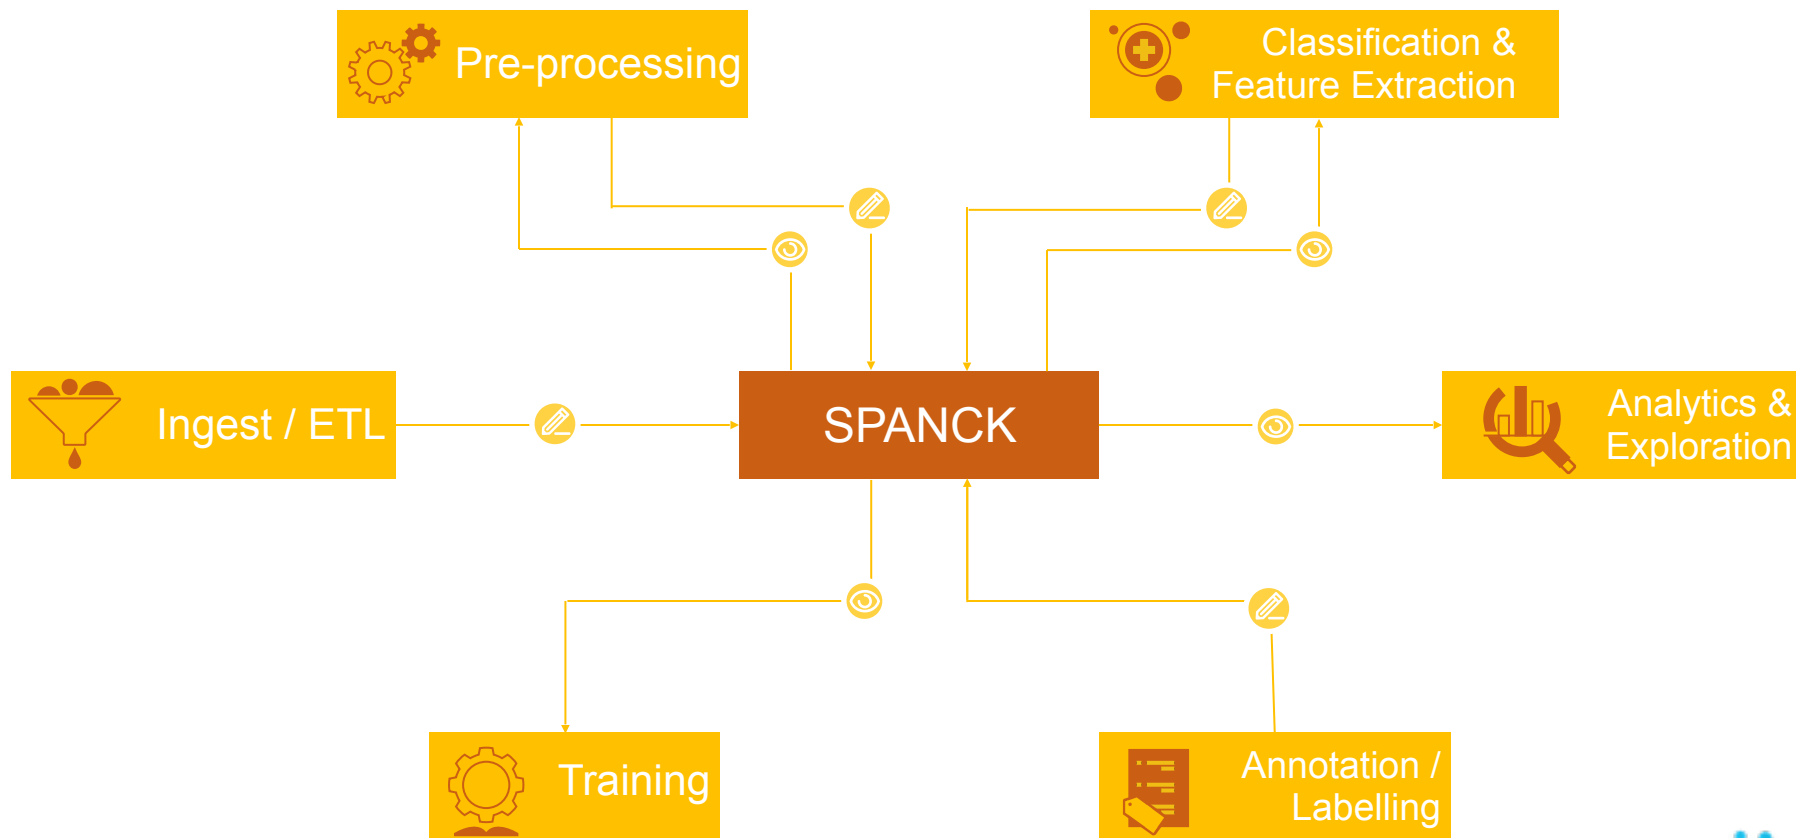
# Sophisticated Analytics

- Visualisation
- Clustering
- Outlier Detection

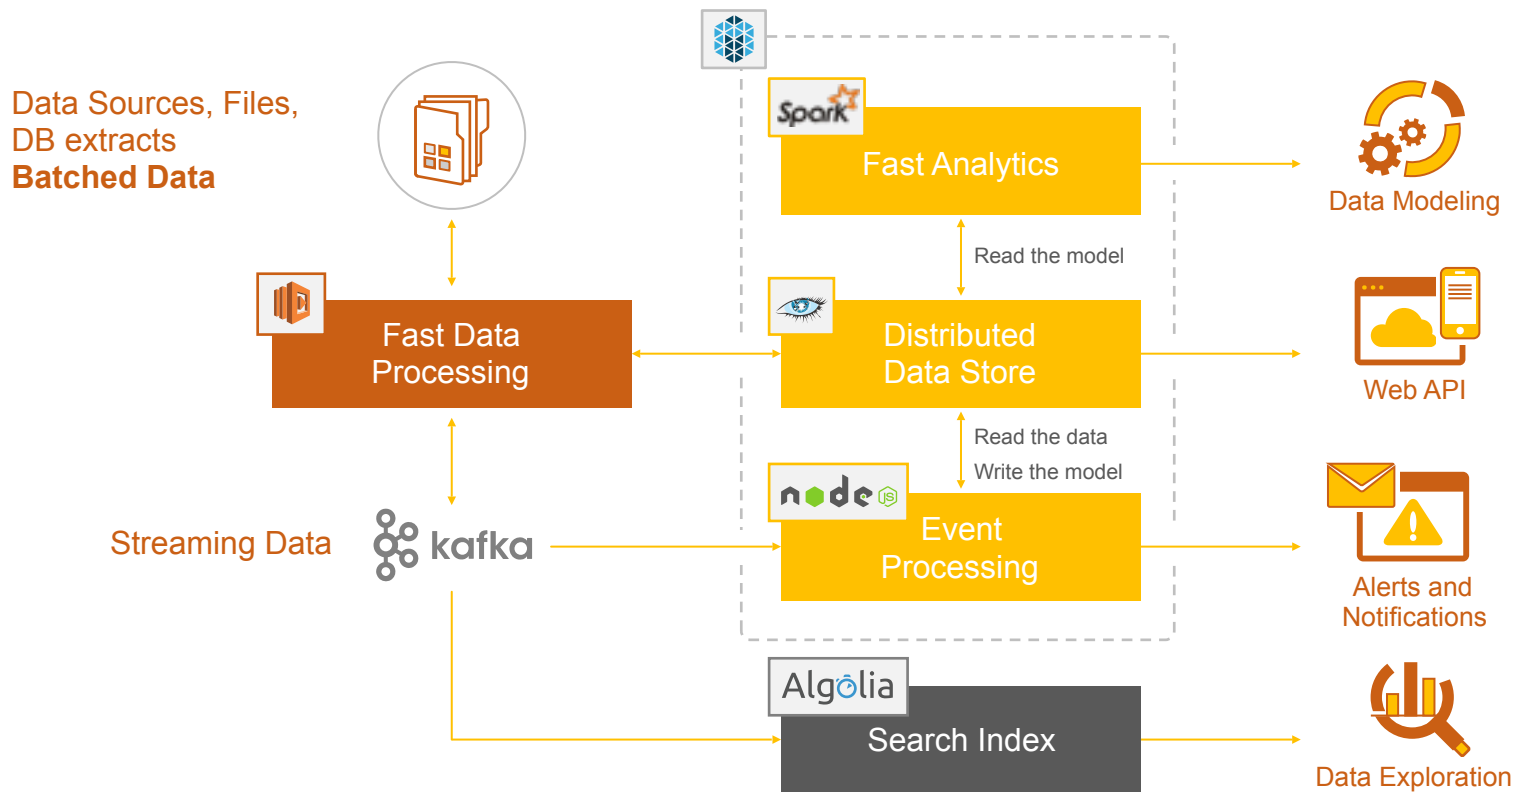
Real time, interactive  
exploration of large visual  
data sets








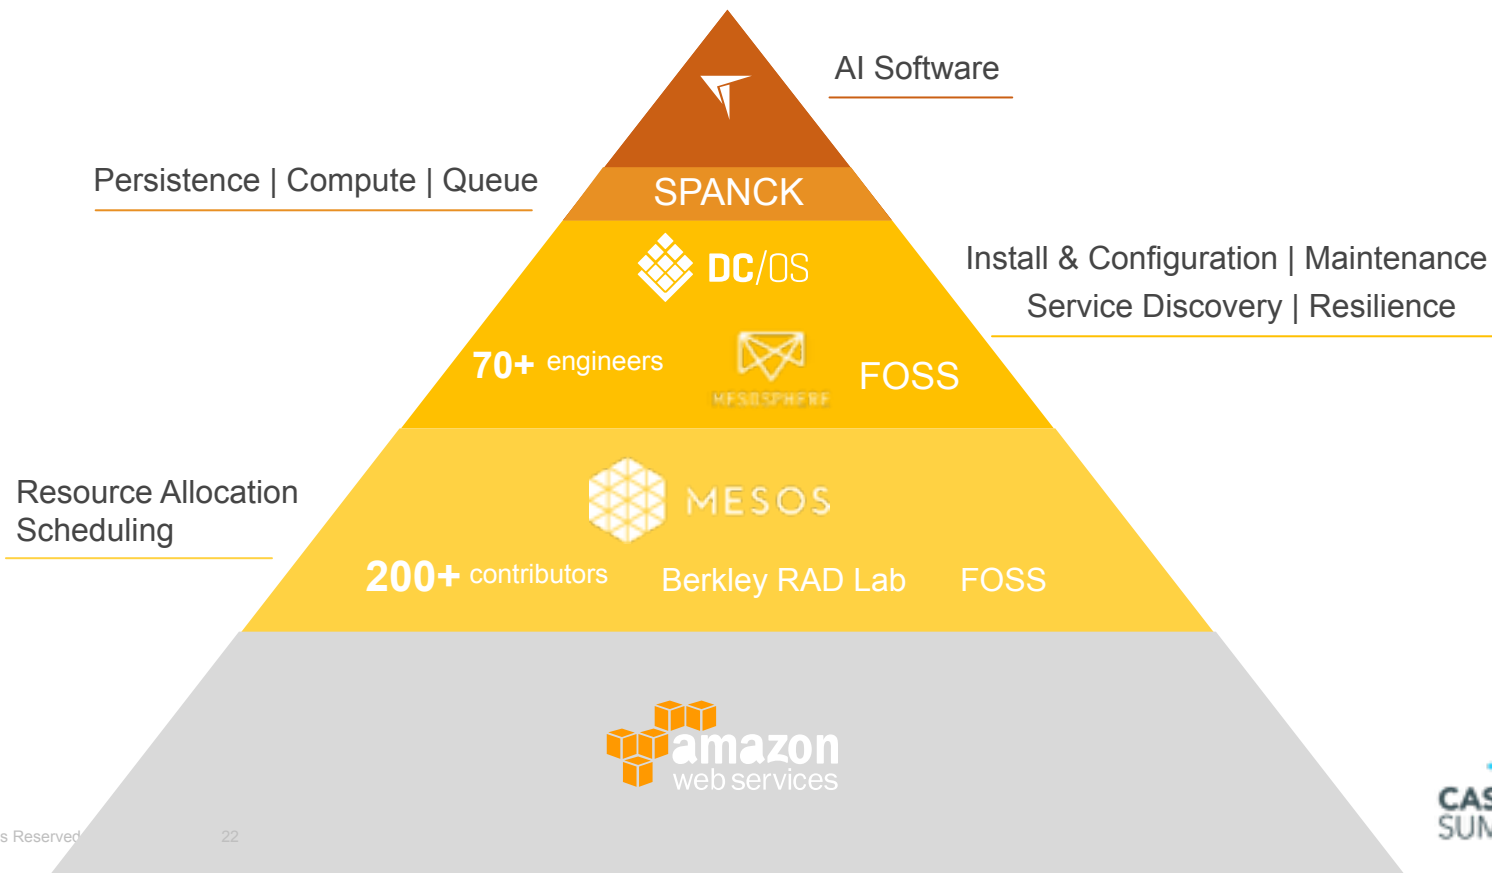
# SPANCK: Spark-Python Algolia Node Cassandra Kafka



# Many Services, 1 Engineer



# Many Services, 1 Engineer




# DCOS Frameworks

## Easy Install

\$ dcos package install kafka --  
options kafka-options.json

\$ dcos package install cassandra --  
options cassandra-options.json

```
{
  "nodes": {
    "count": 12,
    "cpus": 9,
    "mem": 30000,
    "disk": 400000,
    "heap": {
      "size": 8000,
      "new": 1500
    }
  },
  "cassandra": {
    "row_cache_size_in_mb": 12000,
    "row_cache_save_period": 600,
    "commitlog_segment_size_in_mb": 64,
    "concurrent_reads": 24,
    "concurrent_writes": 64,
    "memtable_allocation_type": "offheap_objects",
    "compaction_throughput_mb_per_sec": 128
  }
}
```



# DCOS Frameworks

## Service Discovery

- 0 configuration files
- 0 configuration / orchestration systems
- Works ever time, in every environment

```
function fetchMarathonServiceBaseUrl(serviceName){  
  var dns_url = ['_' + serviceName, "_tcp.marathon.mesos"].join('.');  
  return resolveSrv(dns_url).then(function(addresses){  
    return "http://" + addresses[0].name + ":" + addresses[0].port;  
  })  
}
```

```
fetchCassandraNodes = (mesosHost, authToken) ->  
  mesosDns.fetchMarathonServiceBaseUrl('cassandra').then (baseUrl) ->  
    request(  
      url: baseUrl + "/v1/nodes/connect",  
      headers: headers  
    ).then (result) ->  
      result = JSON.parse(result)  
      return _.pluck(result.nodes, "ip")
```



# DCOS Frameworks

## Maintenance and Admin DCOS Command Line Tools

- Replace a C\* node
- Backup / Restore C\* to AWS S3
- Run C\* repair / cleanup
- Restart C\* nodes
- Replace a Kafka broker
- Rebalance Kafka brokers
- Restart Kafka brokers

```
dcos cassandra --name=cassandra \  
cleanup --key_spaces=dev,test
```

```
dcos cassandra --name=cassandra \  
replace node-4
```

```
dcos kafka broker replace 3
```

# DCOS Frameworks

## Shipping Docker Apps

1 JSON file

1 CLI command

## Apps include

- Spark Drivers
- Web Servers
- Kafka Producers
- Kafka Consumers
- APIs

```
{
  "id": "/my-app",
  "cpus": 1,
  "mem": 2048,
  "instances": 1,
  "env": {
    "NODE_ENV": "dev"
  },
  "container": {
    "type": "DOCKER",
    "docker": {
      "image": "tractableio/my-app:0.1",
      "network": "BRIDGE"
    }
  },
  "healthChecks": [
    {
      "path": "/healthy",
      "maxConsecutiveFailures": 2
    }
  ]
}
```

\*JSON Slightly simplified

# DCOS Cluster in 15 minutes

1.

Deploy DCOS cluster into AWS using CloudFormation template

2.


Install Cassandra & Kafka

3.

Deploy app docker containers into DCOS

4.

Profit!



# Conclusion

- Cassandra's high write speeds allow it to ingest features from Deep Networks
- Cassandra and Spark provides a powerful compute+storage system
- Spark, Cassandra and Kafka can provide a versatile data backbone that supports a range of use cases
- Mesosphere DCOS is a low-effort, high-reward way of running distributed systems