



Advanced DSE analytics client configuration

Jacek Lewandowski


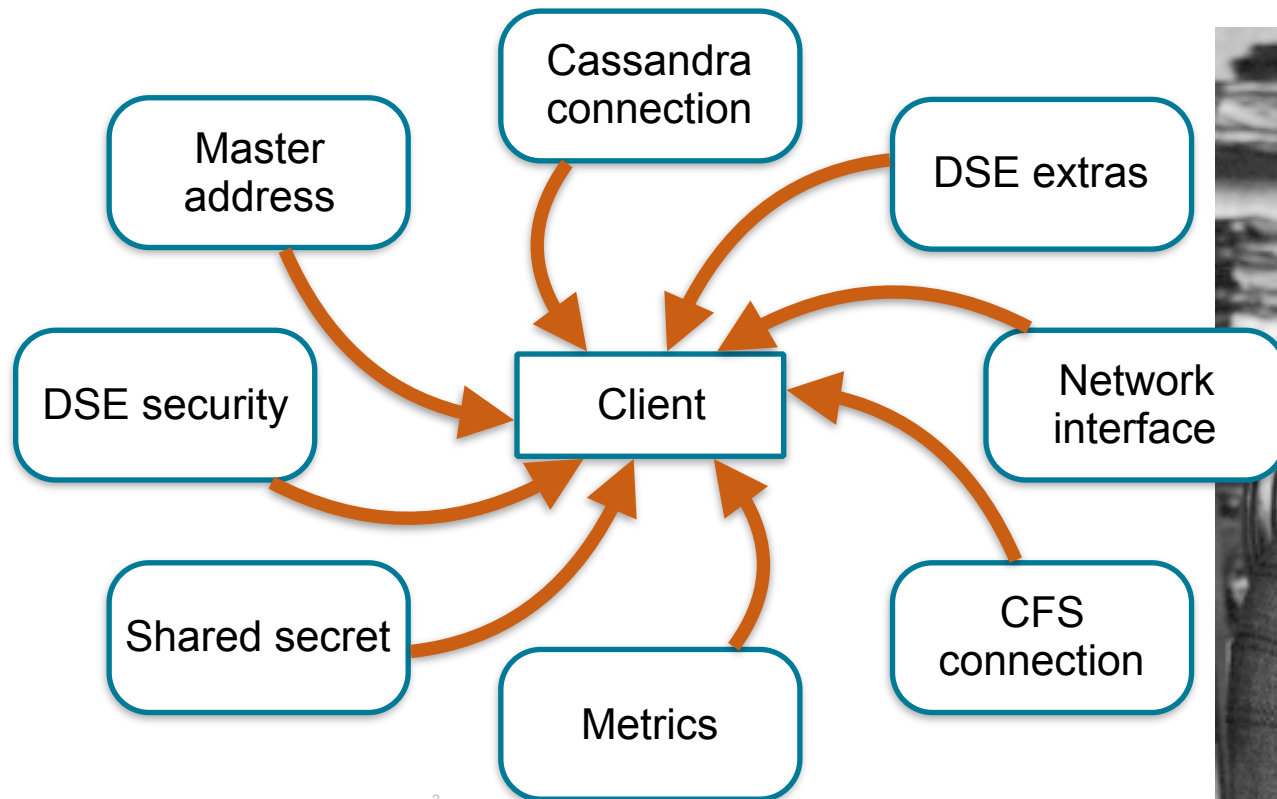
Software engineers in DataStax Analytics Team

jacek.lewandowski@datastax.com

September 8, 2016

- | | |
|---|---|
| 1 | DSE configures client applications automatically |
| 2 | DSE supports Spark HA, even tricky cases |
| 3 | DSE brings useful improvements to Spark |
| 4 | DSE extends usability of Spark shared secret security |
| 5 | No more manual configuration of remote DSE nodes! |

How would you configure your client manually?



No worries, you are using DataStax Enterprise

It does the job for you

Spark

```
$ dse spark
```

Hive

```
$ dse hive
```

Hadoop

```
$ dse hadoop
```

Sqoop

```
$ dse sqoop
```

Pig

```
$ dse pig
```

Client-tool

```
$ dse client-tool
```

No need to provide Cassandra hosts or ports

No need to configure security except providing credentials

No need to specify Spark Master address

No need to specify JobTracker address

It's that simple!

No need to provide any additional configuration to start working

DataStax Enterprise generates the default configuration and lets the experts to overwrite it

Generated configuration

hadoop/dse-core-site.xml
hadoop/dse-mapred-site.xml

Connection configuration for local node

For example: `cassandra.host`

Hadoop static configuration

hadoop/core-site.xml
hadoop/mapred-site.xml

Any additional settings, unrelated to any particular node
Lets you overwrite what was generated by DSE

For example: `cassandra.host`

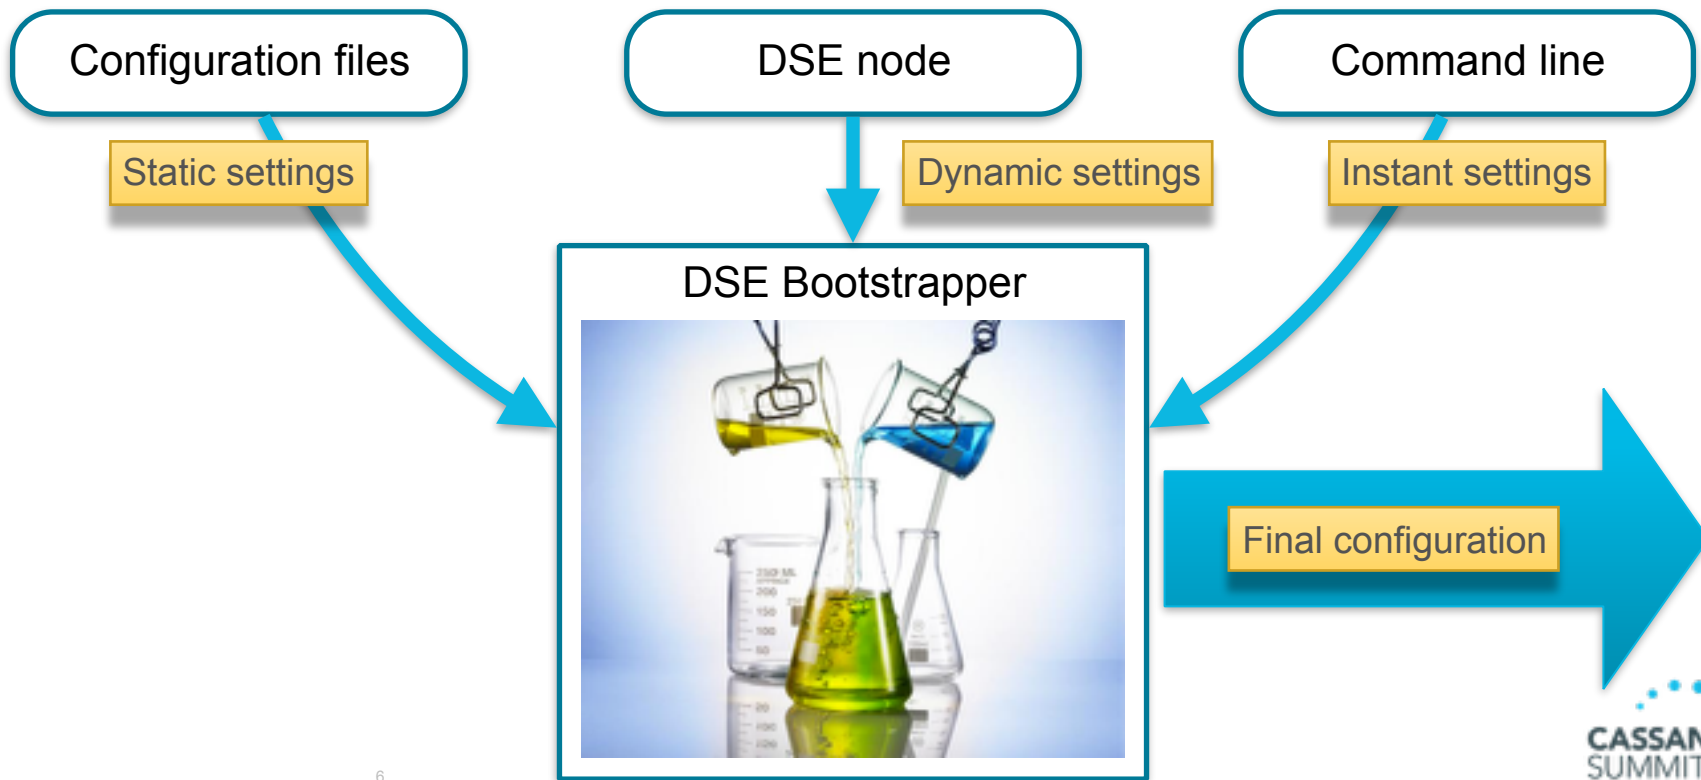
Spark configuration

spark-defaults.properties
command line arguments

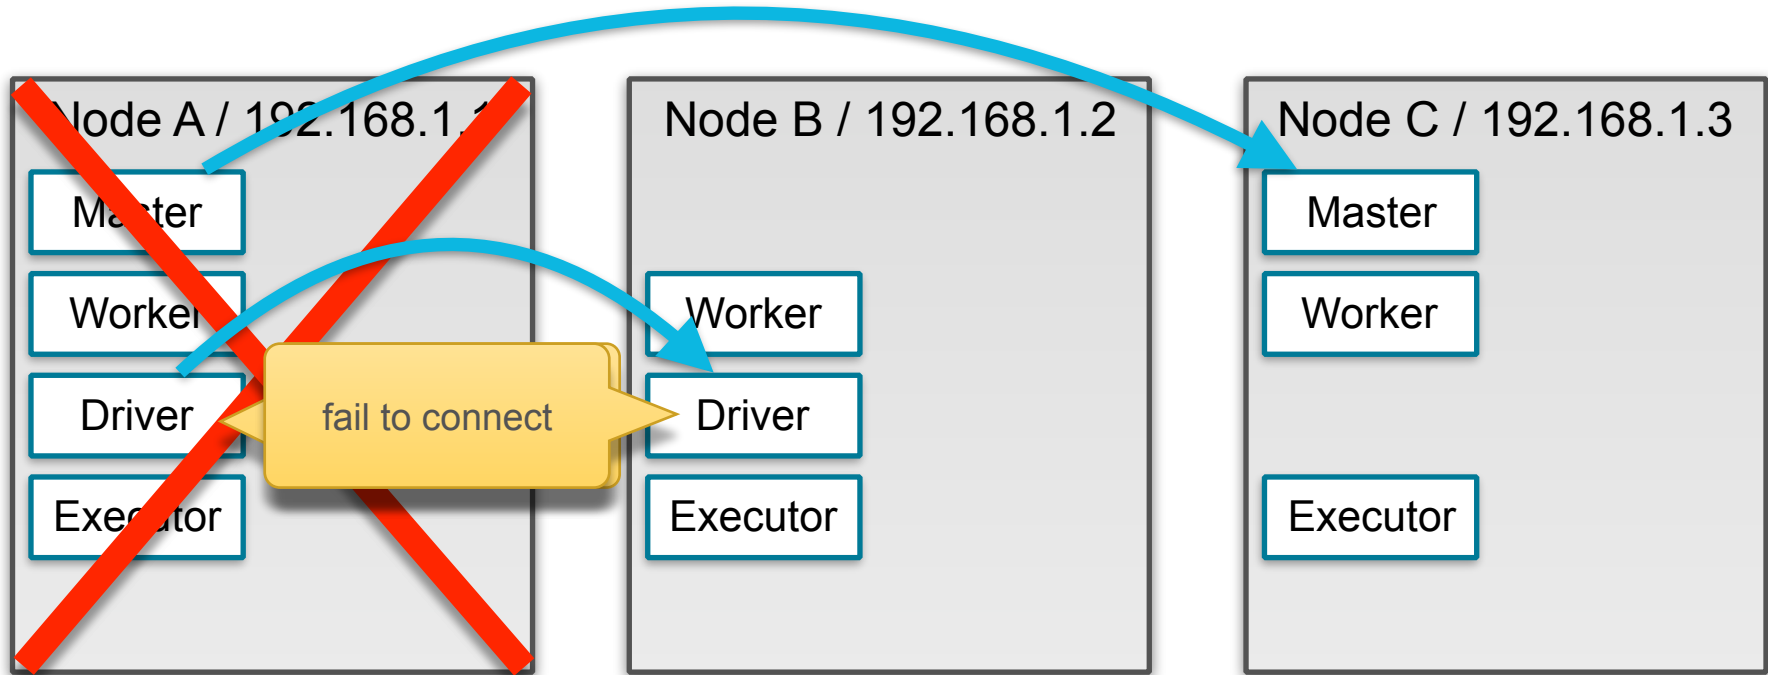
Spark specific settings unrelated to any particular node
Lets you overwrite Hadoop settings just for Spark clients

For example: `spark.hadoop.cassandra.host`

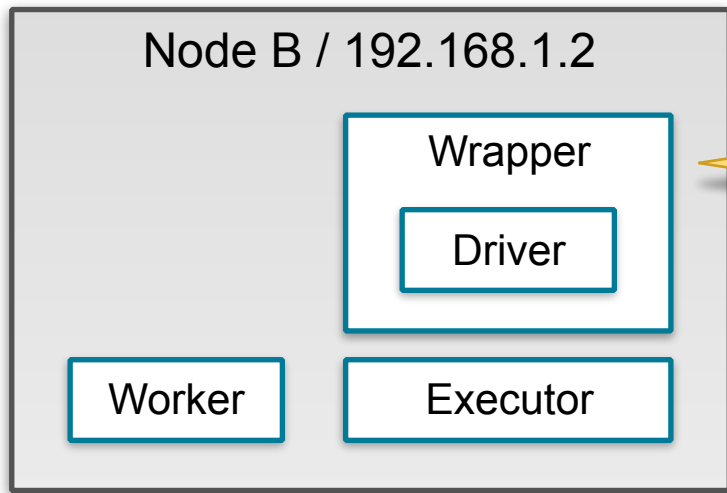
DSE bootstrapper retrieves dynamic settings from a node and merges them with local settings



Now, what would happen if a setting became out of date?



Whenever driver is restarted, DSE updates settings which denotes DSE nodes' addresses



Wrapper updates address which may change

`spark.hadoop.cassandra.host`
`spark.cassandra.connection.host`
`spark.hadoop.fs.default.name`
`spark.hadoop.fs.defaultFS`

These properties may become outdated

`spark.master`

Now, let's talk about confidentiality of settings

It is not guaranteed for command line arguments

```
$ sh test.sh --password=qaz123
```

Any user in the system can see this

```
$ ps
```

PID	TTY	TIME	CMD
8640	ttys000	0:00.03	-bash
8672	ttys000	0:00.00	sh test.sh --password=qaz123
8673	ttys000	0:00.00	sleep 60
8641	ttys001	0:00.01	-bash



This problem surfaces in Spark as well :-)

```
$ bin/spark-submit --deploy-mode cluster --master spark://ursus-major:7077 \  
--properties-file test.properties \  
--class test.Test test-app/test.jar
```

spark.driver.extraJavaOptions = -Ddse.token=12345
spark.cassandra.auth.password = qaz123

Visible in Driver logs, Worker logs, and processes list

```
Launch Command: "/Library/Java/JavaVirtualMachines/jdk-1.8.0_45.jdk/Contents/Home/bin/java" "-cp"  
"/Users/jlewandowski/Downloads/spark-1.6.2-bin-hadoop2.6/conf:/Users/jlewandowski/Downloads/spark-1.6.2-bin-  
hadoop2.6/lib/spark-assembly-1.6.2-hadoop2.6.0.jar:/Users/jlewandowski/Downloads/spark-1.6.2-bin-  
hadoop2.6/lib/datanucleus-api-jdo-3.2.6.jar:/Users/jlewandowski/Downloads/spark-1.6.2-bin-hadoop2.6/lib/datanucleus-core-  
3.2.10.jar:/Users/jlewandowski/Downloads/spark-1.6.2-bin-hadoop2.6/lib/datanucleus-rdbms-3.2.9.jar" "-Xmx1024M" "-  
Xmx1024M" "-Dspark.cassandra.auth.password=qaz123" "-Dakka.loglevel=WARNING" "-Dspark.driver.extraJavaOptions=-  
Ddse.token=12345" "-Dspark.deploy.supervise=false" "-Dspark.submit.deployMode=cluster" "-Dspark.master=spark://ursus-  
major:7077" "-Dspark.rpc.askTimeout=10" "-Dspark.app.name=test.Test" "-  
Dspark.jars=file:/Users/jlewandowski/Downloads/spark-1.6.2-bin-hadoop2.6/test-app/test.jar" "-Ddse.token=12345"  
"org.apache.spark.deploy.worker.DriverWrapper" "spark://Worker@10.0.0.16:62119" "/Users/jlewandowski/Downloads/spark-  
1.6.2-bin-hadoop2.6/work/driver-20160826125838-0000/test.jar" "test.Test"
```

Fortunately DSE Spark passes system properties in a safer way

java options for Driver or Executor

```
ursus-major:driver-20160904114913-0000 jlewandowski$ ls
total 96
driver-20160904114913-0000 wheel 2016 Sep 4 11:49
```

Bytes 0 - 30438 of 30438

```
Launch Command: "/Library/Java/JavaVirtualMachines/jdk1.8.0_45.jdk/Contents/Home/bin/java" "-cp" ... "-Xms1024M" "-Xmx1024M" "-Dguice_include_stack_traces=OFF"
"-Dakka.loglevel=WARNING" "-Dlogback.configurationFile=/Users/jlewandowski/Projects/DataStax/bdp1/resources/spark/conf/logback-spark.xml" "-
Dderby.stream.error.method=com.datastax.bdp.derby.LogbackBridge.getLogger" "org.apache.spark.DseSecureRunner" "org.apache.spark.deploy.worker.DriverWrapper"
"spark://Worker@127.0.0.1:53300" "/var/lib/spark/worker/driver-20160904114913-0000/test.jar" "test.Test"
```

```
[31mWARN [0;39m [32m2016-09-04 11:49:15,288[0;39m com.datastax.driver.core.NettyUtil: Found Netty's native epoll transport, but not running on linux-based
operating system. Using NIO instead.
```

```
[31mWARN [0;39m [32m2016-09-04 11:49:15,934[0;39m org.apache.hadoop.util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using
builtin-java classes where applicable
```

```
[Stage 0:> (0 + 0) / 10]
[Stage 0:> (0 + 5) / 10]
[Stage 0:> (0 + 7) / 10]
```

But this is not enough — system properties and configuration are still exposed in Driver UI

```
$ bin/spark-shell --conf spark.cassandra.auth.password=qaz123 \  
--driver-java-options=-Ddse.token=12345
```

Spark 1.6.2		Jobs	Stages	Storage	Environment	Executors	Spark shell application UI
...							
spark.cassandra.auth.conf.factory	com.datastax.bdp.spark.DseAuthConfFactory						
spark.cassandra.auth.password	qaz123						
spark.cassandra.connection.factory	com.datastax.bdp.spark.DseCassandraConnectionF						
...							
spark.cassandra.sql.pushdown.additionalClasses	org.apache.spark.sql.cassandra.DsePredicateRules						
spark.driver.extraJavaOptions	'-Ddse.token=12345'						
spark.driver.host	127.0.0.1						

DSE Spark addresses this problem too!

`spark.ui.confidentialKeys=password,token`

Just set this property in your Spark configuration



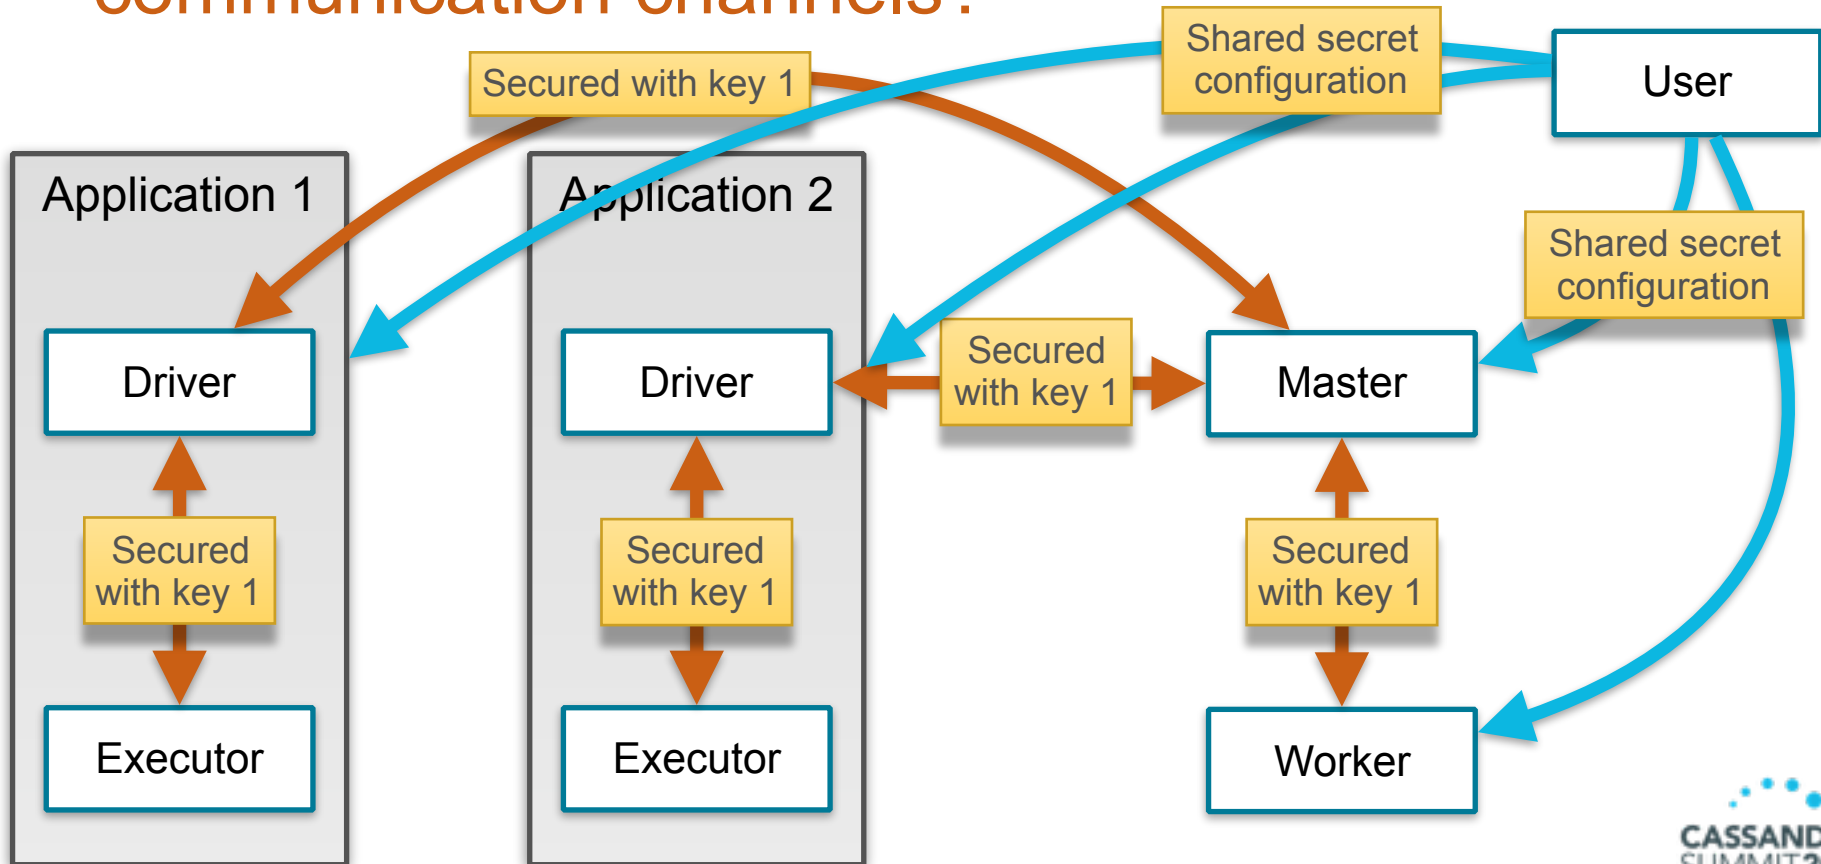
...

spark.cassandra.auth.conf.factory	com.datastax.bdp.spark.DseAuthConfFactory
spark.cassandra.auth.password	*****
spark.cassandra.connection.factory	com.datastax.bdp.spark.DseCassandraConnectionFactory

...

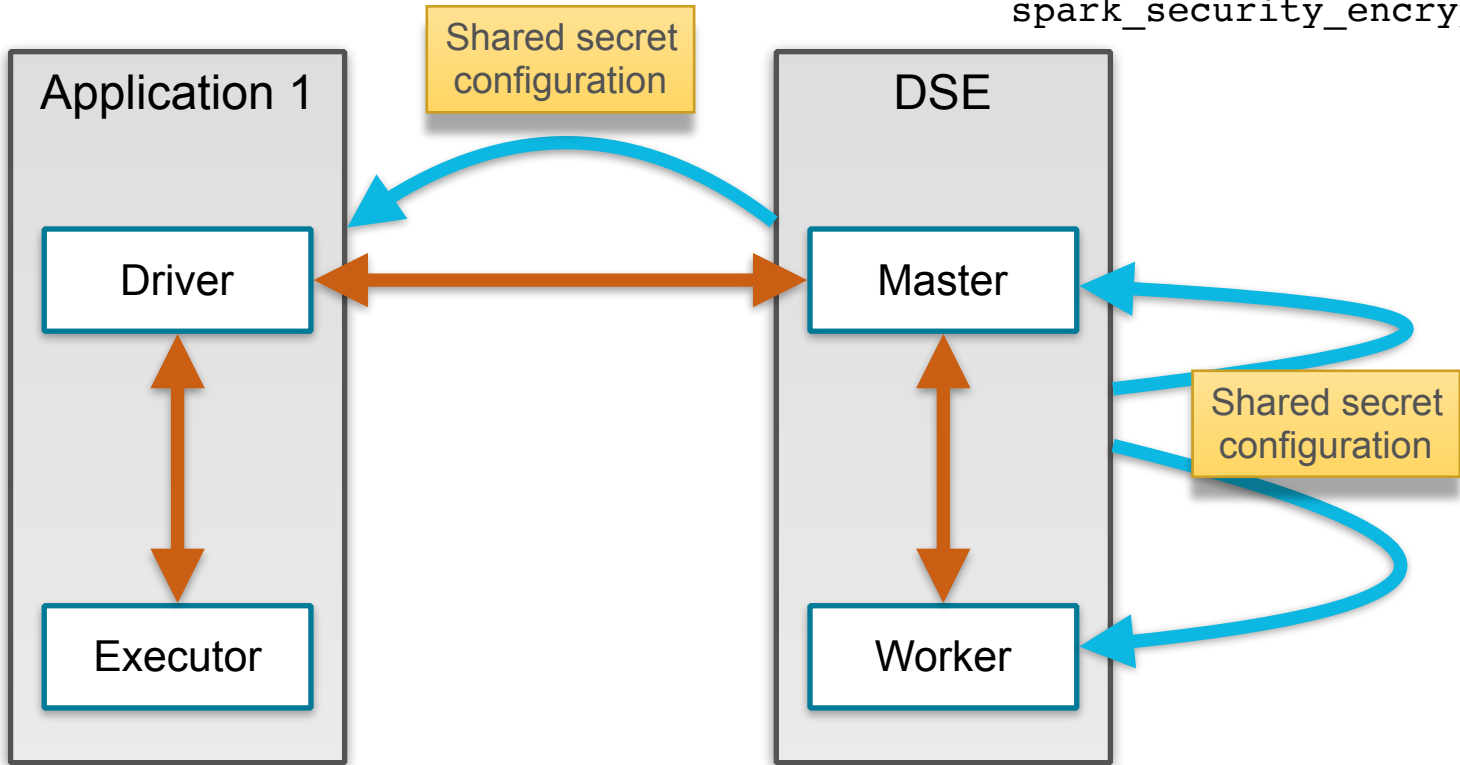
dse.system_memory_in_mb	16384
dse.token	*****
file.encoding	UTF-8

What about confidentiality of Spark communication channels?

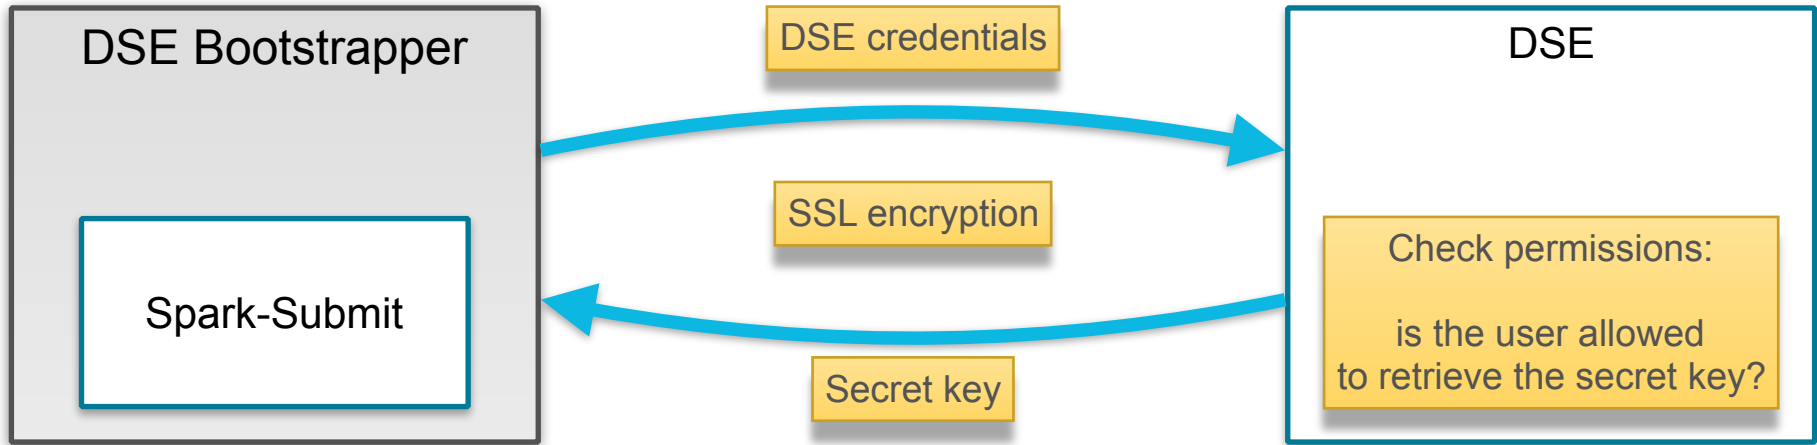


DSE manages shared secret for Spark

```
spark_security_enabled: true  
spark_security_encryption_enabled: true
```



Secret key is provided only to authenticated and authorized users




So, you have your DSE cluster running You want to use it from you Macbook...

DSE installation on unconfigured remote node



Which usually ends up
calling tech support wizards...



DSE can export configuration into a single file ... and import it to any other DSE installation

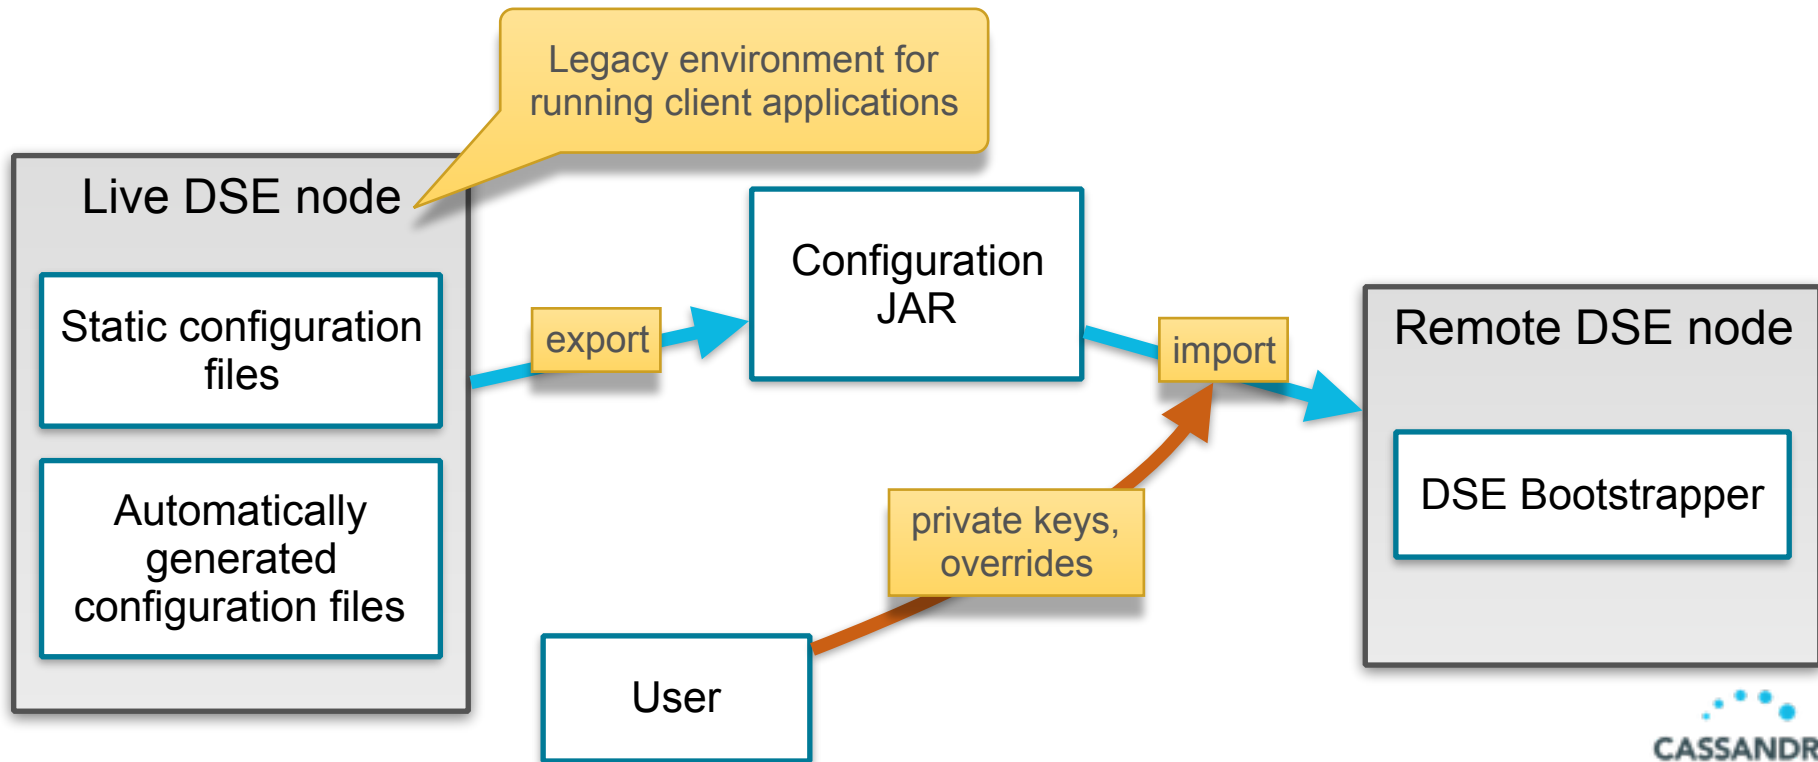
DSE installation
Local running DSE node

```
$ dse client-tool configuration export config.jar
```

DSE installation
Remote, unconfigured DSE node

```
$ dse client-tool configuration import config.jar
```

Import tool does similar job as DSE instance does on startup



So you may still want to call tech support...



...but probably for
a *slightly* different reason



Thank you for your attention

May DSE be with you

- * It configures your client applications automatically
- * It is a provider of HA for Spark and handles even tricky problems
- * It provides tools to configure DSE installations on remote nodes
- * It improves confidentiality of configuration settings
- * It manages distribution of shared secret automatically



Special thanks to Russell Spitzer, Rocco Varela, Piotr Kołaczkowski and the rest of DSE Analytics Team