



Troubleshooting Cassandra

J.B. Langston, Principal Support Engineer

Slides: <http://goo.gl/SdAzOJ>

Troubleshooting Philosophy



Troubleshooting Process

1 *Determine which nodes have problems*

2 *Examine bottlenecks*

3 *Find and understand errors*

4 *Ask what changed*

5 *Determine root cause*

6 *Take corrective action*

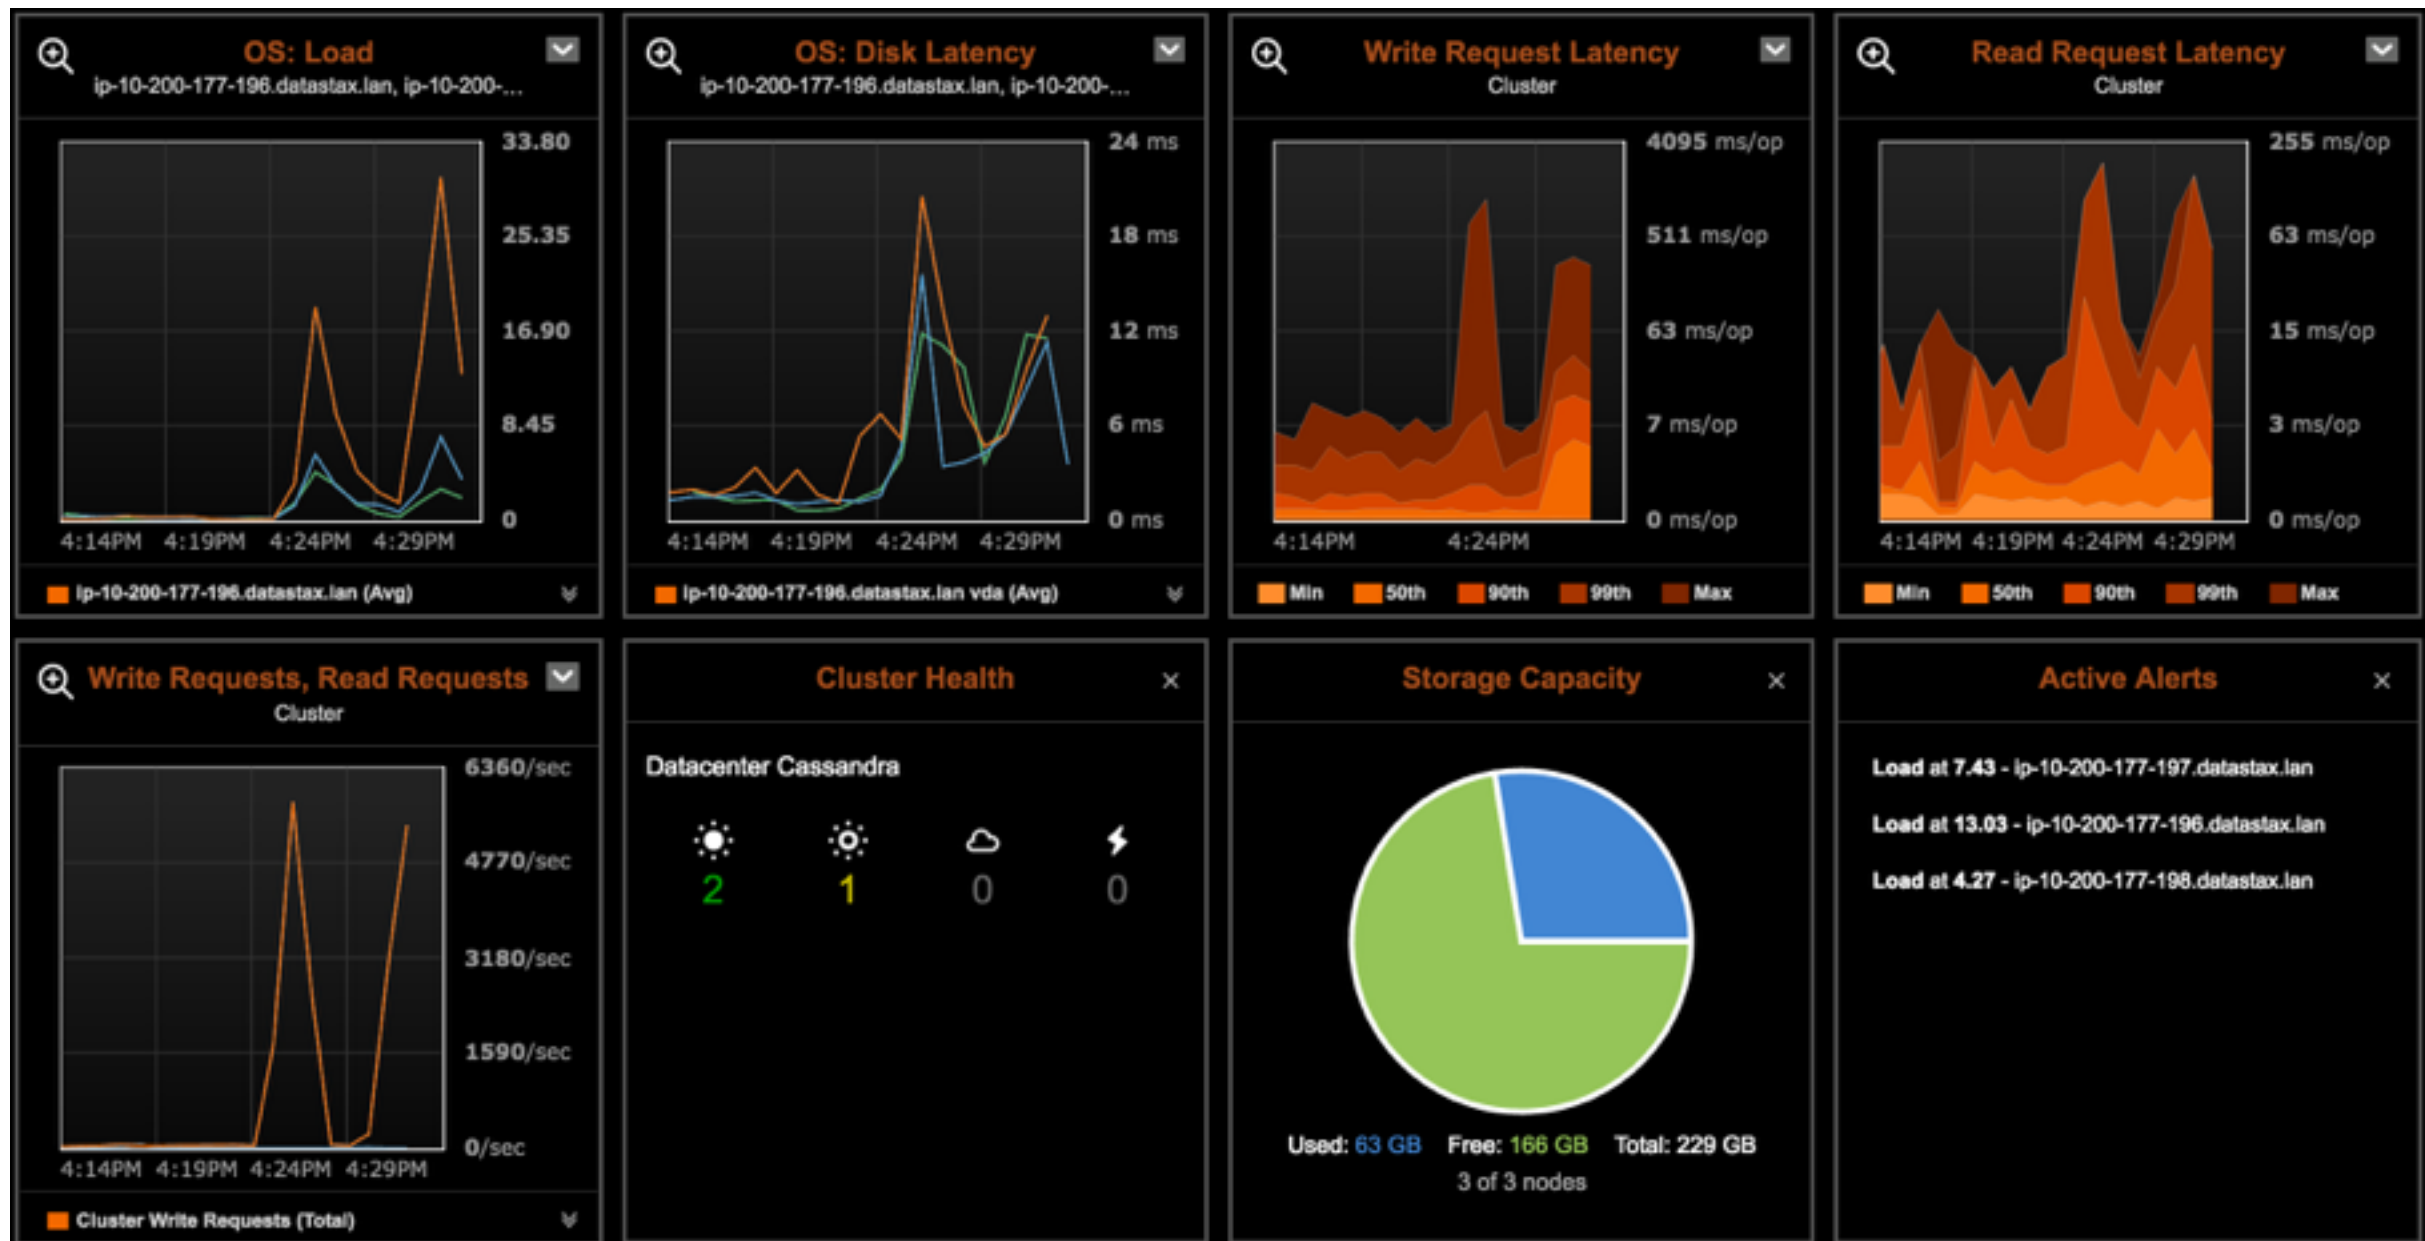
What changed?

- **Did it work before?**
- **Does it work in another environment?**
- **What's different?**
 - Settings
 - Application Code
 - Read/Write Load
 - Data Volume
 - Hardware
 - Network
- **Did you upgrade?**
 - Cassandra
 - Kernel
 - JVM
 - Driver
- **What metrics changed?**
 - OpsCenter
 - Graphite, etc.
- **Change one thing at a time!**
- **Use configuration management**

Tools of the Trade



OpsCenter Metrics



Now with Graphite Integration! (since 6.0)

OpsCenter Diagnostic Tarball

- **Logs**

- system.log
- debug.log
- Spark logs
- OpsCenter logs

- **Configuration**

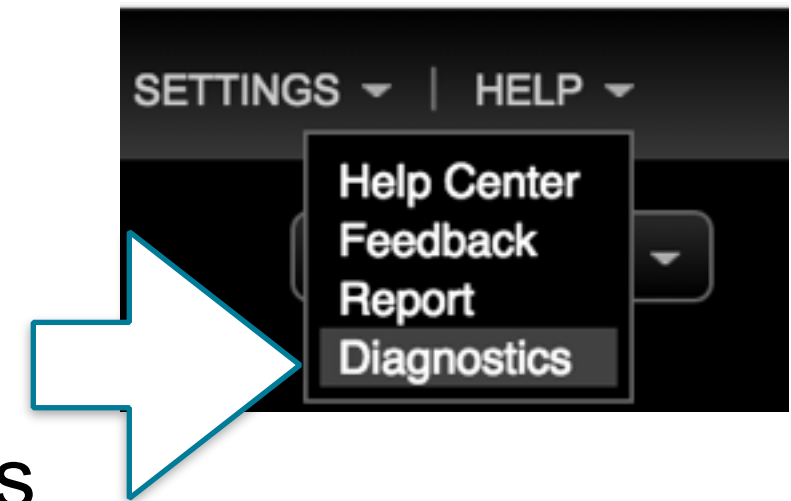
- cassandra.yaml
- cassandra-env.sh
- dse.yaml
- OpsCenter conf/yaml files

- **Schema**

- **Nodetool**

- status
- tablestats
- tpstats
- describecluster
- netstats
- etc...

- **OS metrics**



Cassandra Logs

Default Location

/var/log/cassandra/system.log - INFO and higher
 /var/log/cassandra/debug.log - all enabled messages

Since 3.0 many system.log messages are now in debug.log!

Configuration File

/etc/dse/cassandra/logback.xml

Basic Format

| Level | Thread Type & ID | Date & Time | Source File | Line No. |
|-------|--------------------------|-------------------------|------------------------|----------|
| INFO | [CompactionExecutor:155] | 2015-02-13 02:18:40,986 | CompactionTask.java | :287 |
| WARN | [GossipTasks:1] | 2015-02-17 19:47:37,331 | Gossiper.java | :648 |
| ERROR | [AntiEntropySessions:1] | 2015-02-17 20:32:11,959 | CassandraDaemon.java | :199 |
| DEBUG | [OptionalTasks:1] | 2015-02-20 11:29:14,056 | ColumnFamilyStore.java | :298 |

Configuring Logging Levels

\$ **nodetool getlogginglevels**

show current levels

| Logger Name | Log Level |
|--|-----------|
| ROOT | INFO |
| DroppedAuditEventLogger | INFO |
| SLF4JAuditWriter | INFO |
| com.cryptsoft | OFF |
| com.datastax.bdp.search.solr.metrics.MetricsWriteEventListener | DEBUG |
| com.thinkaurelius.thrift | ERROR |
| org.apache.cassandra | DEBUG |
| org.apache.lucene.index | INFO |
| org.apache.solr.core.CassandraSolrConfig | WARN |
| org.apache.solr.core.RequestHandlers | WARN |
| org.apache.solr.core.SolrCore | WARN |
| org.apache.solr.handler.component | WARN |
| org.apache.solr.search.SolrIndexSearcher | WARN |
| org.apache.solr.update | WARN |

Configuring Logging Levels

```
$ nodetool getlogginglevels
```

```
# show current levels
```

| Logger Name | Log Level |
|--|-----------|
| ROOT | INFO |
| DroppedAuditEventLogger | INFO |
| SLF4JAuditWriter | INFO |
| com.cryptsoft | OFF |
| com.datastax.bdp.search.solr.metrics.MetricsWriteEventListener | DEBUG |
| com.thinkaurelius.thrift | ERROR |
| org.apache.cassandra | DEBUG |
| org.apache.lucene.index | INFO |
| org.apache.solr.core.CassandraSolrConfig | WARN |
| org.apache.solr.core.RequestHandlers | WARN |
| org.apache.solr.core.SolrCore | WARN |
| org.apache.solr.handler.component | WARN |
| org.apache.solr.search.SolrIndexSearcher | WARN |
| org.apache.solr.update | WARN |

```
# all of cassandra
```

```
$ nodetool setlogginglevel org.apache.cassandra TRACE
```

Configuring Logging Levels

```
$ nodetool getlogginglevels
```

| Logger Name | Log Level |
|--|-----------|
| ROOT | INFO |
| DroppedAuditEventLogger | INFO |
| SLF4JAuditWriter | INFO |
| com.cryptsoft | OFF |
| com.datastax.bdp.search.solr.metrics.MetricsWriteEventListener | DEBUG |
| com.thinkaurelius.thrift | ERROR |
| org.apache.cassandra | DEBUG |
| org.apache.lucene.index | INFO |
| org.apache.solr.core.CassandraSolrConfig | WARN |
| org.apache.solr.core.RequestHandlers | WARN |
| org.apache.solr.core.SolrCore | WARN |
| org.apache.solr.handler.component | WARN |
| org.apache.solr.search.SolrIndexSearcher | WARN |
| org.apache.solr.update | WARN |

```

# show current levels

# all of cassandra
$ nodetool setlogginglevel org.apache.cassandra TRACE

# entire package
$ nodetool setlogginglevel org.apache.cassandra.gms TRACE

```

Configuring Logging Levels

```
$ nodetool getlogginglevels
```

```
# show current levels
```

| Logger Name | Log Level |
|--|-----------|
| ROOT | INFO |
| DroppedAuditEventLogger | INFO |
| SLF4JAuditWriter | INFO |
| com.cryptsoft | OFF |
| com.datastax.bdp.search.solr.metrics.MetricsWriteEventListener | DEBUG |
| com.thinkaurelius.thrift | ERROR |
| org.apache.cassandra | DEBUG |
| org.apache.lucene.index | INFO |
| org.apache.solr.core.CassandraSolrConfig | WARN |
| org.apache.solr.core.RequestHandlers | WARN |
| org.apache.solr.core.SolrCore | WARN |
| org.apache.solr.handler.component | WARN |
| org.apache.solr.search.SolrIndexSearcher | WARN |
| org.apache.solr.update | WARN |

```
# all of cassandra
```

```
$ nodetool setlogginglevel org.apache.cassandra TRACE
```

```
# entire package
```

```
$ nodetool setlogginglevel org.apache.cassandra.gms TRACE
```

```
# specific class
```

```
$ nodetool setlogginglevel org.apache.cassandra.service.GCInspector TRACE
```

Configuring Logging Levels

```
$ nodetool getlogginglevels                                # show current levels
```

| Logger Name | Log Level |
|--|-----------|
| ROOT | INFO |
| DroppedAuditEventLogger | INFO |
| SLF4JAuditWriter | INFO |
| com.cryptsoft | OFF |
| com.datastax.bdp.search.solr.metrics.MetricsWriteEventListener | DEBUG |
| com.thinkaurelius.thrift | ERROR |
| org.apache.cassandra | DEBUG |
| org.apache.lucene.index | INFO |
| org.apache.solr.core.CassandraSolrConfig | WARN |
| org.apache.solr.core.RequestHandlers | WARN |
| org.apache.solr.core.SolrCore | WARN |
| org.apache.solr.handler.component | WARN |
| org.apache.solr.search.SolrIndexSearcher | WARN |
| org.apache.solr.update | WARN |

```

# all of cassandra
$ nodetool setlogginglevel org.apache.cassandra TRACE

# entire package
$ nodetool setlogginglevel org.apache.cassandra.gms TRACE

# specific class
$ nodetool setlogginglevel org.apache.cassandra.service.GCInspector TRACE

# reset to defaults
$ nodetool setlogginglevel

```

Overall status



Up or down?

```
$ nodetool status
```

```
Datacenter: Cassandra
```

```
=====
```

```
Status=Up/Down
```

```
|/ State=Normal/Leaving/Joining/Moving
```

| -- | Address | Load | Tokens | Owns | Host ID | Rack |
|----|----------------|---------|--------|------|--------------------------------------|-------|
| UN | 10.200.177.196 | 3.45 GB | 1 | ? | 108af27a-43d8-4814-b617-f8f93ba2bb0e | rack1 |
| UN | 10.200.177.197 | 3.45 GB | 1 | ? | 432bc964-3cd3-4784-9ab7-d7a4a9e063b6 | rack1 |
| UN | 10.200.177.198 | 3.45 GB | 1 | ? | 3c467f89-7cce-485f-bb16-dd782c9a84ec | rack1 |

Note: Non-system keyspaces don't have the same replication settings, effective ownership information meaningless

Lots of nodes? Use grep...

```
$ nodetool status | grep DN
```

| | | | | | | |
|----|----------------|---------|---|---|--------------------------------------|-------|
| DN | 10.200.177.197 | 3.45 GB | 1 | ? | 432bc964-3cd3-4784-9ab7-d7a4a9e063b6 | rack1 |
|----|----------------|---------|---|---|--------------------------------------|-------|

OpsCenter Ring View



OpsCenter Alerts

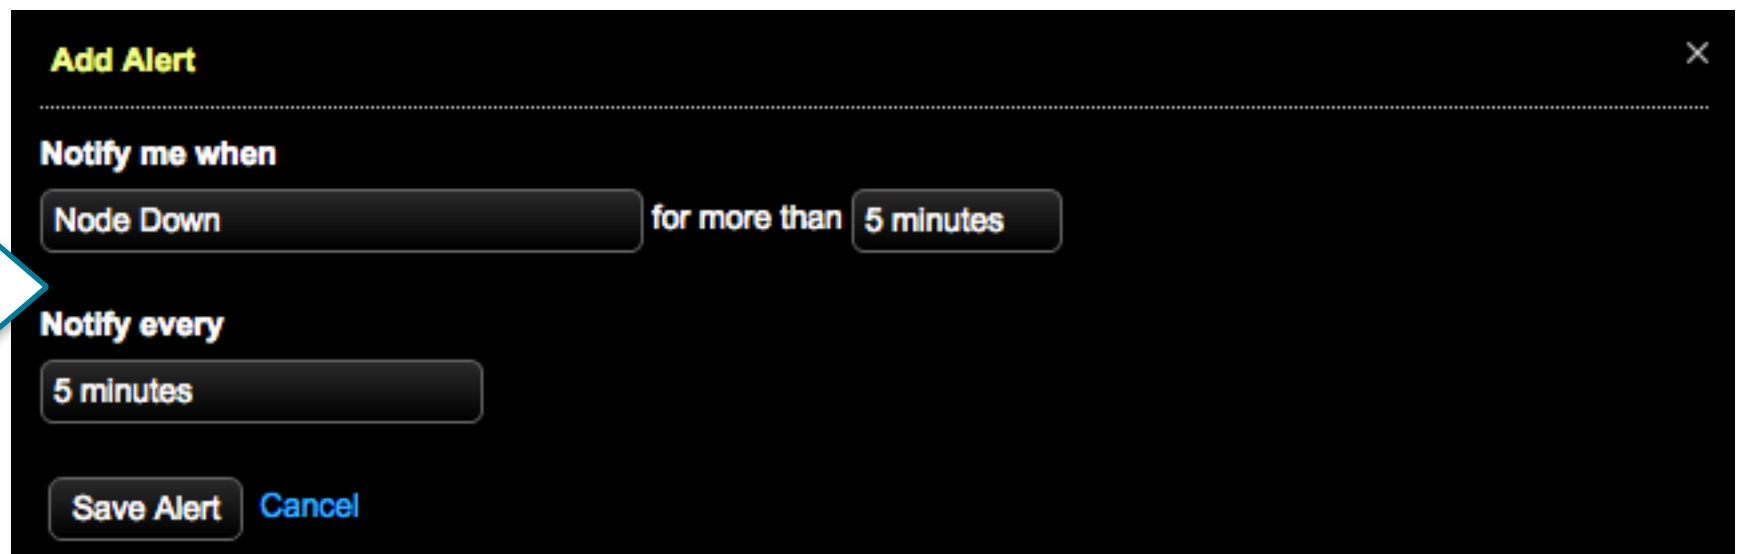
Alert on:

- Nodes Down
- Metrics Exceeding Limit
- Data Balance Issues
- more...

Send Alerts to:

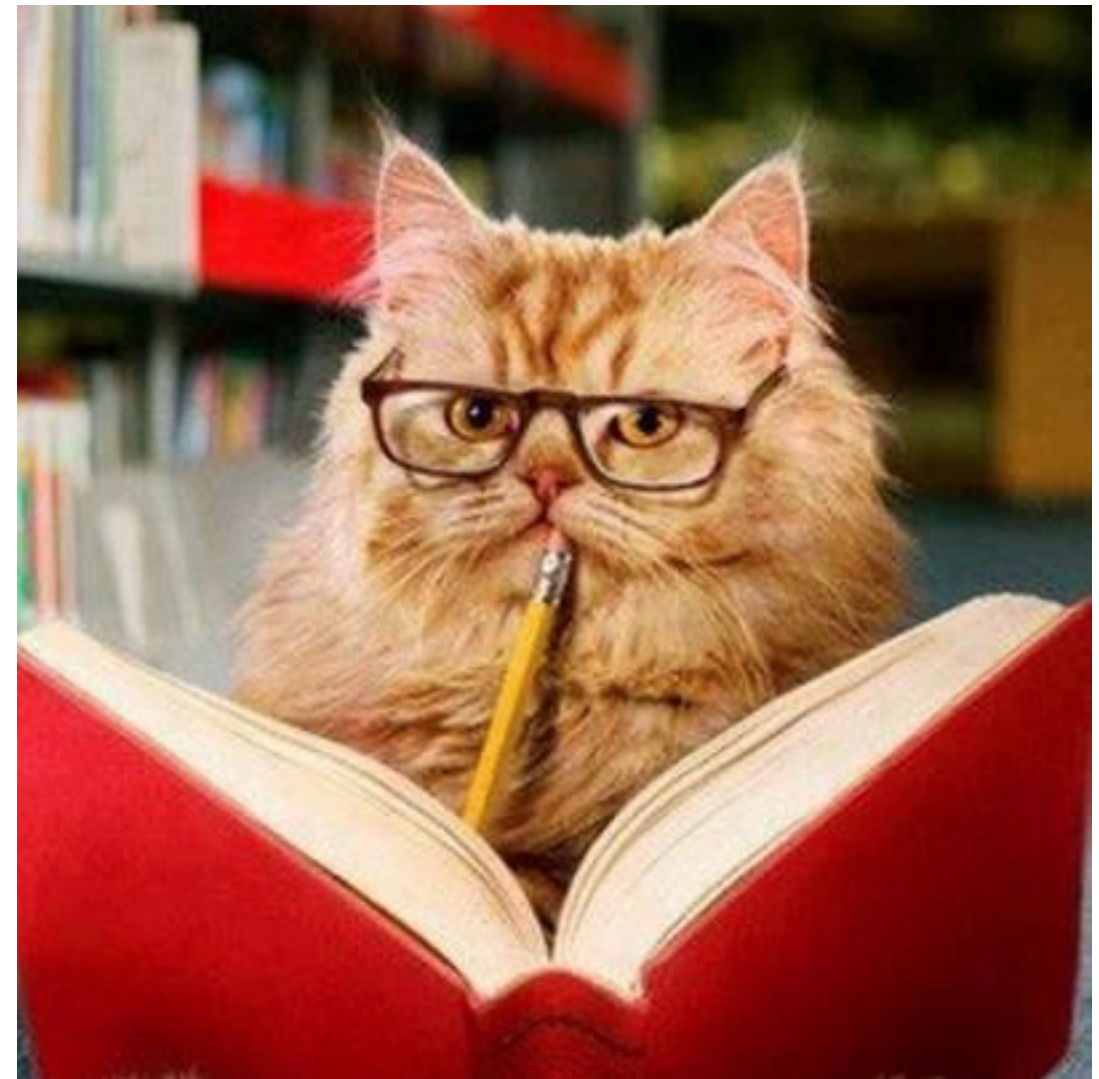
- Email
- HTTP API
- SNMP
- Ex: Hipchat, Slack, Pagerduty

Alerts > Manage Alerts > Add Alert



The screenshot shows a dark-themed 'Add Alert' dialog box. At the top, the title 'Add Alert' is in yellow. Below it, the section 'Notify me when' contains a dropdown menu with 'Node Down' selected and a text input field with '5 minutes'. The section 'Notify every' contains a text input field with '5 minutes'. At the bottom, there are two buttons: 'Save Alert' and 'Cancel'.

Reads and Writes



Latency

Coordinator Latency

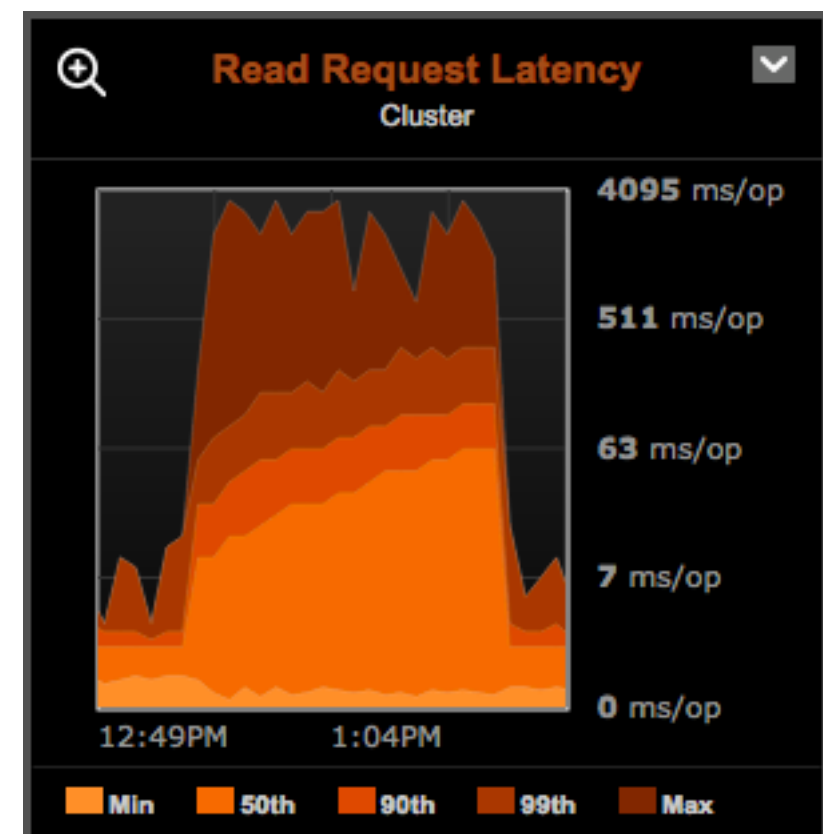
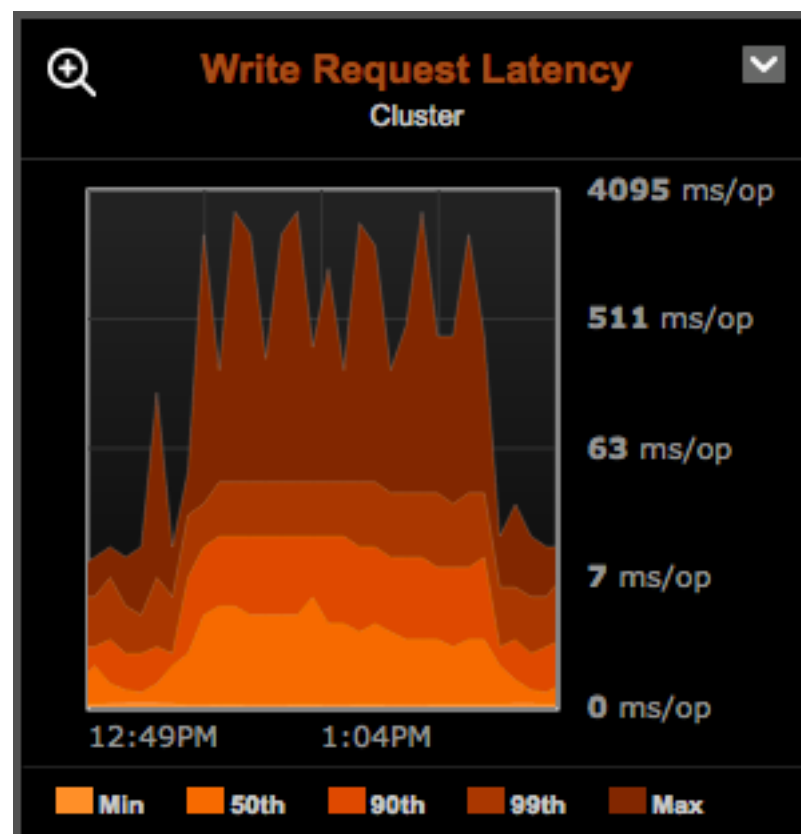
```
$ nodetool proxyhistograms
```

```
proxy histograms
```

| Percentile | Read Latency (micros) | Write Latency (micros) | Range Latency (micros) |
|------------|--------------------------|---------------------------|---------------------------|
| 50% | 1358.10 | 654.95 | 1131.75 |
| 75% | 1955.67 | 1358.10 | 2346.80 |
| 95% | 5839.59 | 4055.27 | 4866.32 |
| 98% | 10090.81 | 7007.51 | 8409.01 |
| 99% | 12108.97 | 10090.81 | 8409.01 |
| Min | 51.01 | 29.52 | 545.79 |
| Max | 1155149.91 | 4966933.18 | 14530.76 |

Common Causes

- Large batches
- Table scans
- Secondary indexes
- IN clause
- Consistency levels
- Network latency
- Cross-DC traffic
- Garbage collection



Local Latency

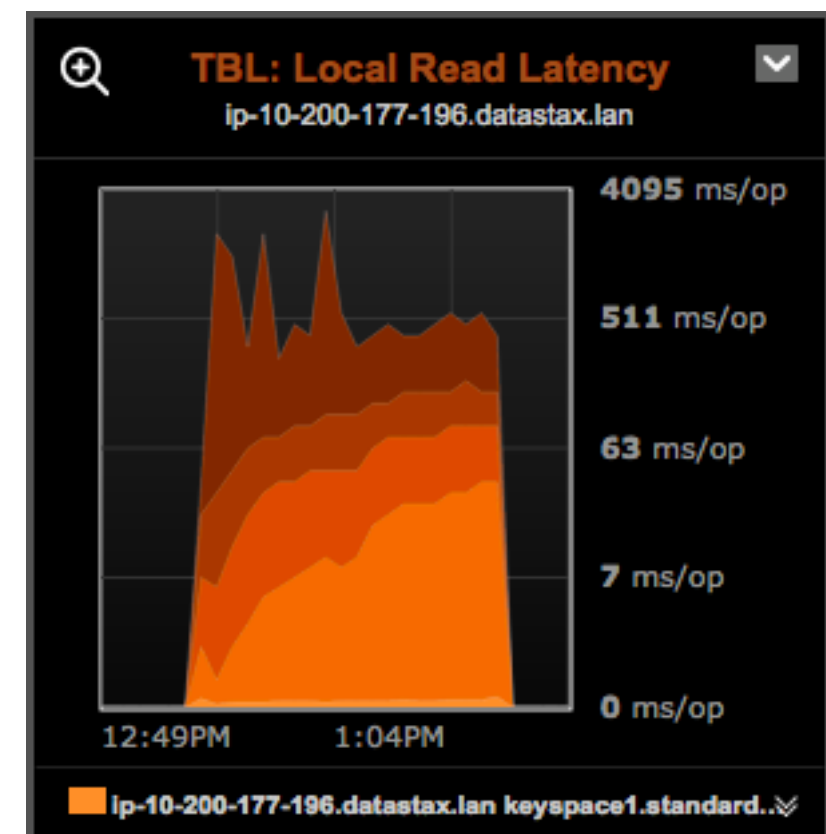
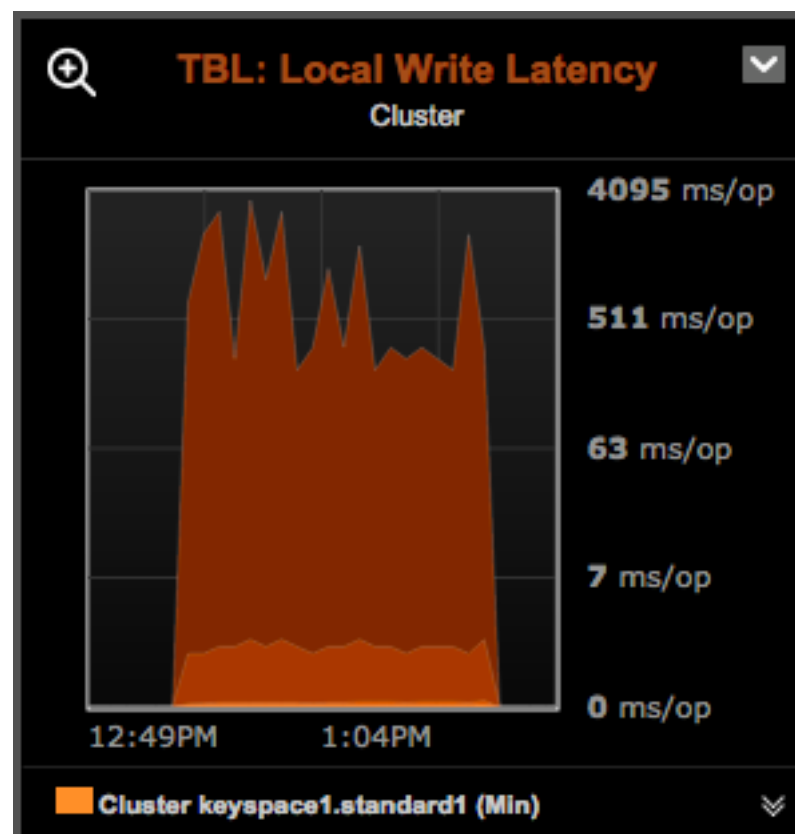
```
$ nodetool tablehistograms keyspace1 standard1
```

keyspace1/standard1 histograms

| Percentile | SSTables | Write Latency (micros) | Read Latency (micros) | Partition Size (bytes) | Cell Count |
|------------|----------|---------------------------|--------------------------|---------------------------|------------|
| 50% | 4.00 | 24.60 | 1629.72 | 14237 | 5 |
| 75% | 4.00 | 35.43 | 3379.39 | 17084 | 5 |
| 95% | 5.00 | 61.21 | 7007.51 | 20501 | 5 |
| 98% | 6.00 | 152.32 | 12108.97 | 20501 | 5 |
| 99% | 6.00 | 454.83 | 14530.76 | 20501 | 5 |
| Min | 1.00 | 4.77 | 182.79 | 536 | 5 |
| Max | 8.00 | 4139110.98 | 1155149.91 | 24601 | 5 |

Common Causes

- I/O bottlenecks
- Lots of sstables
- Garbage collection
- Insufficient threads
- Insufficient cache
- Bloom filter false positive



Read/Write Volume

```
$ nodetool tablestats keyspace1
```

Keyspace: keyspace1

Read Count: 71035

Read Latency: 2.6125789258816075 ms.

Write Count: 2604424

Write Latency: 0.07452477393273357 ms.

Pending Flushes: 0

Table: standard1

SSTable count: 10

Space used (live): 17634915179

Space used (total): 17634915179

Space used by snapshots (total): 0

Off heap memory used (total): 4471240

SSTable Compression Ratio: 0.0

Number of keys (estimate): 835234

Memtable cell count: 8958

Memtable data size: 51012431

Memtable off heap memory used: 0

Memtable switch count: 120

Local read count: 71035

Local read latency: NaN ms

Local write count: 2604424

Local write latency: 0.196 ms

Pending flushes: 0

Bloom filter false positives: 0

Bloom filter false ratio: 0.00000

Bloom filter space used: 3931968

Bloom filter off heap memory used: 3931888

Index summary off heap memory used: 539352

Compression metadata off heap memory used: 0

Compacted partition minimum bytes: 1917

Compacted partition maximum bytes: 9887

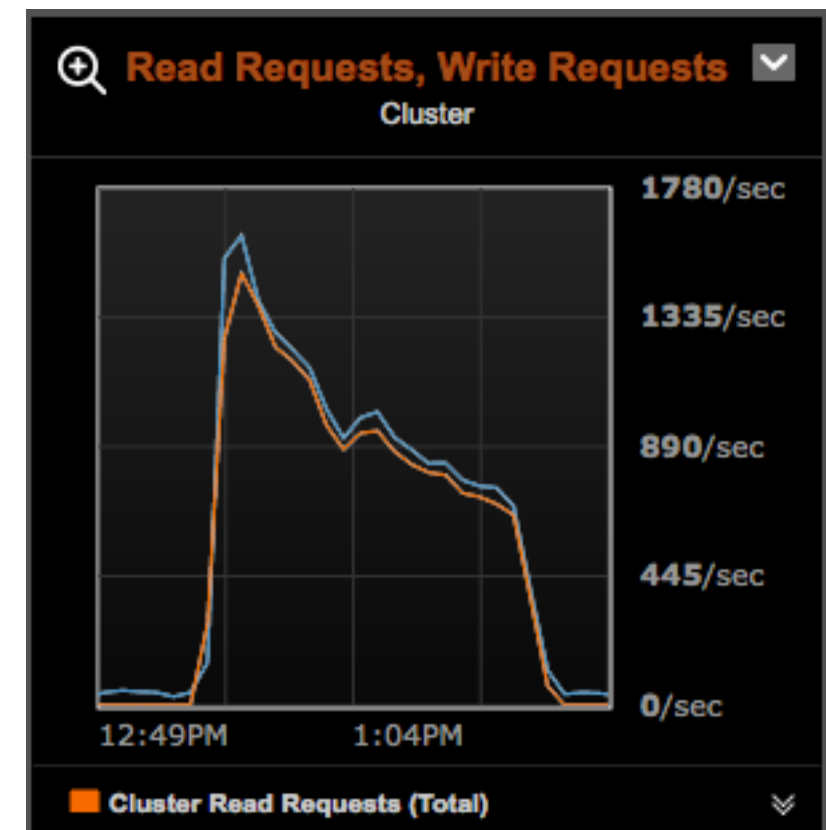
Compacted partition mean bytes: 6143

Average live cells per slice (last five minutes): NaN

Maximum live cells per slice (last five minutes): 0

Average tombstones per slice (last five minutes): NaN

Maximum tombstones per slice (last five minutes): 0



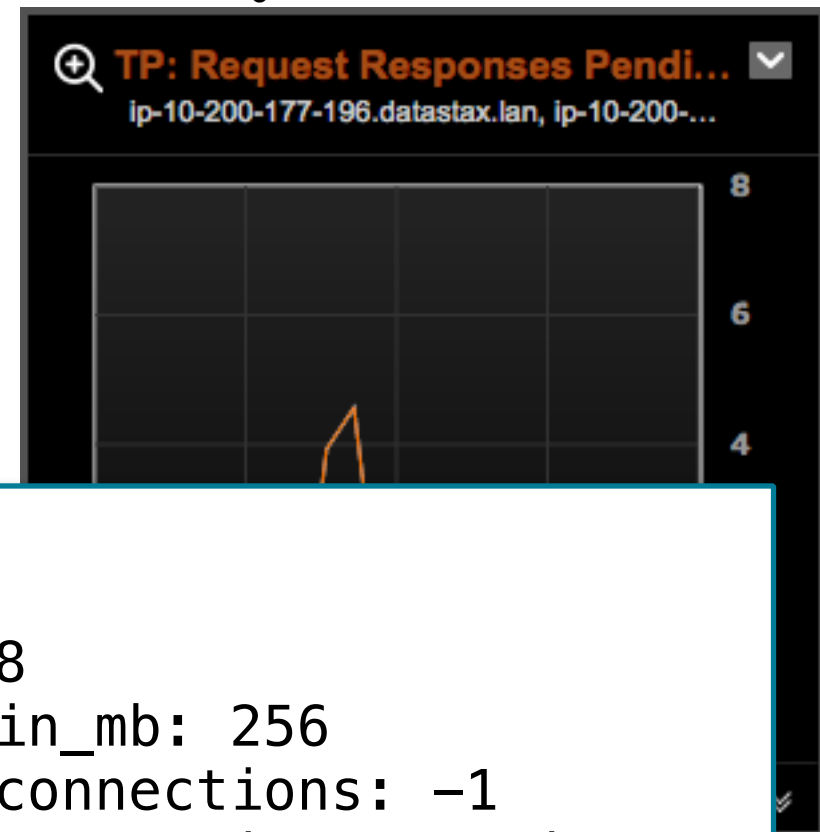
Thread Pools and Timeouts

Native Transport Requests

\$ nodetool tpstats

| Pool Name | Active | Pending | Completed | Blocked | All time blocked |
|----------------------------------|-----------|----------|----------------|----------|------------------|
| MutationStage | 0 | 34 | 1273902 | 0 | 0 |
| ViewMutationStage | 0 | 0 | 0 | 0 | 0 |
| ReadStage | 0 | 0 | 561050 | 0 | 0 |
| RequestResponseStage | 0 | 0 | 894363 | 0 | 0 |
| ReadRepairStage | 0 | 0 | 26007 | 0 | 0 |
| CounterMutationStage | 0 | 0 | 0 | 0 | 0 |
| MiscStage | 0 | 0 | 0 | 0 | 0 |
| CompactionExecutor | 1 | 1 | 4699 | 0 | 0 |
| MemtableReclaimMemory | 0 | 0 | 77 | 0 | 0 |
| PendingRangeCalculator | 0 | 0 | 3 | 0 | 0 |
| GossipStage | 0 | 0 | 21833 | 0 | 0 |
| SecondaryIndexManagement | 0 | 0 | 0 | 0 | 0 |
| HintsDispatcher | 0 | 0 | 87 | 0 | 0 |
| MigrationStage | 0 | 0 | 0 | 0 | 0 |
| MemtablePostFlush | 1 | 1 | 99 | 0 | 0 |
| ValidationExecutor | 0 | 0 | 0 | 0 | 0 |
| Sampler | 0 | 0 | 0 | 0 | 0 |
| MemtableFlushWriter | 1 | 1 | 77 | 0 | 0 |
| InternalResponseStage | 0 | 0 | 38112 | 0 | 0 |
| AntiEntropyStage | 0 | 0 | 0 | 0 | 0 |
| CacheCleanupExecutor | 0 | 0 | 0 | 0 | 0 |
| Native-Transport-Requests | 33 | 0 | 1186750 | 0 | |

| Message type | Dropped |
|------------------|---------|
| READ | 3227 |
| RANGE_SLICE | 0 |
| _TRACE | 0 |
| HINT | 0 |
| MUTATION | |
| COUNTER_MUTATION | |
| BATCH_STORE | |
| BATCH_REMOVE | |
| REQUEST_RESPONSE | |
| PAGED_RANGE | |
| READ_REPAIR | |



cassandra.yaml settings

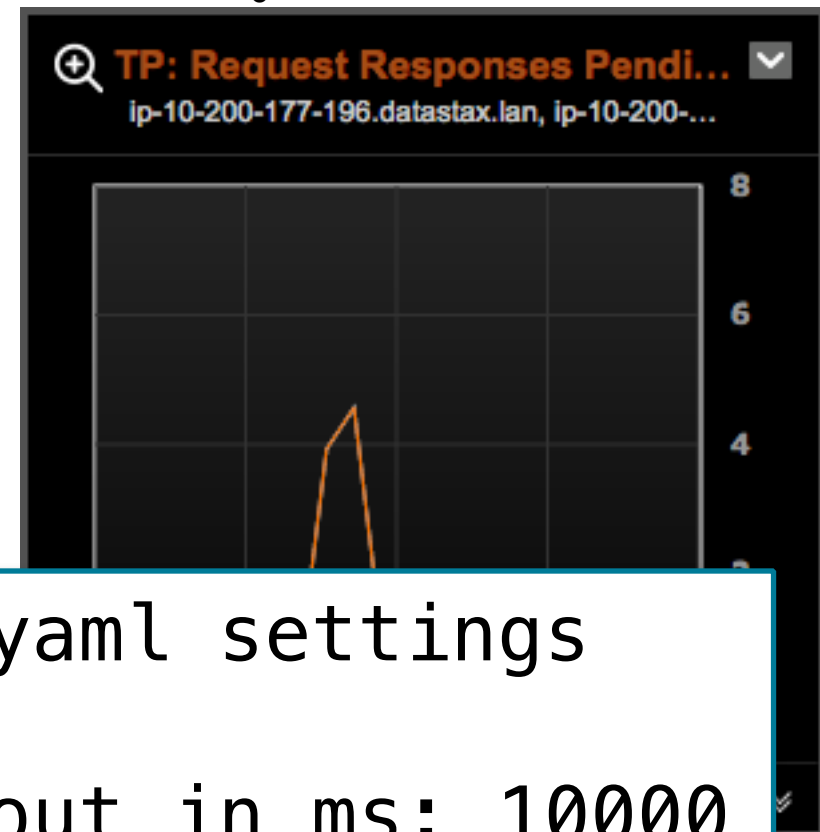
```
native_transport_max_threads: 128
native_transport_max_frame_size_in_mb: 256
native_transport_max_concurrent_connections: -1
native_transport_max_concurrent_connections_per_ip: -1
```


Request Response Threads

\$ nodetool tpstats

| Pool Name | Active | Pending | Completed | Blocked | All time blocked |
|-----------------------------|----------|----------|---------------|----------|------------------|
| MutationStage | 0 | 34 | 1273902 | 0 | 0 |
| ViewMutationStage | 0 | 0 | 0 | 0 | 0 |
| ReadStage | 0 | 0 | 561050 | 0 | 0 |
| RequestResponseStage | 0 | 0 | 894363 | 0 | 0 |
| ReadRepairStage | 0 | 0 | 26007 | 0 | 0 |
| CounterMutationStage | 0 | 0 | 0 | 0 | 0 |
| MiscStage | 0 | 0 | 0 | 0 | 0 |
| CompactionExecutor | 1 | 1 | 4699 | 0 | 0 |
| MemtableReclaimMemory | 0 | 0 | 77 | 0 | 0 |
| PendingRangeCalculator | 0 | 0 | 3 | 0 | 0 |
| GossipStage | 0 | 0 | 21833 | 0 | 0 |
| SecondaryIndexManagement | 0 | 0 | 0 | 0 | 0 |
| HintsDispatcher | 0 | 0 | 87 | 0 | 0 |
| MigrationStage | 0 | 0 | 0 | 0 | 0 |
| MemtablePostFlush | 1 | 1 | 99 | 0 | 0 |
| ValidationExecutor | 0 | 0 | 0 | 0 | 0 |
| Sampler | 0 | 0 | 0 | 0 | 0 |
| MemtableFlushWriter | 1 | 1 | 77 | 0 | 0 |
| InternalResponseStage | 0 | 0 | 38112 | 0 | 0 |
| AntiEntropyStage | 0 | 0 | 0 | 0 | 0 |
| CacheCleanupExecutor | 0 | 0 | 0 | 0 | 0 |
| Native-Transport-Requests | 33 | 0 | 1186750 | 0 | 0 |

| Message type | Dropped |
|-------------------------|-----------|
| READ | 3227 |
| RANGE_SLICE | 0 |
| _TRACE | 0 |
| HINT | 0 |
| MUTATION | 144 |
| COUNTER_MUTATION | 0 |
| BATCH_STORE | 0 |
| BATCH_REMOVE | 0 |
| REQUEST_RESPONSE | 10 |
| PAGED_RANGE | 0 |
| READ_REPAIR | 0 |



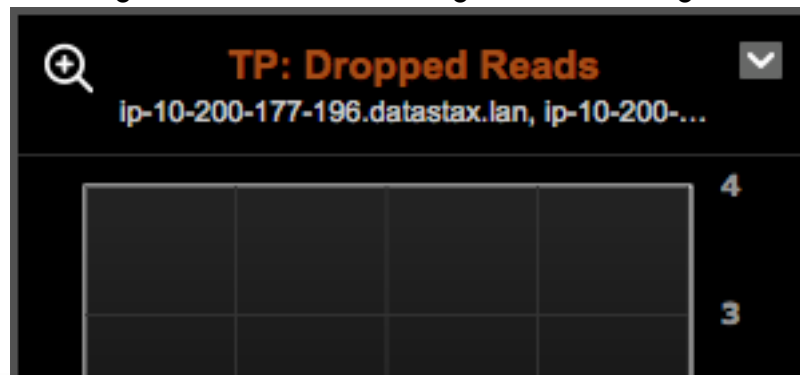
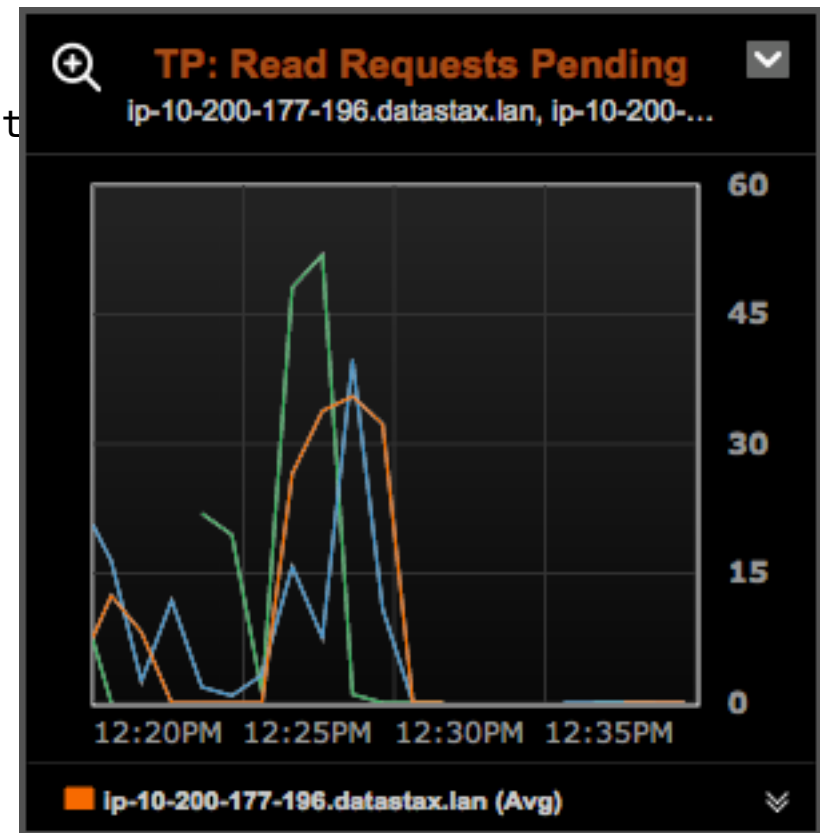
```
# cassandra.yaml settings
request_timeout_in_ms: 10000
```

Read Threads

\$ nodetool tpstats

| Pool Name | Active | Pending | Completed | Blocked | All t |
|---------------------------|-----------|-----------|---------------|----------|-------|
| MutationStage | 0 | 0 | 1214289 | 0 | |
| ViewMutationStage | 0 | 0 | 0 | 0 | |
| ReadStage | 32 | 49 | 508565 | 0 | |
| RequestResponseStage | 0 | 0 | 812868 | 0 | |
| ReadRepairStage | 0 | 0 | 22916 | 0 | |
| CounterMutationStage | 0 | 0 | 0 | 0 | |
| MiscStage | 0 | 0 | 0 | 0 | |
| CompactionExecutor | 0 | 0 | 4766 | 0 | |
| MemtableReclaimMemory | 0 | 0 | 77 | 0 | |
| PendingRangeCalculator | 0 | 0 | 3 | 0 | |
| GossipStage | 0 | 0 | 20991 | 0 | |
| SecondaryIndexManagement | 0 | 0 | 0 | 0 | |
| HintsDispatcher | 0 | 0 | 87 | 0 | |
| MigrationStage | 0 | 0 | 0 | 0 | |
| MemtablePostFlush | 0 | 0 | 99 | 0 | |
| ValidationExecutor | 0 | 0 | 0 | 0 | |
| Sampler | 0 | 0 | 0 | 0 | |
| MemtableFlushWriter | 0 | 0 | 0 | 0 | |
| InternalResponseStage | 0 | 0 | 0 | 0 | |
| AntiEntropyStage | 0 | 0 | 0 | 0 | |
| CacheCleanupExecutor | 0 | 0 | 0 | 0 | |
| Native-Transport-Requests | 24 | | | | |

| Message type | Dropped |
|--------------------|-------------|
| READ | 8802 |
| RANGE_SLICE | 0 |
| _TRACE | 0 |
| HINT | 3 |
| MUTATION | 348 |
| COUNTER_MUTATION | 0 |
| BATCH_STORE | 0 |
| BATCH_REMOVE | 0 |
| REQUEST_RESPONSE | 0 |
| PAGED_RANGE | 0 |
| READ_REPAIR | 0 |



cassandra.yaml settings

concurrent_reads: 32

read_request_timeout_in_ms: 5000

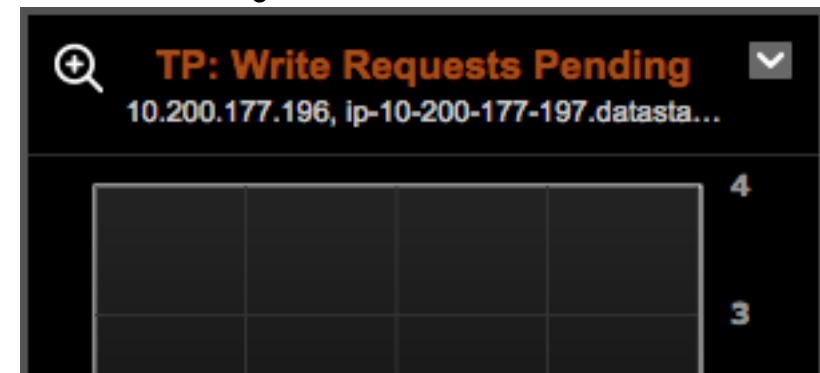
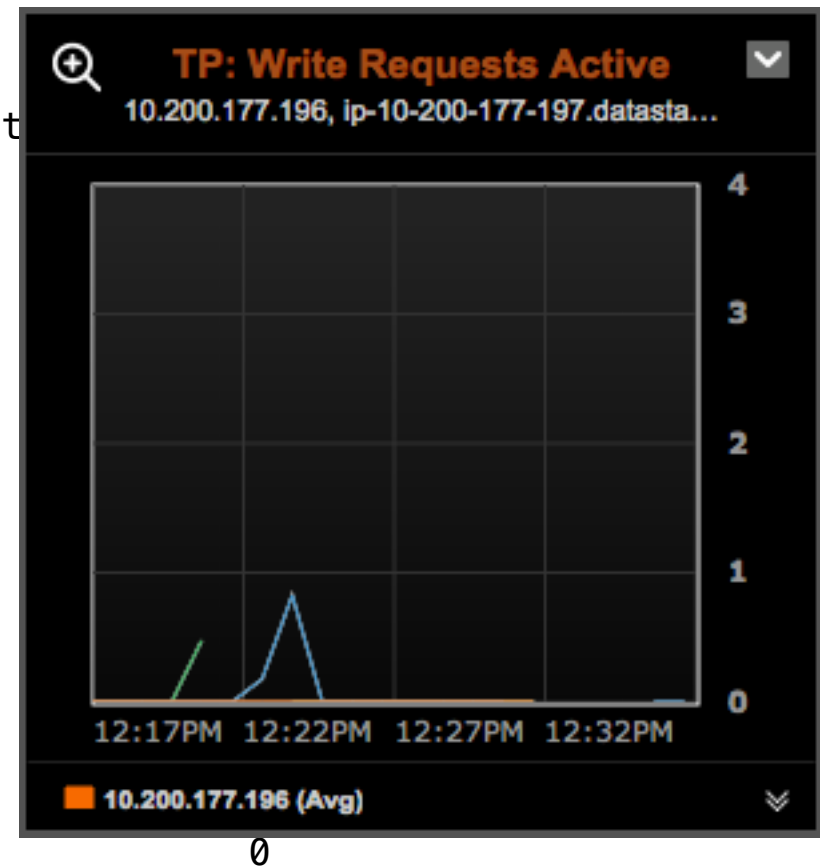
range_request_timeout_in_ms: 10000

Write Threads

\$ nodetool tpstats

| Pool Name | Active | Pending | Completed | Blocked | All t |
|-----------------------------|--------|---------|-----------|---------|-------|
| MutationStage | 0 | 34 | 1273902 | 0 | |
| ViewMutationStage | 0 | 0 | 0 | 0 | |
| ReadStage | 0 | 0 | 561050 | 0 | |
| RequestResponseStage | 0 | 0 | 894363 | 0 | |
| ReadRepairStage | 0 | 0 | 26007 | 0 | |
| CounterMutationStage | 0 | 0 | 0 | 0 | |
| MiscStage | 0 | 0 | 0 | 0 | |
| CompactionExecutor | 1 | 1 | 4699 | 0 | |
| MemtableReclaimMemory | 0 | 0 | 77 | 0 | |
| PendingRangeCalculator | 0 | 0 | 3 | 0 | |
| GossipStage | 0 | 0 | 21833 | 0 | |
| SecondaryIndexManagement | 0 | 0 | 0 | 0 | |
| HintsDispatcher | 0 | 0 | 87 | 0 | |
| MigrationStage | 0 | 0 | 0 | 0 | |
| MemtablePostFlush | 1 | 1 | 99 | 0 | |
| ValidationExecutor | 0 | 0 | 0 | 0 | |
| Sampler | 0 | 0 | 0 | 0 | |
| MemtableFlushWriter | 1 | | | | |
| InternalResponseStage | 0 | | | | |
| AntiEntropyStage | 0 | | | | |
| CacheCleanupExecutor | 0 | | | | |
| Native-Transport-Requests | 33 | | | | |

| Message type | Dropped |
|-------------------------|---------|
| READ | 3227 |
| RANGE_SLICE | 0 |
| _TRACE | |
| HINT | |
| MUTATION | |
| COUNTER_MUTATION | |
| BATCH_STORE | |
| BATCH_REMOVE | |
| REQUEST_RESPONSE | |
| PAGED_RANGE | |
| READ_REPAIR | |



cassandra.yaml settings

```
concurrent_writes: 32
write_request_timeout_in_ms: 2000
counter_write_request_timeout_in_ms: 100
```

SSTables

SSTable Count

```
$ nodetool tablestats keyspace1
```

Keyspace: keyspace1

Read Count: 408439

Read Latency: 48.24464316825768 ms.

Write Count: 437331

Write Latency: 0.1109122129462581 ms.

Pending Flushes: 0

Table: standard1

SSTable count: 20

Space used (live): 37002901216

Space used (total): 37002901216

Space used by snapshots (total): 0

Off heap memory used (total): 9595368

SSTable Compression Ratio: 1.0038997780827315

Number of keys (estimate): 1009214

Memtable cell count: 12226

Memtable data size: 146845620

Memtable off heap memory used: 0

Memtable switch count: 18

Local read count: 408687

Local read latency: 52.956 ms

Local write count: 437471

Local write latency: 0.120 ms

Pending flushes: 0

Bloom filter false positives: 130

Bloom filter false ratio: 0.00007

Bloom filter space used: 3666792

Bloom filter off heap memory used: 4134856

Index summary off heap memory used: 546304

Compression metadata off heap memory used: 4914208

Compacted partition minimum bytes: 447

Compacted partition maximum bytes: 24601

Compacted partition mean bytes: 13783

Average live cells per slice (last five minutes): 1.0

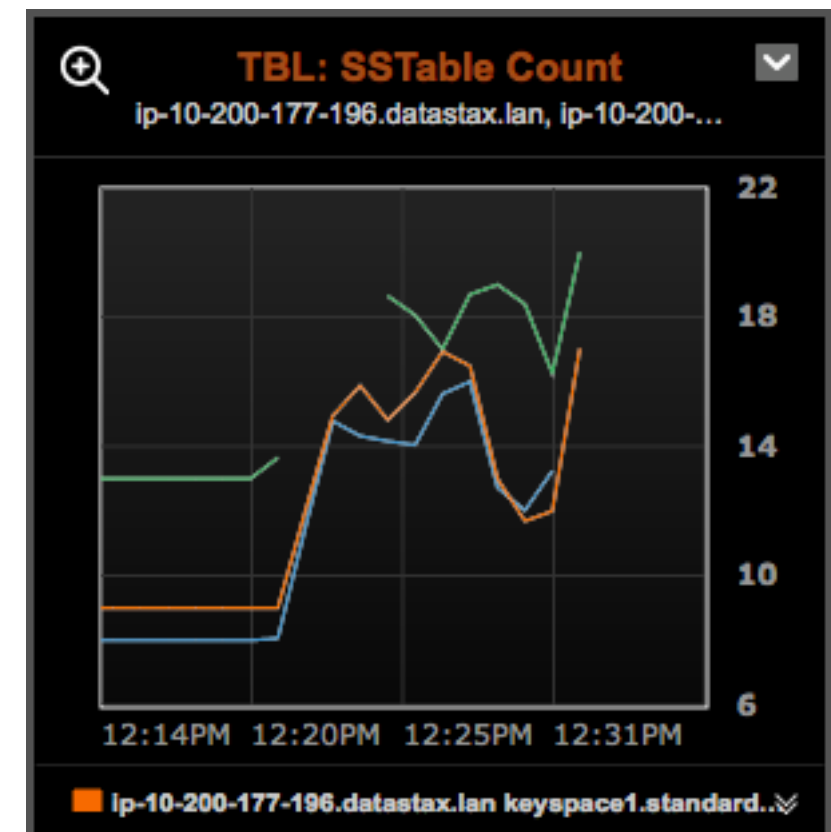
Maximum live cells per slice (last five minutes): 1

Average tombstones per slice (last five minutes): 1.0

Maximum tombstones per slice (last five minutes): 1

Common Causes

- Compaction behind
- Compactions failing
- Flushing too often
- Disk full

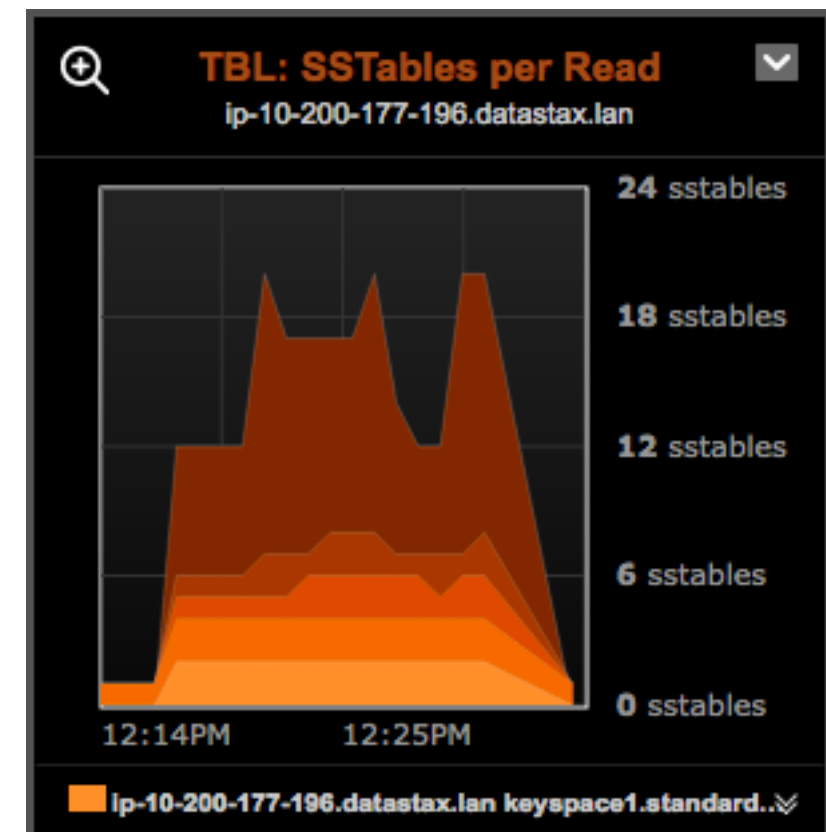


SSTables Per Read

```
$ nodetool tablehistograms keyspace1 standard1
```

keyspace1/standard1 histograms

| Percentile | SSTables | Write Latency (micros) | Read Latency (micros) | Partition Size (bytes) | Cell Count |
|------------|-----------------|---------------------------|--------------------------|---------------------------|------------|
| 50% | 3.00 | 35.43 | 30130.99 | 14237 | 5 |
| 75% | 4.00 | 42.51 | 74975.55 | 17084 | 5 |
| 95% | 6.00 | 88.15 | 129557.75 | 20501 | 5 |
| 98% | 7.00 | 182.79 | 155469.30 | 20501 | 5 |
| 99% | 7.00 | 454.83 | 186563.16 | 20501 | 5 |
| Min | 1.00 | 5.72 | 61.22 | 447 | 5 |
| Max | 20.00 | 2395318.86 | 2874382.63 | 24601 | 5 |



Tombstones

Tombstone Warnings

WARN [SharedPool-Worker-2] 2016-08-21 17:25:00,381 SliceQueryFilter.java:319 - Read **493 live** and **1479 tombstone** cells in **foo.bar** for key: **315847b3-67c4-11e6-9797-1db751d46fa1** (see **tombstone_warn_threshold**). 5000 columns were requested, slices=[-]

ERROR [SharedPool-Worker-3] 2016-08-23 01:12:32,504 SliceQueryFilter.java:284 - Scanned over **100000** tombstones in **foo.bar**; query aborted (see **tombstone_failure_threshold**)

Mitigation

- Avoid implementing queues
- Avoid massive deletes; truncate instead
- Don't insert nulls (be careful with prepared statements)
- Decrease `gc_grace_seconds` if you can

```
# configure gc_grace_seconds in CQL
alter table with gc_grace_seconds = 0; # use with caution!
```

```
# thresholds configurable in cassandra.yaml
tombstone_warn_threshold: 1000
tombstone_failure_threshold: 100000
```


System Resources



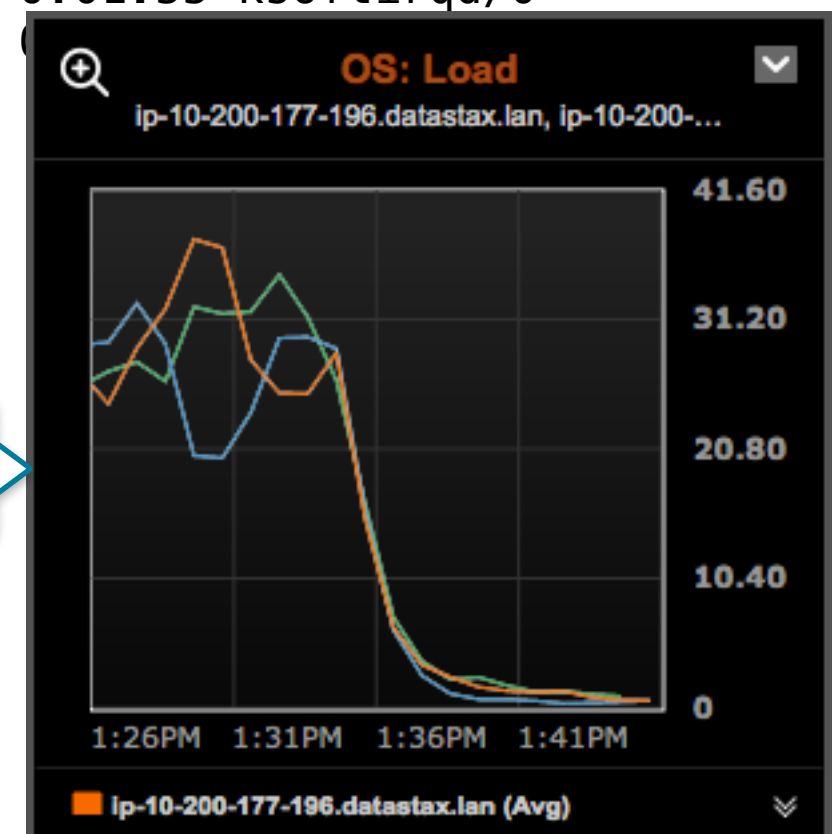
CPU + Disk Utilization = Load Average

```
top - 19:47:11 up 4 days, 5:12, 1 user, load average: 0.12, 0.71, 2.58
Tasks: 108 total, 2 running, 106 sleeping, 0 stopped, 0 zombie
%Cpu(s): 3.6 us, 1.7 sy, 0.0 ni, 94.6 id, 0.0 wa, 0.0 st
KiB Mem: 16434208 total, 8966924 used, 7467284 free, 2476732 cached mem
KiB Swap: 0 total, 0 used, 0 free.
```

1 5 15

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | COMMAND |
|-------|----------|----|-----|---------|--------|-------|---|------|------|----------|-------------|
| 9421 | cassand+ | 20 | 0 | 19.490g | 4.613g | 36832 | S | 23.0 | 29.4 | 39:08.01 | java |
| 14749 | cassand+ | 20 | 0 | 3752964 | 376148 | 15116 | S | 11.3 | 2.3 | 18:58.59 | java |
| 14494 | opscent+ | 20 | 0 | 5051000 | 895292 | 20360 | S | 1.7 | 5.4 | 10:24.98 | java |
| 14641 | automat+ | 20 | 0 | 23632 | 1568 | 1108 | R | 0.3 | 0.0 | 0:00.09 | top |
| 1 | root | 20 | 0 | 33620 | 2888 | 1476 | S | 0.0 | 0.0 | 0:01.82 | init |
| 2 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:00.22 | kthreadd |
| 3 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:01.55 | ksoftirqd/0 |
| 5 | root | 0 | -20 | 0 | 0 | 0 | S | 0.0 | 0.0 | | |

Historic 1 min avg




CPU

Total CPU Utilization

```
top - 19:47:11 up 4 days, 5:12, 1 user, load average: 0.12, 0.71, 2.58
Tasks: 108 total, 2 running, 106 sleeping, 0 stopped, 0 zombie
%Cpu(s): 3.6 us, 1.7 sy, 0.0 ni, 94.6 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 16434208 total, 8966924 used, 7467284 free, 98784 buffers
KiB Swap: 0 total, 0 used, 0 free. 2478732 cached Mem
```

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | COMMAND |
|-------|----------|----|-----|---------|--------|-------|---|------|------|----------|-------------|
| 9421 | cassand+ | 20 | 0 | 19.490g | 4.613g | 36832 | S | 23.0 | 29.4 | 39:08.01 | java |
| 14749 | cassand+ | 20 | 0 | 3752964 | 376148 | 15116 | S | 11.3 | 2.3 | 18:58.59 | java |
| 14494 | opscent+ | 20 | 0 | 5051000 | 895292 | 20360 | S | 1.7 | 5.4 | 10:24.98 | java |
| 14641 | automat+ | 20 | 0 | 23632 | 1568 | 1108 | R | 0.3 | 0.0 | 0:00.09 | top |
| 1 | root | 20 | 0 | 33620 | 2888 | 1476 | S | 0.0 | 0.0 | 0:01.82 | init |
| 2 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:00.22 | kthreadd |
| 3 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:01.55 | ksoftirqd/0 |
| 5 | root | 0 | -20 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:00 | |



Is it Cassandra?

```
top - 19:47:11 up 4 days, 5:12, 1 user, load average: 0.12, 0.71, 2.58
Tasks: 108 total, 2 running, 106 sleeping, 0 stopped, 0 zombie
%Cpu(s): 3.6 us, 1.7 sy, 0.0 ni, 94.6 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 16434208 total, 8966924 used, 7467284 free, 98784 buffers
KiB Swap: 0 total, 0 used, 0 free. 2478732 cached Mem
```

| PID | USER | PR | NI | VIRT | RES | SHR | S | %CPU | %MEM | TIME+ | COMMAND |
|-------|----------|----|-----|---------|--------|-------|---|------|------|----------|--------------|
| 9421 | cassand+ | 20 | 0 | 19.490g | 4.613g | 36832 | S | 23.0 | 29.4 | 39:08.01 | java |
| 14749 | cassand+ | 20 | 0 | 3752964 | 376148 | 15116 | S | 11.3 | 2.3 | 18:58.59 | java |
| 14494 | opscent+ | 20 | 0 | 5051000 | 895292 | 20360 | S | 1.7 | 5.4 | 10:24.98 | java |
| 14641 | automat+ | 20 | 0 | 23632 | 1568 | 1108 | R | 0.3 | 0.0 | 0:00.09 | top |
| 1 | root | 20 | 0 | 33620 | 2888 | 1476 | S | 0.0 | 0.0 | 0:01.82 | init |
| 2 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:00.22 | kthreadd |
| 3 | root | 20 | 0 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:01.55 | ksoftirqd/0 |
| 5 | root | 0 | -20 | 0 | 0 | 0 | S | 0.0 | 0.0 | 0:00.00 | kworker/0:0H |

Per-Thread Utilization

```
$ java -jar sjk-plus-0.4.2.jar ttop -p 22226 -o CPU -n 20
```

```
2016-08-17T18:54:00.188+0000 Process summary
```

```
process cpu=320.99%
```

```
application cpu=303.13% (user=174.83% sys=128.30%)
```

```
other: cpu=17.86%
```

```
heap allocation rate 218mb/s
```

Per-thread utilization

```
[000224] user= 6.39% sys= 4.00% alloc= 14mb/s - Net-Incoming-/10.200.177.196
[000293] user= 7.89% sys= 5.05% alloc= 14mb/s - Net-Incoming-/10.200.177.198
[000305] user= 6.79% sys= 3.17% alloc= 14mb/s - CompactionExecutor:4
[000157] user= 6.19% sys= 4.00% alloc= 6727kb/s - SharedPool-Worker-5
[000153] user= 7.49% sys= 3.51% alloc= 6596kb/s - SharedPool-Worker-2
[000152] user= 7.49% sys= 3.41% alloc= 6436kb/s - SharedPool-Worker-1
[000155] user= 6.39% sys= 3.80% alloc= 6153kb/s - SharedPool-Worker-4
[000154] user= 6.89% sys= 3.29% alloc= 6054kb/s - SharedPool-Worker-3
[000156] user= 6.39% sys= 3.68% alloc= 6030kb/s - SharedPool-Worker-6
[000168] user= 5.89% sys= 3.33% alloc= 5947kb/s - SharedPool-Worker-8
[000181] user= 5.49% sys= 3.28% alloc= 5343kb/s - SharedPool-Worker-9
[000169] user= 6.39% sys= 2.86% alloc= 5202kb/s - SharedPool-Worker-7
[000179] user= 4.80% sys= 2.29% alloc= 4969kb/s - SharedPool-Worker-12
[000158] user= 5.29% sys= 2.43% alloc= 4830kb/s - SharedPool-Worker-11
[000159] user= 4.60% sys= 2.06% alloc= 4734kb/s - SharedPool-Worker-13
[000178] user= 3.80% sys= 2.30% alloc= 4304kb/s - SharedPool-Worker-14
[000180] user= 4.90% sys= 2.45% alloc= 4294kb/s - SharedPool-Worker-10
[000160] user= 3.90% sys= 1.55% alloc= 4160kb/s - SharedPool-Worker-15
[000177] user= 3.30% sys= 1.29% alloc= 3268kb/s - SharedPool-Worker-16
[000247] user= 2.40% sys= 3.48% alloc= 3253kb/s - epollEventLoopGroup-6-2
```

Per-Thread Utilization

```
$ java -jar sjk-plus-0.4.2.jar ttop -p 22226 -o CPU -n 20
```

```
2016-08-17T18:54:00.188+0000 Process summary
```

```
process cpu=320.99%
```

```
application cpu=303.13% (user=174.83% sys=128.30%)
```

```
other: cpu=17.86%
```

```
heap allocation rate 21
```

| | Thread Name |
|----------|---|
| [000224] | user= 6.39% sys= MessagingService-Incoming-/10.200.177.196 |
| [000293] | user= 7.89% sys= MessagingService-Incoming-/10.200.177.198 |
| [000305] | user= 6.79% sys= 3.17% alloc= 14mb/s - CompactionExecutor:4 |
| [000157] | user= 6.19% sys= 4.00% alloc= 6727kb/s - SharedPool-Worker-5 |
| [000153] | user= 7.49% sys= 3.51% alloc= 6596kb/s - SharedPool-Worker-2 |
| [000152] | user= 7.49% sys= 3.41% alloc= 6436kb/s - SharedPool-Worker-1 |
| [000155] | user= 6.39% sys= 3.80% alloc= 6153kb/s - SharedPool-Worker-4 |
| [000154] | user= 6.89% sys= 3.29% alloc= 6054kb/s - SharedPool-Worker-3 |
| [000156] | user= 6.39% sys= 3.68% alloc= 6030kb/s - SharedPool-Worker-6 |
| [000168] | user= 5.89% sys= 3.33% alloc= 5947kb/s - SharedPool-Worker-8 |
| [000181] | user= 5.49% sys= 3.28% alloc= 5343kb/s - SharedPool-Worker-9 |
| [000169] | user= 6.39% sys= 2.86% alloc= 5202kb/s - SharedPool-Worker-7 |
| [000179] | user= 4.80% sys= 2.29% alloc= 4969kb/s - SharedPool-Worker-12 |
| [000158] | user= 5.29% sys= 2.43% alloc= 4830kb/s - SharedPool-Worker-11 |
| [000159] | user= 4.60% sys= 2.06% alloc= 4734kb/s - SharedPool-Worker-13 |
| [000178] | user= 3.80% sys= 2.30% alloc= 4304kb/s - SharedPool-Worker-14 |
| [000180] | user= 4.90% sys= 2.45% alloc= 4294kb/s - SharedPool-Worker-10 |
| [000160] | user= 3.90% sys= 1.55% alloc= 4160kb/s - SharedPool-Worker-15 |
| [000177] | user= 3.30% sys= 1.29% alloc= 3268kb/s - SharedPool-Worker-16 |
| [000247] | user= 2.40% sys= 3.48% alloc= 3253kb/s - epollEventLoopGroup-6-2 |

JVM Utilization

```
$ java -jar sjk-plus-0.4.2.jar ttop -p 22226 -o CPU -n 20
```

```
2016-08-17T18:54:00.188+0000 Process summary
```

```
process cpu=320.99%
```

```
application cpu=303.28.30%)
```

```
other: cpu=17.86%
```

```
heap allocation rat
```

JVM (GC, etc.)

```
[000224] user= 6.39% sys= 4.00% alloc= 29mb/s - MessagingService-Incoming-/10.200.177.196
[000293] user= 7.89% sys= 5.05% alloc= 29mb/s - MessagingService-Incoming-/10.200.177.198
[000305] user= 6.79% sys= 3.17% alloc= 14mb/s - CompactionExecutor:4
[000157] user= 6.19% sys= 4.00% alloc= 6727kb/s - SharedPool-Worker-5
[000153] user= 7.49% sys= 3.51% alloc= 6596kb/s - SharedPool-Worker-2
[000152] user= 7.49% sys= 3.41% alloc= 6436kb/s - SharedPool-Worker-1
[000155] user= 6.39% sys= 3.80% alloc= 6153kb/s - SharedPool-Worker-4
[000154] user= 6.89% sys= 3.29% alloc= 6054kb/s - SharedPool-Worker-3
[000156] user= 6.39% sys= 3.68% alloc= 6030kb/s - SharedPool-Worker-6
[000168] user= 5.89% sys= 3.33% alloc= 5947kb/s - SharedPool-Worker-8
[000181] user= 5.49% sys= 3.28% alloc= 5343kb/s - SharedPool-Worker-9
[000169] user= 6.39% sys= 2.86% alloc= 5202kb/s - SharedPool-Worker-7
[000179] user= 4.80% sys= 2.29% alloc= 4969kb/s - SharedPool-Worker-12
[000158] user= 5.29% sys= 2.43% alloc= 4830kb/s - SharedPool-Worker-11
[000159] user= 4.60% sys= 2.06% alloc= 4734kb/s - SharedPool-Worker-13
[000178] user= 3.80% sys= 2.30% alloc= 4304kb/s - SharedPool-Worker-14
[000180] user= 4.90% sys= 2.45% alloc= 4294kb/s - SharedPool-Worker-10
[000160] user= 3.90% sys= 1.55% alloc= 4160kb/s - SharedPool-Worker-15
[000177] user= 3.30% sys= 1.29% alloc= 3268kb/s - SharedPool-Worker-16
[000247] user= 2.40% sys= 3.48% alloc= 3253kb/s - epollEventLoopGroup-6-2
```

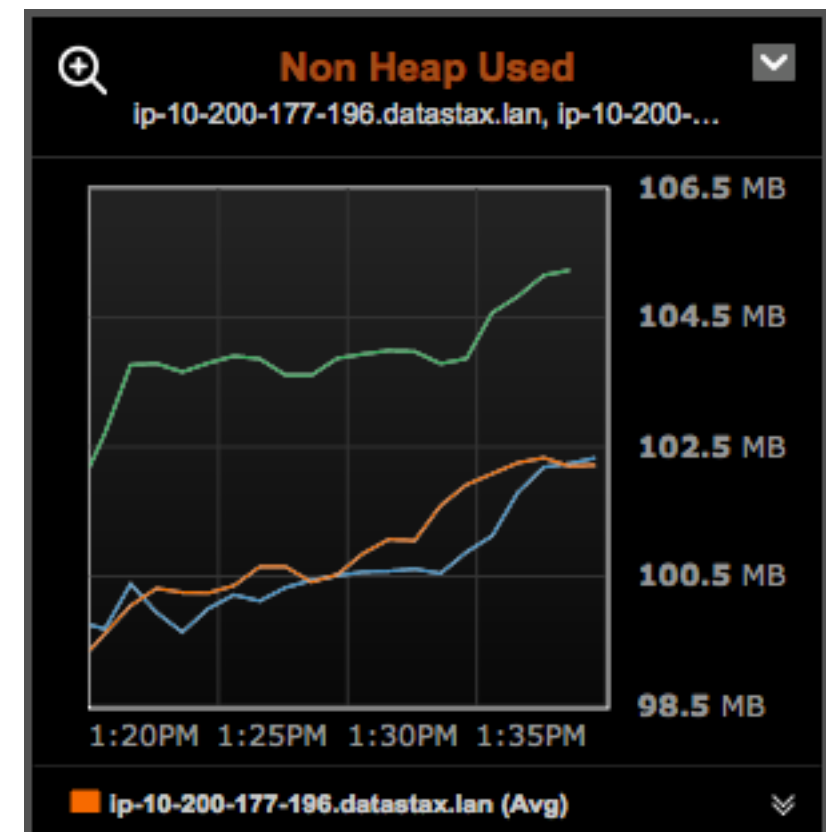
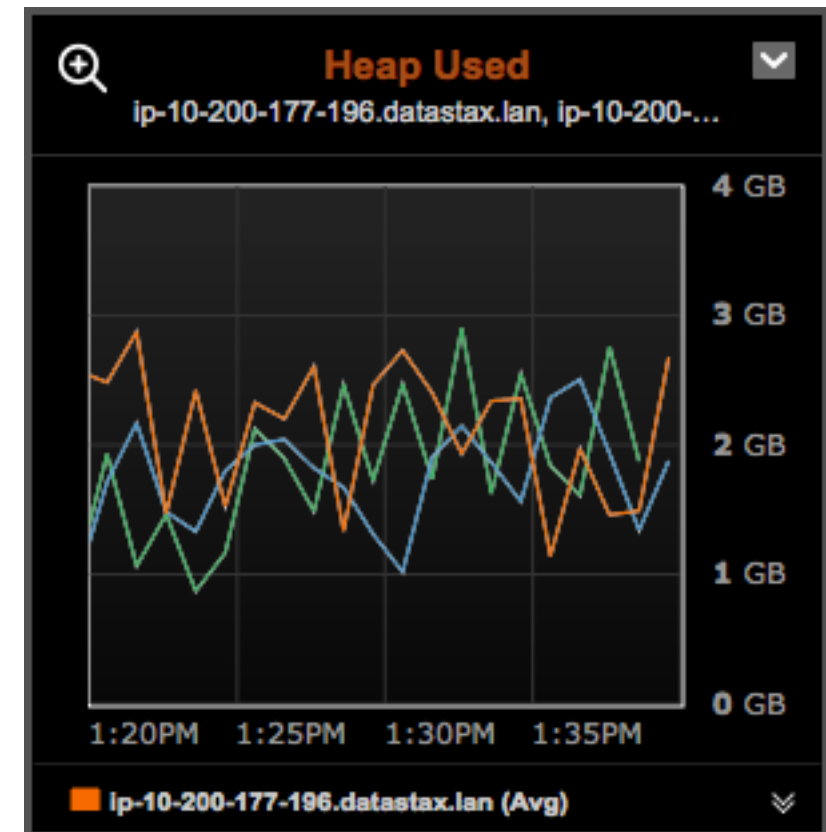

Memory

Heap/Off-Heap Usage

\$ nodetool info

```

ID : 432bc964-3cd3-4784-9ab7-d7a4a9e063b6
Gossip active : true
Thrift active : true
Native Transport active: true
Load : 13.41 GB
Generation No : 1471459228
Uptime (seconds) : 97288
Heap Memory (MB) : 1065.20 / 4012.00
Off Heap Memory (MB) : 3.49
Data Center : Cassandra
Rack : rack1
Exceptions : 0
Key Cache : entries 9002, size 760.92 KB,
capacity 100 MB, 19723 hits, 29234 requests, 0.675 recent hit
rate, 14400 save period in seconds
Row Cache : entries 0, size 0 bytes, capacity 0
bytes, 0 hits, 0 requests, NaN recent hit rate, 0 save period
in seconds
Counter Cache : entries 0, size 0 bytes, capacity 50
MB, 0 hits, 0 requests, NaN recent hit rate, 7200 save period
in seconds
Token : -3074457345618258603
  
```



Caches

Cassandra Caches

\$ nodetool info

```
ID : 6817e9ca-e79d-4fed-946e-7318bcfd5343
Gossip active : true
Thrift active : true
Native Transport active: true
Load : 22.68 MB
Generation No : 1426523950
Uptime (seconds) : 1557
Heap Memory (MB) : 270.85 / 1842.00
Off Heap Memory (MB) : 0.11
Data Center : us-east
Rack : 1b
Exceptions : 0
Key Cache : entries 156962, size 12.83 MB, capacity 100 MB, 649 hits, 713 requests,
              0.910 recent hit rate, 14400 save period in seconds
Row Cache : entries 0, size 0 bytes, capacity 0 bytes, 0 hits, 0 requests,
              NaN recent hit rate, 0 save period in seconds
Counter Cache : entries 0, size 0 bytes, capacity 50 MB, 0 hits, 0 requests,
                  NaN recent hit rate, 7200 save period in seconds
Token : 80372383360720788
```

Cache Size vs Capacity

\$ nodetool info

| | |
|-------------------------|---|
| ID | : 6817e9ca-e79d-4fed-946e-7318bcfd5343 |
| Gossip active | : true |
| Thrift active | : true |
| Native Transport active | : true |
| Load | : 22.68 MB |
| Generation No | : 1426523950 |
| Uptime (seconds) | : 1557 |
| Heap Memory (MB) | : 270.85 / 1842.00 |
| Off Heap Memory (MB) | : 0.11 |
| Data Center | : us-east |
| Rack | : 1b |
| Exceptions | : 0 |
| Key Cache | : entries 156962, size 12.83 MB, capacity 100 MB , 649 hits, 713 requests, 0.910 recent hit rate, 14400 save period in seconds |
| Row Cache | : entries 0, size 0 bytes, capacity 0 bytes , 0 hits, 0 requests, NaN recent hit rate, 0 save period in seconds |
| Counter Cache | : entries 0, size 0 bytes, capacity 50 MB , 0 hits, 0 requests, NaN recent hit rate, 7200 save period in seconds |
| Token | : 80372383360720788 |

In Use

Available

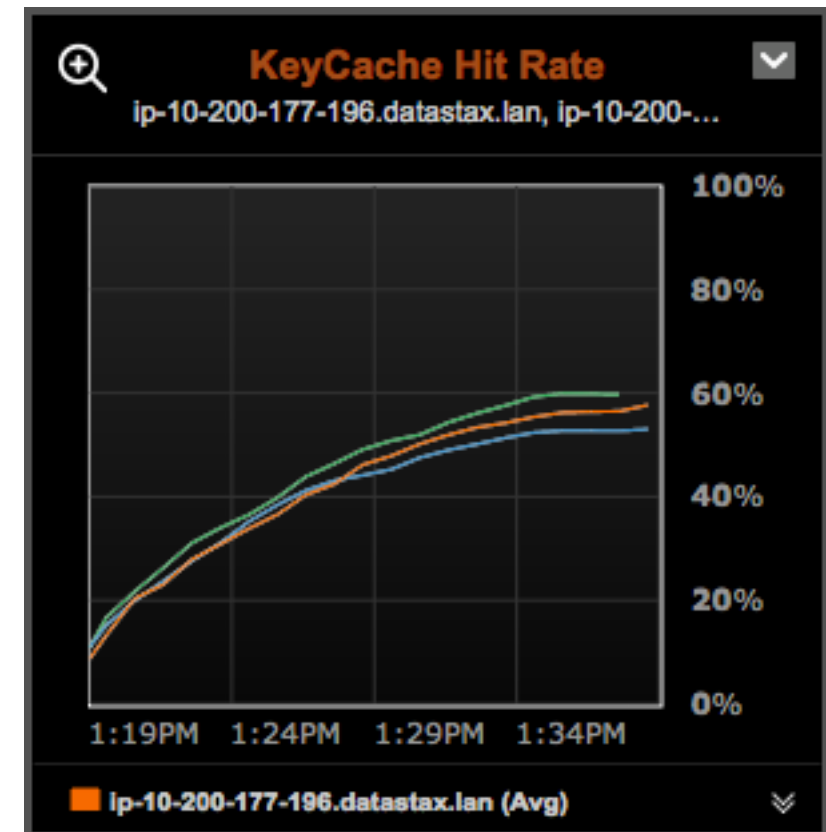
Cache Hit Rate

\$ nodetool info

```

ID : 6817e9ca-e79d-4fed-946e-7318bcfd5343
Gossip active : true
Thrift active : true
Native Transport active: true
Load : 22.68 MB
Generation No : 1426523950
Uptime (seconds) : 1557
Heap Memory (MB) : 270.85 / 1842.00
Off Heap Memory (MB) : 0.11
Data Center :
Rack :
Exceptions :
Key Cache : entries 156962, size 12.83 MB, capacity 100 MB, 649 hits, 713 requests,
0.910 recent hit rate, 14400 save period in seconds
Row Cache : entries 0, size 0 bytes, capacity 0 bytes, 0 hits, 0 requests,
NaN recent hit rate, 0 save period in seconds
Counter Cache : entries 0, size 0 bytes, capacity 50 MB, 0 hits, 0 requests,
NaN recent hit rate, 7200 save period in seconds
Token : 80372383360720788
    
```

Hit Rate



cassandra.yaml Cache Settings

\$ nodetool info

```
ID : 6817e9ca-e79d-4fed-946e-7318bcfd5343
Gossip active : true
Thrift active : true
Native Transport active: true
Load : 22.68 MB
Generation No : 1426523950
Uptime (seconds) : 1557
Heap Memory (MB) : 270.85 / 1842.00
Off Heap Memory (MB) : 0.11
Data Center : us-east
Rack : 1b
Exceptions : 0
Key Cache : entries 156962, size 12.83 MB, capacity 100 MB, 649 hits, 713 requests,
           0.910 recent hit rate, 14400 save period in seconds
Row Cache : entries 0, size 0 bytes, capacity 0 bytes, 0 hits, 0 requests,
           NaN recent hit rate, 0 save period in seconds
Counter Cache : entries 0, size 0 bytes, capacity 50 MB, 0 hits, 0 requests,
              NaN recent hit rate, 7200 save period in seconds
Token : 80372383360720788
```

cassandra.yaml settings

key_cache_size_in_mb: 100

row_cache_size_in_mb: 0

counter_cache_size_in_mb: 50

cql per-table configuration

alter table foo with caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'};

CQL Cache Settings

```
$ nodetool info
```

```
ID : 6817e9ca-e79d-4fed-946e-7318bcfd5343
Gossip active : true
Thrift active : true
Native Transport active: true
Load : 22.68 MB
Generation No : 1426523950
Uptime (seconds) : 1557
Heap Memory (MB) : 270.85 / 1842.00
Off Heap Memory (MB) : 0.11
Data Center : us-east
Rack : 1b
Exceptions : 0
Key Cache : entries 156962, size 12.83 MB, capacity 100 MB, 649 hits, 713 requests,
           0.910 recent hit rate, 14400 save period in seconds
Row Cache : entries 0, size 0 bytes, capacity 0 bytes, 0 hits, 0 requests,
           NaN recent hit rate, 0 save period in seconds
Counter Cache : entries 0, size 0 bytes, capacity 50 MB, 0 hits, 0 requests,
              NaN recent hit rate, 7200 save period in seconds
Token : 80372383360720788
```

```
# cassandra.yaml settings
```

```
key_cache_size_in_mb: 100
```

```
row_cache_size_in_mb: 0
```

```
counter_cache_size_in_mb: 50
```

```
# cql per-table configuration
```

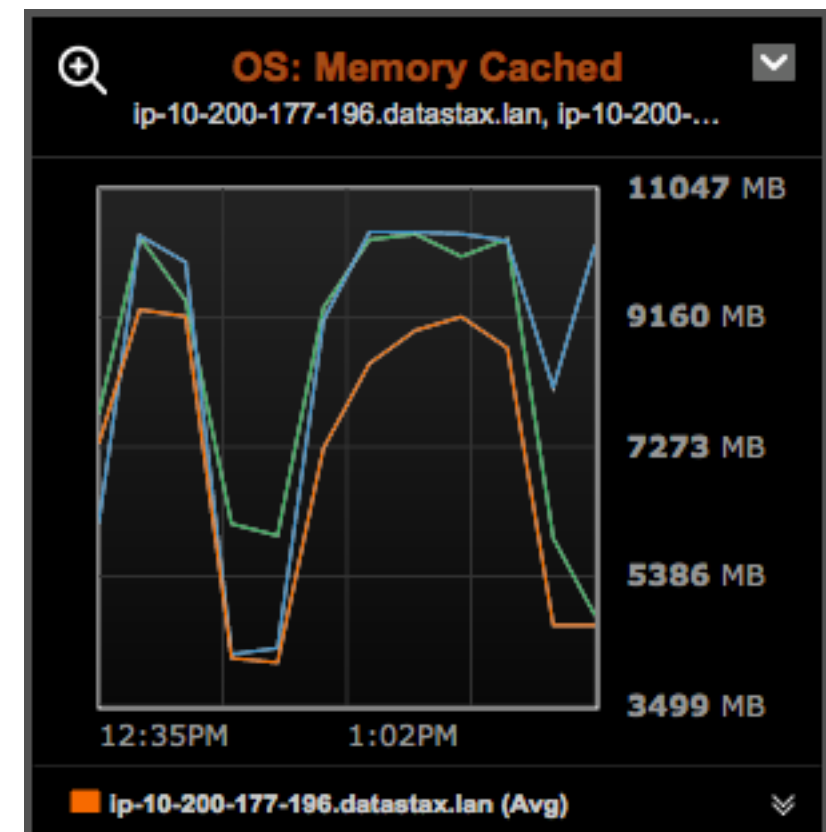
```
alter table foo with caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'};
```


Linux Page Cache

\$ free -h

| | total | used | free | shared | buffers | cached |
|--------------------|-------|------|------|--------|---------|--------|
| Mem: | 15G | 6.3G | 9.4G | 416K | 101M | 1.4G |
| -/+ buffers/cache: | | 4.8G | 10G | | | |
| Swap: | 0B | 0B | 0B | | | |

Don't forget about me!



Large Data Structures

Bloom Filters

```
$ nodetool tablestats keyspace1
```

Keyspace: keyspace1

Read Count: 4433559

Read Latency: 0.16309469277390917 ms.

Write Count: 15528537

Write Latency: 0.03023585988815302 ms.

Pending Flushes: 0

Table: standard1

SSTable count: 5

Space used (live): 3464993787

Space used (total): 3464993787

Space used by snapshots (total): 0

Off heap memory used (total): 19517188

SSTable Compression Ratio: 0.0

Number of keys (estimate): 10196252

Memtable cell count: 123604

Memtable data size: 34015680

Memtable off heap memory used: 0

Memtable switch count: 38

Local read count: 4433559

Local read latency: 0.179 ms

Local write count: 15528537

Local write latency: 0.034 ms

Pending flushes: 0

Bloom filter false positives: 2

Bloom filter false ratio: 0.00000

Bloom filter space used: 17159136

Bloom filter off heap memory used: 17159096

Index summary off heap memory used: 2358092

Compression metadata off heap memory used: 0

Compacted partition minimum bytes: 180

Compacted partition maximum bytes: 258

Compacted partition mean bytes: 258

Average live cells per slice (last five minutes): 1.0

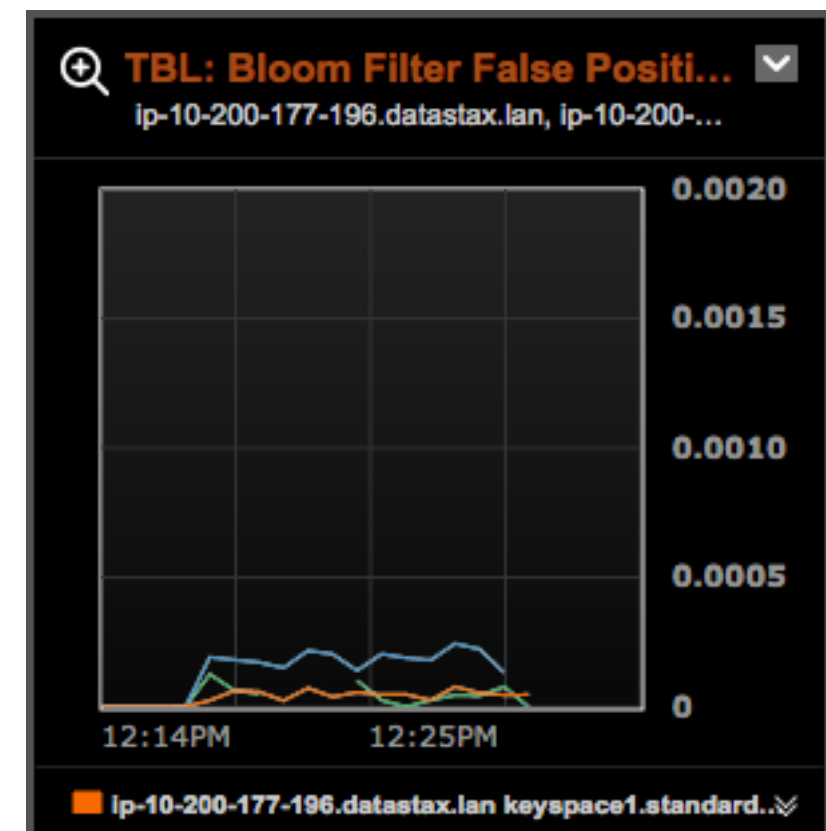
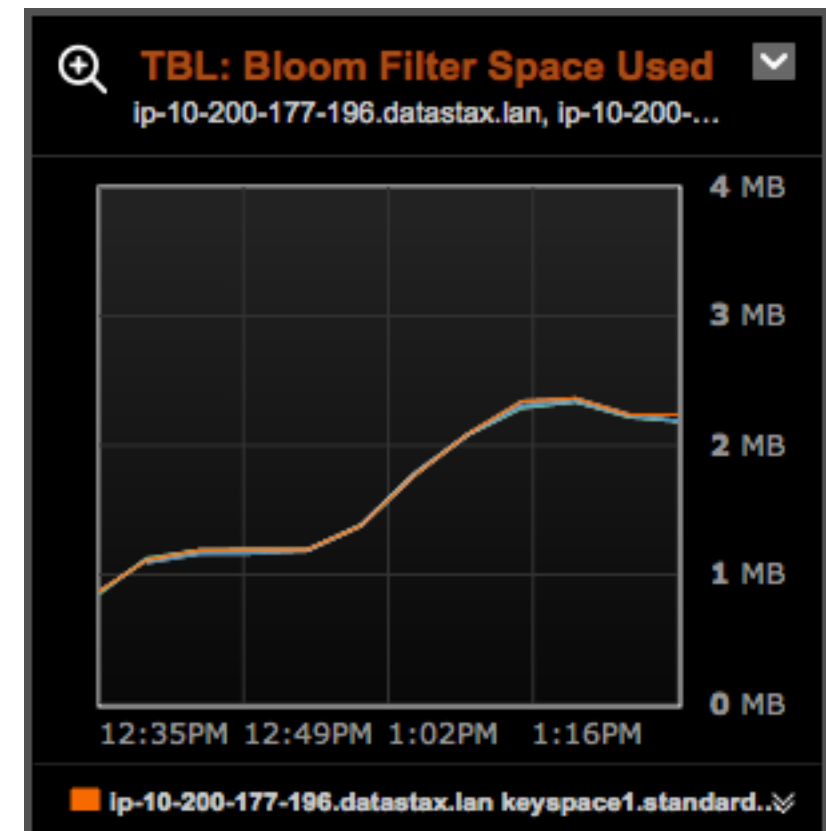
Maximum live cells per slice (last five minutes): 1

Average tombstones per slice (last five minutes): 1.0

Maximum tombstones per slice (last five minutes): 1

CQL configuration

alter table foo with bloom_filter_fp_chance = 0.01;



Index Summaries

```
$ nodetool tablestats keyspace1
```

```
Keyspace: keyspace1
```

```
Read Count: 4433559
```

```
Read Latency: 0.16309469277390917 ms.
```

```
Write Count: 15528537
```

```
Write Latency: 0.03023585988815302 ms.
```

```
Pending Flushes: 0
```

```
Table: standard1
```

```
SSTable count: 5
```

```
Space used (live): 3464993787
```

```
Space used (total): 3464993787
```

```
Space used by snapshots (total): 0
```

```
Off heap memory used (total): 19517188
```

```
SSTable Compression Ratio: 0.0
```

```
Number of keys (estimate): 10196252
```

```
Memtable cell count: 123604
```

```
Memtable data size: 34015680
```

```
Memtable off heap memory used: 0
```

```
Memtable switch count: 38
```

```
Local read count: 4433559
```

```
Local read latency: 0.179 ms
```

```
Local write count: 15528537
```

```
Local write latency: 0.034 ms
```

```
Pending flushes: 0
```

```
Bloom filter false positives: 2
```

```
Bloom filter false ratio: 0.00000
```

```
Bloom filter space used: 17159136
```

```
Bloom filter off heap memory used: 17159096
```

```
Index summary off heap memory used: 2358092
```

```
Compression metadata off heap memory used: 0
```

```
Compacted partition minimum bytes: 180
```

```
Compacted partition maximum bytes: 258
```

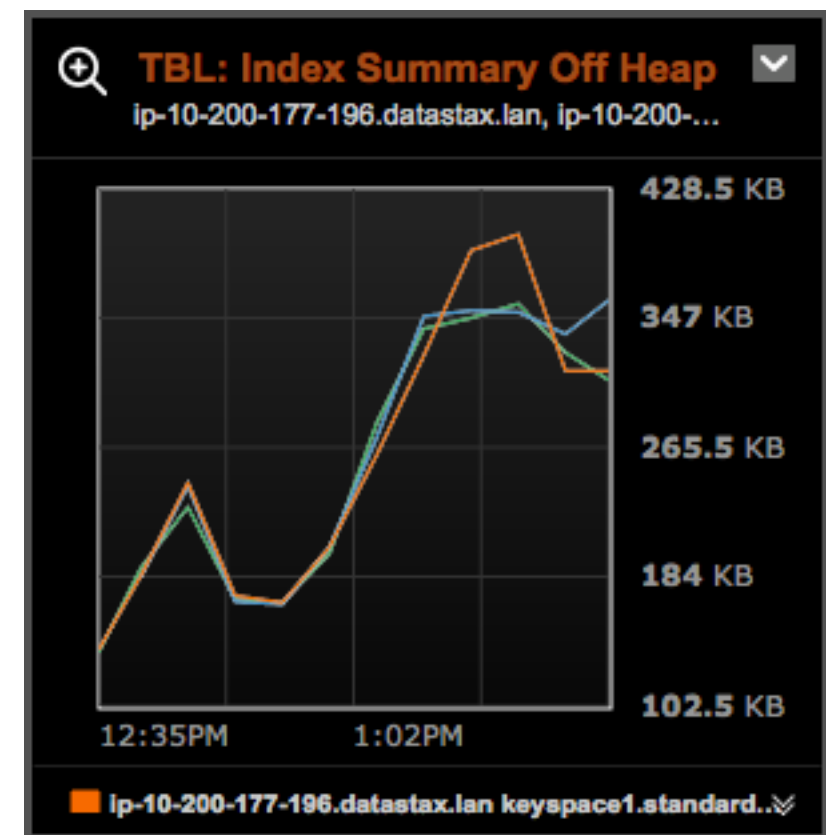
```
Compacted partition mean bytes: 258
```

```
Average live cells per slice (last five minutes): 1.0
```

```
Maximum live cells per slice (last five minutes): 1
```

```
Average tombstones per slice (last five minutes): 1.0
```

```
Maximum tombstones per slice (last five minutes): 1
```



```
# cassandra.yaml settings
```

```
index_summary_capacity_in_mb: 400
```

```
index_summary_resize_interval_in_minutes: 60
```

Compression Metadata

```
$ nodetool tablestats keyspace1
```

Keyspace: keyspace1

Read Count: 408439

Read Latency: 48.24464316825768 ms.

Write Count: 437331

Write Latency: 0.1109122129462581 ms.

Pending Flushes: 0

Table: standard1

SSTable count: 20

Space used (live): 37002901216

Space used (total): 37002901216

Space used by snapshots (total): 0

Off heap memory used (total): 9595368

SSTable Compression Ratio: 1.0038997780827315

Number of keys (estimate): 1009214

Memtable cell count: 12226

Memtable data size: 146845620

Memtable off heap memory used: 0

Memtable switch count: 18

Local read count: 408687

Local read latency: 52.956 ms

Local write count: 437471

Local write latency: 0.120 ms

Pending flushes: 0

Bloom filter false positives: 130

Bloom filter false ratio: 0.00007

Bloom filter space used: 3666792

Bloom filter off heap memory used: 4134856

Index summary off heap memory used: 546304

Compression metadata off heap memory used: 4914208

Compacted partition minimum bytes: 447

Compacted partition maximum bytes: 24601

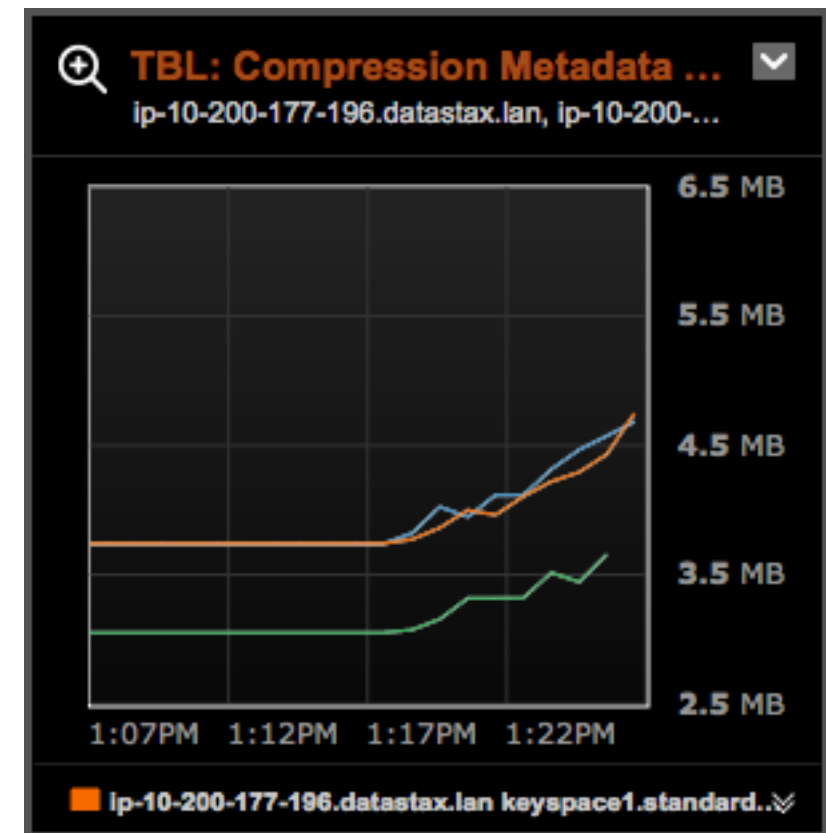
Compacted partition mean bytes: 13783

Average live cells per slice (last five minutes): 1.0

Maximum live cells per slice (last five minutes): 1

Average tombstones per slice (last five minutes): 1.0

Maximum tombstones per slice (last five minutes): 1



CQL configuration

```
alter table foo with compression = {'sstable_compression': ''};
```

Disk

I/O Utilization

\$ iostat -xd 5

Linux 3.13.0-92-generic (ip-10-200-177-196.datastax.lan)

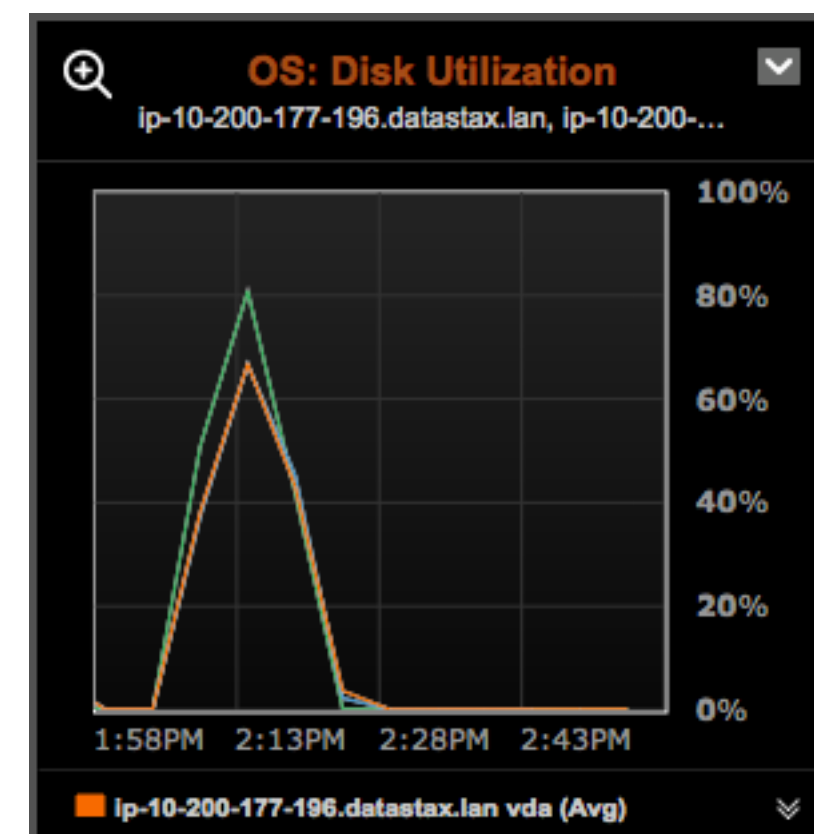
08/18/2016

_x86_64_

(4 CPU)

| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
|---------|--------|--------|------|--------|--------|----------|----------|----------|-------|---------|---------|-------|-------|
| vda | 0.00 | 1.95 | 7.94 | 5.33 | 691.72 | 2027.59 | 409.76 | 0.10 | 7.56 | 1.58 | 16.46 | 0.75 | 0.99 |
| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
| vda | 0.00 | 28.00 | 0.00 | 79.40 | 0.00 | 31832.80 | 801.83 | 0.97 | 12.21 | 0.00 | 12.21 | 0.80 | 6.32 |
| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
| vda | 0.00 | 63.80 | 0.00 | 208.60 | 0.00 | 86608.80 | 830.38 | 3.77 | 18.06 | 0.00 | 18.06 | 1.55 | 32.40 |

% Utilization



I/O Utilization

\$ iostat -xd 5

Linux 3.13.0-92-generic (ip-10-200-177-196.datastax.lan)

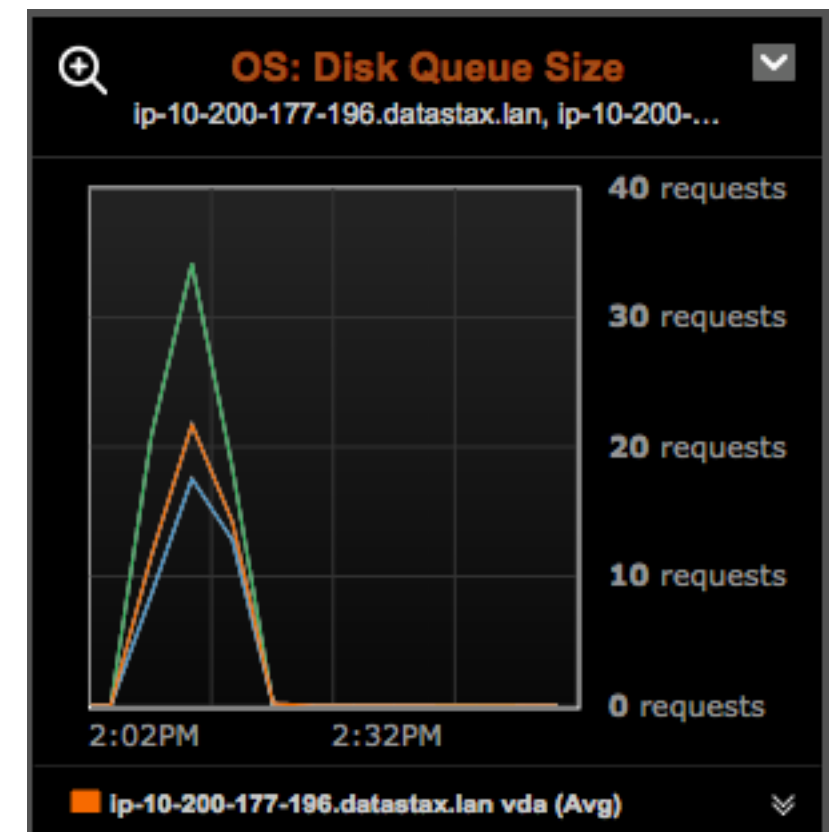
08/18/2016

_x86_64_

(4 CPU)

| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
|---------|--------|--------|------|--------|--------|----------|----------|----------|-------|---------|---------|-------|-------|
| vda | 0.00 | 1.95 | 7.94 | 5.33 | 691.72 | 2027.59 | 409.76 | 0.10 | 7.56 | 1.58 | 16.46 | 0.75 | 0.99 |
| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
| vda | 0.00 | 28.00 | 0.00 | 79.40 | 0.00 | 31832.80 | 801.83 | 0.97 | 12.21 | 0.00 | 12.21 | 0.80 | 6.32 |
| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
| vda | 0.00 | 63.80 | 0.00 | 208.60 | 0.00 | 86608.80 | 830.38 | 3.77 | 18.06 | 0.00 | 18.06 | 1.55 | 32.40 |

Queue Size



I/O Utilization

\$ iostat -xd 5

Linux 3.13.0-92-generic (ip-10-200-177-196.datastax.lan)

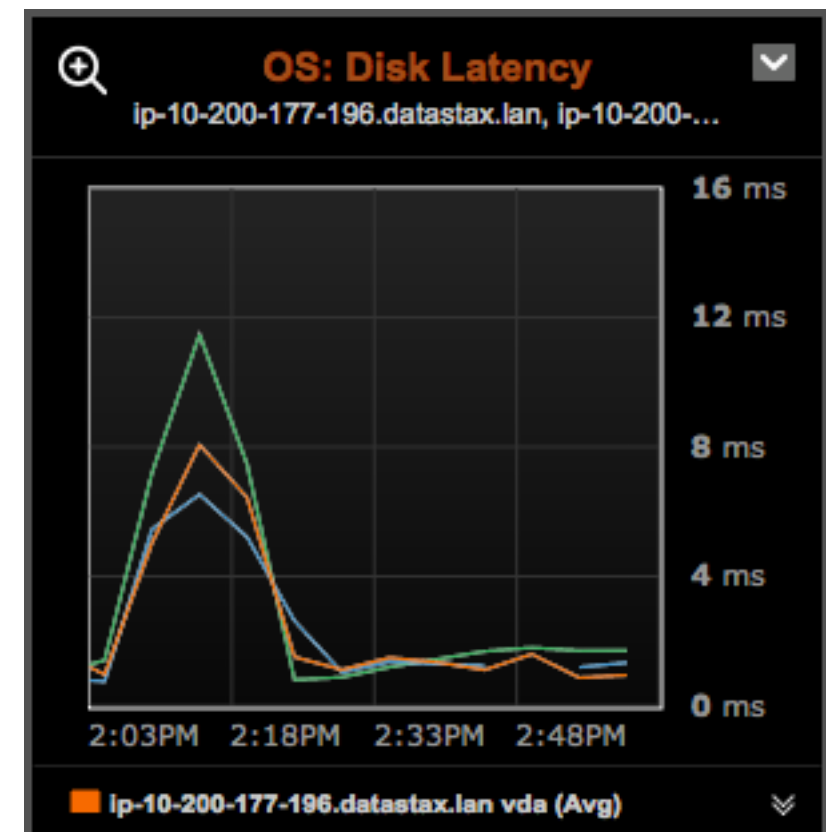
08/18/2016

_x86_64_

(4 CPU)

| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
|---------|--------|--------|------|--------|--------|----------|----------|----------|-------|---------|---------|-------|-------|
| vda | 0.00 | 1.95 | 7.94 | 5.33 | 691.72 | 2027.59 | 409.76 | 0.10 | 7.56 | 1.58 | 16.46 | 0.75 | 0.99 |
| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
| vda | 0.00 | 28.00 | 0.00 | 79.40 | 0.00 | 31832.80 | 801.83 | 0.97 | 12.21 | 0.00 | 12.21 | 0.80 | 6.32 |
| Device: | rrqm/s | wrqm/s | r/s | w/s | rkB/s | wkB/s | avgrq-sz | avgqu-sz | await | r_await | w_await | svctm | %util |
| vda | 0.00 | 63.80 | 0.00 | 208.60 | 0.00 | 86608.80 | 830.38 | 3.77 | 18.06 | 0.00 | 18.06 | 1.55 | 32.40 |

Average Wait



Space Used per Node

\$ nodetool status

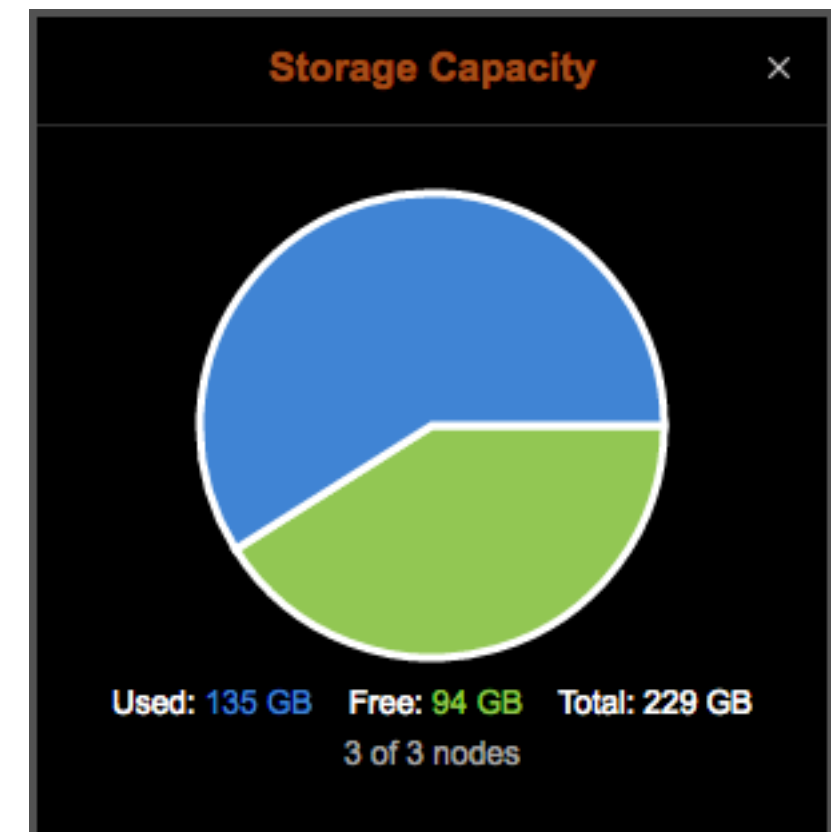
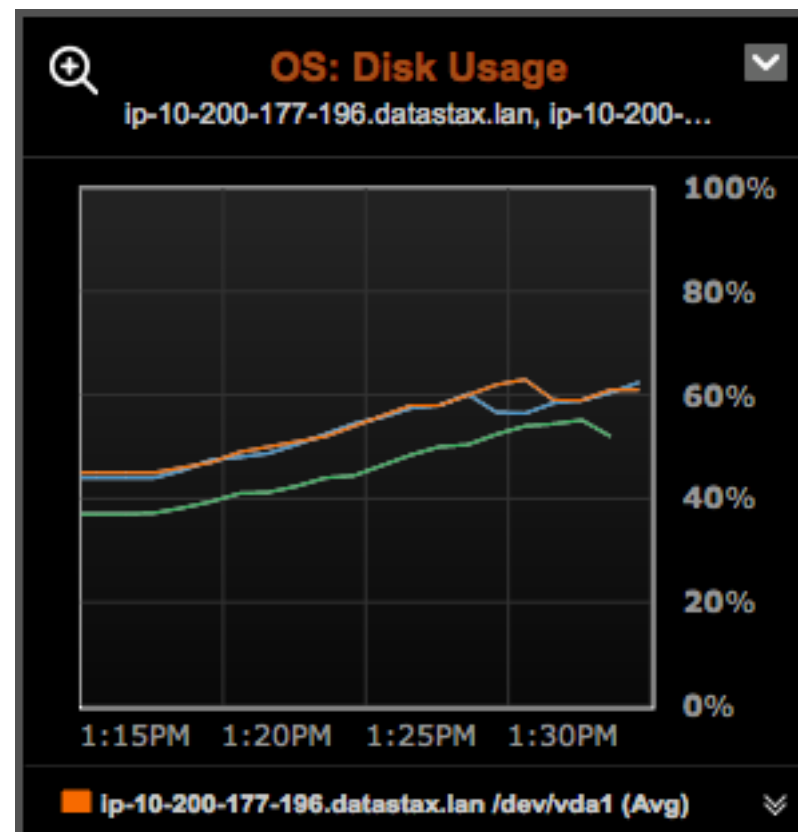
Datacenter: Cassandra

=====

Status=Up/Down

|/ State=Normal/Leaving/Joining/Moving

| -- | Address | Load | Tokens | Owns | Host ID | Rack |
|----|----------------|----------|--------|------|--------------------------------------|-------|
| UN | 10.200.177.196 | 36.02 GB | 256 | ? | fbb34d33-5f93-4f24-84e0-bd9cc42b1860 | rack1 |
| UN | 10.200.177.197 | 35.73 GB | 256 | ? | 0373dd12-25ef-4d97-a1c1-a4014dbd6eaa | rack1 |
| UN | 10.200.177.198 | 29.85 GB | 256 | ? | a9d22c58-2554-40c4-b4b1-37c7e96d00b4 | rack1 |



Space Used Per Table

\$ nodetool tablestats keyspace1

Keyspace: keyspace1

Read Count: 4433559

Read Latency: 0.16309469277390917 ms.

Write Count: 15528537

Write Latency: 0.03023585988815302 ms.

Pending Flushes: 0

Table: standard1

SSTable count: 5

Space used (live): 3464993787

Space used (total): 3464993787

Space used by snapshots (total): 0

Off heap memory used (total): 19517188

SSTable Compression Ratio: 0.0

Number of keys (estimate): 10196252

Memtable cell count: 123604

Memtable data size: 34015680

Memtable off heap memory used: 0

Memtable switch count: 38

Local read count: 4433559

Local read latency: 0.179 ms

Local write count: 15528537

Local write latency: 0.034 ms

Pending flushes: 0

Bloom filter false positives: 2

Bloom filter false ratio: 0.00000

Bloom filter space used: 17159136

Bloom filter off heap memory used: 17159096

Index summary off heap memory used: 2358092

Compression metadata off heap memory used: 0

Compacted partition minimum bytes: 180

Compacted partition maximum bytes: 258


Compacted partition mean bytes: 258

Average live cells per slice (last five minutes): 1.0

Maximum live cells per slice (last five minutes): 1

Average tombstones per slice (last five minutes): 1.0

Maximum tombstones per slice (last five minutes): 1



Uneven Data Distribution

Disk Usage Per Node

\$ nodetool status

Datacenter: Cassandra

=====

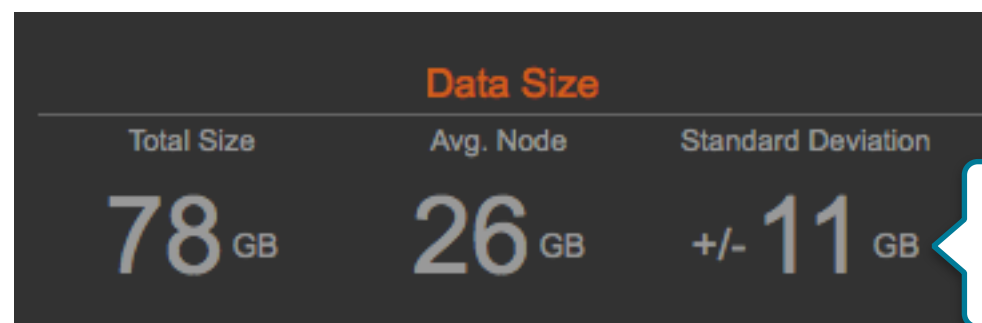
Status=Up/Down

|/ State=Normal/Leaving/Joining/Moving

| -- | Address | Load | Tokens | Owns | Host ID | Rack |
|----|----------------|----------|--------|------|--------------------------------------|-------|
| UN | 10.200.177.196 | 31.85 GB | 256 | ? | fb34d33-5f93-4f24-84e0-bd9cc42b1860 | rack1 |
| UN | 10.200.177.197 | 13.63 GB | 256 | ? | 0373dd12-25ef-4d97-a1c1-a4014dbd6eaa | rack1 |
| UN | 10.200.177.198 | 32.95 GB | 256 | ? | a9d22c58-2554-40c4-b4b1-37c7e96d00b4 | rack1 |

Note: They don't have the same replication settings, effective ownership information means

Very Uneven!



Standard deviation

Ownership Percentage

```
$ nodetool status keyspace1
```

```
Datacenter: Cassandra
```

```
=====
```

```
Status=Up/Down
```

```
|/ State=Normal/Leaving/Joining/Moving
```

| -- | Address | Load | Tokens | Owns (effective) | Host ID | Rack |
|----|----------------|----------|--------|------------------|--------------------------------------|-------|
| UN | 10.200.177.196 | 32.71 GB | 256 | 67.1% | fbb34d33-5f93-4f24-84e0-bd9cc42b1860 | rack1 |
| UN | 10.200.177.197 | 31.82 GB | 256 | 69.1% | 0373dd12-25ef-4d97-a1c1-a4014dbd6eaa | rack1 |
| UN | 10.200.177.198 | 32.95 GB | 256 | 63.8% | a9d22c58-2554-40c4-b4b1-37c7e96d00b4 | rack1 |

Must specify keyspace

- Should be roughly equal
- May vary with vnodes up to 10%
- Adds up to 100% times RF

Unbalanced Racks

\$ nodetool status

Datacenter: Cassandra

=====

Status=Up/Down

|/ State=Normal/Leaving/Joining/Moving

| -- | Address | Load | Tokens | Owns | Host ID | Rack |
|----|----------------|---------|--------|------|--------------------------------------|-------|
| UN | 10.200.177.196 | 3.45 GB | 256 | ? | 108af27a-43d8-4814-b617-f8f93ba2bb0e | rack1 |
| UN | 10.200.177.197 | 3.45 GB | 256 | ? | 432bc964-3cd3-4784-9ab7-d7a4a9e063b6 | rack2 |
| UN | 10.200.177.198 | 3.45 GB | 256 | ? | 3c467f89-7cce-485f-bb16-dd782c9a84ec | rack2 |

Ensure same # in each rack

Note: Non-system keyspaces don't have the same replication settings, effective ownership information is meaningless

recommended snitch in cassandra.yaml
 endpoint_snitch: **GossipingPropertyFileSnitch**

cassandra-rackdc.properties setting
 dc=Cassandra
 dc_suffix=Cassandra
rack=rack1

Inconsistent Token Count

\$ nodetool status

Datacenter: Cassandra

=====

Status=Up/Down

|/ State=Normal/Leaving/Joining/Moving

| -- | Address | Load | Tokens | Owens | Host ID | Rack |
|----|----------------|---------|--------|-------|--------------------------------------|-------|
| UN | 10.200.177.196 | 3.45 GB | 256 | ? | 108af27a-43d8-4814-b617-f8f93ba2bb0e | rack1 |
| UN | 10.200.177.197 | 3.45 GB | 256 | ? | 432bc964-3cd3-4784-9ab7-d7a4a9e063b6 | rack2 |
| UN | 10.200.177.198 | 3.45 GB | 128 | ? | 3c467f89-7cce-485f-bb16-dd782c9a84ec | rack2 |

Note: Non-sys
meaningless

Must be the same within a DC

ings, effective ownership information

cassandra.yaml setting
num_tokens: 256

Uneven Token Distribution*

\$ **nodetool ring**

* If not using vnodes

Datacenter: Cassandra

=====

| Address | Rack | Status | State | Load | Owns | Token |
|----------------|-------|--------|--------|----------|------|----------------------|
| 10.200.177.196 | rack1 | Up | Normal | 13.41 GB | ? | 3074457345618258602 |
| 10.200.177.197 | rack1 | Up | Normal | 4.41 GB | ? | -9223372036854775808 |
| 10.200.177.198 | rack1 | Up | Normal | 4.4 GB | ? | -3074457345618258603 |
| | | | | | | 3074457345618258602 |

to generate evenly spaced tokens

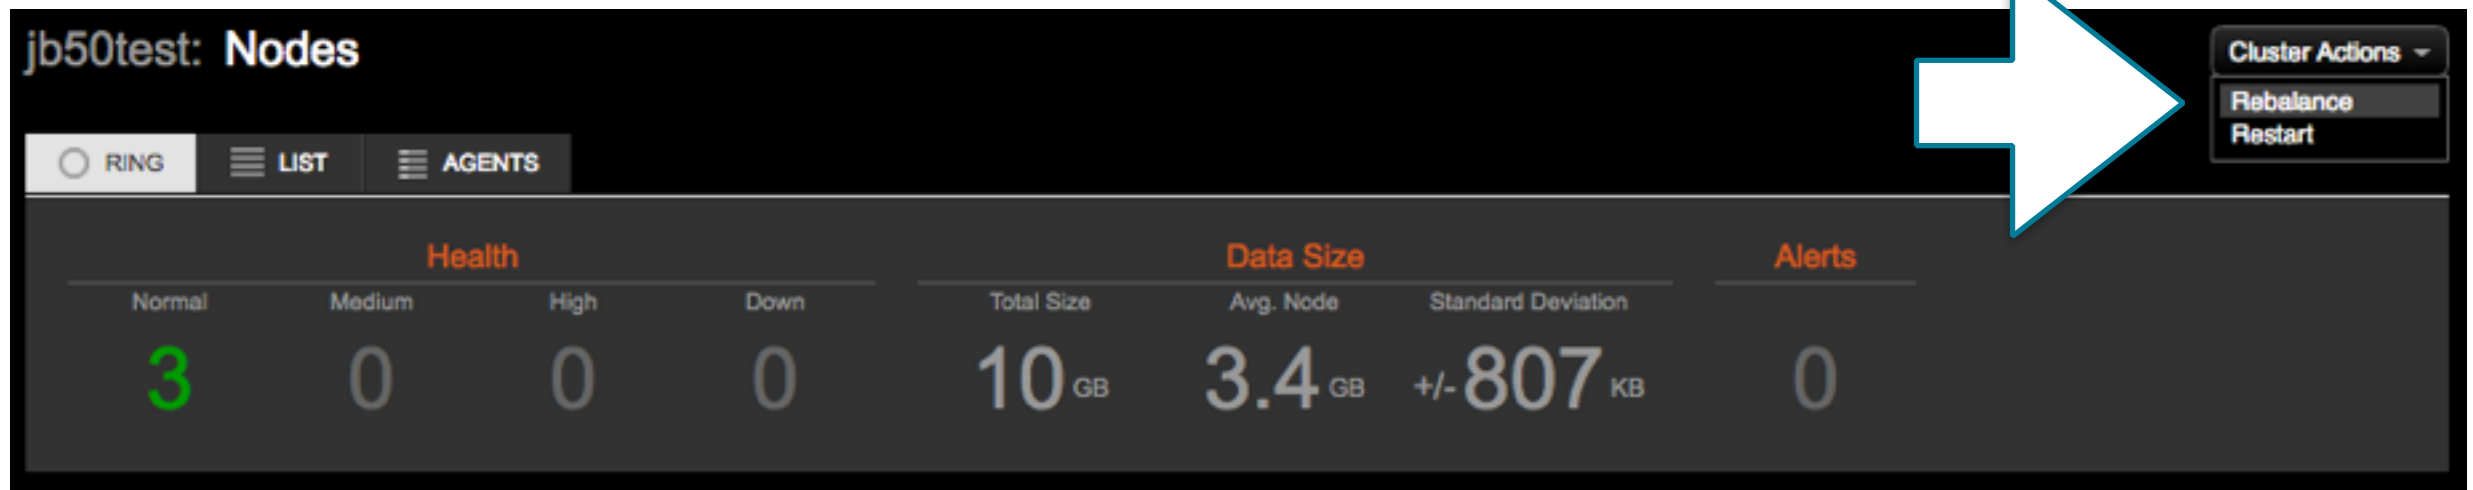
\$ token-generator 3 5

cassandra.yaml setting – make sure it's set*
initial_token: -9223372036854775808

to change after the node is bootstrapped

\$ nodetool move 3074457345618258602

Must be evenly spaced

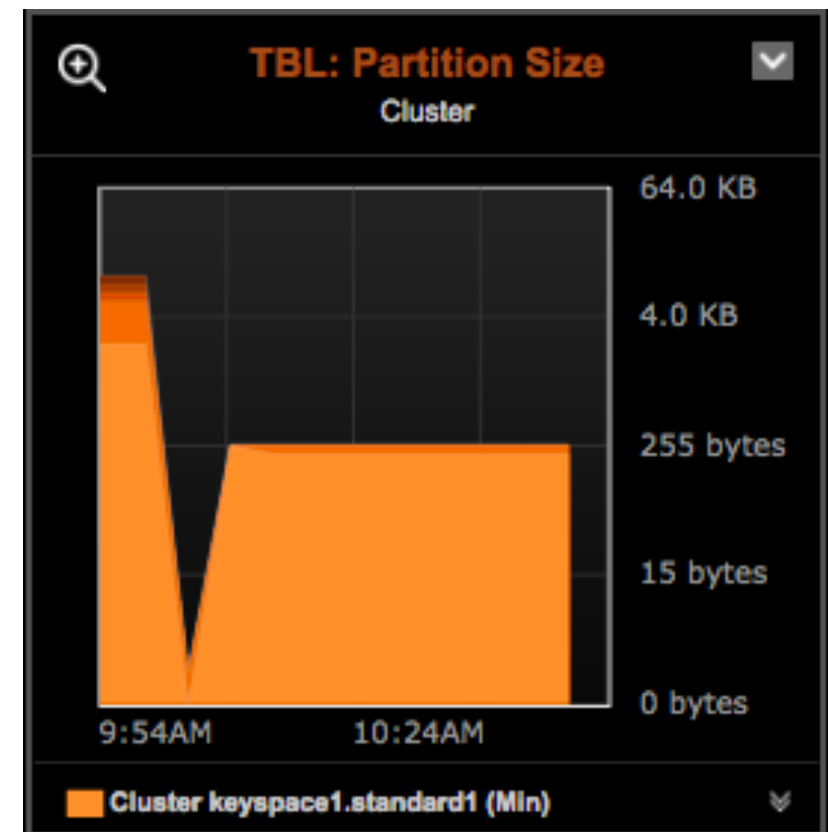
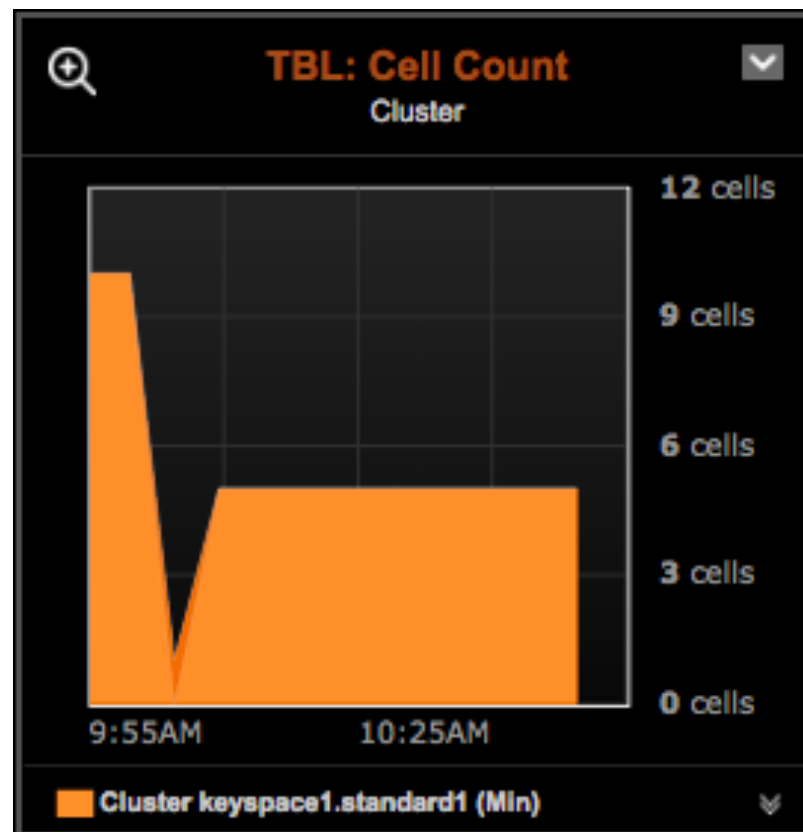


Partition Size and Cell Count

\$ nodetool tablehistograms keyspace1 standard1

foo/bar histograms

| Percentile | SSTables | Write Latency (micros) | Read Latency (micros) | Partition Size (bytes) | Cell Count |
|------------|----------|---------------------------|--------------------------|---------------------------|------------|
| 50% | 3.00 | 124.00 | 924.00 | 29521 | 149 |
| 75% | 3.00 | 215.00 | 1331.00 | 61214 | 310 |
| 95% | 3.00 | 642.00 | 2299.00 | 219342 | 924 |
| 98% | 3.00 | 1109.00 | 3311.00 | 379022 | 1331 |
| 99% | 3.00 | 1331.00 | 3973.00 | 454826 | 1916 |
| Min | 0.00 | 43.00 | 51.00 | 1332 | 11 |
| Max | 3.00 | 2759.00 | 42510.00 | 2346799 | 6866 |



Compactions - large partitions

INFO [CompactionExecutor:633] 2015-09-05 07:34:55,115 CompactionTask.java:141 -
Compacting [SSTableReader(path='/data/ax/user/ax-user-ka-13684-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13674-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13670-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13687-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13680-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13685-Data.db')]

WARN [CompactionExecutor:633] 2015-09-05 07:45:25,016 SSTableWriter.java:240 -
Compacting large partition ax/user:root (2503794228 bytes)

INFO [CompactionExecutor:633] 2015-09-05 07:45:25,114 CompactionTask.java:274 -
Compacting to [/data/ax/user/ax-user-ka-13691,].
400,997,778 bytes to 276,263,497 (~68% of original) in 629,997ms = 0.418201MB/s.
21 total partitions merged to 5. Partition merge counts were {3:3, 6:2, }

Warning!

Compactions - keyspace/table:key

INFO [CompactionExecutor:633] 2015-09-05 07:34:55,115 CompactionTask.java:141 -
Compacting [SSTableReader(path='/data/ax/user/ax-user-ka-13684-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13674-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13670-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13687-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13680-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13685-Data.db')]

WARN [CompactionExecutor:633] 2015-09-05 07:45:25,016 SSTableWriter.java:240 -
Compacting large partition **ax/user:root** (2503794228 bytes)

INFO [CompactionExecutor:633] 2015-09-05 07:45:25,016 CompactionTask.java:274 -
Compacted 6 sstables
400,997,778 bytes to 276,263,497 (~68% of original) in 629,997ms = 0.418201MB/s.
21 total partitions merged to 5. Partition merge counts were {3:3, 6:2, }

keyspace/table:key

Compactions - partition size

INFO [CompactionExecutor:633] 2015-09-05 07:34:55,115 CompactionTask.java:141 -
Compacting [SSTableReader(path='/data/ax/user/ax-user-ka-13684-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13674-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13670-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13687-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13680-Data.db'),
SSTableReader(path='/data/ax/user/ax-user-ka-13685-Data.db')]

WARN [CompactionExecutor:633] 2015-09-05 07:45:25,016 SSTableWriter.java:240 -
Compacting large partition ax/user:root (**2503794228 bytes**)

INFO [CompactionExecutor:633] 2015-09-05 07:45:25,016 CompactionTask.java:274 -
Compacted 6 sstables to [/data/ax/user/ax-user-ka-13684-Data.db, ...].
400,997,778 bytes to 276,263,497 (~68% of original) in 629,997ms = 0.418201MB/s.
21 total partitions merged to 5. Partition merge counts were {3:3, 6:2, }

Size

Compacted Partition Sizes

```
$ nodetool tablestats keyspace1
```

```
Keyspace: keyspace1
```

```
Read Count: 4433559
```

```
Read Latency: 0.16309469277390917 ms.
```

```
Write Count: 15528537
```

```
Write Latency: 0.03023585988815302 ms.
```

```
Pending Flushes: 0
```

```
Table: standard1
```

```
SSTable count: 5
```

```
Space used (live): 3464993787
```

```
Space used (total): 3464993787
```

```
Space used by snapshots (total): 0
```

```
Off heap memory used (total): 19517188
```

```
SSTable Compression Ratio: 0.0
```

```
Number of keys (estimate): 10196252
```

```
Memtable cell count: 123604
```

```
Memtable data size: 34015680
```

```
Memtable off heap memory used: 0
```

```
Memtable switch count: 38
```

```
Local read count: 4433559
```

```
Local read latency: 0.179 ms
```

```
Local write count: 15528537
```

```
Local write latency: 0.034 ms
```

```
Pending flushes: 0
```

```
Bloom filter false positives: 2
```

```
Bloom filter false ratio: 0.00000
```

```
Bloom filter space used: 17159136
```

```
Bloom filter off heap memory used: 17159096
```

```
Index summary off heap memory used: 2358092
```

```
Compression metadata off heap memory used: 0
```

```
Compacted partition minimum bytes: 180
```

```
Compacted partition maximum bytes: 258
```

```
Compacted partition mean bytes: 258
```

```
Average live cells per slice (last five minutes): 1.0
```

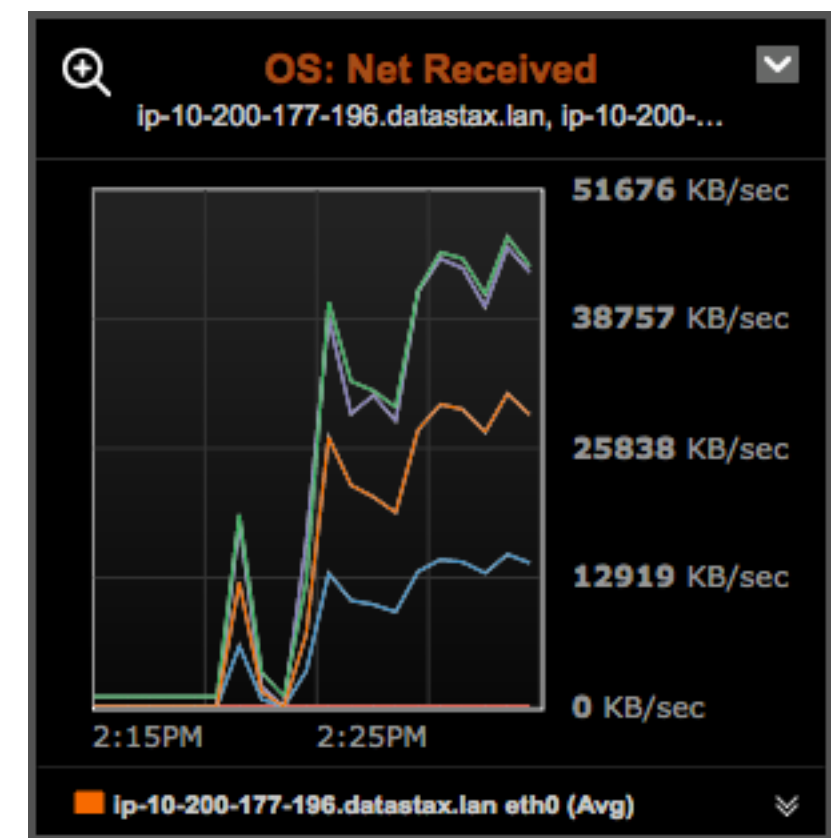
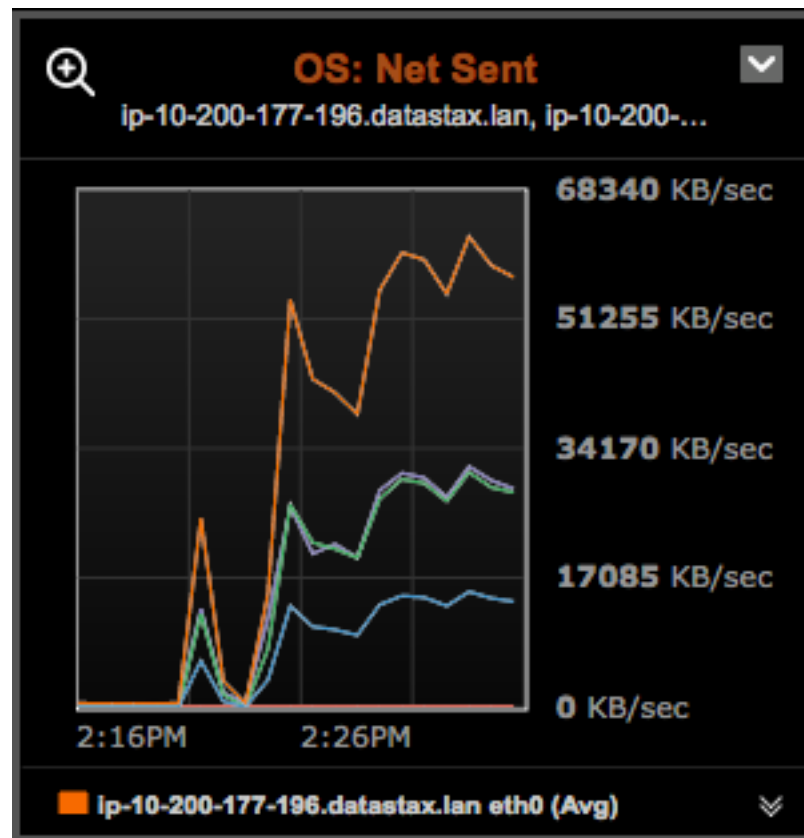
```
Maximum live cells per slice (last five minutes): 1
```

```
Average tombstones per slice (last five minutes): 1.0
```

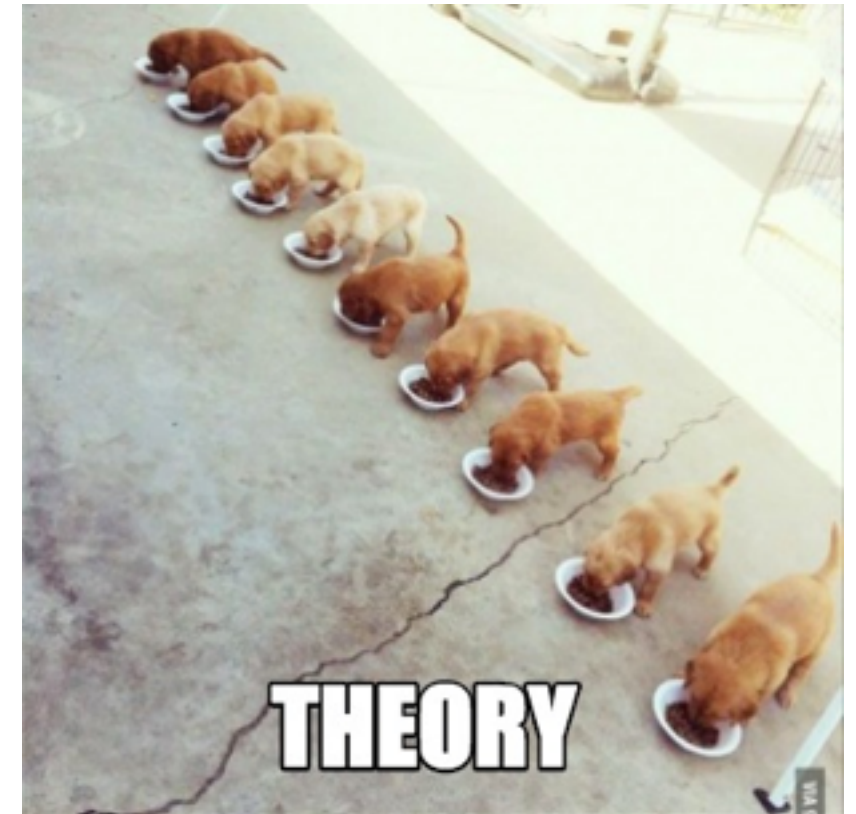
```
Maximum tombstones per slice (last five minutes): 1
```

Network

Network Utilization



Background Processes



Flushing

Memtable Size

```
$ nodetool tablestats keyspace1
```

Keyspace: keyspace1

Read Count: 4433559

Read Latency: 0.16309469277390917 ms.

Write Count: 15528537

Write Latency: 0.03023585988815302 ms.

Pending Flushes: 0

Table: standard1

SSTable count: 5

Space used (live): 3464993787

Space used (total): 3464993787

Space used by snapshots (total): 0

Off heap memory used (total): 19517188

SSTable Compression Ratio: 0.0

Number of keys (estimate): 10196252

Memtable cell count: 123604

Memtable data size: 34015680

Memtable off heap memory used: 0

Memtable switch count: 38

Local read count: 4433559

Local read latency: 0.179 ms

Local write count: 15528537

Local write latency: 0.034 ms

Pending flushes: 0

Bloom filter false positives: 2

Bloom filter false ratio: 0.00000

Bloom filter space used: 17159136

Bloom filter off heap memory used: 17159096

Index summary off heap memory used: 2358092

Compression metadata off heap memory used: 0

Compacted partition minimum bytes: 180

Compacted partition maximum bytes: 258

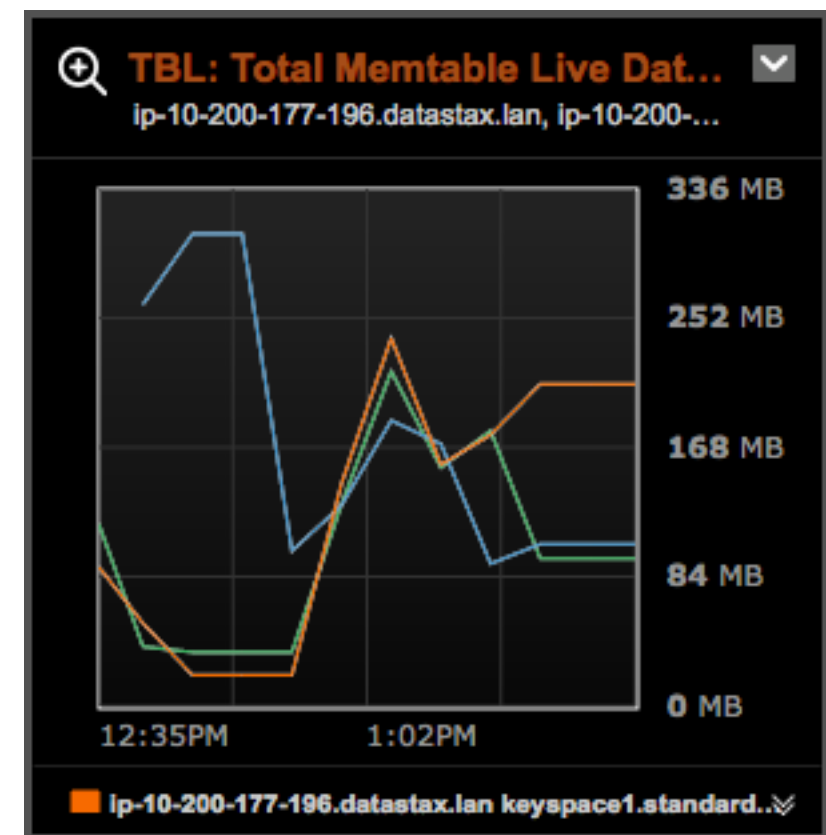
Compacted partition mean bytes: 258

Average live cells per slice (last five minutes): 1.0

Maximum live cells per slice (last five minutes): 1

Average tombstones per slice (last five minutes): 1.0

Maximum tombstones per slice (last five minutes): 1



cassandra.yaml settings

memtable_heap_space_in_mb: 2048

memtable_offheap_space_in_mb: 2048

memtable_allocation_type: heap_buffers

Enqueueing Flushes

Thread

Keyspace

Table

DEBUG [**SlabPoolCleaner**] 2016-08-25 19:16:57,771 ColumnFamilyStore.java:1186 -
Flushing largest CFS(Keyspace='**keyspace1**', ColumnFamily='**standard1**') to free up room.
Used total: **0.64**/0.00, live: **0.33**/0.00, flushing: **0.31**/0.00, this: **0.31**/0.31

Total % used

Cleanup threshold

This memtable

DEBUG [SlabPoolCleaner] 2016-08-25 19:16:57,771 ColumnFamilyStore.java:845 -
Enqueueing flush of standard1: **321247911** (31%) on-heap, 0 (0%) off-heap

Size of this memtable

percent of total size to trigger flush
memtable_cleanup_threshold: 0.33

Enqueueing Flushes

Different thread!

DEBUG [**COMMIT-LOG-ALLOCATOR**] 2016-08-26 14:58:28,166 ColumnFamilyStore.java:
845 - Enqueueing flush of standard1: 280740184 (**27%**) on-heap, 0 (0%) off-heap

Less than cleanup threshold

total commitlog size in cassandra.yaml
commitlog_total_space_in_mb: 8192

Compaction

Compactions

```
$ nodetool compactionstats
```

pending tasks: 3

Pending

Keyspace/table

| action type | keyspace | table | completed | total | unit | progress |
|-------------|-----------|-----------|-----------|------------|-------|----------|
| Compaction | keyspace1 | standard1 | 508710090 | 1283148921 | bytes | 39.65% |
| Compaction | keyspace1 | standard1 | 54302737 | 1275330433 | bytes | 4.26% |

Active compaction remaining time : 0h01m58s

Progress

configure throughput in cassandra.yaml (0 to unthrottle)
compaction_throughput_mb_per_sec: 16

configure throughput temporarily in nodetool
nodetool setcompactionthroughput 16

configure number of compactors in cassandra.yaml
concurrent_compactors: 1




change compaction strategy in CQL
alter table foo with compaction = {
'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy',
'max_threshold': '32', 'min_threshold': '4'}

Compactions in OpsCenter

jb50test: Activities Cluster Actions

GENERAL EVENT LOG

| ACTIVITY | HOSTNAME | DESCRIPTION | STATUS |
|---------------------|--------------------------------|-----------------------------------|------------------------------------|
| Cassandra Operation | ip-10-200-177-196.datastax.lan | Compaction of keyspace1.standard1 | Running <div><div></div></div> 49% |
| Cassandra Operation | ip-10-200-177-197.datastax.lan | Compaction of keyspace1.standard1 | Running <div><div></div></div> 43% |
| Cassandra Operation | ip-10-200-177-198.datastax.lan | Compaction of keyspace1.standard1 | Running <div><div></div></div> 47% |

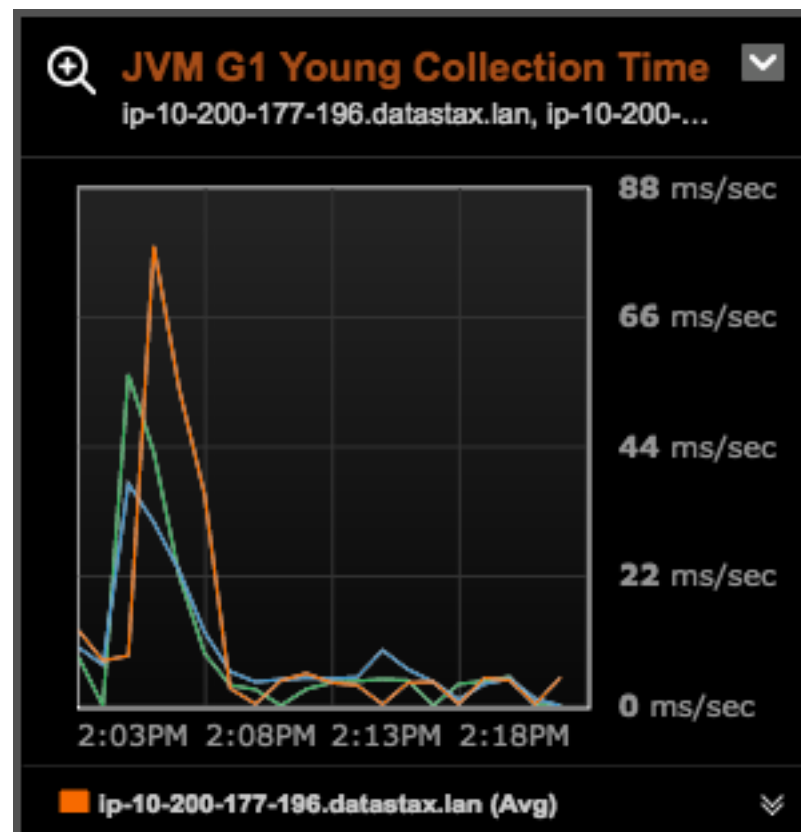
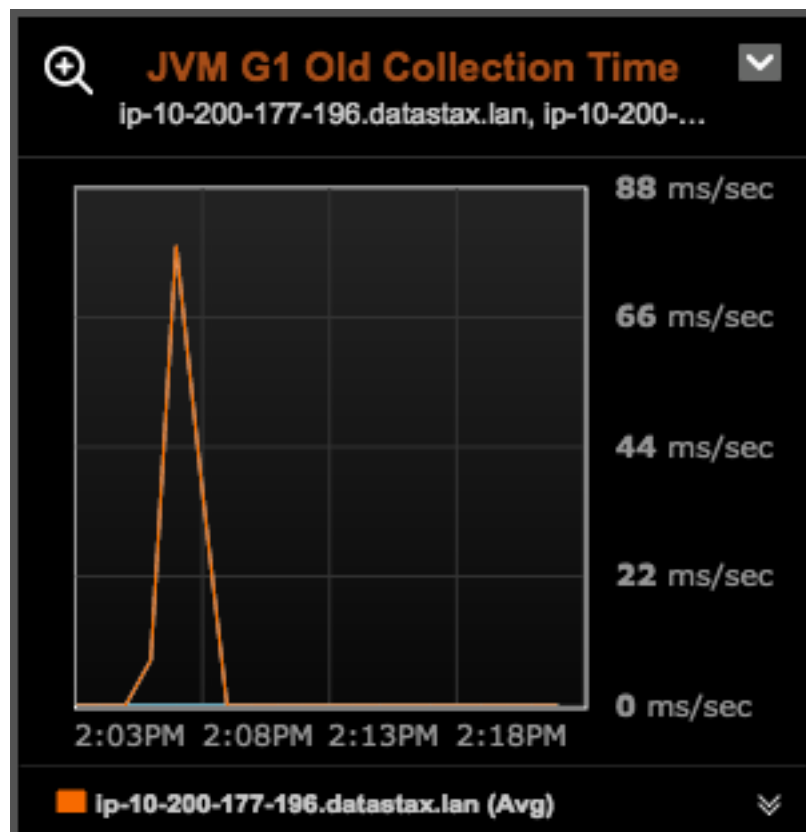


Garbage Collection

Garbage Collections

INFO [Service Thread] **2016-08-30 18:06:18,444** GCInspector.java:284 - G1 Young Generation GC in **609ms**. G1 Eden Space: 990904320 -> 0; G1 Old Gen: 1693531744 -> 2522955360; G1 Survivor Space: 4194304 -> 9437184;

WARN [Service Thread] **2016-08-30 18:07:32,376** GCInspector.java:282 - G1 Old Generation GC in **1233ms**. G1 Old Gen: 2651134656 -> 1064769688; G1 Survivor Space: 7340032 -> 0;



Heap Allocation Rate

```
$ java -jar sjk-plus-0.4.2.jar ttop -p 22226 -o ALLOC -n 20
```

2016-08-17T18:54:00.188+0000 Process summary

process cpu=320.99%

application cpu=303.13% (user=174.83% sys=128.30%)

other: cpu=17.86%

heap allocation rate 218mb/s

| | | | | |
|----------|-------------|------------|-----------------|---|
| [000224] | user= 6.39% | sys= 4.00% | alloc= 29mb/s | - MessagingService-Incoming-/10.200.177.196 |
| [000293] | user= 7.89% | sys= 5.05% | alloc= 29mb/s | - MessagingService-Incoming-/10.200.177.198 |
| [000305] | user= 6.79% | sys= 3.17% | alloc= 14mb/s | - CompactionExecutor:4 |
| [000157] | user= 6.19% | sys= 4.00% | alloc= 6727kb/s | - SharedPool-Worker-5 |
| [000153] | user= 7.49% | sys= 3.51% | alloc= 6596kb/s | - SharedPool-Worker-2 |
| [000152] | user= 7.49% | sys= 3.41% | alloc= 6436kb/s | - SharedPool-Worker-1 |
| [000155] | user= 6.39% | sys= 3.80% | alloc= 6153kb/s | - SharedPool-Worker-4 |
| [000154] | user= 6.89% | sys= 3.29% | alloc= 6054kb/s | - SharedPool-Worker-3 |
| [000156] | user= 6.39% | sys= 3.68% | alloc= 6030kb/s | - SharedPool-Worker-6 |
| [000168] | user= 5.89% | sys= 3.33% | alloc= 5947kb/s | - SharedPool-Worker-8 |
| [000181] | user= 5.49% | sys= 3.28% | alloc= 5343kb/s | - SharedPool-Worker-9 |
| [000169] | user= 6.39% | sys= 2.86% | alloc= 5202kb/s | - SharedPool-Worker-7 |
| [000179] | user= 4.80% | sys= 2.29% | alloc= 4969kb/s | - SharedPool-Worker-12 |
| [000158] | user= 5.29% | sys= 2.43% | alloc= 4830kb/s | - SharedPool-Worker-11 |
| [000159] | user= 4.60% | sys= 2.06% | alloc= 4734kb/s | - SharedPool-Worker-13 |
| [000178] | user= 3.80% | sys= 2.30% | alloc= 4304kb/s | - SharedPool-Worker-14 |
| [000180] | user= 4.90% | sys= 2.45% | alloc= 4294kb/s | - SharedPool-Worker-10 |
| [000160] | user= 3.90% | sys= 1.55% | alloc= 4160kb/s | - SharedPool-Worker-15 |
| [000177] | user= 3.30% | sys= 1.29% | alloc= 3268kb/s | - SharedPool-Worker-16 |
| [000247] | user= 2.40% | sys= 3.48% | alloc= 3253kb/s | - epollEventLoopGroup-6-2 |

Per-Thread Heap Allocation Rate

```
$ java -jar sjk-plus-0.4.2.jar ttop -p 22226 -o ALLOC -n 20
```

```
2016-08-17T18:54:00.188+0000 Process summary
```

```
process cpu=320.99%
```

```
application cpu=303.13% (user=174.83% sys=128.30%)
```

```
other: cpu=17.86%
```

```
heap allocation rate 218mb/s
```

| | | | | |
|----------|-------------|------------|----------------------|--|
| [000224] | user= 6.39% | sys= 4.00% | alloc= 29mb/s | – MessagingService-Incoming-/10.200.177.196 |
| [000293] | user= 7.89% | sys= 5.05% | alloc= 29mb/s | – MessagingService-Incoming-/10.200.177.198 |
| [000305] | user= 6.79% | sys= 3.17% | alloc= 14mb/s | – CompactionExecutor:4 |
| [000157] | user= 6.19% | sys= 4.00% | alloc= 6727kb/s | – SharedPool-Worker-5 |
| [000153] | user= 7.49% | sys= 3.51% | alloc= 6596kb/s | – SharedPool-Worker-2 |
| [000152] | user= 7.49% | sys= 3.41% | alloc= 6436kb/s | – SharedPool-Worker-1 |
| [000155] | user= 6.39% | sys= 3.80% | alloc= 6153kb/s | – SharedPool-Worker-4 |
| [000154] | user= 6.89% | sys= 3.29% | alloc= 6054kb/s | – SharedPool-Worker-3 |
| [000156] | user= 6.39% | sys= 3.68% | alloc= 6030kb/s | – SharedPool-Worker-6 |
| [000168] | user= 5.89% | sys= 3.33% | alloc= 5947kb/s | – SharedPool-Worker-8 |
| [000181] | user= 5.49% | sys= 3.28% | alloc= 5343kb/s | – SharedPool-Worker-9 |
| [000169] | user= 6.39% | sys= 2.86% | alloc= 5202kb/s | – SharedPool-Worker-7 |
| [000179] | user= 4.80% | sys= 2.29% | alloc= 4969kb/s | – SharedPool-Worker-12 |
| [000158] | user= 5.29% | sys= 2.43% | alloc= 4830kb/s | – SharedPool-Worker-11 |
| [000159] | user= 4.60% | sys= 2.06% | alloc= 4734kb/s | – SharedPool-Worker-13 |
| [000178] | user= 3.80% | sys= 2.30% | alloc= 4304kb/s | – SharedPool-Worker-14 |
| [000180] | user= 4.90% | sys= 2.45% | alloc= 4294kb/s | – SharedPool-Worker-10 |
| [000160] | user= 3.90% | sys= 1.55% | alloc= 4160kb/s | – SharedPool-Worker-15 |
| [000177] | user= 3.30% | sys= 1.29% | alloc= 3268kb/s | – SharedPool-Worker-16 |
| [000247] | user= 2.40% | sys= 3.48% | alloc= 3253kb/s | – epollEventLoopGroup-6-2 |

Repairs & Streaming

nodetool netstats - repair id

Mode: NORMAL

Repair 028763b0-cc1e-11e4-a20c-a1d01a3fbf30

/54.174.19.98

Receiving 6 files, 117949006 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-162-Data.db

851792/17950738 bytes(4%) received from /54.174.19.98

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

3786324/46561942 bytes(8%) sent to /54.174.19.98

Repair 020ed850-cc1e-11e4-a20c-a1d01a3fbf30

/54.174.245.247

Receiving 4 files, 93304584 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-161-Data.db

6094594/46561942 bytes(13%) received from /54.174.245.247

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

34195028/46561942 bytes(73%) sent to /54.174.245.247

Repair 018c88f0-cc1e-11e4-a20c-a1d01a3fbf30

/54.153.39.203 (using /172.31.10.65)

Receiving 3 files, 49959102 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-160-Data.db

9371380/46561942 bytes(20%) received from /54.153.39.203

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-159-Data.db

2533414/2533414 bytes(100%) received from /54.153.39.203

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-158-Data.db

1147584/1147584 bytes(100%) sent to /54.153.39.203

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

46561942/46561942 bytes(100%) sent to /54.153.39.203

Read Repair Statistics:

Attempted: 39576

Mismatch (Blocking): 0

Mismatch (Background): 746

| Pool Name | Active | Pending | Completed |
|-----------|--------|---------|-----------|
| Commands | n/a | 58 | 2545817 |
| Responses | n/a | 0 | 2833081 |

nodetool netstats - sstable names

Mode: NORMAL

Repair 028763b0-cc1e-11e4-a20c-a1d01a3fbf30

/54.174.19.98

Receiving 6 files, 117949006 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-162-Data.db

851792/17950738 bytes(4%) received from /54.174.19.98

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

3786324/46561942 bytes(8%) sent to /54.174.19.98

Repair 020ed850-cc1e-11e4-a20c-a1d01a3fbf30

/54.174.245.247

Receiving 4 files, 93304584 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-161-Data.db

6094594/46561942 bytes(13%) received from /54.174.245.247

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

34195028/46561942 bytes(73%) sent to /54.174.245.247

Repair 018c88f0-cc1e-11e4-a20c-a1d01a3fbf30

/54.153.39.203 (using /172.31.10.65)

Receiving 3 files, 49959102 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-160-Data.db

9371380/46561942 bytes(20%) received from /54.153.39.203

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-159-Data.db

2533414/2533414 bytes(100%) received from /54.153.39.203

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-158-Data.db

1147584/1147584 bytes(100%) sent to /54.153.39.203

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

46561942/46561942 bytes(100%) sent to /54.153.39.203

Read Repair Statistics:

Attempted: 39576

Mismatch (Blocking): 0

Mismatch (Background): 746

| Pool Name | Active | Pending | Completed |
|-----------|--------|---------|-----------|
| Commands | n/a | 58 | 2545817 |
| Responses | n/a | 0 | 2833081 |

nodetool netstats - streaming progress

Mode: NORMAL

Repair 028763b0-cc1e-11e4-a20c-a1d01a3fbf30

/54.174.19.98

Receiving 6 files, 117949006 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-162-Data.db

851792/17950738 bytes(4%) received from /54.174.19.98

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

3786324/46561942 bytes(8%) sent to /54.174.19.98

Repair 020ed850-cc1e-11e4-a20c-a1d01a3fbf30

/54.174.245.247

Receiving 4 files, 93304584 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-161-Data.db

6094594/46561942 bytes(13%) received from /54.174.245.247

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

34195028/46561942 bytes(73%) sent to /54.174.245.247

Repair 018c88f0-cc1e-11e4-a20c-a1d01a3fbf30

/54.153.39.203 (using /172.31.10.65)

Receiving 3 files, 49959102 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-160-Data.db

9371380/46561942 bytes(20%) received from /54.153.39.203

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-tmp-jb-159-Data.db

2533414/2533414 bytes(100%) received from /54.153.39.203

Sending 2 files, 47709526 bytes total

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-158-Data.db

1147584/1147584 bytes(100%) sent to /54.153.39.203

/var/lib/cassandra/data/Keyspace1/Standard1/Keyspace1-Standard1-jb-157-Data.db

46561942/46561942 bytes(100%) sent to /54.153.39.203

Read Repair Statistics:

Attempted: 39576

Mismatch (Blocking): 0

Mismatch (Background): 746

| Pool Name | Active | Pending | Completed |
|-----------|--------|---------|-----------|
| Commands | n/a | 58 | 2545817 |
| Responses | n/a | 0 | 2833081 |

Repairs - unique id

INFO [AntiEntropySessions:1] 2015-03-16 17:24:44,097 RepairSession.java (line 246)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] new session: will sync /54.174.19.98, /54.153.108.157, /54.153.39.203, /54.174.245.247 on range (3074457345618258602,6148914691236517205] for Keyspace1.[Counter3, Standard1, SuperCounter1, Super1, Counter1]

INFO [RepairJobTask:2] 2015-03-16 17:24:48,102 RepairJob.java (line 161)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] requesting merkle trees for Standard1 (to [/54.153.108.157, /54.153.39.203, /54.174.245.247, /54.174.19.98])

INFO [AntiEntropyStage:1] 2015-03-16 17:24:52,749 RepairSession.java (line 166)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Received merkle tree for Standard1 from /54.174.19.98

INFO [RepairJobTask:4] 2015-03-16 17:24:52,757 Differencer.java (line 67)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Endpoints /54.153.39.203 and /54.174.245.247 are consistent for Standard1

INFO [RepairJobTask:6] 2015-03-16 17:24:55,438 Differencer.java (line 74)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Endpoints /54.153.108.157 and /54.174.19.98 have 6222 range(s) out of sync for Standard1

INFO [RepairJobTask:6] 2015-03-16 17:24:55,441 StreamingRepairTask.java (line 64)
[streaming task #562f46d0-cc01-11e4-a84d-098a653a7013] Performing streaming repair of 6222 ranges with /54.153.108.157

INFO [StreamReceiveTask:1] 2015-03-16 17:25:01,524 StreamingRepairTask.java (line 92)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] streaming task succeed, returning response to /54.174.19.98

INFO [AntiEntropyStage:1] 2015-03-16 17:25:01,805 RepairSession.java (line 223)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Standard1 is fully synced

INFO [AntiEntropySessions:1] 2015-03-16 17:25:01,806 RepairSession.java (line 284)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] session completed successfully

Repairs - session started

INFO [AntiEntropySessions:1] 2015-03-16 17:24:44,097 RepairSession.java (line 246)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] **new session**: will sync /54.174.19.98, /54.153.108.157, /54.153.39.203, /54.174.245.247 on range (3074457345618258602,6148914691236517205] for Keyspace1.[Counter3, Standard1, SuperCounter1, Super1, Counter1]

INFO [RepairJobTask:2] 2015-03-16 17:24:48,102 RepairJob.java (line 161)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] requesting merkle trees for Standard1 (to [/54.153.108.157, /54.153.39.203, /54.174.245.247, /54.174.19.98])

INFO [AntiEntropyStage:1] 2015-03-16 17:24:52,749 RepairSession.java (line 166)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Received merkle tree for Standard1 from /54.174.19.98

INFO [RepairJobTask:4] 2015-03-16 17:24:52,757 Differencer.java (line 67)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Endpoints /54.153.39.203 and /54.174.245.247 are consistent for Standard1

INFO [RepairJobTask:6] 2015-03-16 17:24:55,438 Differencer.java (line 74)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Endpoints /54.153.108.157 and /54.174.19.98 have 6222 range(s) out of sync for Standard1

INFO [RepairJobTask:6] 2015-03-16 17:24:55,441 StreamingRepairTask.java (line 64)
[streaming task #562f46d0-cc01-11e4-a84d-098a653a7013] Performing streaming repair of 6222 ranges with /54.153.108.157

INFO [StreamReceiveTask:1] 2015-03-16 17:25:01,524 StreamingRepairTask.java (line 92)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] streaming task succeed, returning response to /54.174.19.98

INFO [AntiEntropyStage:1] 2015-03-16 17:25:01,805 RepairSession.java (line 223)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Standard1 is fully synced

INFO [AntiEntropySessions:1] 2015-03-16 17:25:01,806 RepairSession.java (line 284)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] session completed successfully

Repairs - session complete

INFO [AntiEntropySessions:1] 2015-03-16 17:24:44,097 RepairSession.java (line 246)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] new session: will sync /54.174.19.98, /54.153.108.157, /54.153.39.203, /54.174.245.247 on range (3074457345618258602,6148914691236517205] for Keyspace1.[Counter3, Standard1, SuperCounter1, Super1, Counter1]

INFO [RepairJobTask:2] 2015-03-16 17:24:48,102 RepairJob.java (line 161)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] requesting merkle trees for Standard1 (to [/54.153.108.157, /54.153.39.203, /54.174.245.247, /54.174.19.98])

INFO [AntiEntropyStage:1] 2015-03-16 17:24:52,749 RepairSession.java (line 166)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Received merkle tree for Standard1 from /54.174.19.98

INFO [RepairJobTask:4] 2015-03-16 17:24:52,757 Differencer.java (line 67)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Endpoints /54.153.39.203 and /54.174.245.247 are consistent for Standard1

INFO [RepairJobTask:6] 2015-03-16 17:24:55,438 Differencer.java (line 74)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Endpoints /54.153.108.157 and /54.174.19.98 have 6222 range(s) out of sync for Standard1

INFO [RepairJobTask:6] 2015-03-16 17:24:55,441 StreamingRepairTask.java (line 64)
[streaming task #562f46d0-cc01-11e4-a84d-098a653a7013] Performing streaming repair of 6222 ranges with /54.153.108.157

INFO [StreamReceiveTask:1] 2015-03-16 17:25:01,524 StreamingRepairTask.java (line 92)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] streaming task succeed, returning response to /54.174.19.98

INFO [AntiEntropyStage:1] 2015-03-16 17:25:01,805 RepairSession.java (line 223)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] Standard1 is fully synced

INFO [AntiEntropySessions:1] 2015-03-16 17:25:01,806 RepairSession.java (line 284)
[repair #562f46d0-cc01-11e4-a84d-098a653a7013] **session completed successfully**

Errors



Exception Count - nodetool info

\$ nodetool info

```

ID : 108af27a-43d8-4814-b617-f8f93ba2bb0e
Gossip active : false
Thrift active : false
Native Transport active: false
Load : 49.56 GB
Generation No : 0
Uptime (seconds) : 263824
Heap Memory (MB) : 2698.60 / 4012.00
Off Heap Memory (MB) : 4.10
Data Center : Cassandra
Rack : rack1
Exceptions : 5
Key Cache : entries 12298, size 1.03 MB, capacity 100 MB, 2592495 hits, 7854163 requests
0.330 recent hit rate, 14400 save period in seconds
Row Cache : entries 0, size 0 bytes, capacity 0 bytes, 0 hits, 0 requests, NaN recent hit
rate, 0 save period in seconds
Counter Cache : entries 0, size 0 bytes, capacity 50 MB, 0 hits, 0 requests, NaN recent hit
rate, 7200 save period in seconds
Token : -9223372036854775808
  
```

Exceptions

java.io.EOFException

```
at java.io.DataInputStream.readFully(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
  (CacheService.java:356)
at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:386)
at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
  (CassandraDaemon.java:480)
```


Exceptions – stack trace

java.io.EOFException

```
at java.io.DataInputStream.readFully(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
  (CacheService.java:356)
at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:386)
at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
  (CassandraDaemon.java:480)
```

Exceptions – organization

java.io.EOFException

```
at java.io.DataInputStream.readFully(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
  (CacheService.java:356)
at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:386)
at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
  (CassandraDaemon.java:480)
```


Exceptions – subsystem

java.io.EOFException

```
at java.io.DataInputStream.readFully(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
  (CacheService.java:356)
at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:386)
at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
  (CassandraDaemon.java:480)
```

Exceptions – class

java.io.EOFException

```
at java.io.DataInputStream.readFully(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
(CacheService.java:356)
at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
(ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
(ColumnFamilyStore.java:386)
at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
(CassandraDaemon.java:480)
```

Exceptions – methods

java.io.EOFException

```
at java.io.DataInputStream.readFully(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
  (CacheService.java:356)
at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
  (ColumnFamilyStore.java:386)
at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
  (CassandraDaemon.java:480)
```

Exceptions – source files

```
java.io.EOFException
  at java.io.DataInputStream.readFully(DataInputStream.java:197)
  at java.io.DataInputStream.readFully(DataInputStream.java:169)
  at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
  at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
    (CacheService.java:356)
  at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
  at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
  at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
    (ColumnFamilyStore.java:415)
  at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
    (ColumnFamilyStore.java:386)
  at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
  at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
  at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
  at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
  at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
  at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
  at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
  at org.apache.cassandra.service.CassandraDaemon.activate
    (CassandraDaemon.java:480)
```


Exceptions – line numbers

```
java.io.EOFException
  at java.io.DataInputStream.readFully(DataInputStream.java:197)
  at java.io.DataInputStream.readFully(DataInputStream.java:169)
  at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
  at org.apache.cassandra.service.CacheService$KeyCacheSerializer.deserialize
    (CacheService.java:356)
  at org.apache.cassandra.cache.AutoSavingCache.loadSaved(AutoSavingCache.java:119)
  at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
  at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
    (ColumnFamilyStore.java:415)
  at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
    (ColumnFamilyStore.java:386)
  at org.apache.cassandra.db.Keyspace.initCf(Keyspace.java:309)
  at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
  at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
  at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
  at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
  at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
  at com.datastax.bdp.server.DseDaemon.setup(DseDaemon.java:376)
  at org.apache.cassandra.service.CassandraDaemon.activate
    (CassandraDaemon.java:480)
```

Exceptions – nested exceptions

```
org.apache.thrift.transport.TTransportException: ...  
  at org.apache.thrift.transport.TIOStreamTransport.read  
  at com.datastax.bdp.transport.server.TPreviewableTransport.readUntilEof  
  at com.datastax.bdp.transport.server.TPreviewableTransport.preview  
  at com.datastax.bdp.transport.server.TNegotiatingServerTransport.open  
  at com.datastax.bdp.transport.server.TNegotiatingServerTransport$...  
  at com.datastax.bdp.transport.server.TNegotiatingServerTransport$...  
  at org.apache.cassandra.thrift.CustomTThreadPoolServer$WorkerProcess.run  
  at java.util.concurrent.ThreadPoolExecutor.runWorker  
  at java.util.concurrent.ThreadPoolExecutor$Worker.run  
  at java.lang.Thread.run
```

Caused by: java.net.SocketException: Connection reset

```
  at java.net.SocketInputStream.read  
  at java.net.SocketInputStream.read  
  at java.io.BufferedInputStream.fill  
  at java.io.BufferedInputStream.read1  
  at java.io.BufferedInputStream.read  
  at org.apache.thrift.transport.TIOStreamTransport.read  
  ... 9 more
```

Exceptions – error message

```
org.apache.thrift.transport.TTransportException: ...
  at org.apache.thrift.transport.TIOStreamTransport.read
  at com.datastax.bdp.transport.server.TPreviewableTransport.readUntilEof
  at com.datastax.bdp.transport.server.TPreviewableTransport.preview
  at com.datastax.bdp.transport.server.TNegotiatingServerTransport.open
  at com.datastax.bdp.transport.server.TNegotiatingServerTransport$...
  at com.datastax.bdp.transport.server.TNegotiatingServerTransport$...
  at org.apache.cassandra.thrift.CustomThreadPoolServer$WorkerProcess.run
  at java.util.concurrent.ThreadPoolExecutor.runWorker
  at java.util.concurrent.ThreadPoolExecutor$Worker.run
  at java.lang.Thread.run
Caused by: java.net.SocketException: Connection reset
  at java.net.SocketInputStream.read
  at java.net.SocketInputStream.read
  at java.io.BufferedInputStream.fill
  at java.io.BufferedInputStream.read1
  at java.io.BufferedInputStream.read
  at org.apache.thrift.transport.TIOStreamTransport.read
  ... 9 more
```

Google Tips

Do

- Use exception **and** several package+class+method names
- Use quotation marks around individual elements
- Use “site:” to limit search to relevant web sites
 - **Mailing list:** site:www.mail-archive.com/
user@cassandra.apache.org
 - **JIRA:** site:issues.apache.org/jira/browse/CASSANDRA
 - **StackOverflow:** site:stackoverflow.com cassandra
- Narrow or broaden as necessary

Don't

- Include source file + line number
- Include specific numbers and strings

Exceptions – search terms

java.io.EOFException

at **java.io.DataInputStream.readFully**(DataInputStream.java:197)
at java.io.DataInputStream.readFully(DataInputStream.java:169)
at org.apache.cassandra.utils.ByteBufferUtil.read(ByteBufferUtil.java:395)
at **org.apache.cassandra.service.CacheService\$KeyCacheSerializer.deserialize**
(CacheService.java:356)
at **org.apache.cassandra.cache.AutoSavingCache.loadSaved**(AutoSavingCache.java:119)
at org.apache.cassandra.db.ColumnFamilyStore.<init>(ColumnFamilyStore.java:261)
at **org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore**
(ColumnFamilyStore.java:415)
at org.apache.cassandra.db.ColumnFamilyStore.createColumnFamilyStore
(ColumnFamilyStore.java:386)
at **org.apache.cassandra.db.Keyspace.initCf**(Keyspace.java:309)
at org.apache.cassandra.db.Keyspace.<init>(Keyspace.java:266)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:110)
at org.apache.cassandra.db.Keyspace.open(Keyspace.java:88)
at org.apache.cassandra.db.SystemKeyspace.checkHealth(SystemKeyspace.java:536)
at org.apache.cassandra.service.CassandraDaemon.setup(CassandraDaemon.java:246)
at **com.datastax.bdp.server.DseDaemon.setup**(DseDaemon.java:376)
at org.apache.cassandra.service.CassandraDaemon.activate
(CassandraDaemon.java:480)

Exceptions – search terms

org.apache.thrift.transport.TTransportException: ...

at org.apache.thrift.transport.TIOStreamTransport.read

at **com.datastax.bdp.transport.server.TPreviewableTransport.readUntilEof**

at com.datastax.bdp.transport.server.TPreviewableTransport.preview

at com.datastax.bdp.transport.server.TNegotiatingServerTransport.open

at com.datastax.bdp.transport.server.TNegotiatingServerTransport\$...

at com.datastax.bdp.transport.server.TNegotiatingServerTransport\$...

at org.apache.cassandra.thrift.CustomThreadPoolServer\$WorkerProcess.run

at java.util.concurrent.ThreadPoolExecutor.runWorker

at java.util.concurrent.ThreadPoolExecutor\$Worker.run

at java.lang.Thread.run

Caused by: **java.net.SocketException: Connection reset**

at java.net.SocketInputStream.read

at java.net.SocketInputStream.read

at java.io.BufferedInputStream.fill

at java.io.BufferedInputStream.read1

at java.io.BufferedInputStream.read

at **org.apache.thrift.transport.TIOStreamTransport.read**

... 9 more

Thanks!