

# Introduction to Cassandra.yaml

& friends

**THE LAST PICKLE**

# Hi, I'm Edward Capriolo.

@edwardcapriolo

<https://www.linkedin.com/in/edwardcapriolo>

<http://www.slideshare.net/edwardcapriolo>

Consultant

The Last Pickle

Cassandra user since (v 0.6.5)

White Plains, NY USA



improve Apache  
based solutions.



**THE LAST PICKLE**



**IN MY DAY THE SCHEMA AND  
CONFIG WAS IN XML**

**TO CHANGE SCHEMA YOU HAD TO MANUALLY COPY  
THE FILE TO ALL NODES. UP HILL BOTH WAYS.**

[imggenerator.net](http://imggenerator.net)

This talk is the ‘gateway’ talk...

Many ‘picklers’ (TLP staff) are covering some points I will quickly cover over in depth in other talks.

# Section Overview

1. Key configuration settings
2. Configuration outside of the yam1
3. Multi-system configuration settings
4. Advanced settings
5. Exotic settings

# Basic setup

1. `$ wget <apache-cassandra*.tar.gz>`
2. `$ tar -xf <apache-cassandra*.tar.gz>`
3. `$ apache-cassandra*/bin/cassandra`

Result:

Web scale distributed storage



Drop Mic.

Well almost...

We have to do a bit of configuration.



# Before we dive into config

```
cqlsh> CREATE KEYSPACE test WITH replication =  
{'class': 'SimpleStrategy', 'replication_factor' : 1};  
cqlsh> USE test;  
cqlsh:test> CREATE COLUMNFAMILY trip (src varchar,  
... dest varchar, PRIMARY KEY (src,dest));
```

```
cqlsh:test> INSERT INTO trip (src, dest) VALUES ('ny', 'ca');
```

```
cqlsh:test> SELECT * FROM trip;
```


src	dest
ny	ca

```
cqlsh:test> INSERT INTO trip (src, dest) VALUES ('fl', 'ca');  
cqlsh:test> SELECT * FROM trip;
```

src	dest
fl	ca
ny	ca


# Single Data Center

4 Nodes at Replication Factor 3



# Multiple Data Center

DC1: 4 Nodes at Replication Factor 3  
DC2: 2 Nodes at Replication Factor 2



# Where does the data go?

```
data_file_directories:
```

```
– /var/lib/cassandra/data
```

1. User data is stored in all listed directories
2. Do: fast seek'ing storage (SSD)
3. Do: ample free space (30% overhead)
4. Don't: Store on a SAN



# Commit log storage

commitlog\_directory:

- /var/lib/cassandra/commitlog

1. Stores unflushed mutations (write/deletes)
2. Don't: Assume these are log4j type logs
3. Do: use a dedicated disk if possible
4. Do: provide at least 10GB (write velocity)



Ok we now where  
(most of) the data goes...

How do clients connect?

# Default port binding

```
1 $ netstat -nl
2 tcp      0      0 127.0.0.1:7000      0.0.0.0:*           LISTEN
3 tcp6     0      0 127.0.0.1:9042     :::*                LISTEN
```

1. Cassandra does not bind to 0.0.0.0
2. 127.0.0.1 not web scale
3. 7000 is the “Storage Port” inter node traffic
4. 9042 is the “Native Port” client traffic

# Native transport

```
start_native_transport: true (default)
native_transport_port: 9042 (default)
listen_address: localhost
```

1. Change listen\_address to a client-reachable address
2. Do: consider transport security
3. Do: consider network routing performance
4. Don't: put nodes on a public network. EVAR



Outside the `yaml` file...

# cassandra-env.sh (& friends)

1. JVM and startup params defined outside the YAML
2. Newer version of c\* use jvm.options

# Memory usage

```
#MAX_HEAP_SIZE="1G"  
#HEAP_NEWSIZE="100M"
```

1.  $\max(\min(1/2 \text{ ram}, 1024\text{MB}), \min(1/4 \text{ ram}, 8\text{GB}))$
2. Do: set lower when experimenting with workstation
3. Do: leave ample free memory for disk cache

# JMX

```
1 JMX_PORT="7199"
2 if [ "$LOCAL_JMX" = "yes" ]; then
3   JVM_OPTS="$JVM_OPTS -Dcassandra.jmx.local.port=$JMX_PORT -XX:+DisableExplicitGC" ; else
4   JVM_OPTS="$JVM_OPTS -Dcom.sun.management.jmxremote.port=$JMX_PORT"
5   JVM_OPTS="$JVM_OPTS -Dcom.sun.management.jmxremote.rmi.port=$JMX_PORT"
6   JVM_OPTS="$JVM_OPTS -Dcom.sun.management.jmxremote.ssl=false"
7   JVM_OPTS="$JVM_OPTS -Dcom.sun.management.jmxremote.authenticate=true"
8   JVM_OPTS="$JVM_OPTS -Dcom.sun.management.jmxremote.password.file=/etc/cassandra/jmxremote.password"
```

1. bin/nodetool uses JMX to administer Cassandra
2. All management tools require password if set

Check out Nate's talk on Securing Cassandra to learn more



# Multi-node configurations

# Phi convict threshold

```
# phi_convict_threshold: 8
```

1. Threshold for failure detector
2. False positives make nodes appear down to peers
3. Do: Raise for flaky WAN networks 10 - 12

# Defining network topology

```
# endpoint_snitch: SimpleSnitch
```

1. Snitch with config data determines topology
2. Do: use SimpleSnitch for single switch/LAN
3. Consider: Multi DC to start

# Gossiping Property File Snitch

```
conf/cassandra-rackdc.properties  
dc=dc1  
rack=rack1
```

1. Information is propagated around the cluster
2. DC may not be physical but is a replication unit
3. Rack has impact on replication copies
4. Don't: Change rack unless you understand the impact



# Internode communications

```
internode_compression: all | dc | none  
inter_dc_tcp_nodelay: false
```

## 1. WAN can benefit from reduced size

```
server_encryption_options:  
  internode_encryption: none  
internode_authenticator:  
o.a.c.auth.AllowAllInternodeAuthenticator
```

## 2. Settings which server nodes use to communicate

# Broadcast address


```
broadcast_address: 1.2.3.4  
listen_on_broadcast_address: false  
broadcast_rpc_address: 1.2.3.4
```

1. Gossip a specific address (not bind address)
2. Useful in NAT and cloud environments

# Broadcast address

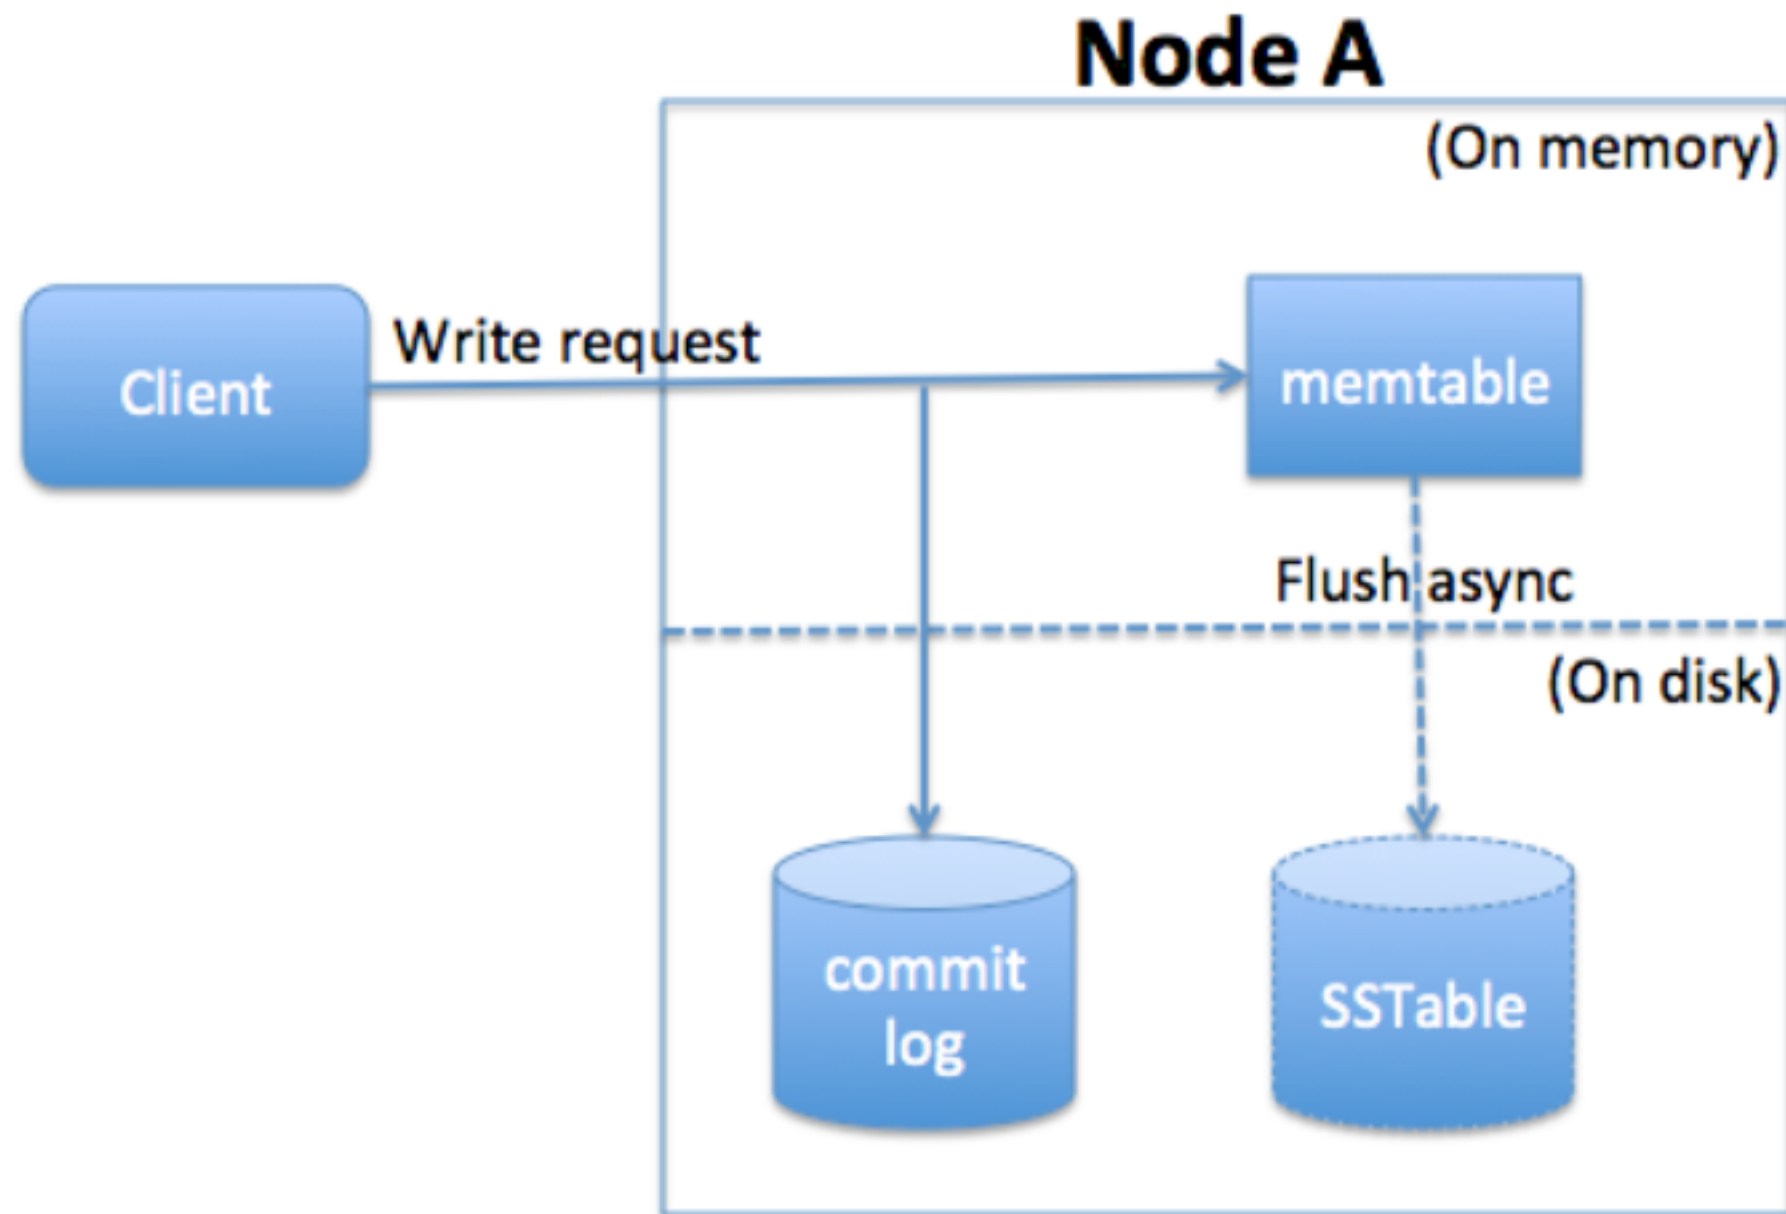
DC1: 4 Nodes at Replication Factor 3

DC2: 2 Nodes at Replication Factor 2



Advanced settings

# Write path



<http://www.toadworld.com/platforms/nosql/w/wiki/11621.an-introduction-to-apache-cassandra>

# Memtables

```
#memtable_flush_writers: 1
```

## 1. Default One per data directory

```
# memtable_cleanup_threshold  
defaults to 1 /  
(memtable_flush_writers +  
#memtable_cleanup_threshold: 0.11
```

2.  $1 / (1 + 1) = .5$

# .5 of what you ask?

```
#If omitted, both set to 1/4 the heap  
#memtable_heap_space_in_mb: 2048  
#memtable_offheap_space_in_mb: 2048
```

1. Depending on the next setting dictates how much of each memory type is used

```
#heap_buffers: on heap nio buffers  
#offheap_buffers: off heap nio buffers  
#offheap_objects: off heap objects  
#memtable_allocation_type: heap_buffers
```

2. Based on column value buffers vs objects may be better

# Trickle fsync

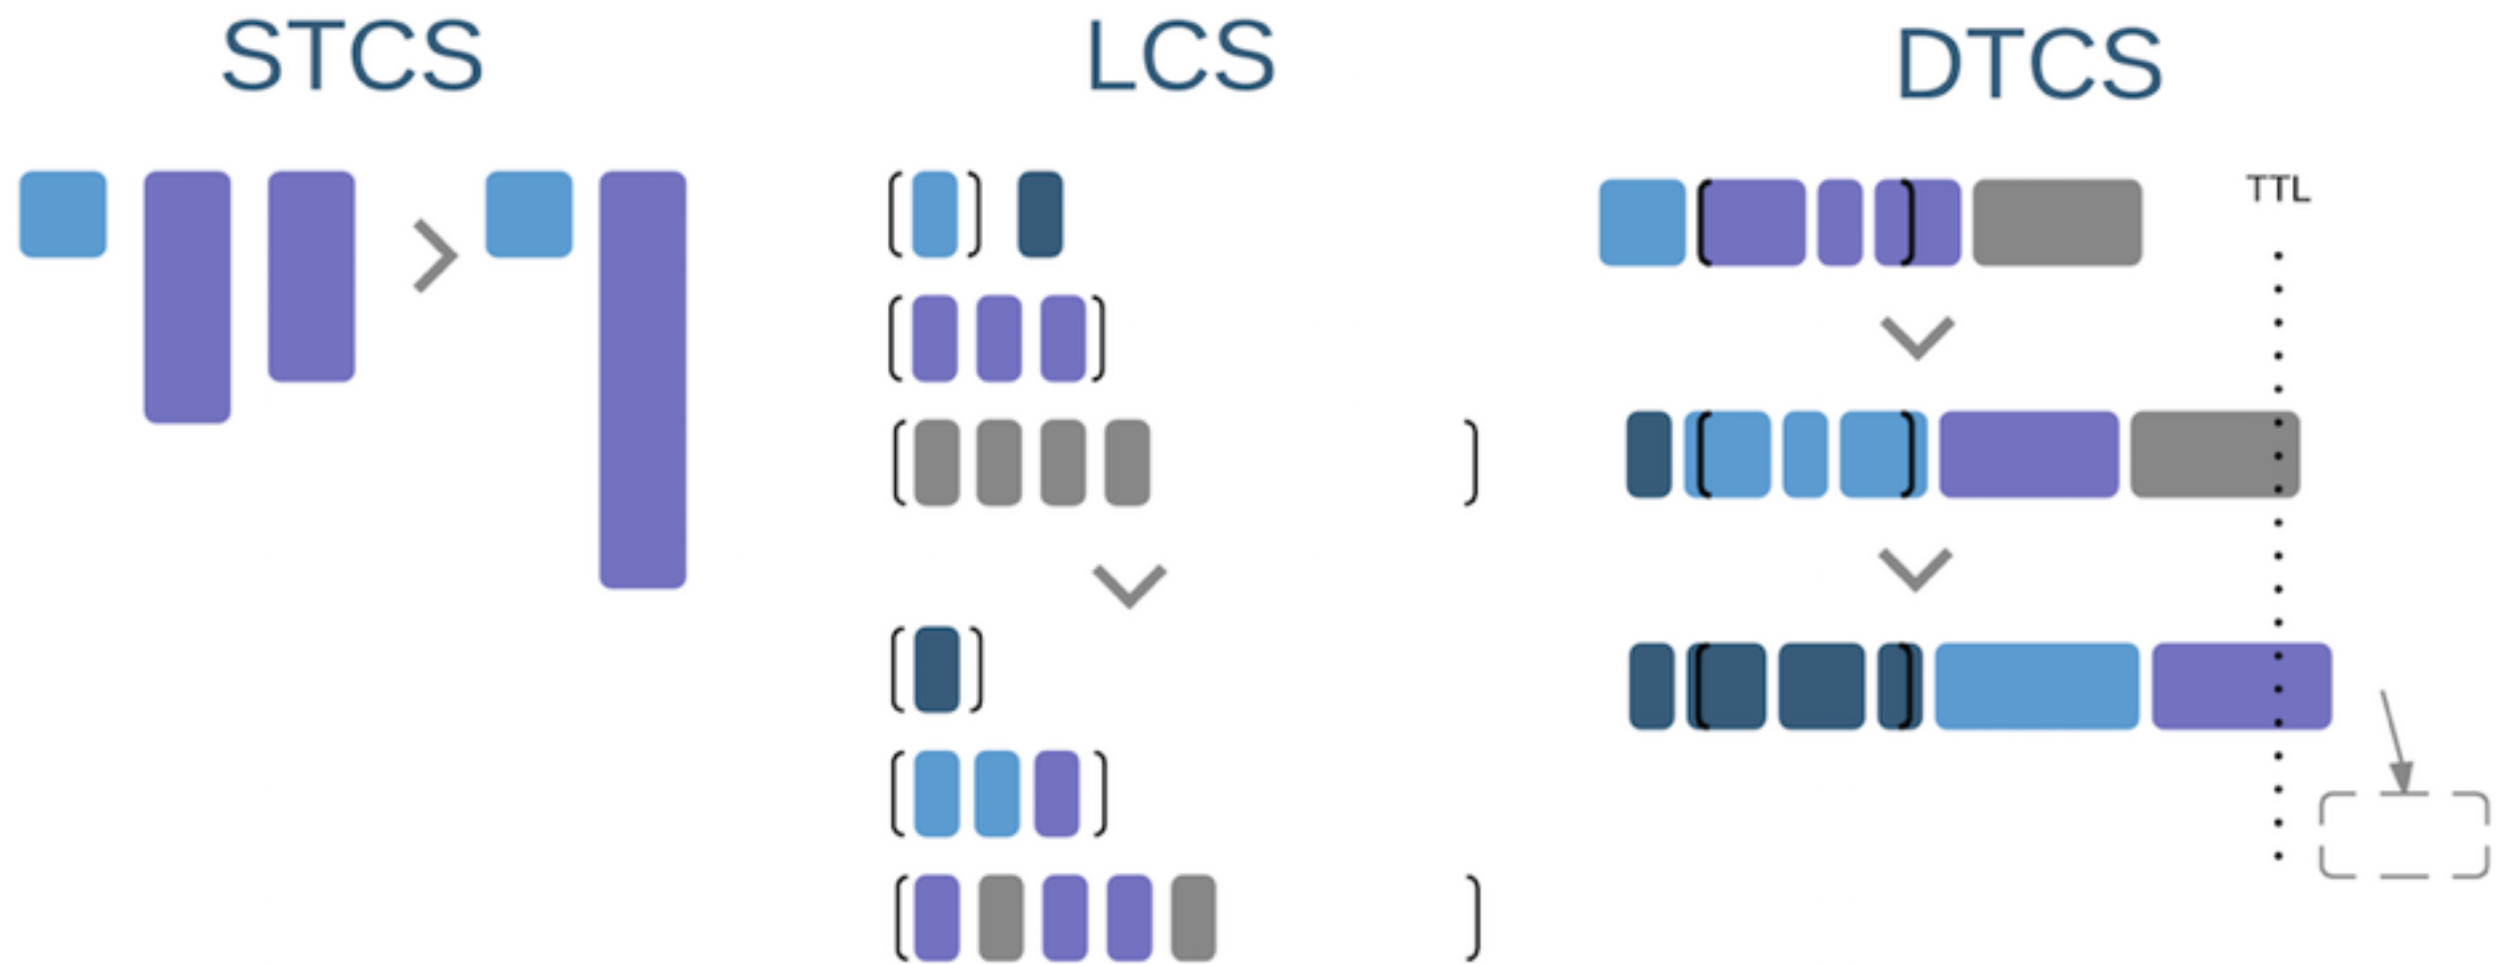
```
trickle_fsync: false  
trickle_fsync_interval_in_kb: 10240
```

1. Optimization to periodically f-sync large files
2. Designed to prevent latency spikes in read path





# Compaction



<https://www.instaclustr.com/blog/2016/01/27/apache-cassandra-compaction/>

# Compaction

```
concurrent_compactors: 1  
compaction_throughput_mb_per_sec: 16
```

1. Control resources used by compaction
2. Compaction throughput can be changed at runtime
3. Generally `concurrent_compactors < 8` and `> 1`

# Disk Failure settings

```
disk_failure_policy: stop  
commit_failure_policy: stop
```

1. stop\_paranoid: shut down gossip and client transports even for single-sstable errors, kill the JVM for errors during startup
2. die: shut down gossip and Thrift and kill the JVM, so the node can be replaced

# Hints

```
hinted_handoff_enabled: true
max_hint_window_in_ms: 10800000
hinted_handoff_throttle_in_kb: 1024
max_hints_delivery_threads: 2
hints_directory: /var/lib/cassandra/hints
hints_flush_period_in_ms: 10000
max_hints_file_size_in_mb: 128
hints_compression: LZ4Compressor
```

1. Hints recently redesigned, again again
2. Don't: tune high and overwhelming recovering node
3. Don't: tune low and have out of sync data

# Disk optimization strategy

```
#disk_optimization_strategy:  ssd
```

1. Tip for those with rota



Exotic settings

# Auto bootstrap

`auto_bootstrap : true(hidden variable)`

1. “Bootstrapping” here means: Should the node joining attempt to acquire data from other nodes or startup empty
2. Can be used when bringing on new datacenter
3. Can be used when streaming/ join issues

# Backup\*Ish options

```
incremental_backups: false  
snapshot_before_compaction: false  
auto_snapshot: true
```

1. Enable with external backup like tools
2. Creates hard link files operator must clean up
3. Enabling and not cleaning will cause disk fill up
4. Truncate/drop makes snapshot



# Per operation default timeouts

```
read_request_timeout_in_ms: 5000  
write_request_timeout_in_ms: 2000  
request_timeout_in_ms: 10000
```

1. Each operation type has different timeout
2. Applied on the coordinator not the client
3. Previously was only global rpc\_timeout

# Commit Log sync

```
commitlog_sync: periodic  
commitlog_sync_period_in_ms: 10000  
commitlog_segment_size_in_mb: 32
```

1. Alternative batch mode blocks ack to clients
2. Commit logs persist until Memtable's flush

# Thanks!

@edwardcapriolo

**THE LAST PICKLE**