



# Develop scalable applications with DataStax Drivers for Apache Cassandra and DataStax Enterprise

Alex Popescu & Bulat Shakirzyanov

- |   |                                 |
|---|---------------------------------|
| 1 | Overview of DataStax drivers    |
| 2 | Load balancing                  |
| 3 | Fault tolerance                 |
| 4 | Address resolution              |
| 5 | Sneak peek at upcoming features |



# DataStax Drivers

Smart clients for Apache Cassandra and DataStax Enterprise

# DataStax Drivers



# Goals of DataStax Drivers

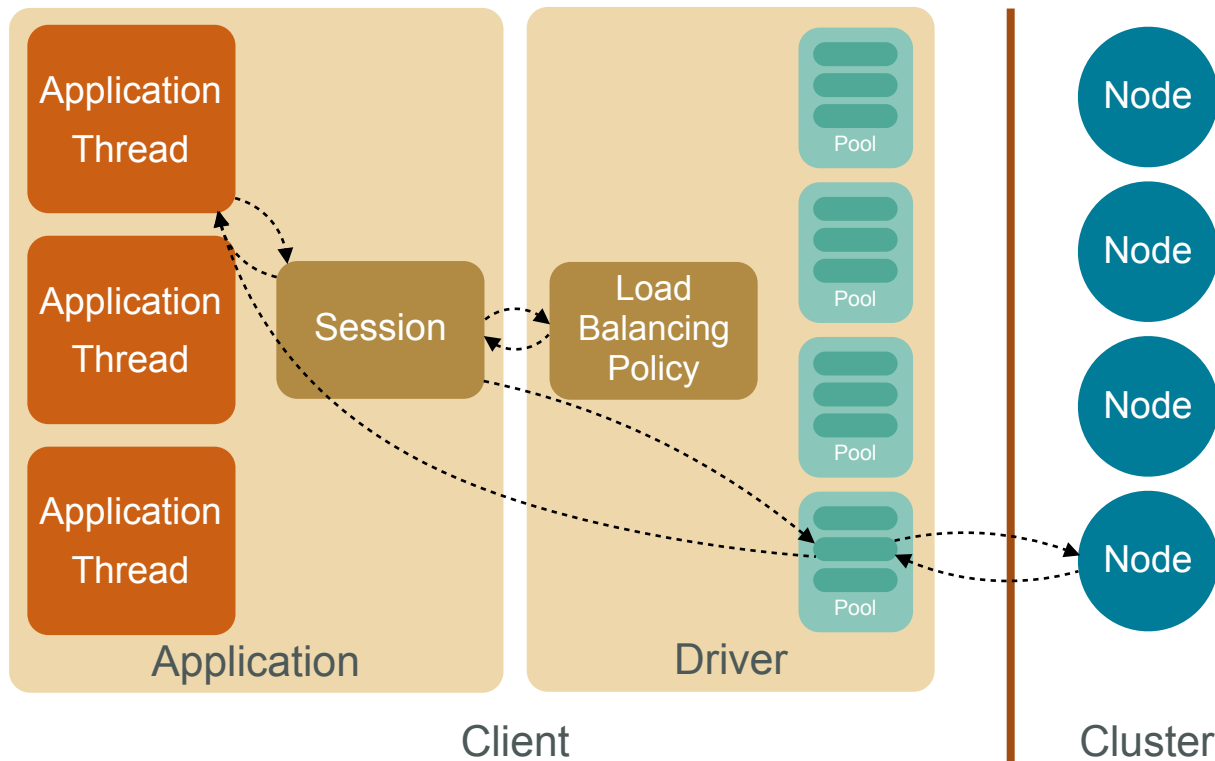
- Consistent set of features across languages
  - Asynchronous execution
  - Automatic cluster discovery
  - Connection pools and automatic reconnection
  - Load balancing
  - Fault tolerant
  - Address resolution
- Flexible to the core
- Consistent terminology
  - Cluster -> Session -> PreparedStatement & Statement -> Future or ResultSet




# Load Balancing

Principles and Implementations

# Load Balancing




# DataCenter Aware Balancing





# Token Aware Balancing







# Fault Tolerance

Sources of Failure and Error Handling

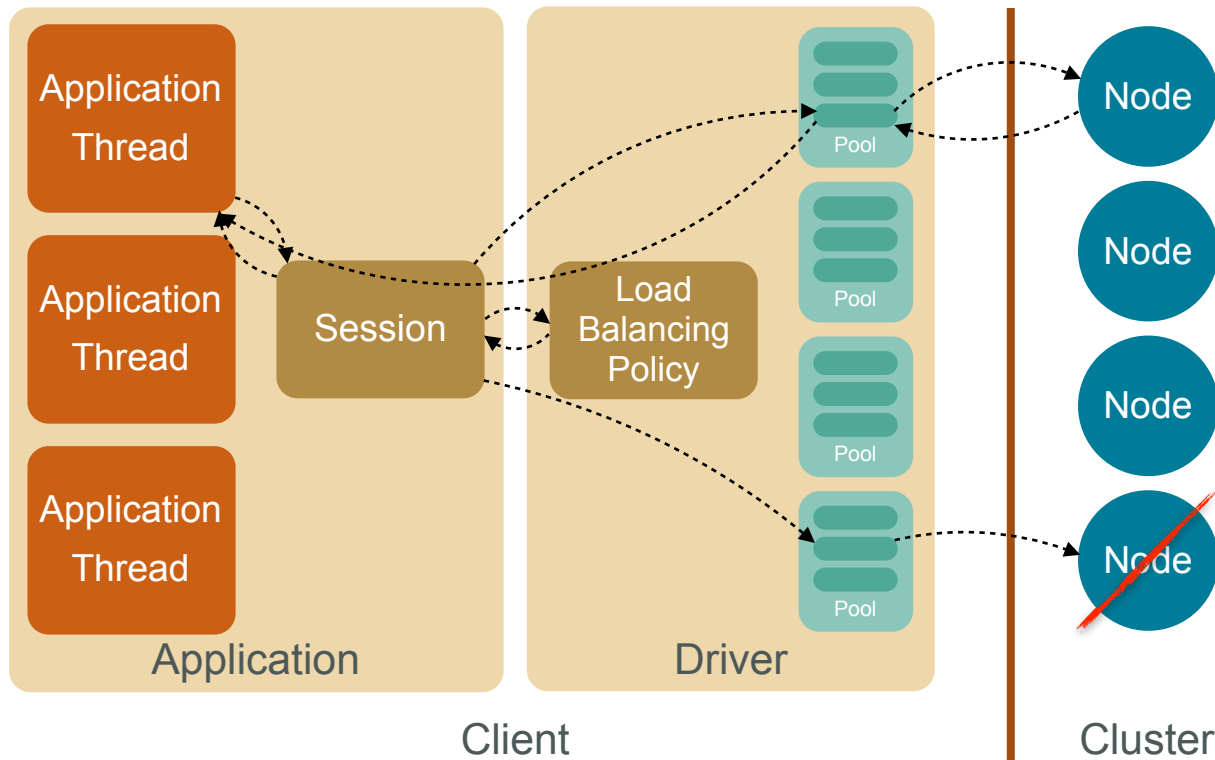
# Fault Tolerance



# Possible Failures



# Automatic Retry of Server Errors

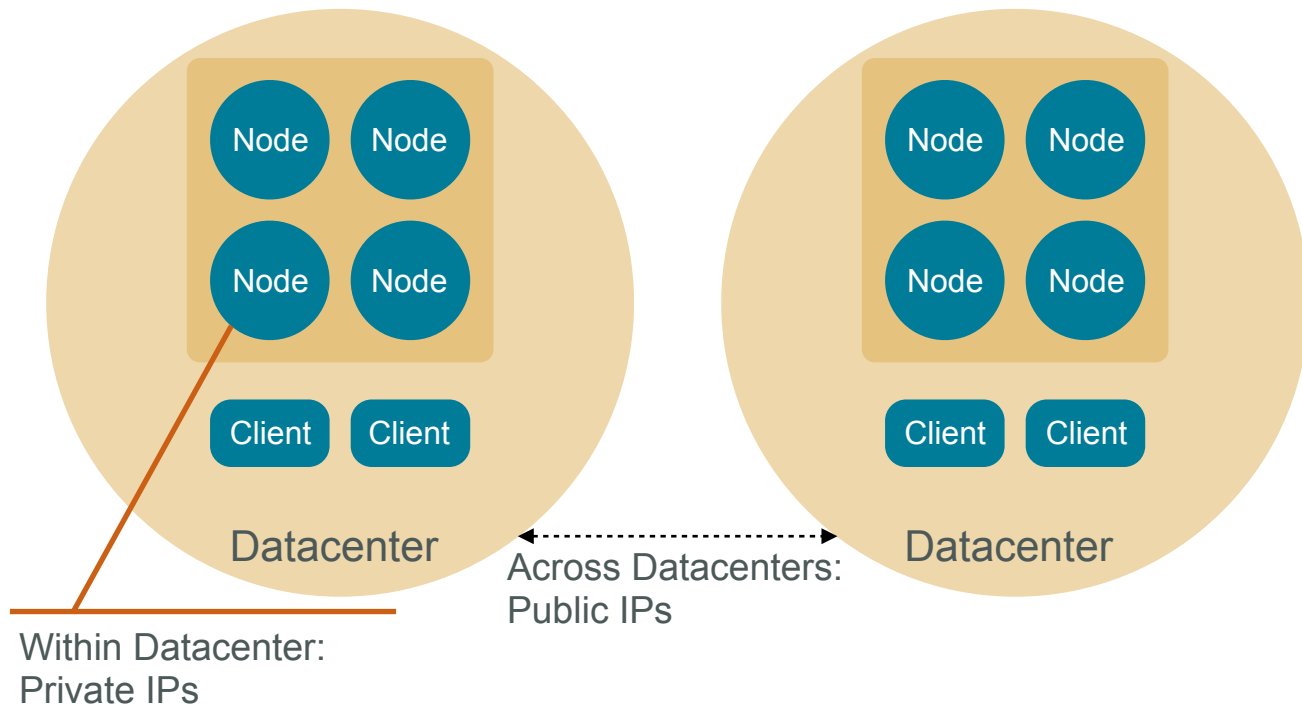




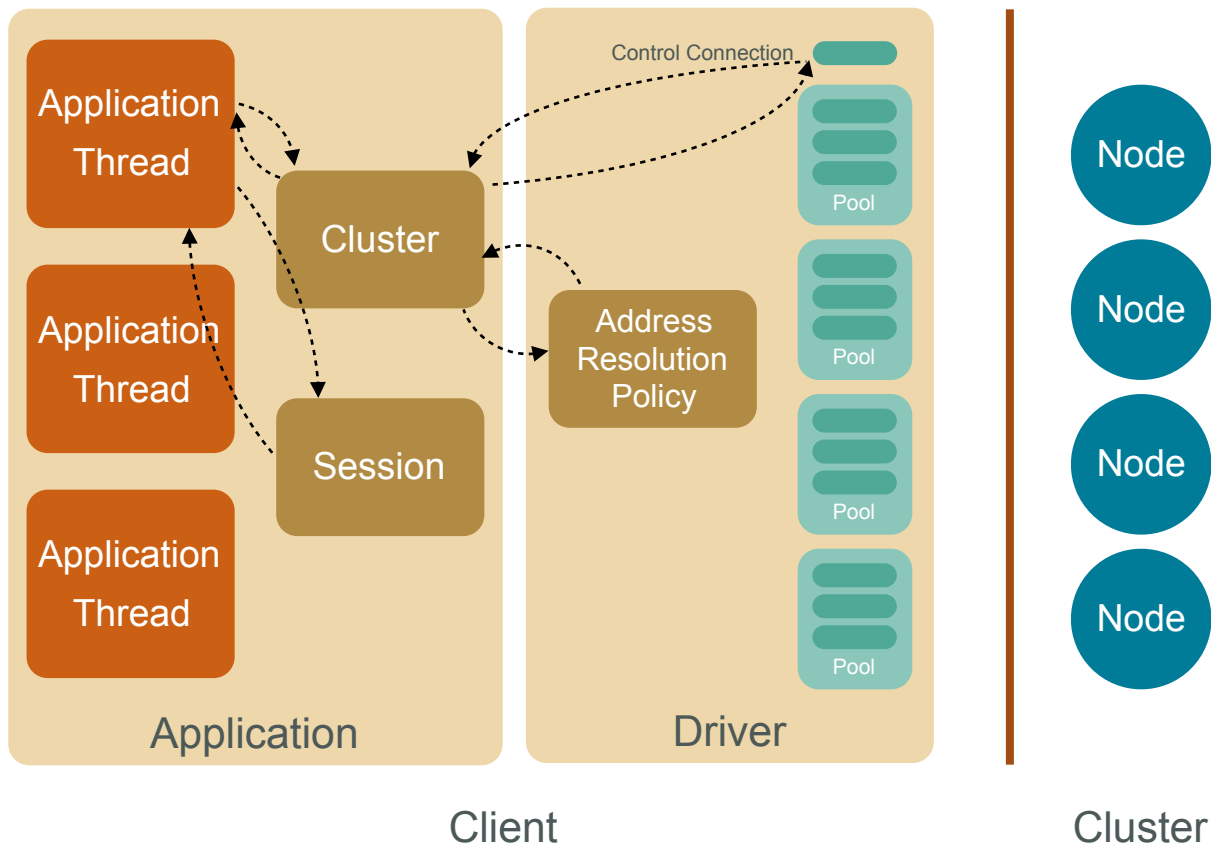
# Address Resolution

Topology Aware Client

# Multiple Addresses

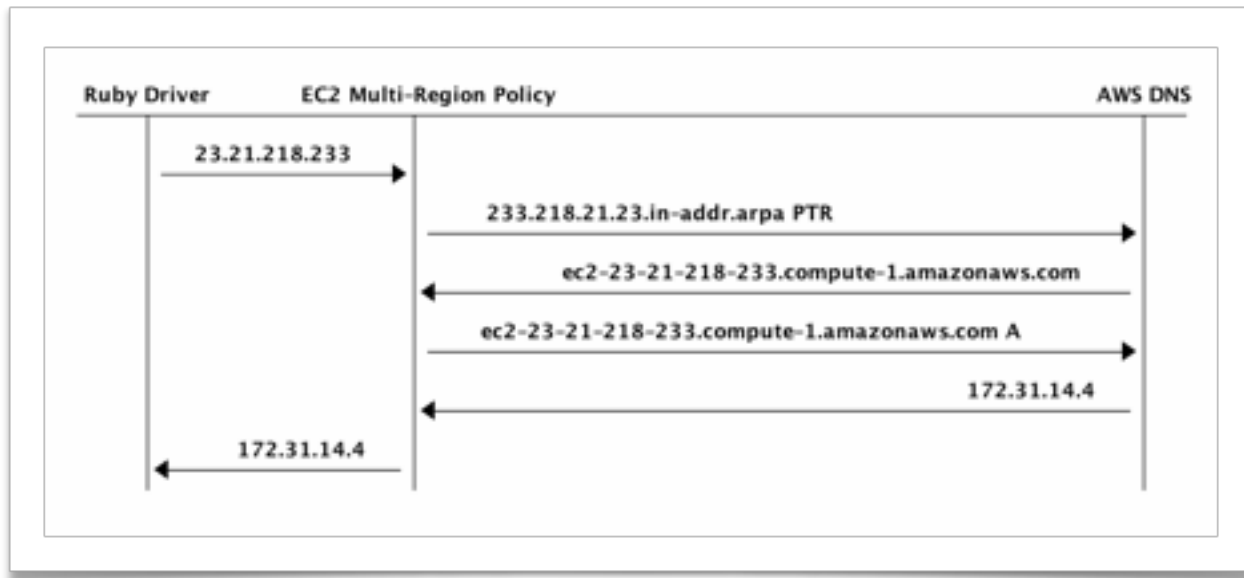


# Address Resolution





# EC2 Multi-Region Address Resolution





# One more thing...

Sneak peek of new and upcoming driver features

# Sneak peek: Slow query log

```
Cluster cluster = ...  
QueryLogger queryLogger = QueryLogger.builder(cluster)  
    .withConstantThreshold(...)   
    .withMaxQueryStringLength(...)   
    .build();  
cluster.register(queryLogger);
```

[Copy](#)


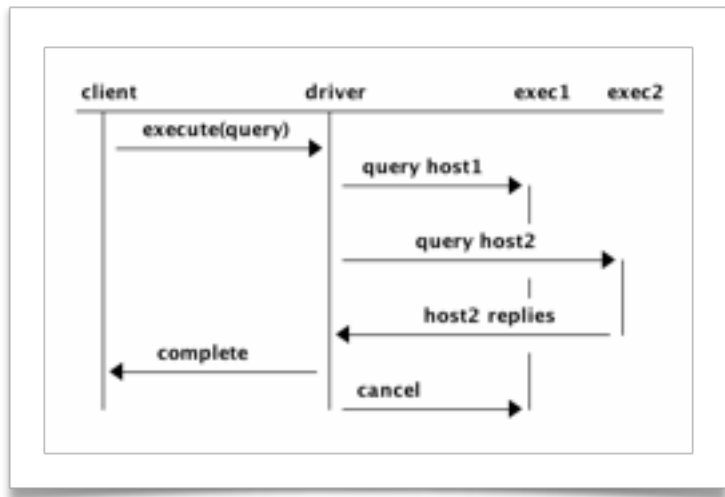
```
<logger name="com.datastax.driver.core.QueryLogger.SLOW">  
  <level value="DEBUG"/>  
</logger>
```

[Copy](#)

# Sneak peek: Execution profiles

```
ep1 = ExecutionProfile(load_balancing_policy=TokenAwarePolicy(DCAwareRoundRobinPolicy(local_dc='dc1')))  
ep2 = ExecutionProfile(load_balancing_policy=TokenAwarePolicy(DCAwareRoundRobinPolicy(local_dc='dc2')),  
                        row_factory=tuple_factory, request_timeout=None) # target dc2, return tuples, never timeout  
session = Cluster(execution_profiles={EXEC_PROFILE_DEFAULT: ep1, 'other-dc': ep2}).connect()
```

# Sneak peak: Speculative retries



..

# CASSANDRA SUMMIT 2016

## Q&A

Thank you!

