



Effective Testing in DSE

Predrag Knežević

predrag.knezevic@datastax.com

[@pedjaknezevic](#)

Why Taking Care?

- Automated testing for quality control of shipped product/service
- Number of tests and total testing times increase over time
- Shorter delivery cycles → continuous testing

- Run tests on each pre-merge check, but
- Keep feedback cycles short
- Ensure repeatable test execution anywhere

DSE Build Facts


- Junit based test infrastructure
- December 2014 (DSE 4.6)
 - Ant based build system
 - ~5h for running all tests on Jenkins, with a rather complicated job layout
- July 2016 (DSE 4.7+)
 - Gradle based build system
 - 40-60mins for running all tests on Jenkins
 - 16 hours of total testing time
 - The number of tests doubled!
 - Repeatable test execution across all machines
 - Simple setup

Why Moving to Gradle?

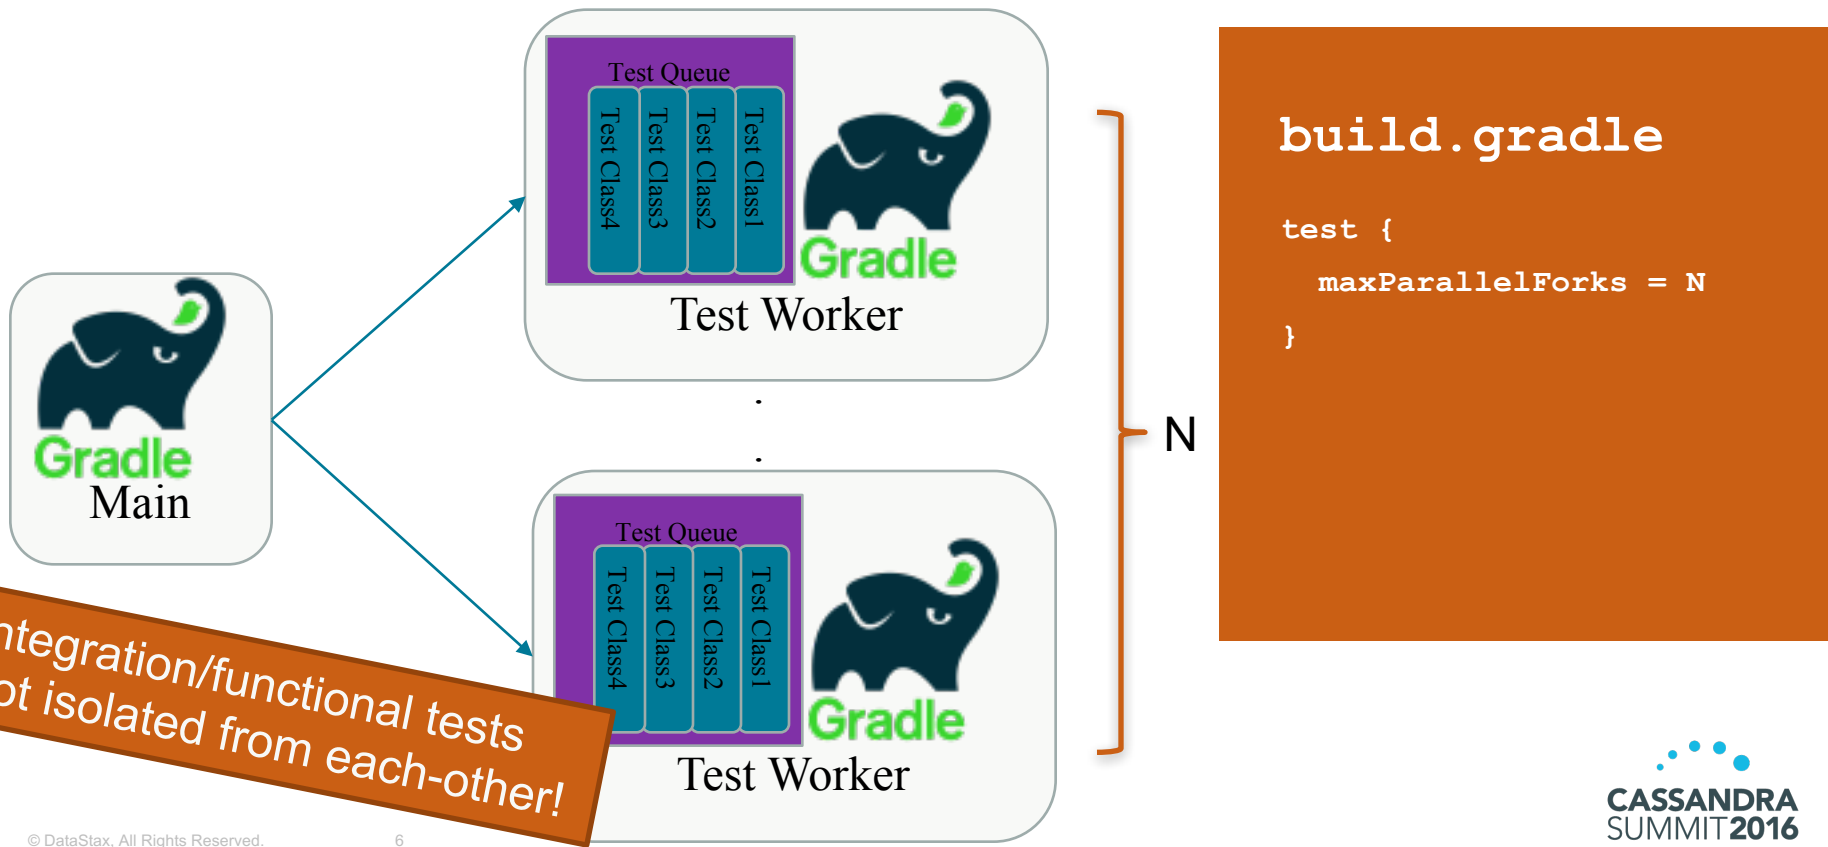
- Built-in support for parallel test execution
- Readable - build scripts based on Groovy (easy learning for Java devs)
- Repeatable builds/environment setup across machines
- Powerful dependency management
- Sane conventions, but configurable when needed
- Easy project modularization
- Excellent Eclipse/IntelliJ support
- Easy extendable through plugins or additional Java/Groovy code living in the script or project
- All Ant tasks still available

Running Tests

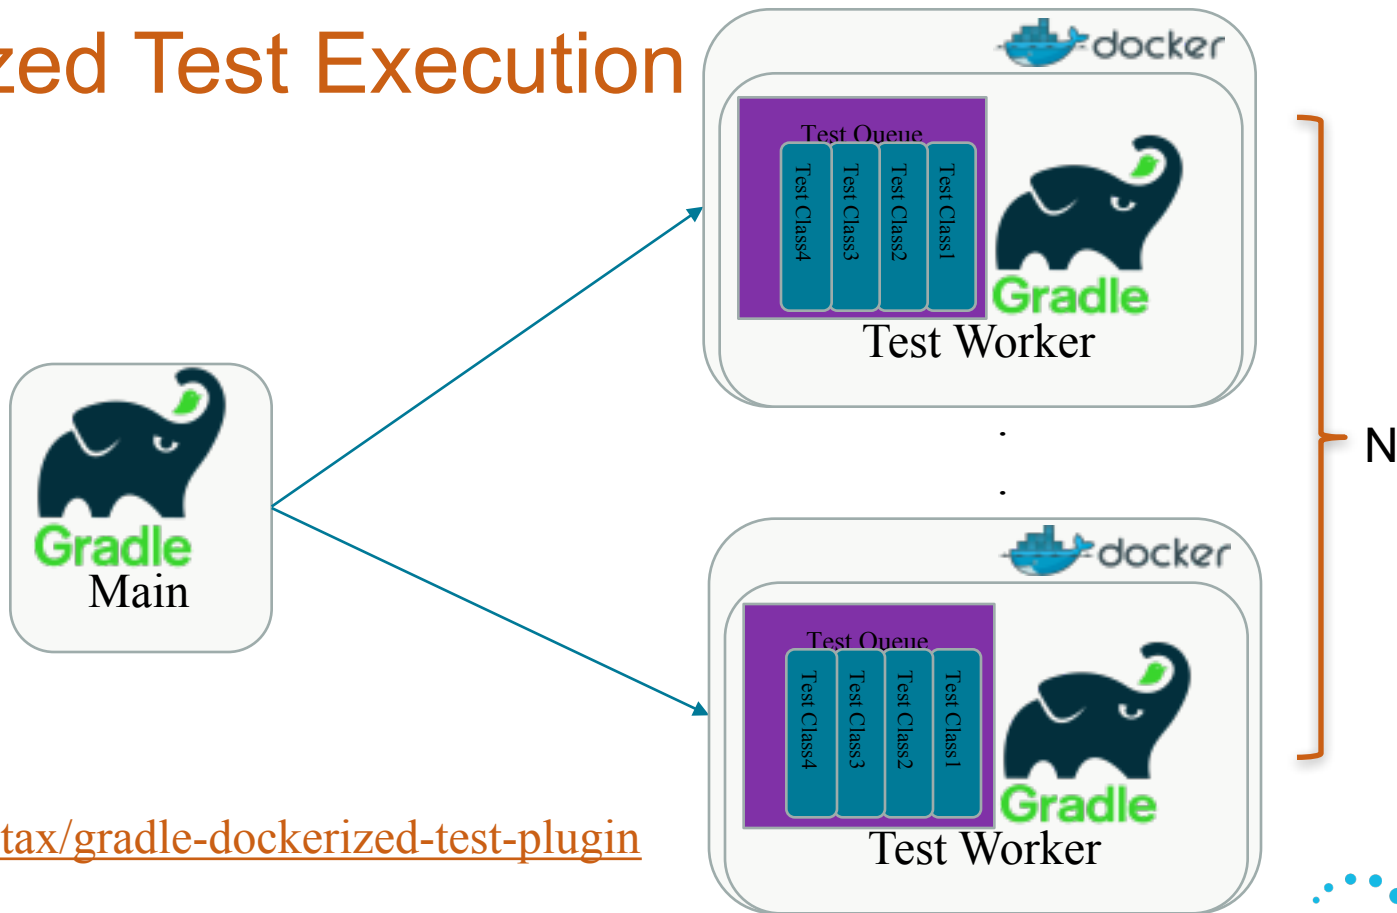
- All: `gradlew test`
- Single: `gradlew test -Dtest.single=FooTest`
- By default sequential execution
 - Low resources usage on modern multicore hardware
 - Long test round duration



Parallel Test Execution



Dockerized Test Execution



github.com/datastax/gradle-dockerized-test-plugin

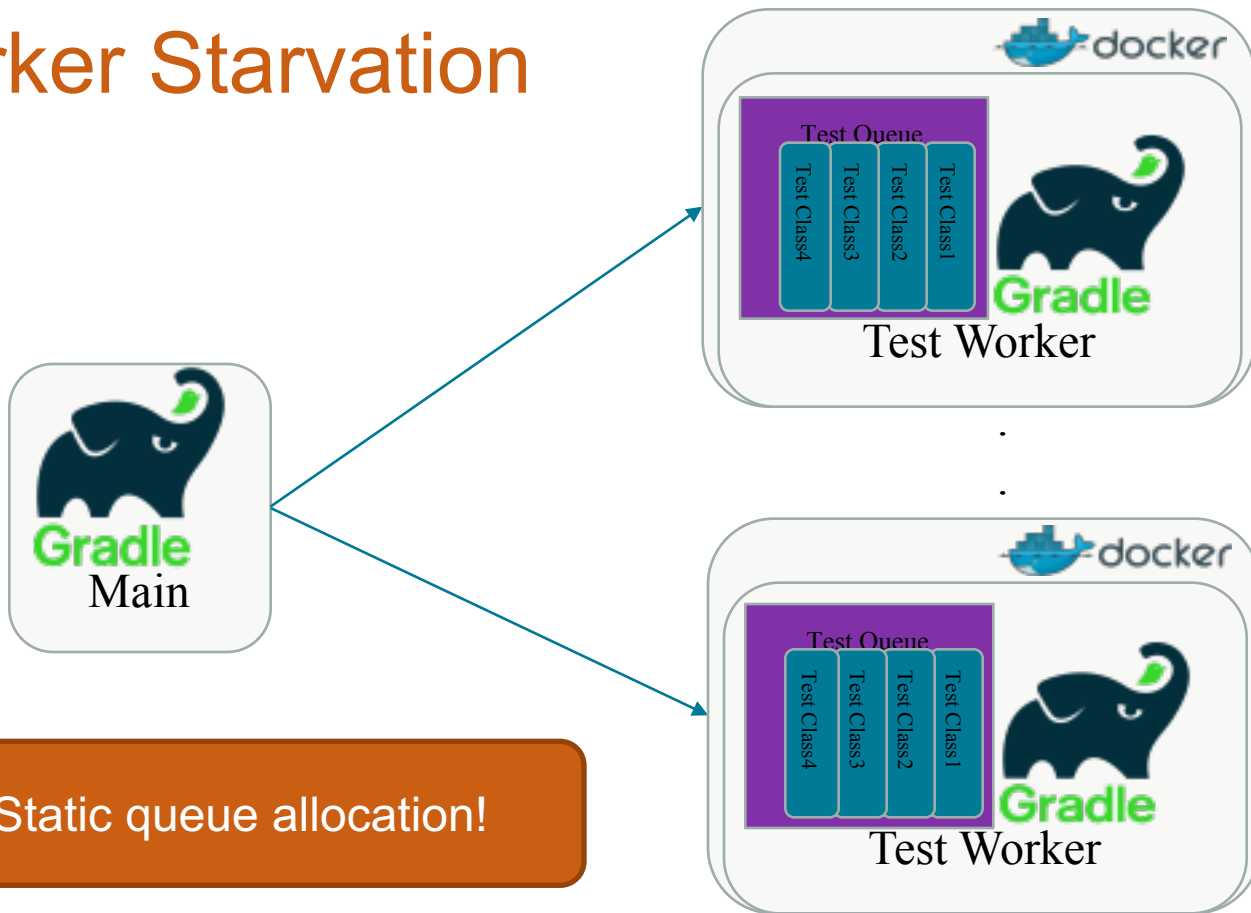
Dockerized Test Execution (2)

- Install Docker locally (native or boot2docker)
- No changes on production or test code required
- Test environment
 - Consistent across all machines (dev + CI)
→ no more “it works on my machine”
 - Managed as code (Dockerfiles) within project
 - Easy machine bootstrapping
 - Fully isolated
- Easy testing against several and/or appropriate environment


```
build.gradle
```

```
test {  
    maxParallelForks = N  
    docker {  
        image = 'test-image'  
    }  
}
```


Worker Starvation




Improved Queue Management



The duration of test round depends on the order of the test classes in the queue!

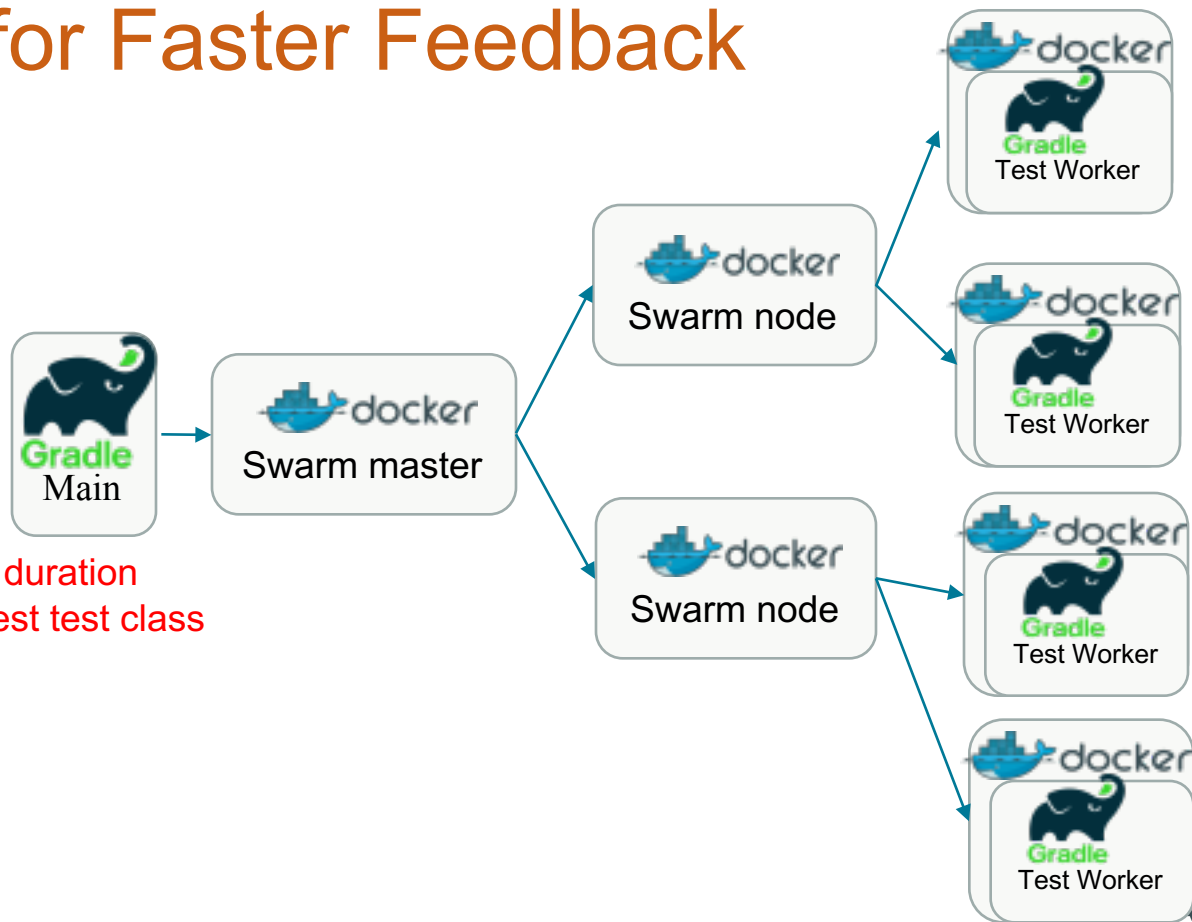
Test Scheduling

- NP-hard
- Longest Processing Time (LPT) algorithm
- History access required



More Workers for Faster Feedback

- Powerful multi-core machine
- Docker Swarm
 - Virtual Docker engine
 - Test workers run on a Swarm node
 - Cluster can shrink or grow



- Lower bound for the test round duration is equal to the duration of slowest test class

Split Slow Test Classes for More Throughput

- Manual
 - Group by common fixture
 - Extreme: one test per class
- Bytecode manipulations (auto split annotated class)
- Grouping tests classes into JUnit suites forces their sequential execution!

```
class FooTest {  
    @Test  
    void bar1() {  
    }  
  
    @Test  
    void bar2() {  
    }  
}
```



```
class Foo1Test {  
    @Test  
    void bar1() {  
    }  
}  
  
class Foo2Test {  
    @Test  
    void bar2() {  
    }  
}
```

No Embedded DSE Nodes

- Running cluster within single test worker JVM possible only by using separate classpath loaders
 - Still requires a good amount of hacking to make it work decently
 - Node shutdowns might still be problematic
 - Thread can hang around
 - Some objects cannot be garbage collected → memory leaks
- Standalone nodes enable reuse across test classes

Remote Code Execution

- MobilityRPC library (<http://github.com/npgall/mobility-rpc>)
- Remote JUnit Testrunner (<http://github.com/datastax/remote-junit-runner>)
 - Useful for integration tests requiring application context, but application cannot be easily embedded into test JVM
 - Support any existing JUnit runner on the remote side (Parametrized, Spock, Suites)

```
@RunWith(Remote.class)
public class RemoteTest {

    @Test
    public void foo() {
        // test
    }
}
```

Injecting Faults

- JBoss Byteman (<http://byteman.jboss.org>)
- Inject at runtime
 - Exceptions
 - Delays
 - Arbitrary side-effects
- Enables/simplifies testing of edge/uncommon cases

Logging

- Logback (<http://logback.qos.ch/>) based infrastructure
- Log file per a test case
 - Send all log messages to the single logback server asynchronously (reactor-logback adapter) keeping DSE nodes responsive
 - Route log statements to the proper file using SiftingAdapter and appropriate discriminator
 - Use JUnit rules to mark the beginning and the end of a test and propagate this information to the logback discriminator
 - Write thread dumps in case of a failure
- Scan log files for known issues and fail tests if they occur (e.g. Netty memory leaks)
- Turn DEBUG messages on to help later digging in a case of failure

Log Event Sender

```
<appender name="SOCKET" class="com.datastax.bdp.logback.SocketAppender">
  <remoteHost>${logbackServer}</remoteHost>
  <port>12345</port>
  <reconnectionDelay>1 seconds</reconnectionDelay>
  <nodeId>${nodeid}</nodeId>
  <eventDelayLimit>120 seconds</eventDelayLimit>
</appender>

<appender name="ASYNC" class="reactor.logback.AsyncAppender">
  <includeCallerData>true</includeCallerData>
  <appender-ref ref="SOCKET"/>
</appender>

<root level="DEBUG">
  <appender-ref ref="ASYNC"/>
</root>
```

Log Event Router

```
<appender name="SIFT" class="ch.qos.logback.classic.sift.SiftingAppender">
  <discriminator class="com.datastax.bdp.test.ng.LogFileDiscriminator">
    <key>TEST_LOG_FILE</key>
    <defaultValue>system.log</defaultValue>
  </discriminator>
  <sift>
    <appender name="FILE-${TEST_LOG_FILE}" class="ch.qos.logback.core.FileAppender">
      <filter class="com.datastax.bdp.test.ng.ForbiddenLogEventsDetector">
        <logFile>${TEST_LOG_FILE}</logFile>
        <detector class="com.datastax.bdp.test.ng.NettyLeakDetector"/>
      </filter>
      <encoder>
        <pattern>%X{nodeid} %5p [%t] %d{ISO8601} %F \ (line %line\) %m%n</pattern>
        <immediateFlush>false</immediateFlush>
      </encoder>
      <file>${LOG_DIR}/${TEST_LOG_FILE}</file>
      <append>true</append>
    </appender>
  </sift>
</appender>
```

Test Start/End Detection

```
@Rule
public TestRule watcher = new TestWatcher() {

    private String logFile;

    @Override
    protected void starting(Description description) {
        logFile = description.getClassName()+"/"+description.getMethodName()+".log";
        logToFile(logFile);
    }

    protected void failed(Throwable e, Description description) {
        threadDumpLogger.error("Test {}.{} failed, thread dump:\n{}\n", description.getClassName(),
            description.getMethodName(), getThreadDumps());
        logToFile(description.getClassName()+"/after.log");
    }

    protected void finished(Description description) {
        logToFile(description.getClassName()+"/after.log");
        ForbiddenLogEventsDetector.checkForIssues(logFile);
    }
};
```

Resources

- Gradle (gradle.org)
- Docker (docker.com)
- Gradle Dockerized Test Plugin (github.com/datastax/gradle-dockerized-test-plugin)
- MobilityRPC (<http://github.com/npgall/mobility-rpc>)
- Remote JUnit Testrunner (<http://github.com/datastax/remote-junit-runner>)
- JBoss Byteman (<http://byteman.jboss.org>)
- Logback (<http://logback.qos.ch/>)
- Project Reactor Addons (github.com/reactor/reactor-addons), Logback adapter



CASSANDRA SUMMIT 2016

Thank you!

Questions?