



Ameesh Divatia

The Promise and Perils of Encrypting Cassandra Data

## DATA BREACH STATISTICS

DATA RECORDS LOST OR STOLEN SINCE 2013

4,762,376,968

ONLY 4% of breaches were "Secure Breaches" where encryption was used and the stolen data was rendered useless.

DATA RECORDS ARE LOST OR STOLEN AT THE FOLLOWING FREQUENCY



EVERY DAY

3,575,358

Records



EVERY HOUR

148,973

Records



EVERY MINUTE

2,483

Records



EVERY SECOND

41

Records

# Perils of Cloud Hosted Databases

## Cloud hosted database services being adopted by enterprise IT

- Where is the data stored?


### Threats

- Stolen cloud administrator credentials
- Violations of compliance regulations


### Need

- Complete control of data by enterprise IT at all times
  - Encrypt all data that is uploaded to the cloud

## How does enterprise IT deal with this?



# Cassandra Security



## Use Case – Apache Cassandra

- Raw customer data stored in multi-tenant clusters as clear text
- Encryption available as an option only at rest

## Vulnerabilities

- Data in memory is in the clear
- Encryption key is in the clear in memory risking the entire data store

# Anatomy of a Hack

baffle

## Threat Scenario

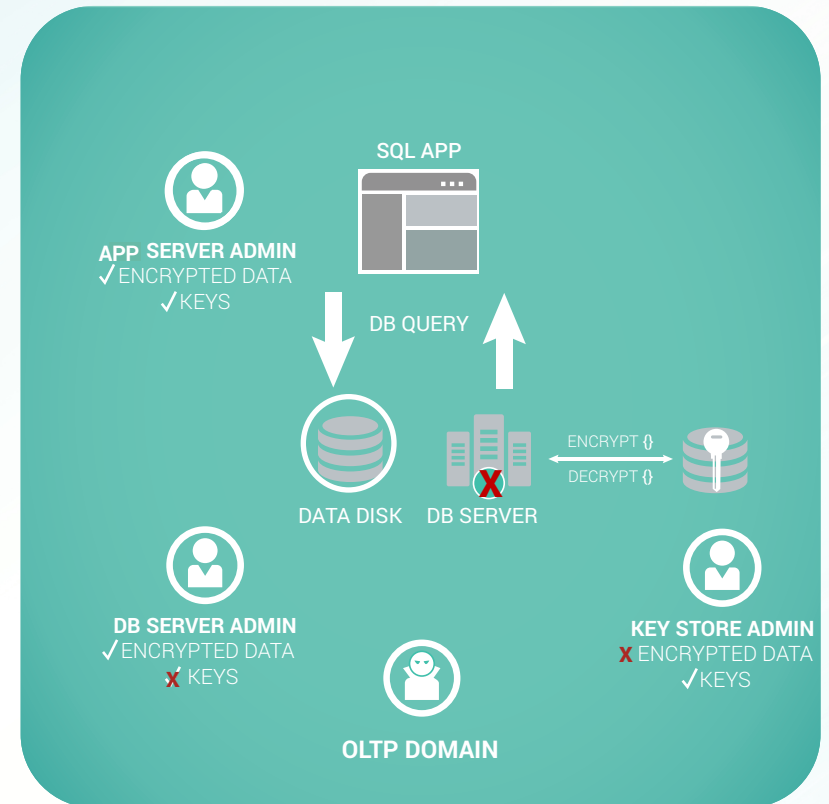
- Database admin has access to data and keys used for 'at rest' security

## Vulnerabilities

- Administrator's credentials are stolen by a hacker
- Hacker has access to the data as well as the key

## Need

- Separation of keys and data into separate domains
- Never decrypting data in a single compute instance
- Operate on encrypted data**





# Encryption Adoption Challenges

baffle

**Complexity:** Perception that once the data is encrypted, the customer loses control

**Key management:** What if the keys are lost, stolen or stagnant?



# Encryption Adoption Challenges

baffle

**Performance degradation:** Perception that application performance will slow down

**Workflow impact:** Changes the current paradigm of app to database communication



# A Practical Approach

baffle

## Driver level insertion for data protection

- Transparent application operation

## Seamless key management with Amazon KMS

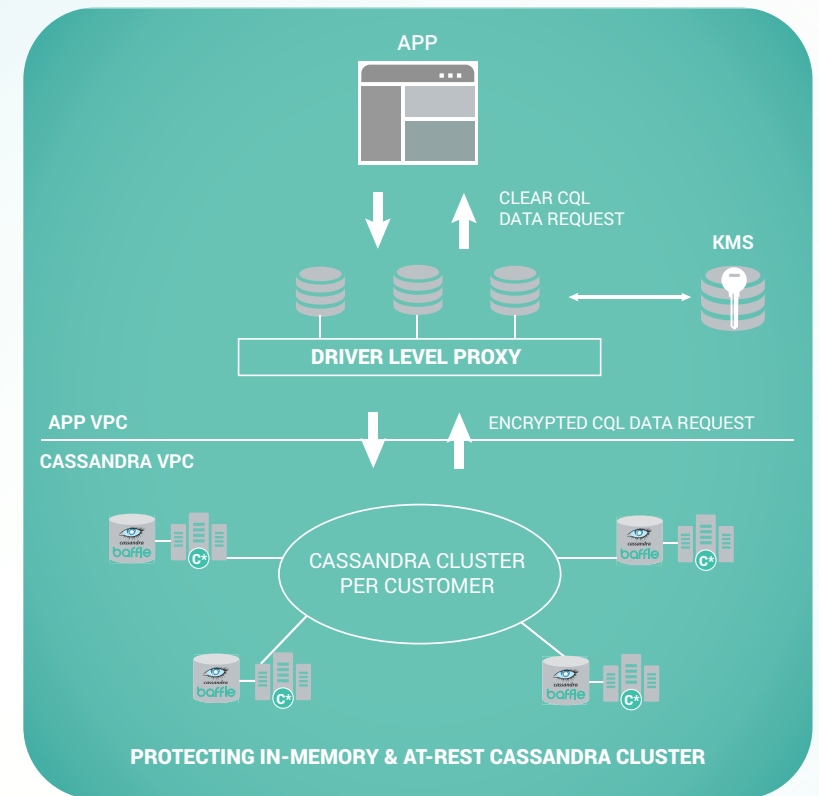
- Generate, use, store, rotate and retire keys

## Performance at scale using massive parallelization

- Elastic load balancing in response to demand

## Frictionless integration into enterprise application workflows

- Active monitoring of service components to ensure high availability and failover





# Selective Privacy Example

baffle

Employee Name	Job Title	Base Pay	Overtime Pay	Other Pay	Benefits	Total Pay & Benefits	Year	Notes	Agency	Status
David Shinn	Deputy Chief 3	129150.01		0	342802.63	38780.04471952.6	510732.68	2014	San Francisco	PT
Amy P Hart	Asst Med Examiner	318835.49	10712.95	60563.54	89540.23	390112	479652.21	2014	San Francisco	FT
William J Coaker Jr.	Chief Investment Officer	257340		0	82313.7	96570.66339653.7	436224.36	2014	San Francisco	PT
Gregory P Suhr	Chief of Police	307450.04		0	19266.72	91302.46326716.8	418019.22	2014	San Francisco	FT
Joanne M Hayes-White	Chief, Fire Department	302068		0	24165.44	91201.66326233.4	417435.1	2014	San Francisco	FT
Ellen G Moffatt	Asst Med Examiner	270222.04	6009.22	67956.2	71580.48344187.5	415767.94		2014	San Francisco	FT
John L Martin	Dept Head V	311298.55		0	0	89772.32311298.6	401070.87	2014	San Francisco	FT



employee\_data

Column encryption



Select key



employee\_name

Select encryption type



Select key



job\_title

Select encryption type



Select key



base\_pay

Encryption using specified key

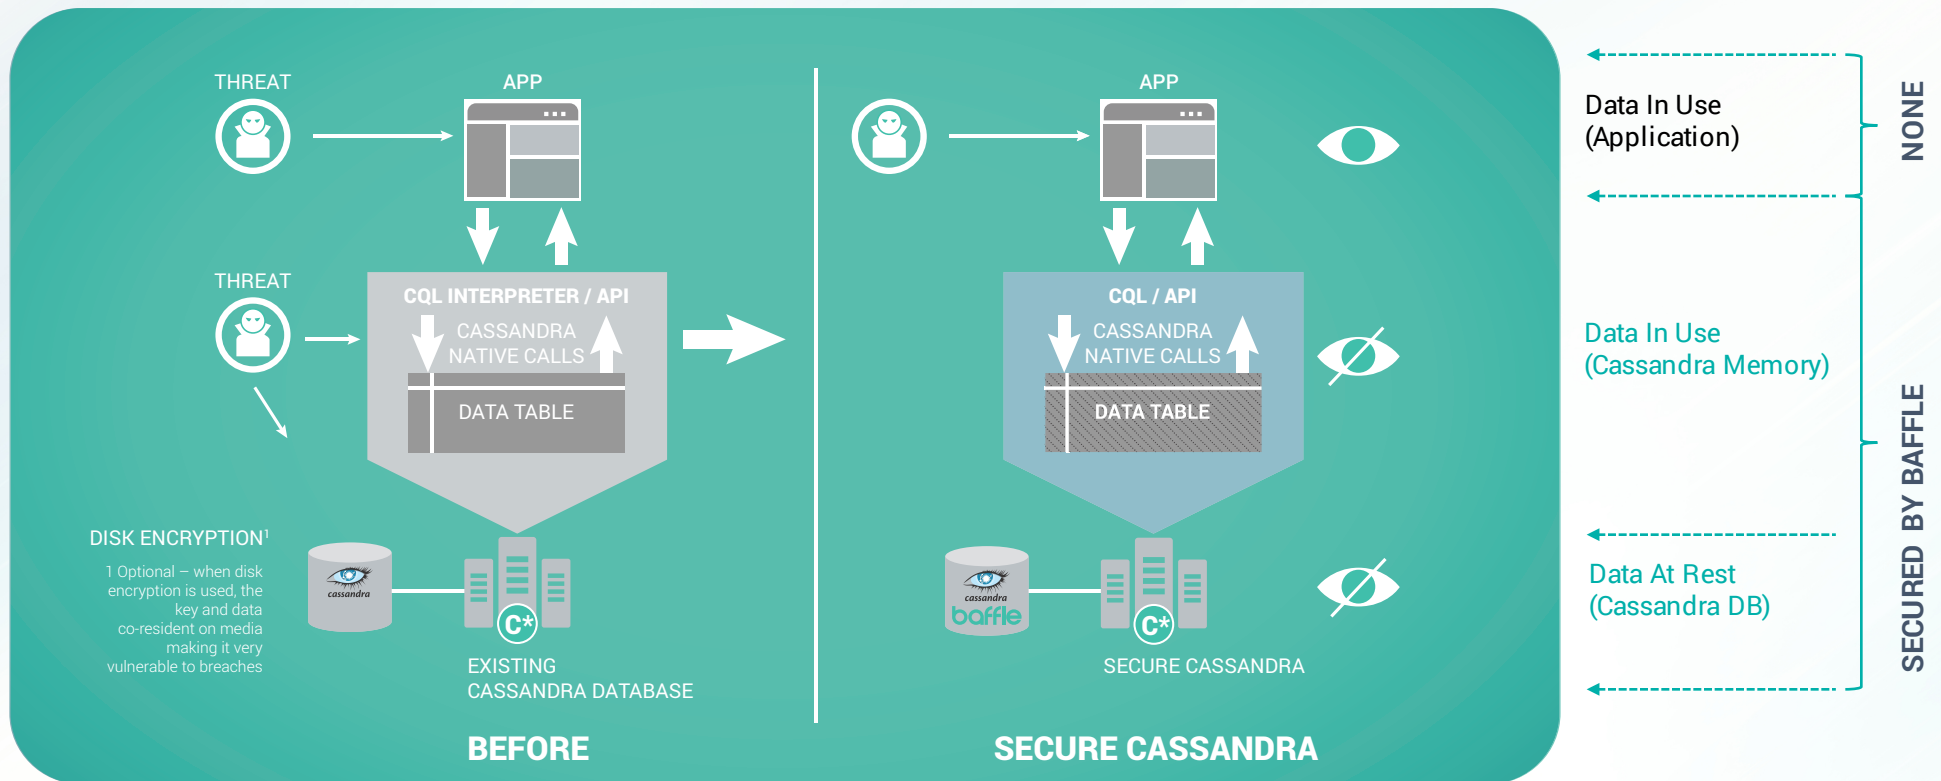


acme\_key p01



# Secure Cassandra – Before and After

baffle



# Driver Modifications









Added couple of pipeline stages for encryption / decryption in the java driver

```
Connection.java
```

```
pipeline.addLast("messageDecryptor", new MessageDecryptor(transformDB));
```

```
pipeline.addLast("messageEncryptor", new MessageEncryptor.Encryptor(transformDB));
```

# Tiers of Security

PARAMETER	STOCK CASSANDRA	SECURE CASSANDRA
PROTECTION		
Data on disk	 Encrypted	 Encrypted
Data in Cassandra server memory	 Clear	 Encrypted
Data in App memory	 Clear	 Clear
Vulnerability	No data Protection (App, DB Server, Disk)	Data Protection below CQL Interface (DB Server & Disk)
IMPACT		
Driver	Unchanged	Added pipeline stages
DB Binaries	Unchanged	Add support for custom types
OPERATIONS ON ENCRYPTED DATA	None	Search, Sort, Aggregate

# Cassandra Modifications



Added custom types for encrypted values with serializers/ deserializers

Added custom functions for operations such as compares and aggregates on the encrypted values

Custom Integer Type:

Db/marshal/EncIntType.java

```
public class EncIntType extends AbstractType<BlindInt>
{
    public static final EncIntType instance = new EncIntType();
    private EncIntType()
    {
        super(ComparisonType.CUSTOM);
    }
    ...
}
public synchronized int compareCustom(ByteBuffer o1, ByteBuffer o2)
{
```



# Conclusion



Cloud hosted databases have *perils* for stored data

Practical approach to protecting Cassandra data

- Driver level modifications to enable encryption
- Support custom types and functions for encrypted operations

**The *promise* of keeping data encrypted...always! is attainable**

- No threats to availability
- No impact on stability
- Minimal impact on performance



**CASSANDRA  
SUMMIT 2016**

**Thank You**

**baffle**

**Make data breaches irrelevant**

[www.baffle.io](http://www.baffle.io)