

# Real World Repairs @ Netflix

Vinay Chella

Cassandra MVP, Cloud Data Architect





# NETFLIX











# Vinay Chella

Cloud Data Architect

Cassandra MVP

Cloud Database Engineering @ Netflix







# Agenda

- What is Repair and Why?
- How do we Repair @ Netflix?
- Use cases
- Pain points
- Alternatives to Repair
- When to/ not to repair
- What is missing in C\*

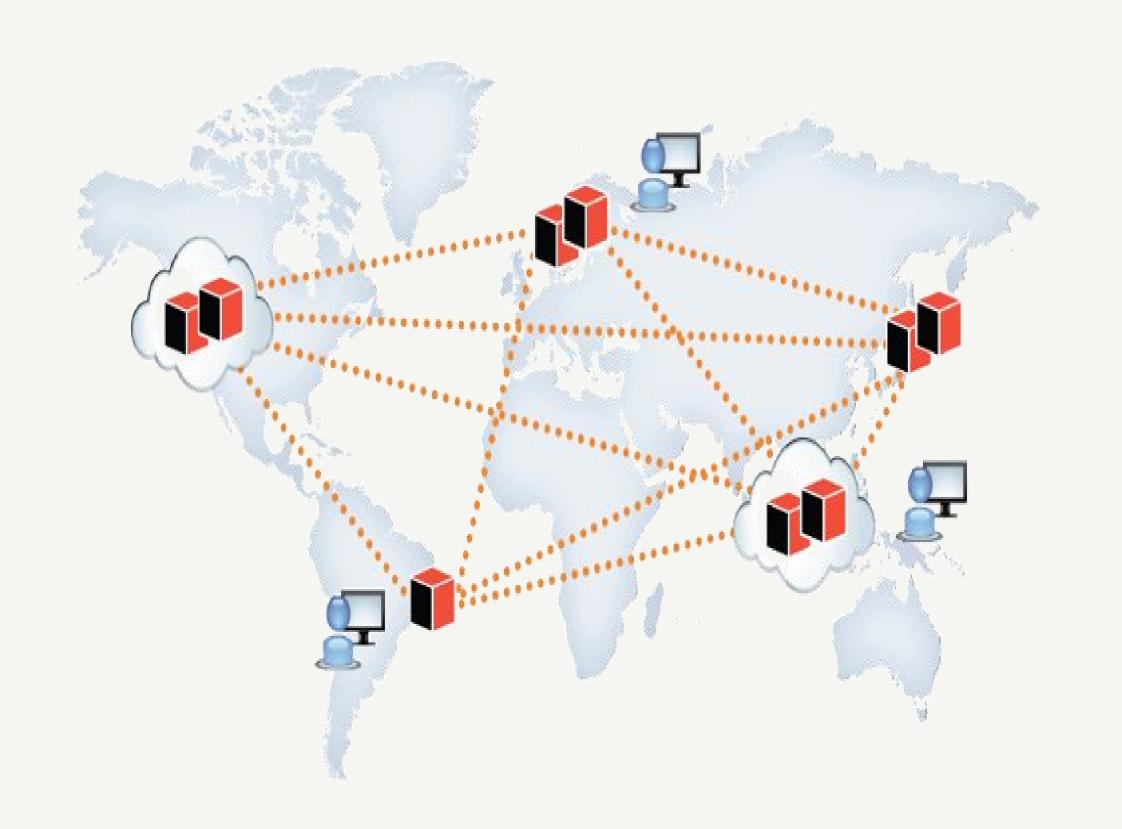






# Why is Repair Important?

- Entropy is inevitable
- Inconsistent data might impact your business
- Data Resurrection

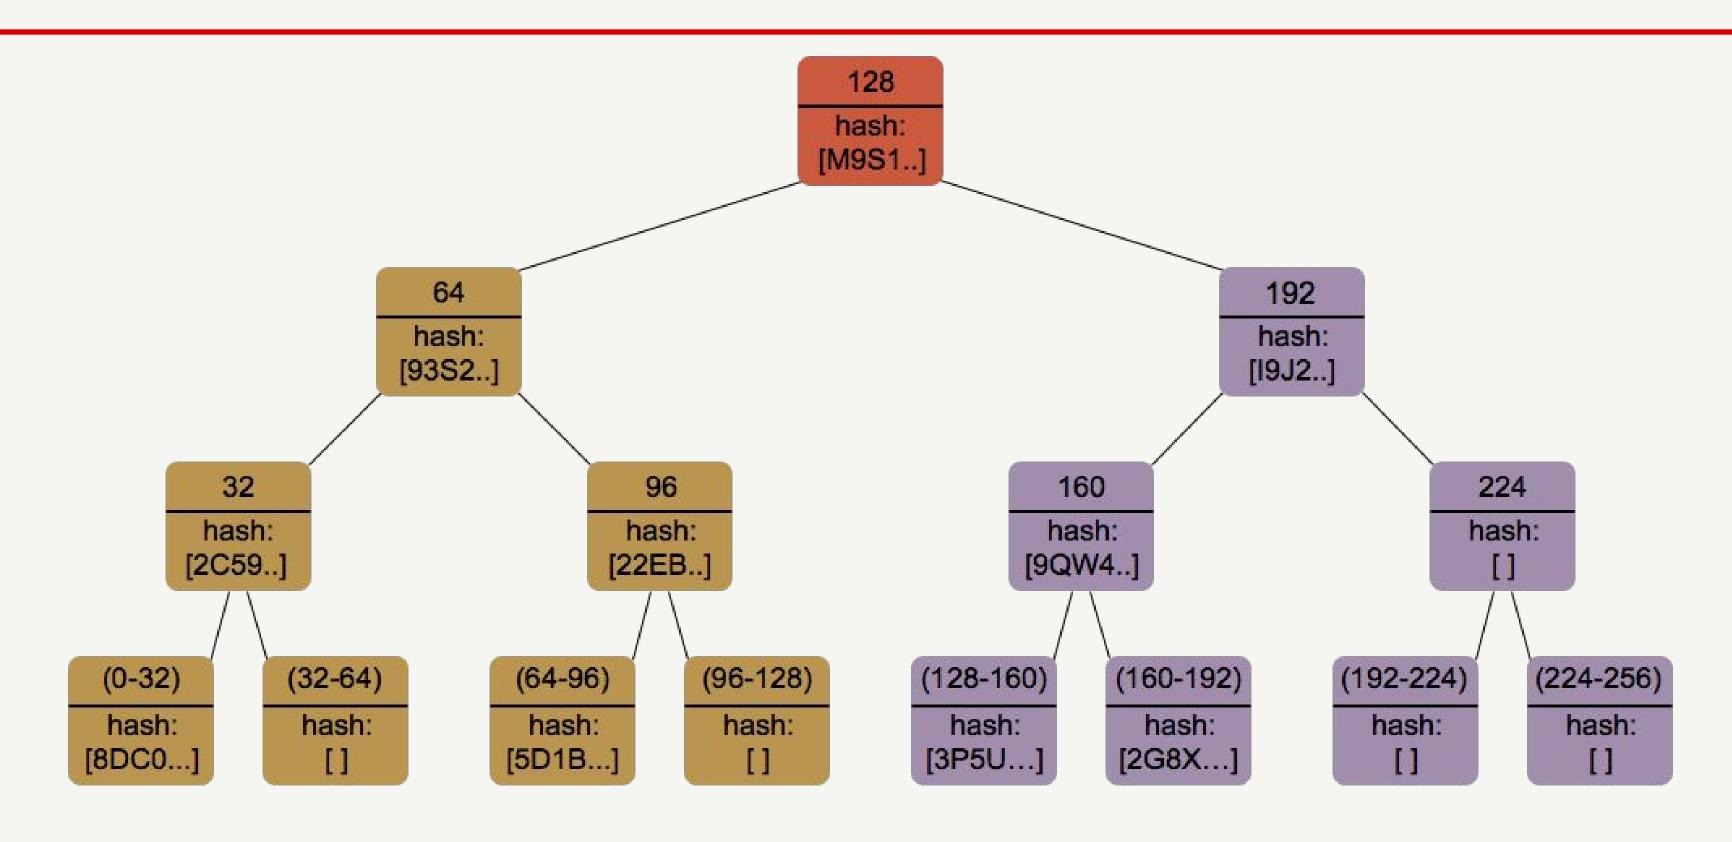








# What is Repair



Row key: jack	Row key: jill	Row key: terry	Row key: misty
Row token: 5	Row token: 7	Row token: 10	Row token: 20
hash: 8DC0	hash: 5D1B	hash: 3P5U	hash: @G8X







## Repair in Netflix world

#### Case 1

- ~300 TB cluster size
- ~1 TB per node
- SSDs
- 288 Nodes Multi Region Cluster
- Critical data, No TTLs

#### Case 2

- ~180 TB cluster size
- ~3.6 TB per node
- HDDs
- 48 Nodes Island cluster
- No TTLs, long lived







# How do we Repair @ Netflix?

- Primary Range (-pr)
- Subrange (-st and -et)
- On every *node* in every *dc* sequentially
- Parallel (-par) & SEQUENTIAL









# How do we Repair @ Netflix? Cont...

- More granular Repair
  - -Selective tables only
- On schedule basis
  - -Jenkins
- Repair progress tracking





# Impacts of UnRepaired C\*/ Use-cases

- Incorrect bookmark index
- Losing the ratings
- Recommendations impact
- Broken recently watched



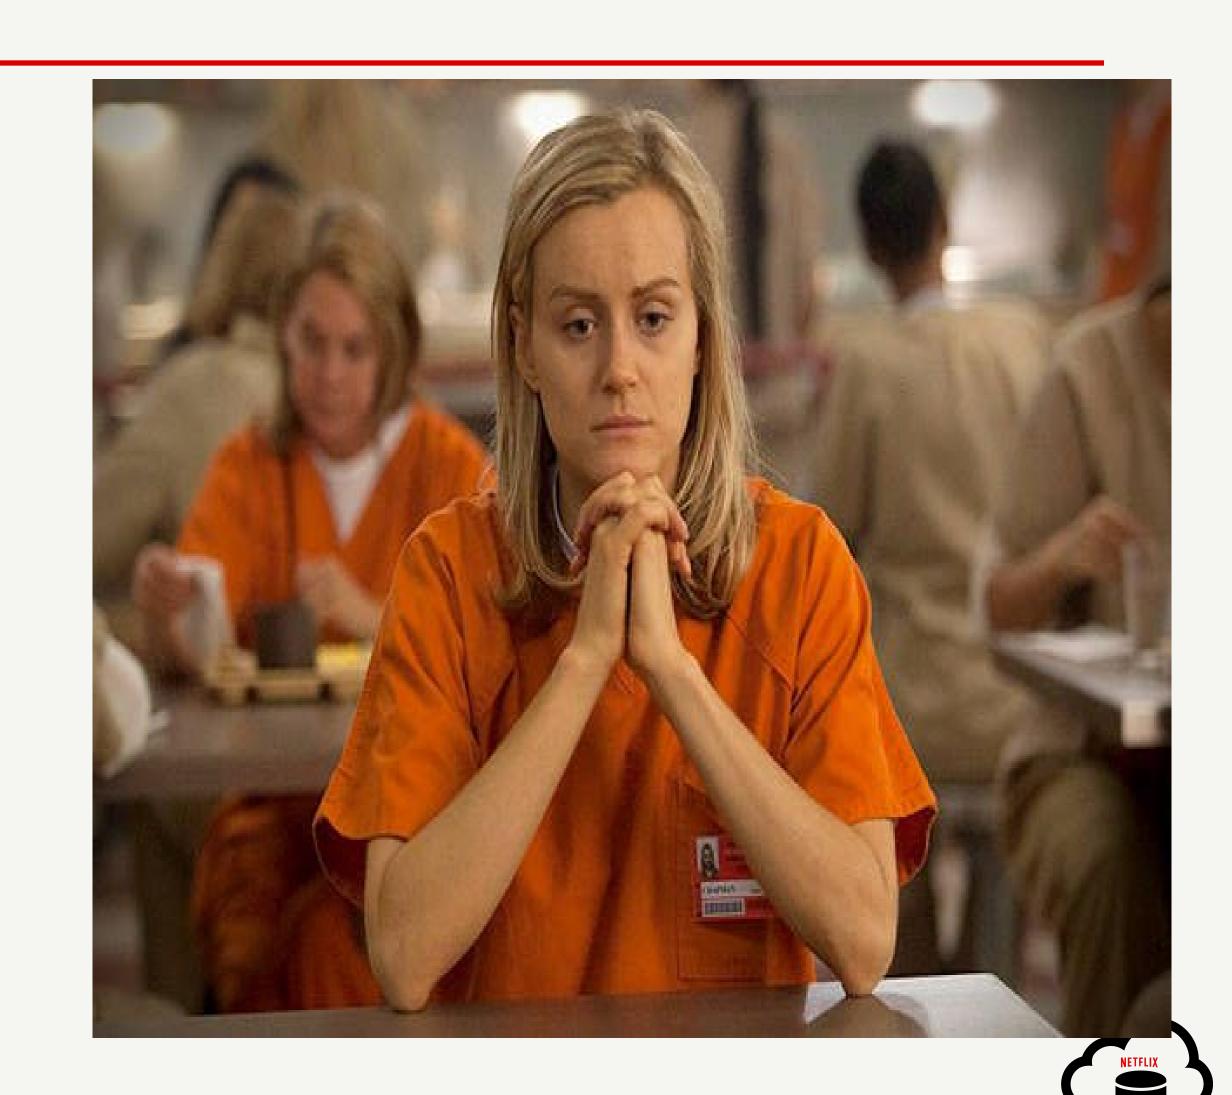






## Pain Points

- Over streaming
- Stream timeouts
- Stuck repairs
- Tracking/ resuming repair
- Compactions bottlenecks
- Disk I/O
- CPU usage
- Latencies/SLA
- Wide partitions









How do we live with it?







Well, we don't, we fix it.







# Over streaming

- Subrange repair
  - -Techniques
  - -Algorithm









# Subrange Repair

- Get subranges
  <a href="https://github.com/pauloricardomg/cassandra-list-subra">https://github.com/pauloricardomg/cassandra-list-subra</a>
  <a href="mailto:nges">nges</a>
  - Thrift: describe\_splits\_ex
- Repair those ranges using <a href="https://github.com/BrianGallew/cassandra range repair">https://github.com/BrianGallew/cassandra range repair</a>







## Subrange Repair

#### Regular Repair

- Size per Node 950 GB
  - Time to repair ~11 Hours
    - Primary Range
    - SERIAL
- Size per Node 550 GB
  - Time to repair ~24 Hours
  - Primary Range
  - SERIAL
  - Wide partitions



### Subrange Repair

- Size per Node 950 GB
  - Time to repair ~2.5 Hours
  - Primary Range
  - Threads #: 5
  - Split size 64K
- Size per Node **550 GB** 
  - Time to repair ~5.5 Hours
  - Primary Range
  - Threads # 10
  - Split size 32K
  - Wide partitions

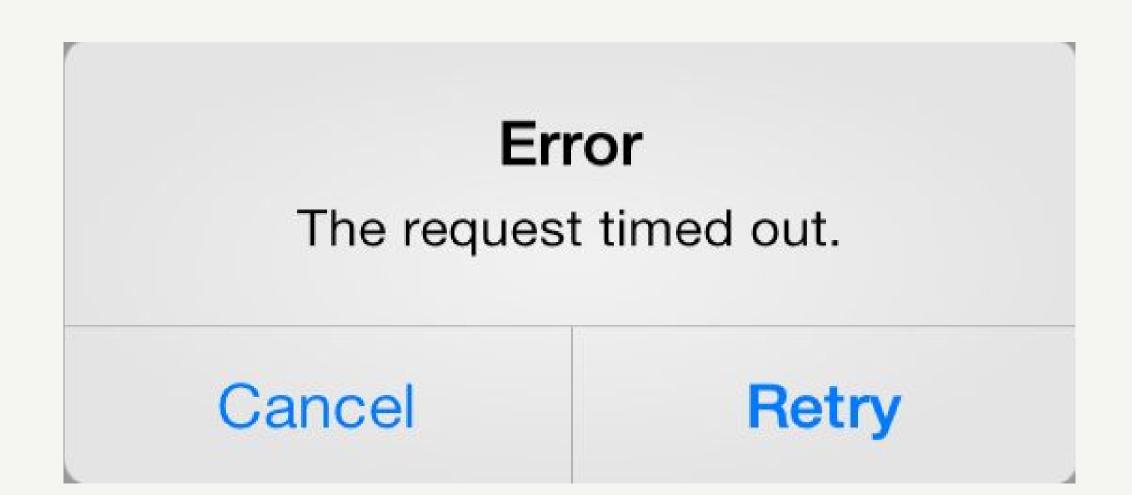




## To overcome stream timeouts

#### YAML Tunings

- streaming\_socket\_timeout\_in\_ms
- stream\_throughput\_outbound\_megabits\_per\_sec
- inter\_dc\_stream\_throughput\_outbound\_megabits\_per\_sec



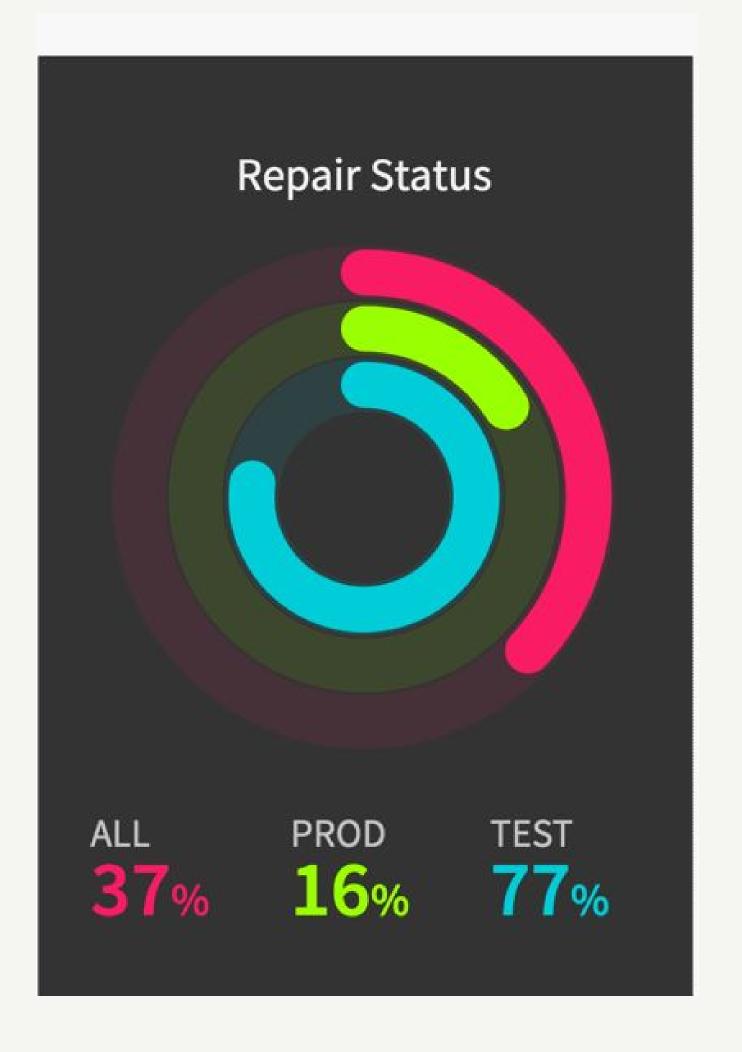






## To overcome resume issue

- Track it
  - -Record the progress
  - -Record the status
- Resume-ability
  - -Resume from failed or paused range









# To fix stuck repairs

- Use Repair Tracking data
- Timeouts for long running repair
  - JMX notifications
- Self healable repair (Automate it !!)
  - Grep for repair logs
  - nodetool compactionstats
  - nodetool netstats
  - JMX: forceTerminateAllRepairSessions()







# Compaction bottlenecks

- Why compactions?
  - -Validation phase
- Tune compaction settings
  - compaction\_throughput\_mb\_per\_sec



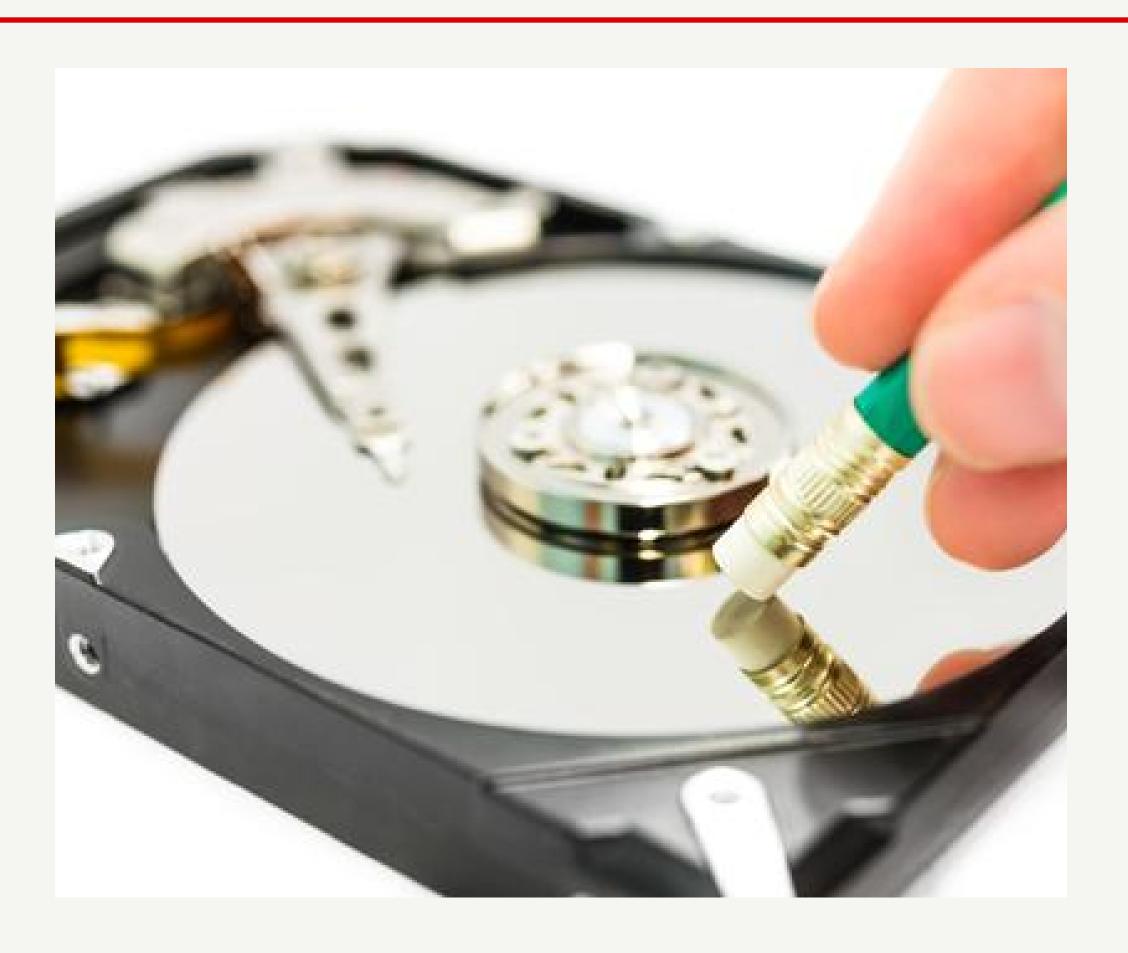






# Disk I/O issues

• SSD









# High CPU Usage?

- Throttled streaming
  - nodetool setstreamthroughput
- Throttling subrange repair
  - -Keysplit
  - -#threads









# To minimize impacts on latencies

- Running Repairs during off-peak hours
- Switch to SERIAL repair
  - -Instead of PARALLEL
- If Subrange
  - -Reduce #threads
  - -Smaller subrange







# Wide partition issues

Reduce the partition size









# Other ways to Repair

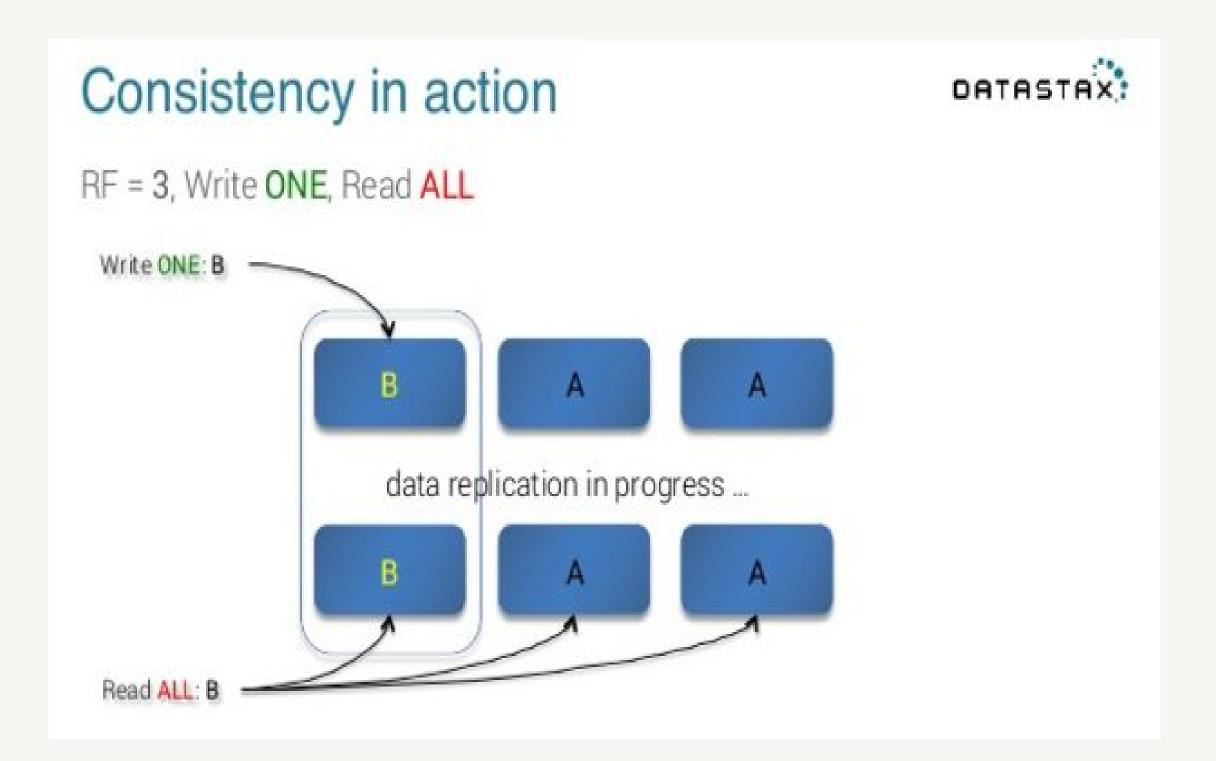






# Other ways to Repair: Tickler

- Row Tickler
- All Row Tickler









# When to Repair

- Multi-Region clusters
- Low consistent read and writes
- Frequent node outages
- Low read-repair chance
- Flaky networks



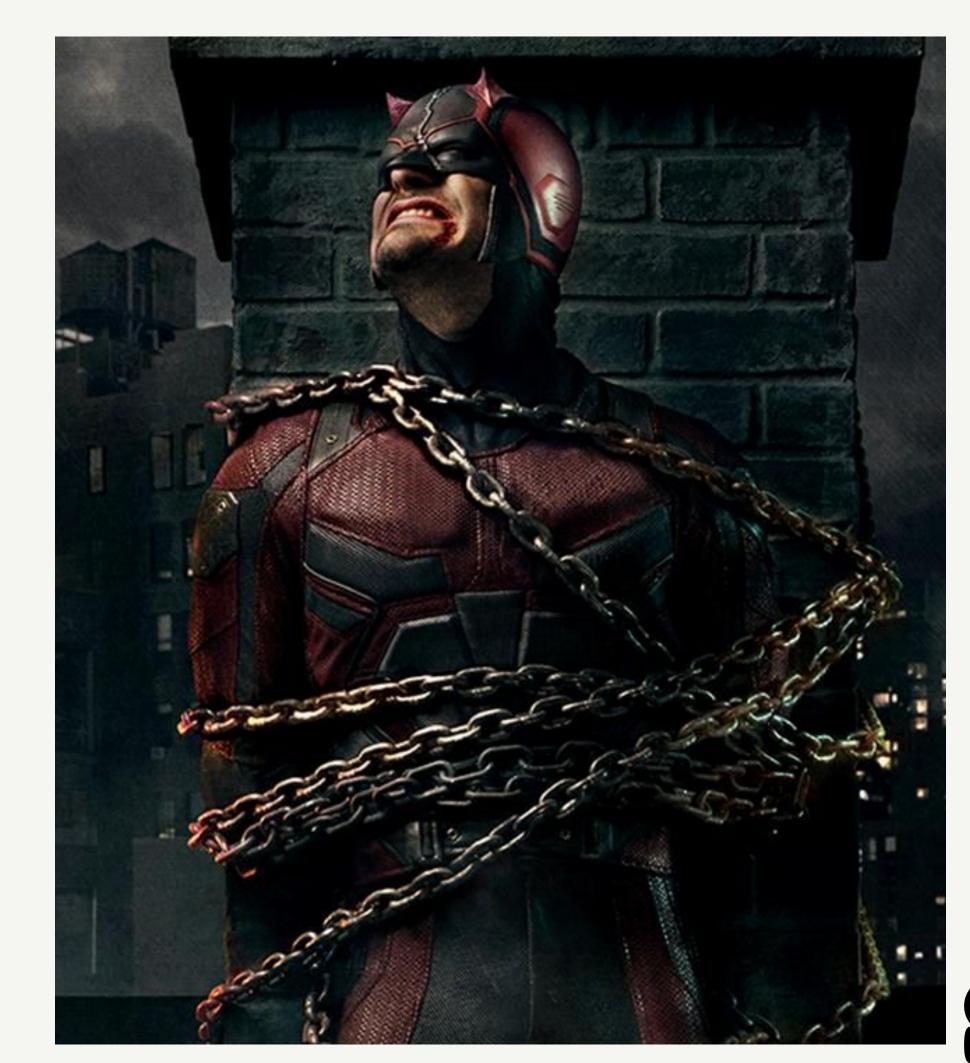






# When not to Repair

- Short TTIs
- Highly consistent read and writes
- High read repair chance settings in island clusters









# What is missing in C\*

- Metrics and insights
- Incremental Repair bug fixes
- CASSANDRA-10070 Automatic repair
  scheduling
- Mutation-based Repairs?

# Photo Not Available







# Q&A





