

Real world tales of repair

THE LAST PICKLE

CASSANDRA SUMMIT - SEPTEMBER 2016

Alexander Dejanovski
@alexanderdeja

Consultant
www.thelastpickle.com

Datastax MVP for Apache Cassandra

About The Last Pickle

We help people deliver and improve Apache Cassandra based solutions.

With staff in 5 countries and over 50 years combined experience in Apache Cassandra.





What and why ?

Full repair

Incremental repair

How to make it work

Automated repairs

What is repair ?

A maintenance operation that (briefly)
restores strong consistency throughout the
cluster



Why do we need repair ?

- Eventual consistency
- Downtime / failure recovery
- Safe deletes



Tombstones need repair too

Missing tombstones can lead to zombie data
(repair within `gc_grace_seconds`)

What and why ?

Full repair

Incremental repair

How to make it work

Automated repairs

How does anti-entropy repair works ?

Reads all data

How does anti-entropy repair works ?

Reads all data
Calculates hashes

How does anti-entropy repair works ?

Reads all data
Calculates hashes
Compares hashes

How does anti-entropy repair works ?

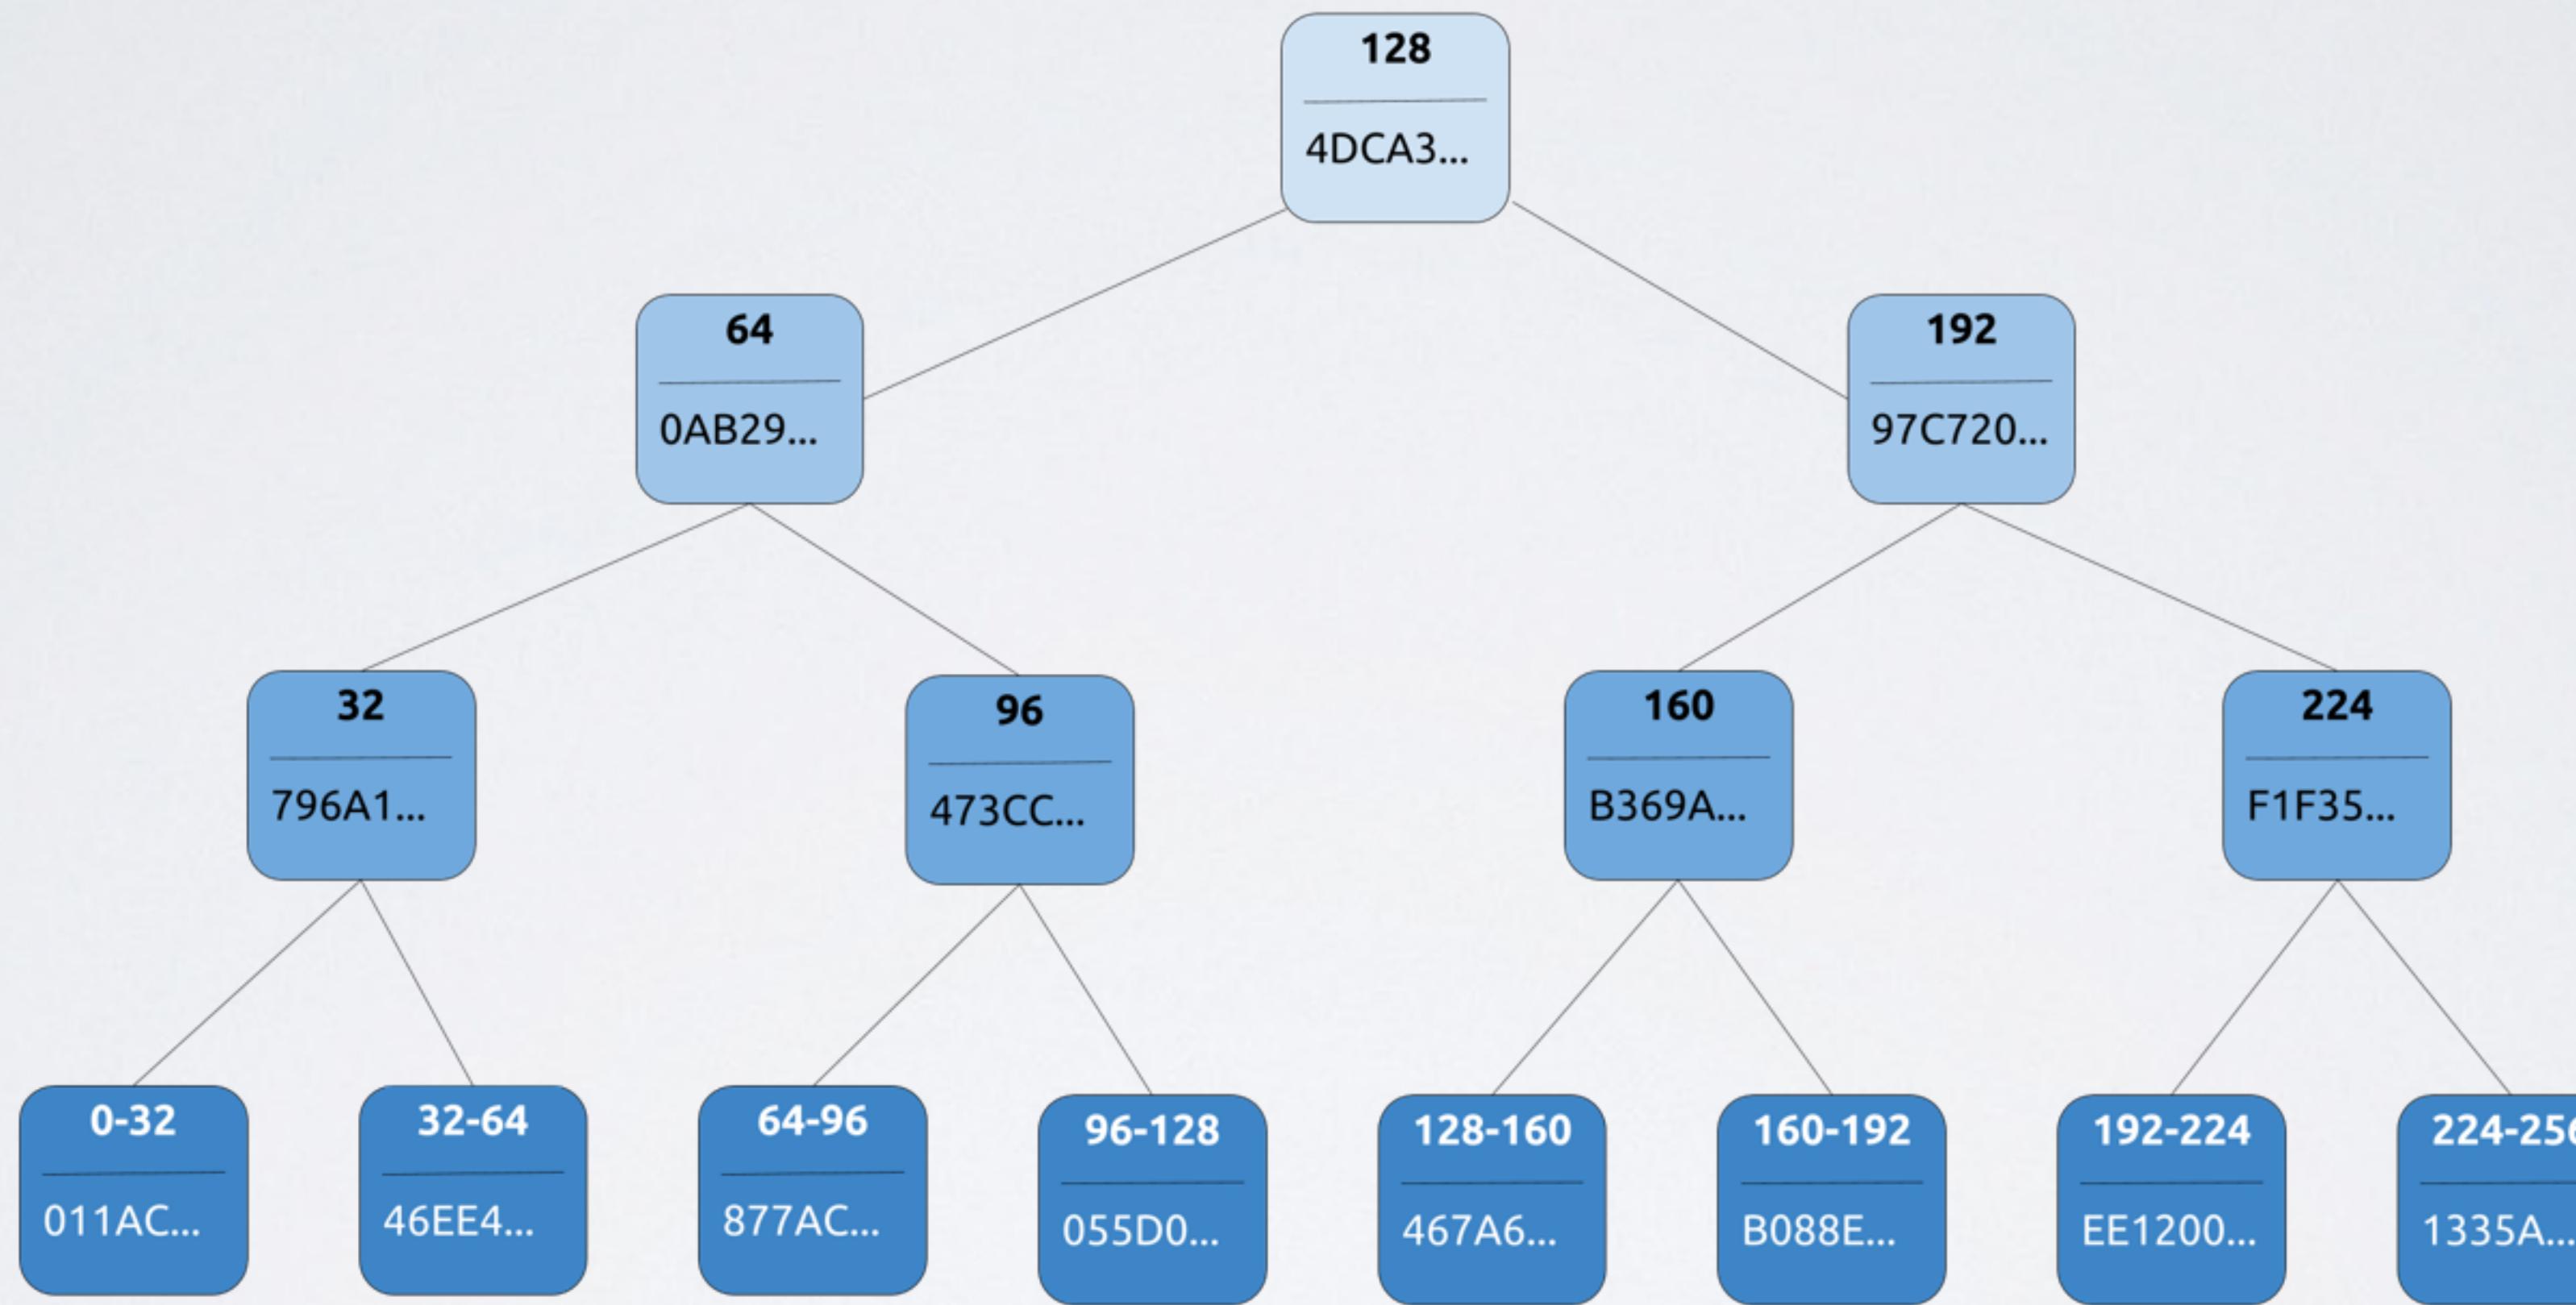
Reads all data

Calculates hashes

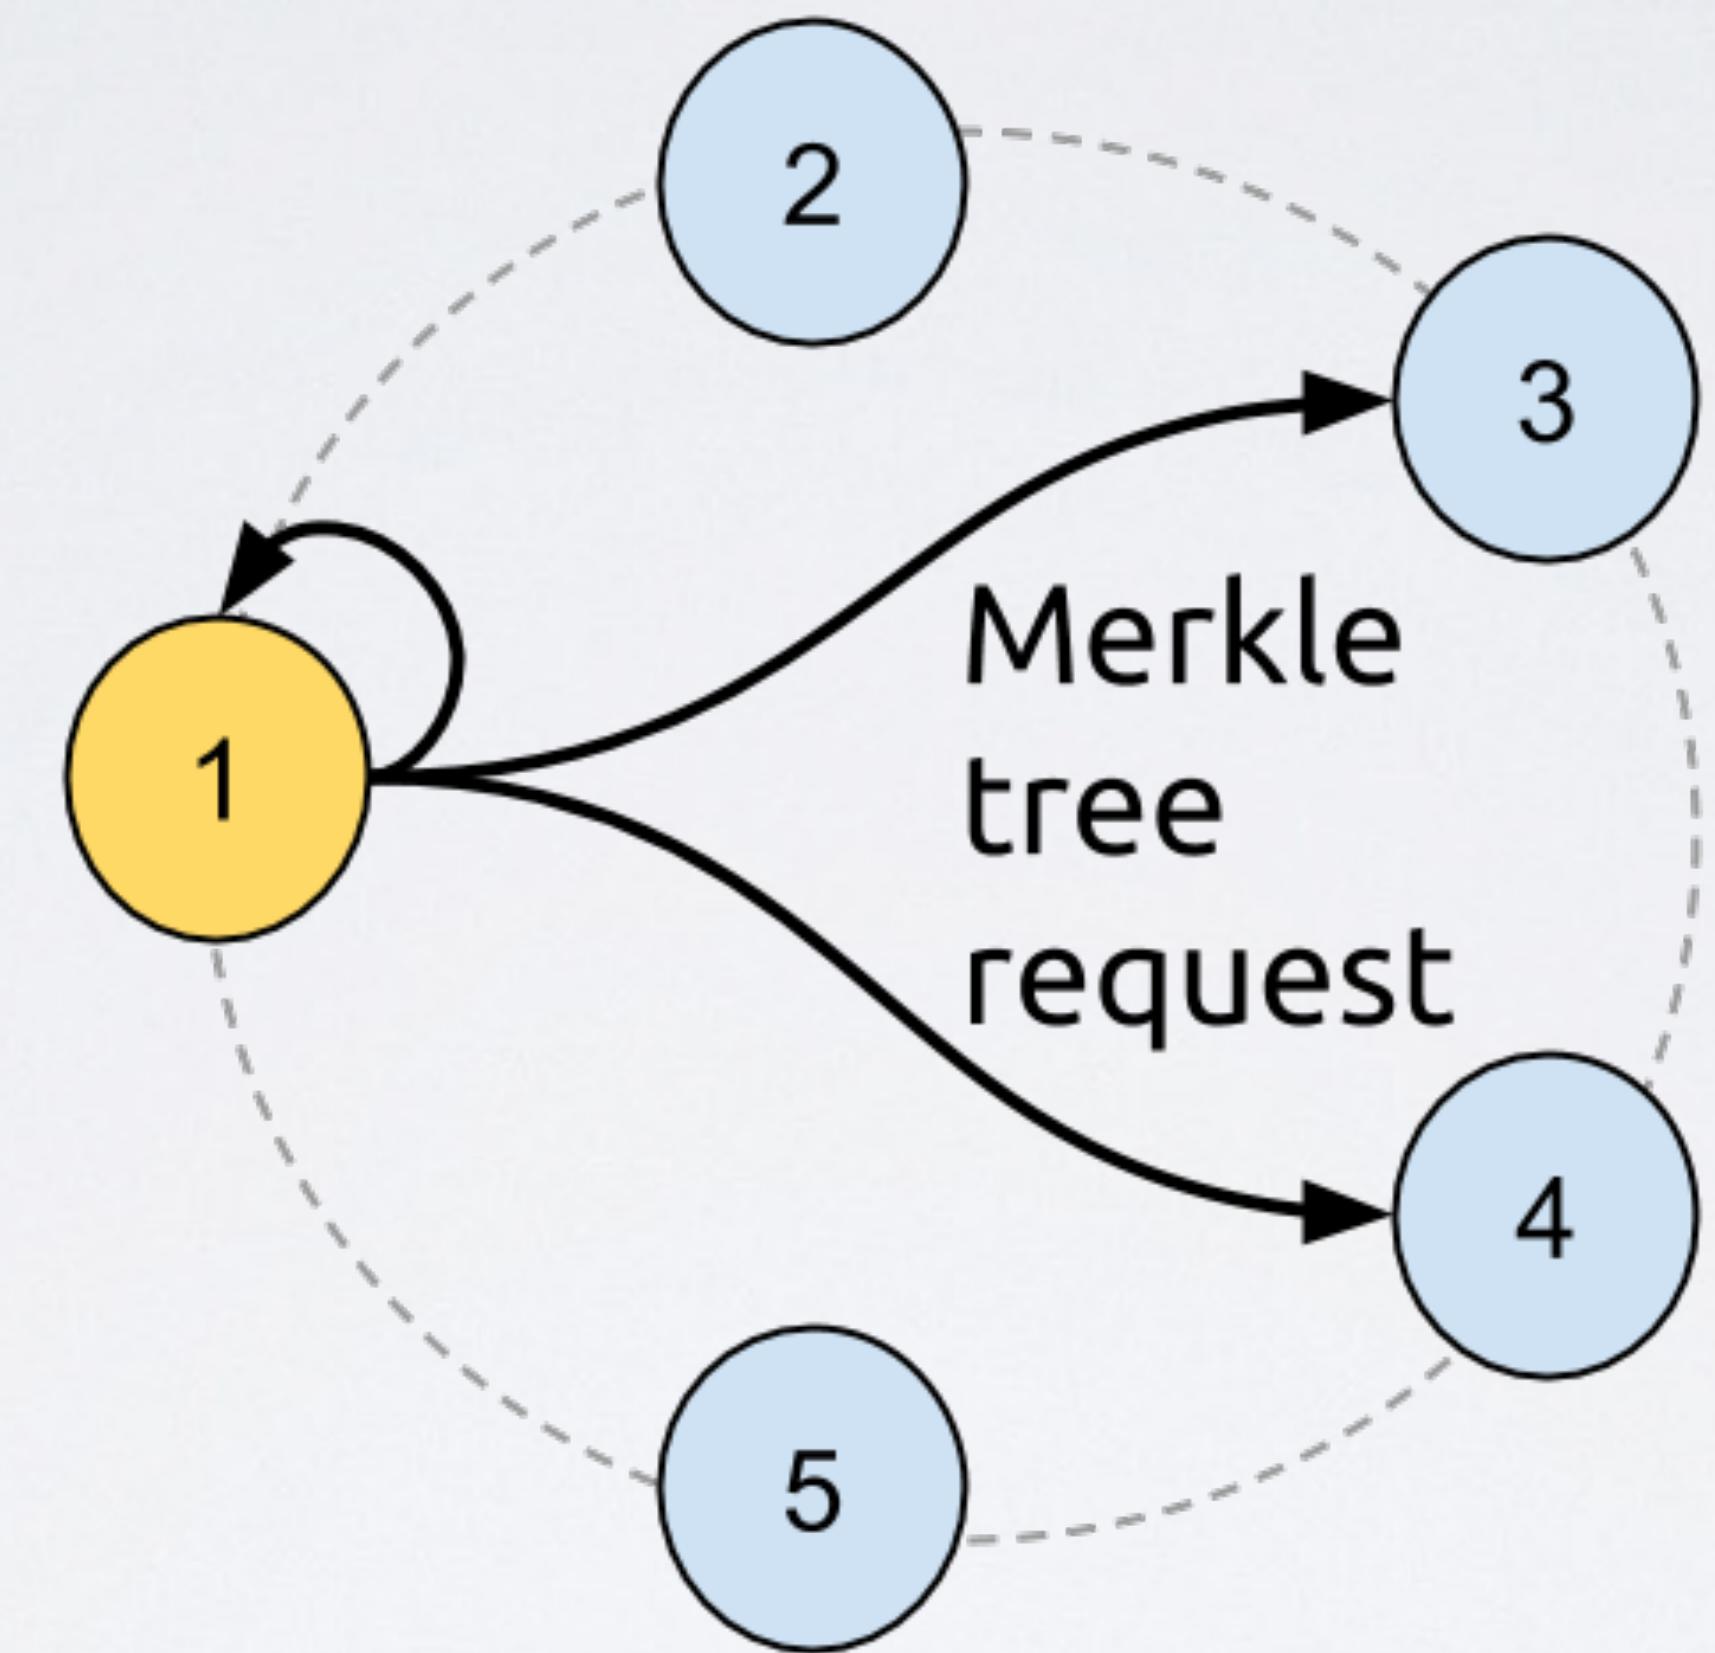
Compares hashes

Streams mismatching partitions

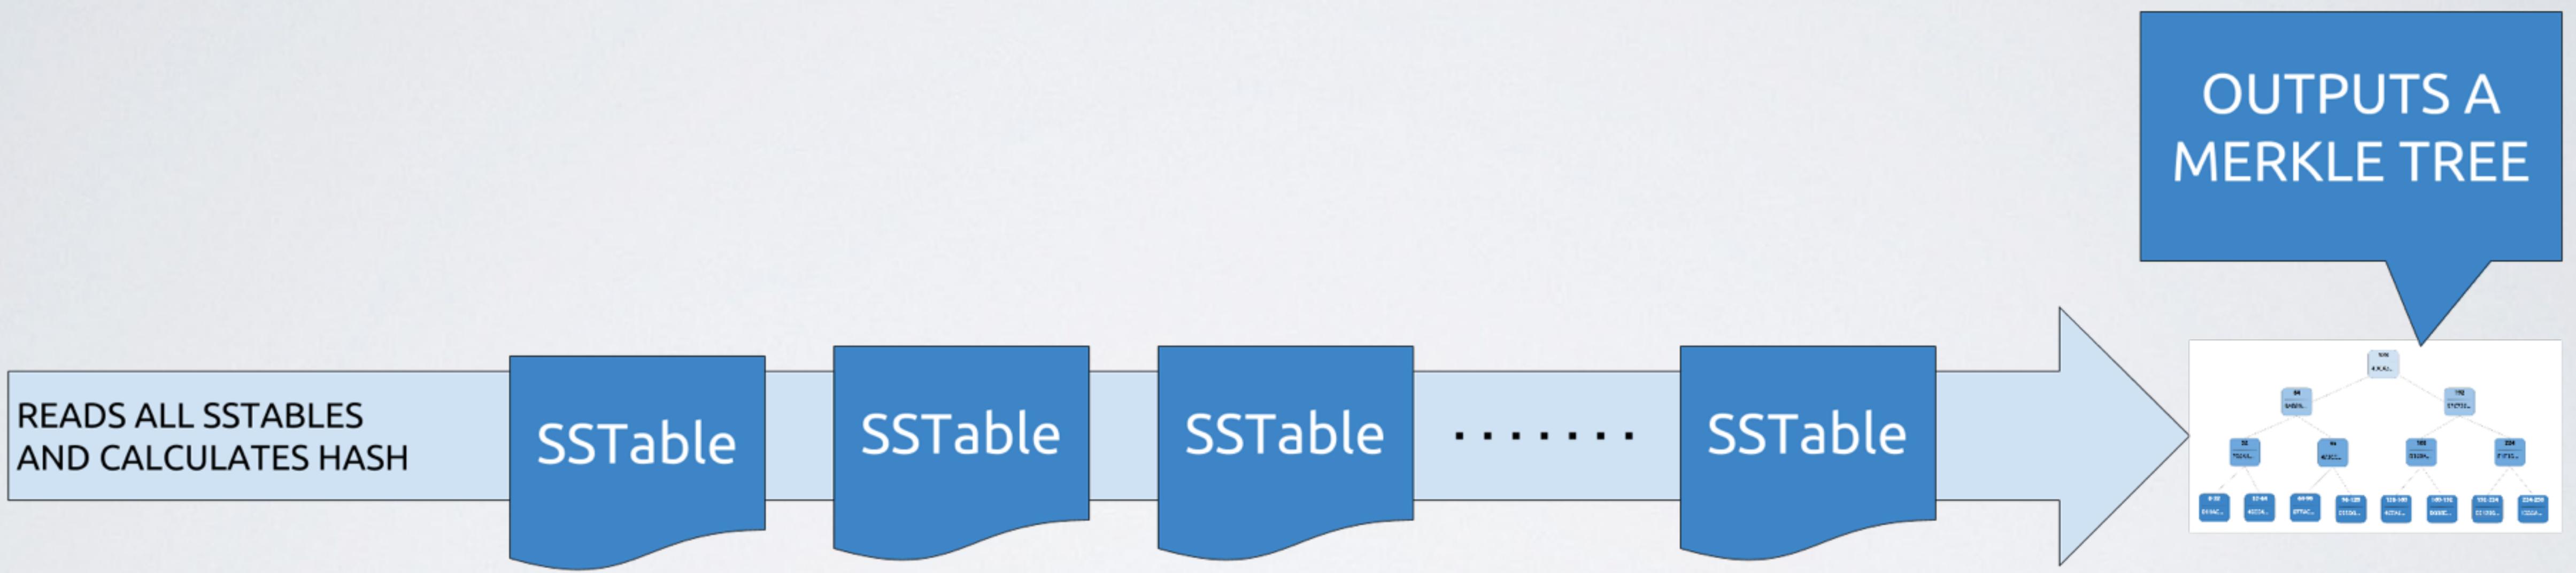
How does anti-entropy repair works ?



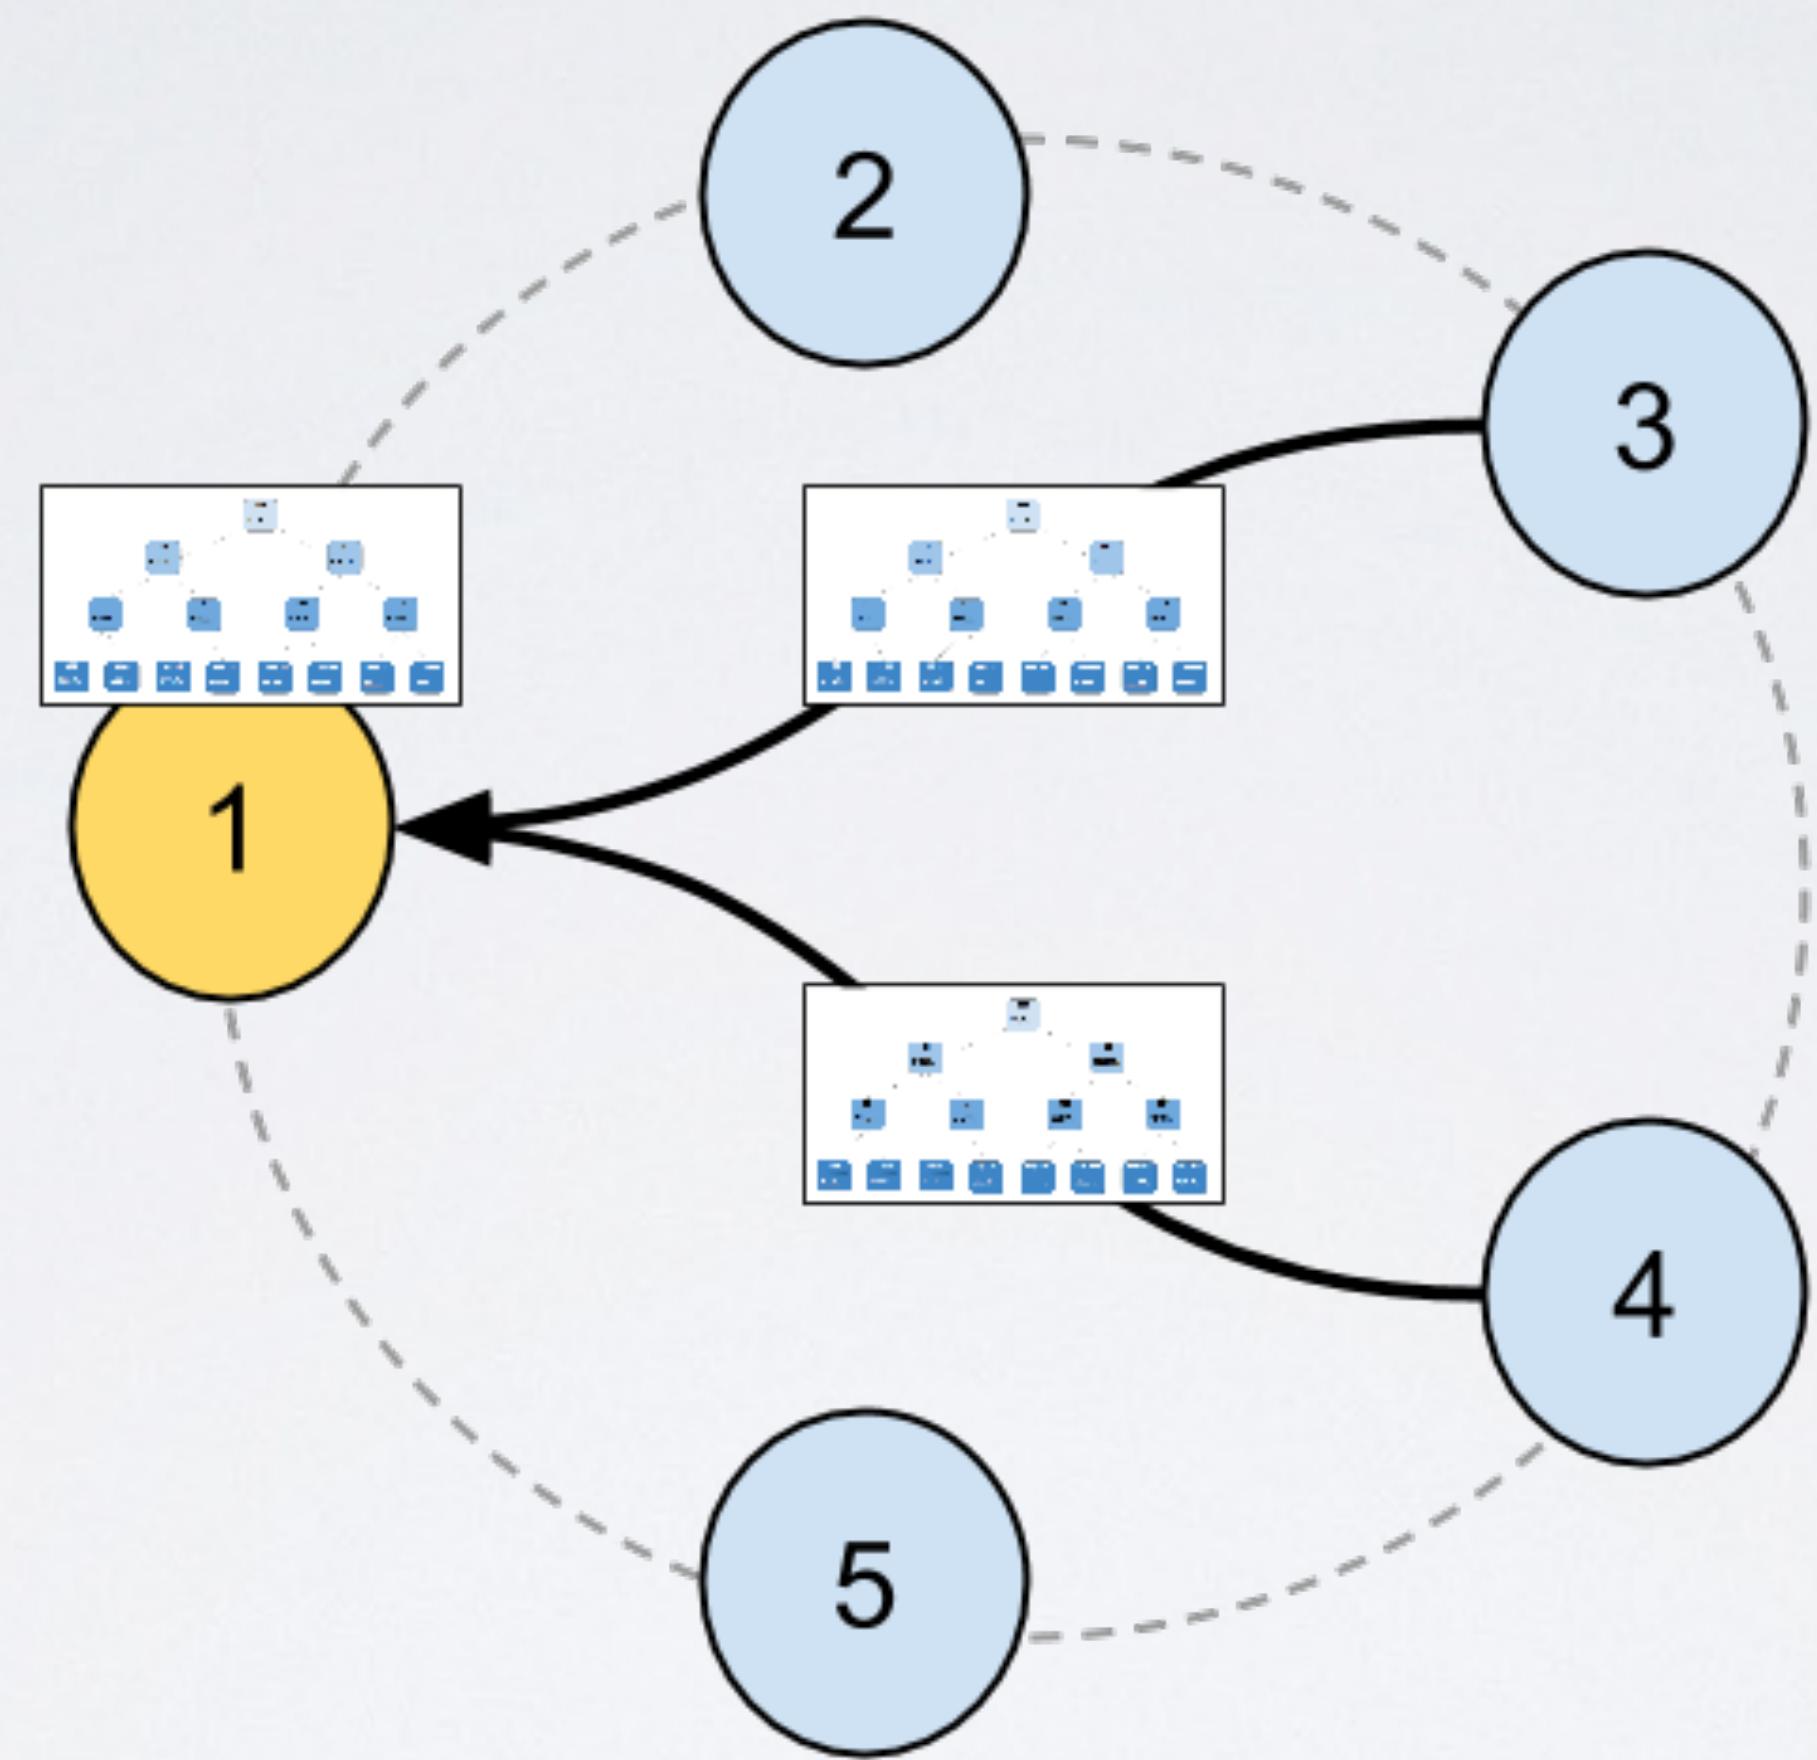
Merkle tree is requested to all replicas



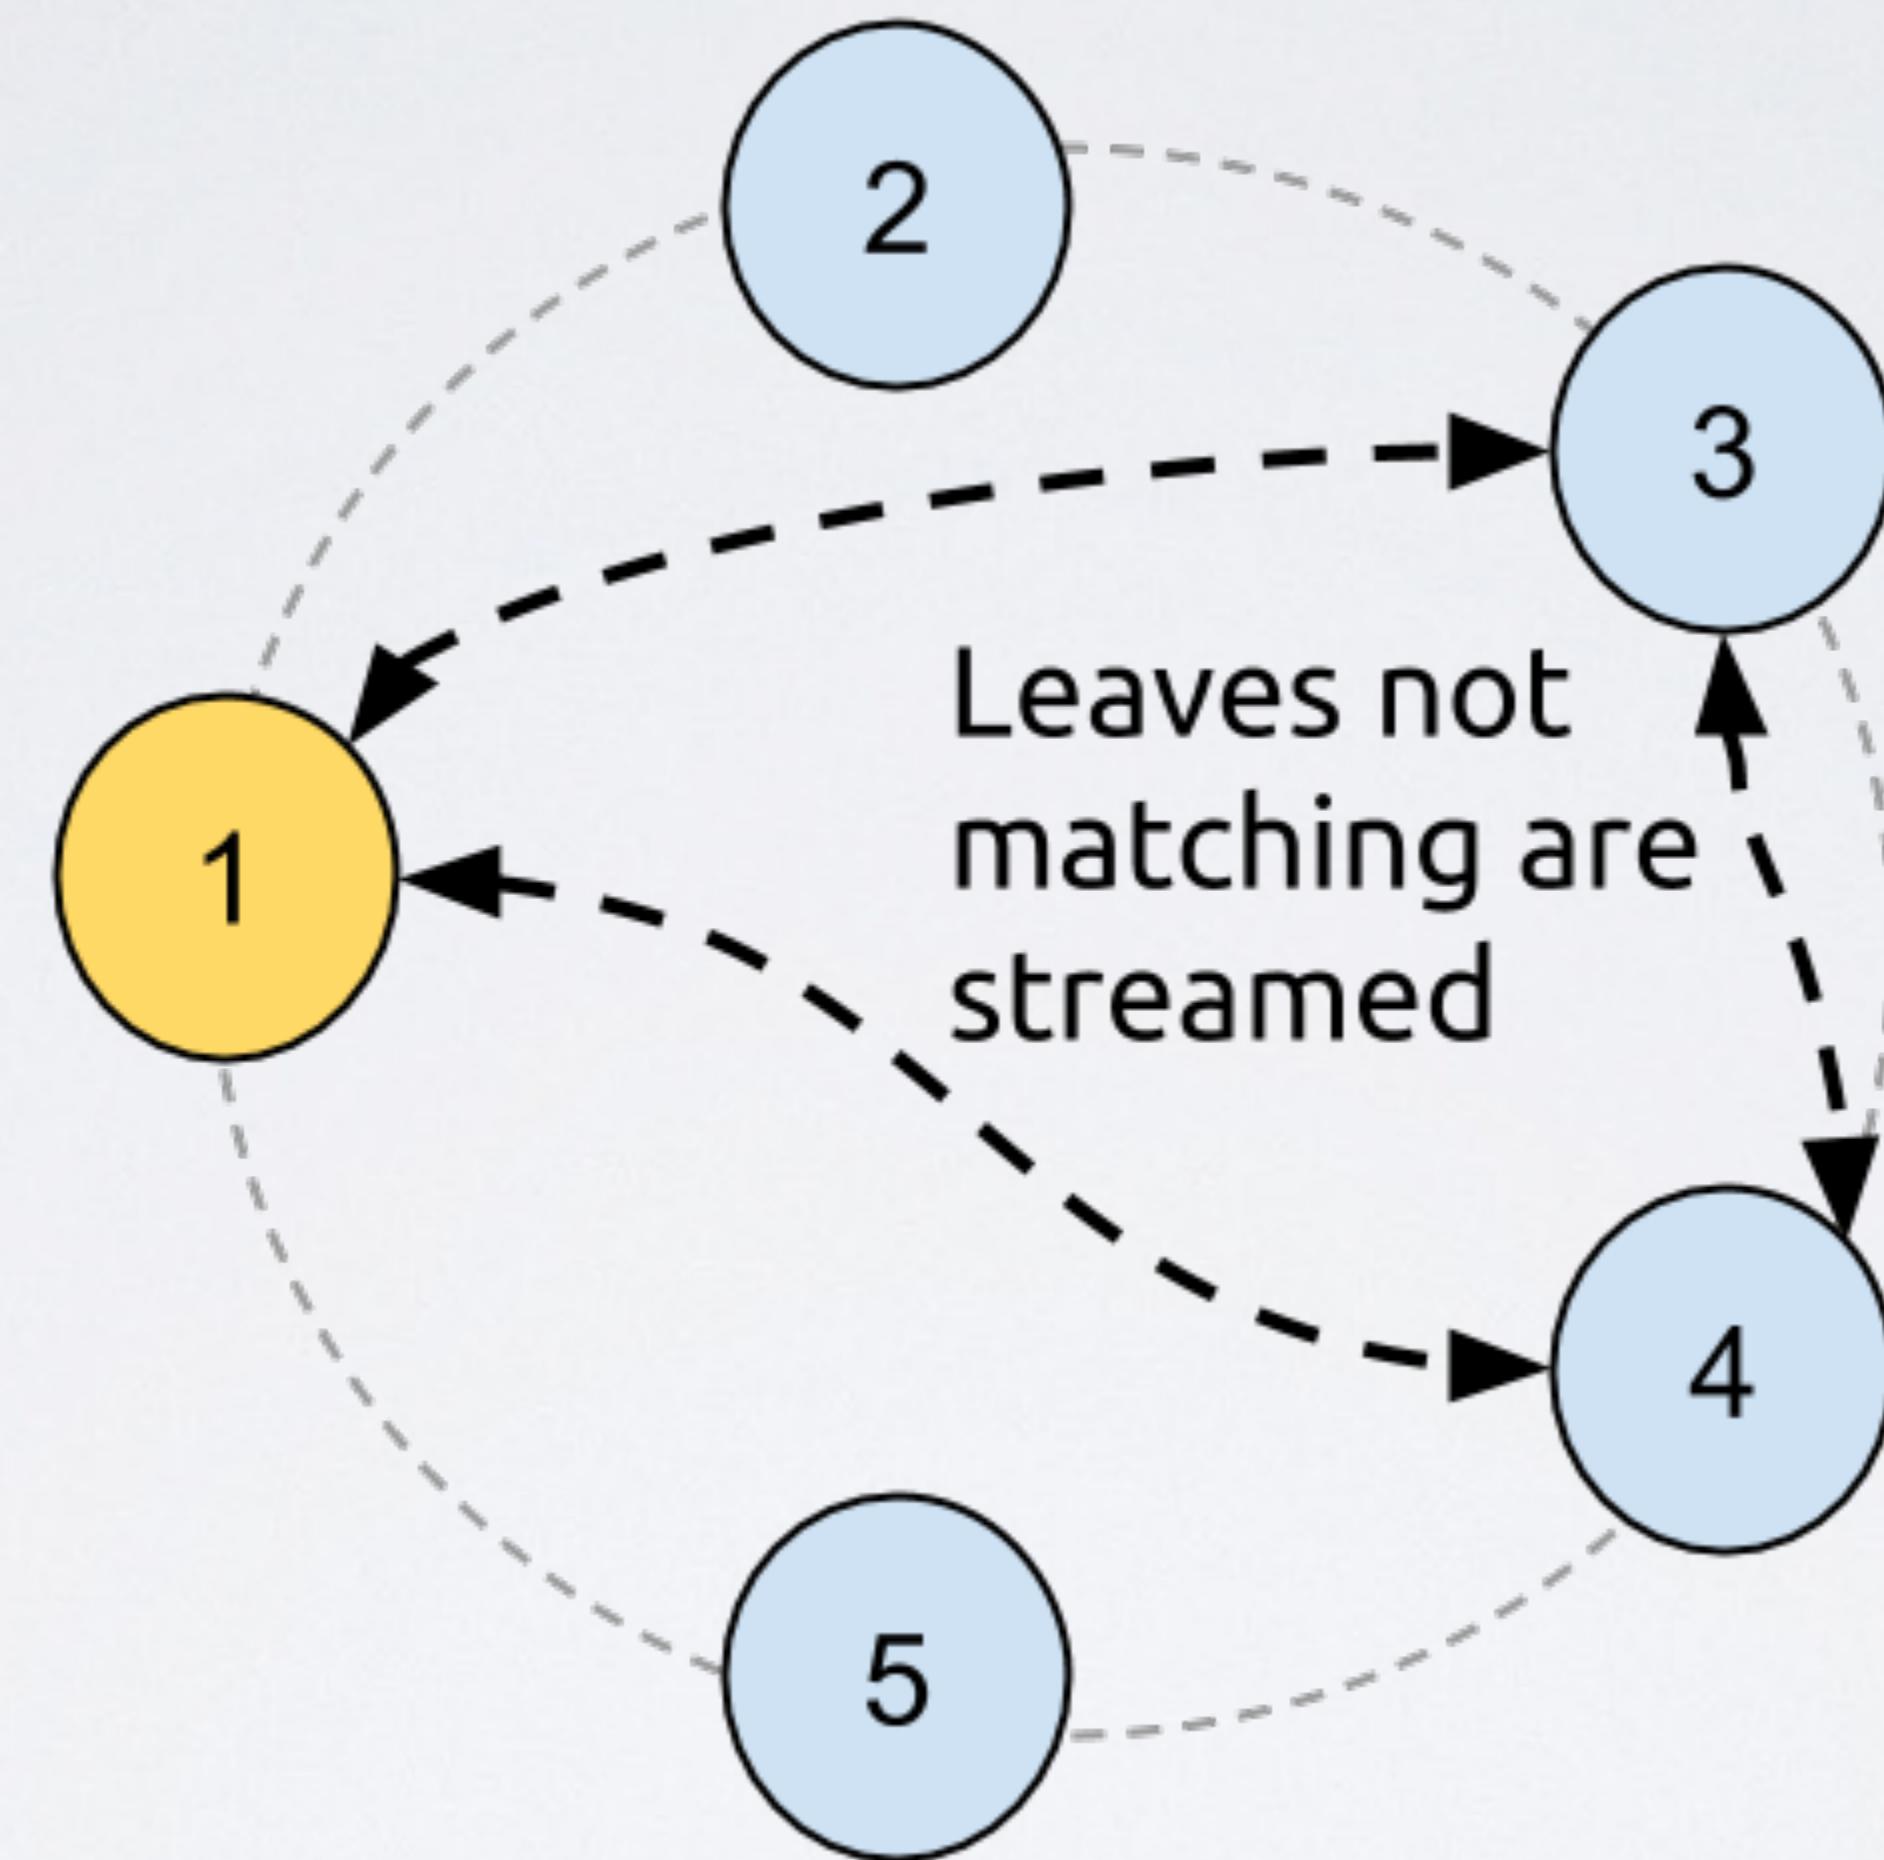
Validation compaction



Merkle tree comparison



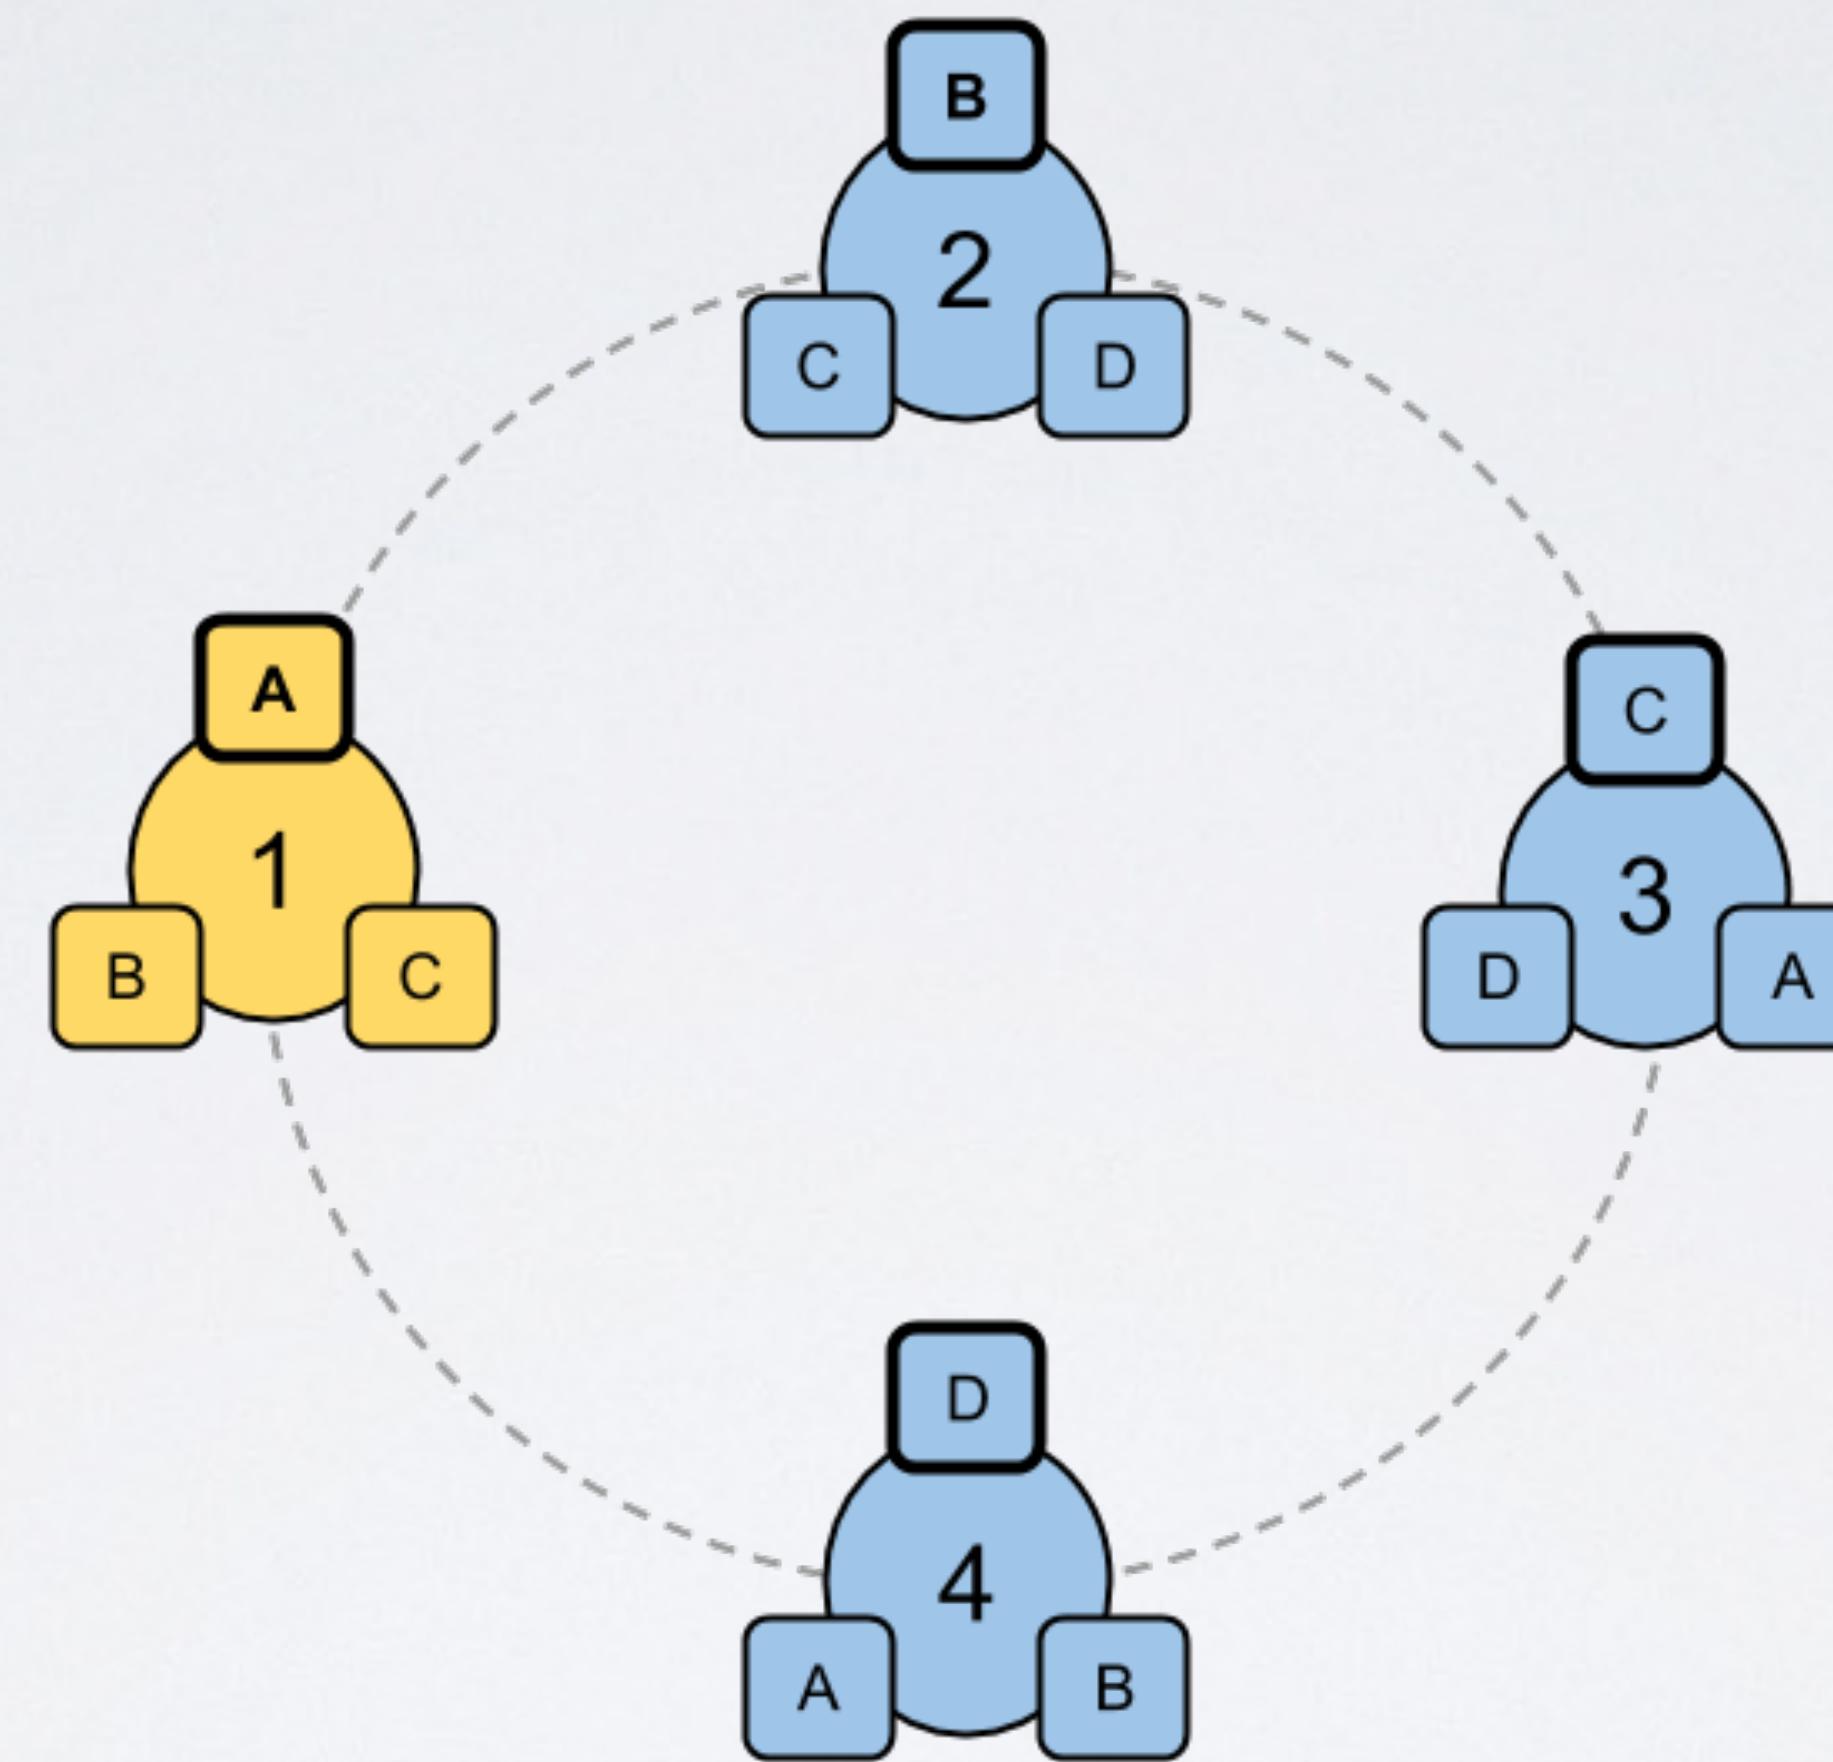
Streaming



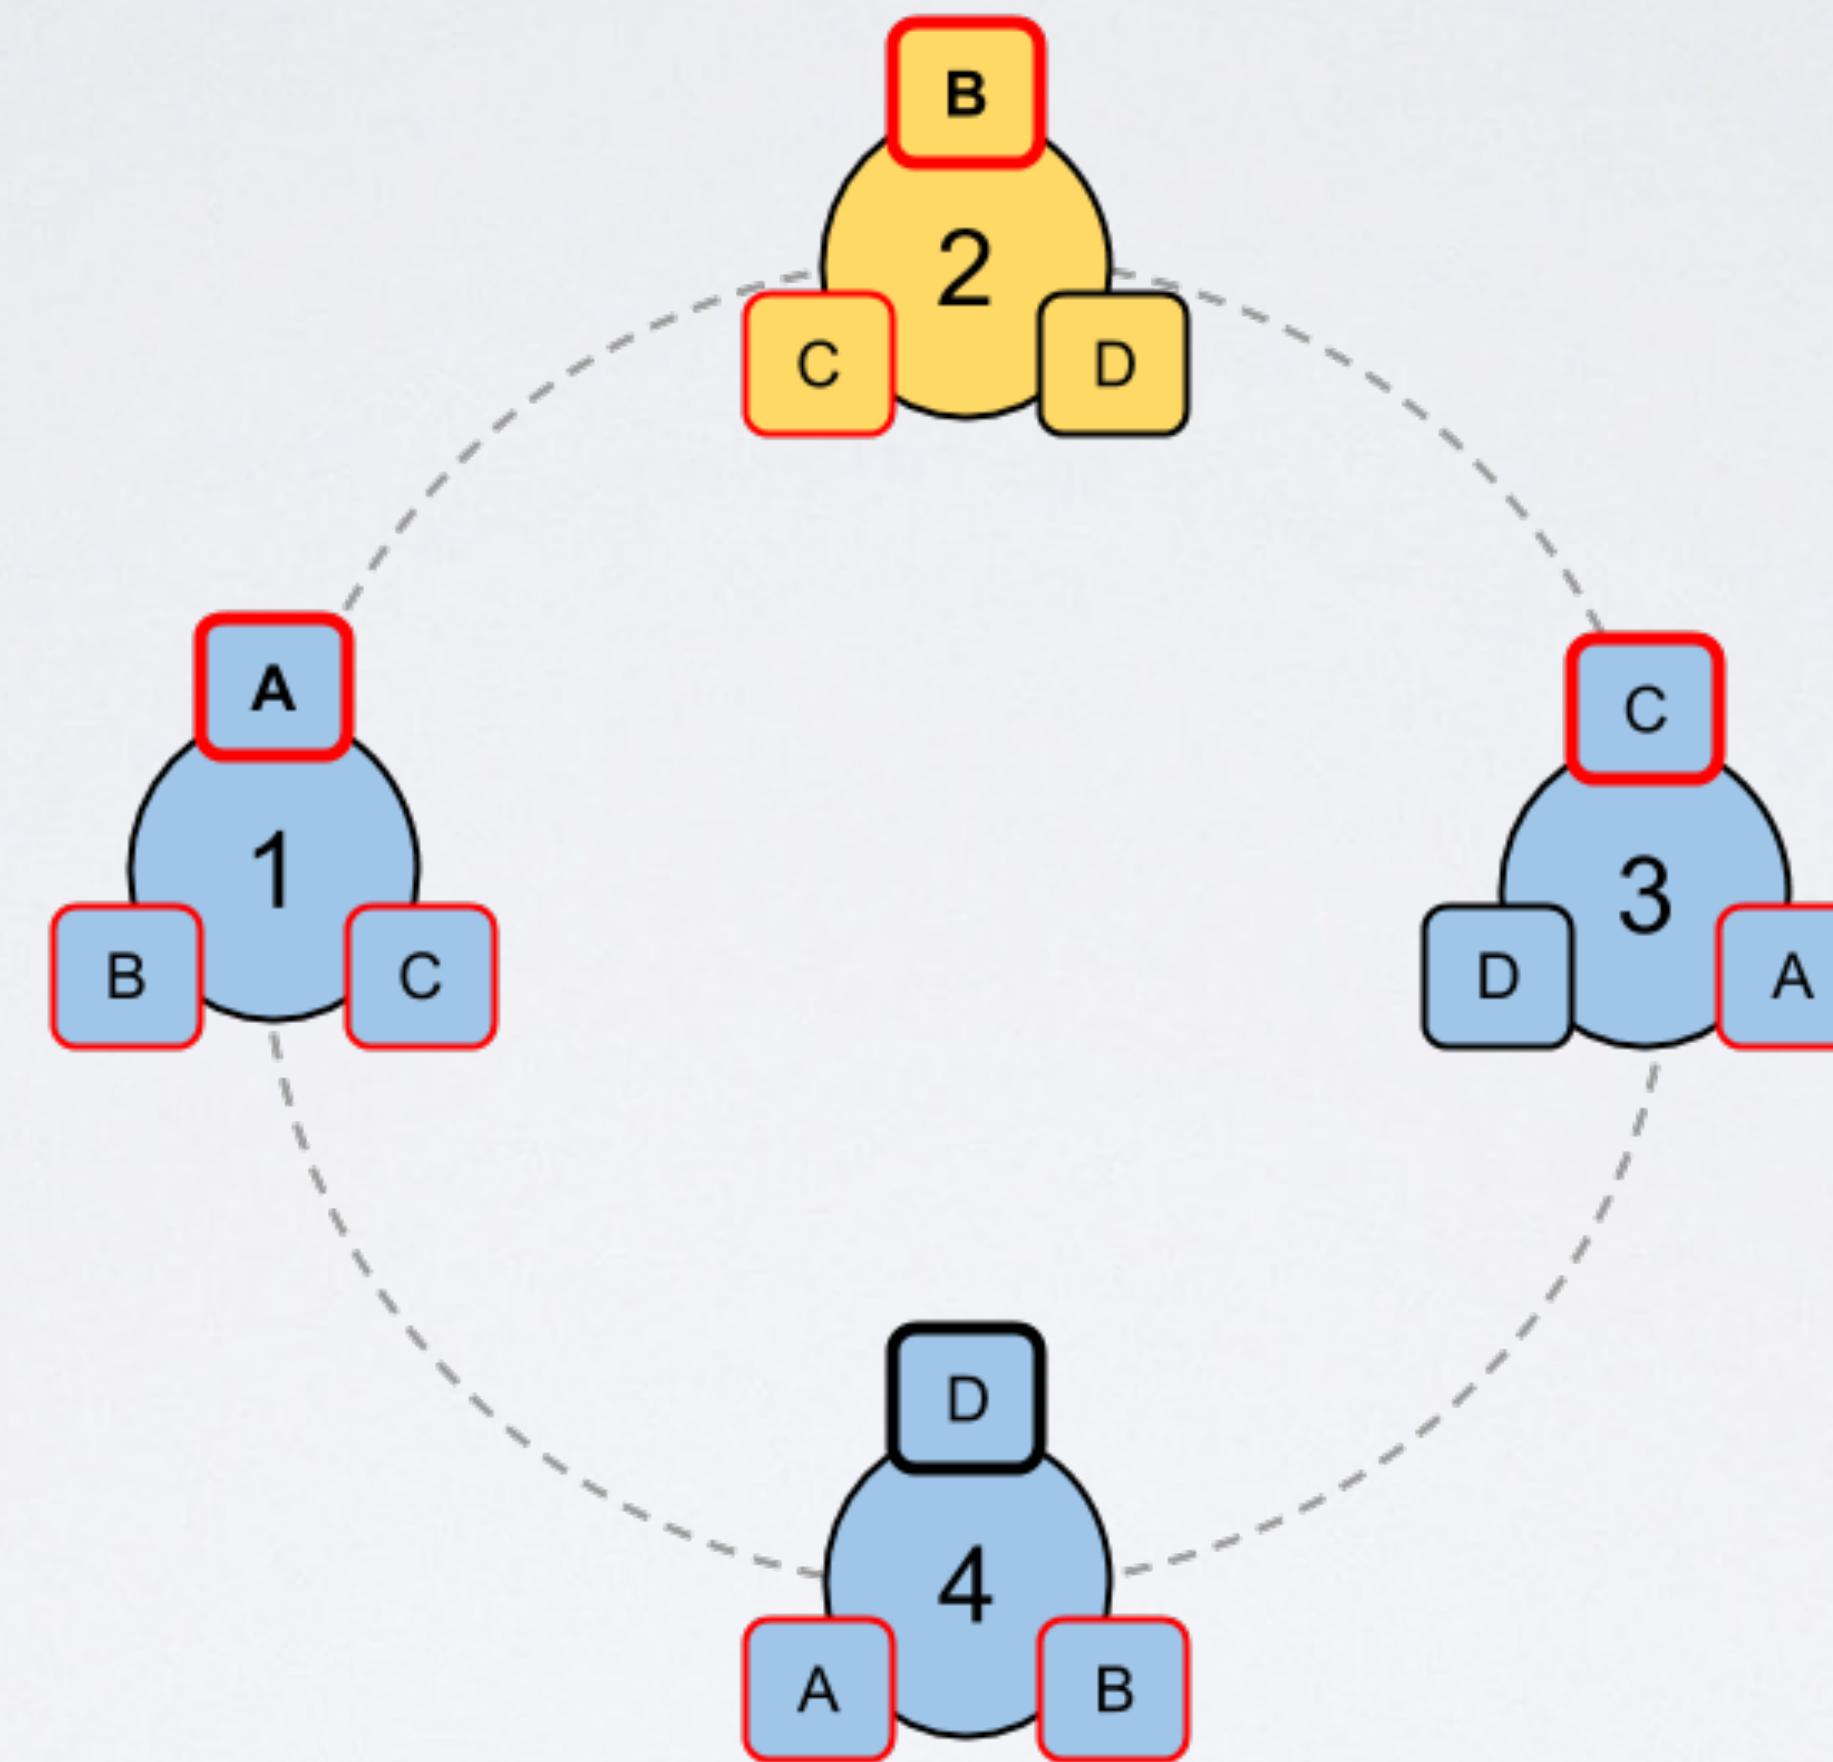
How do we run repair ?

nodetool repair

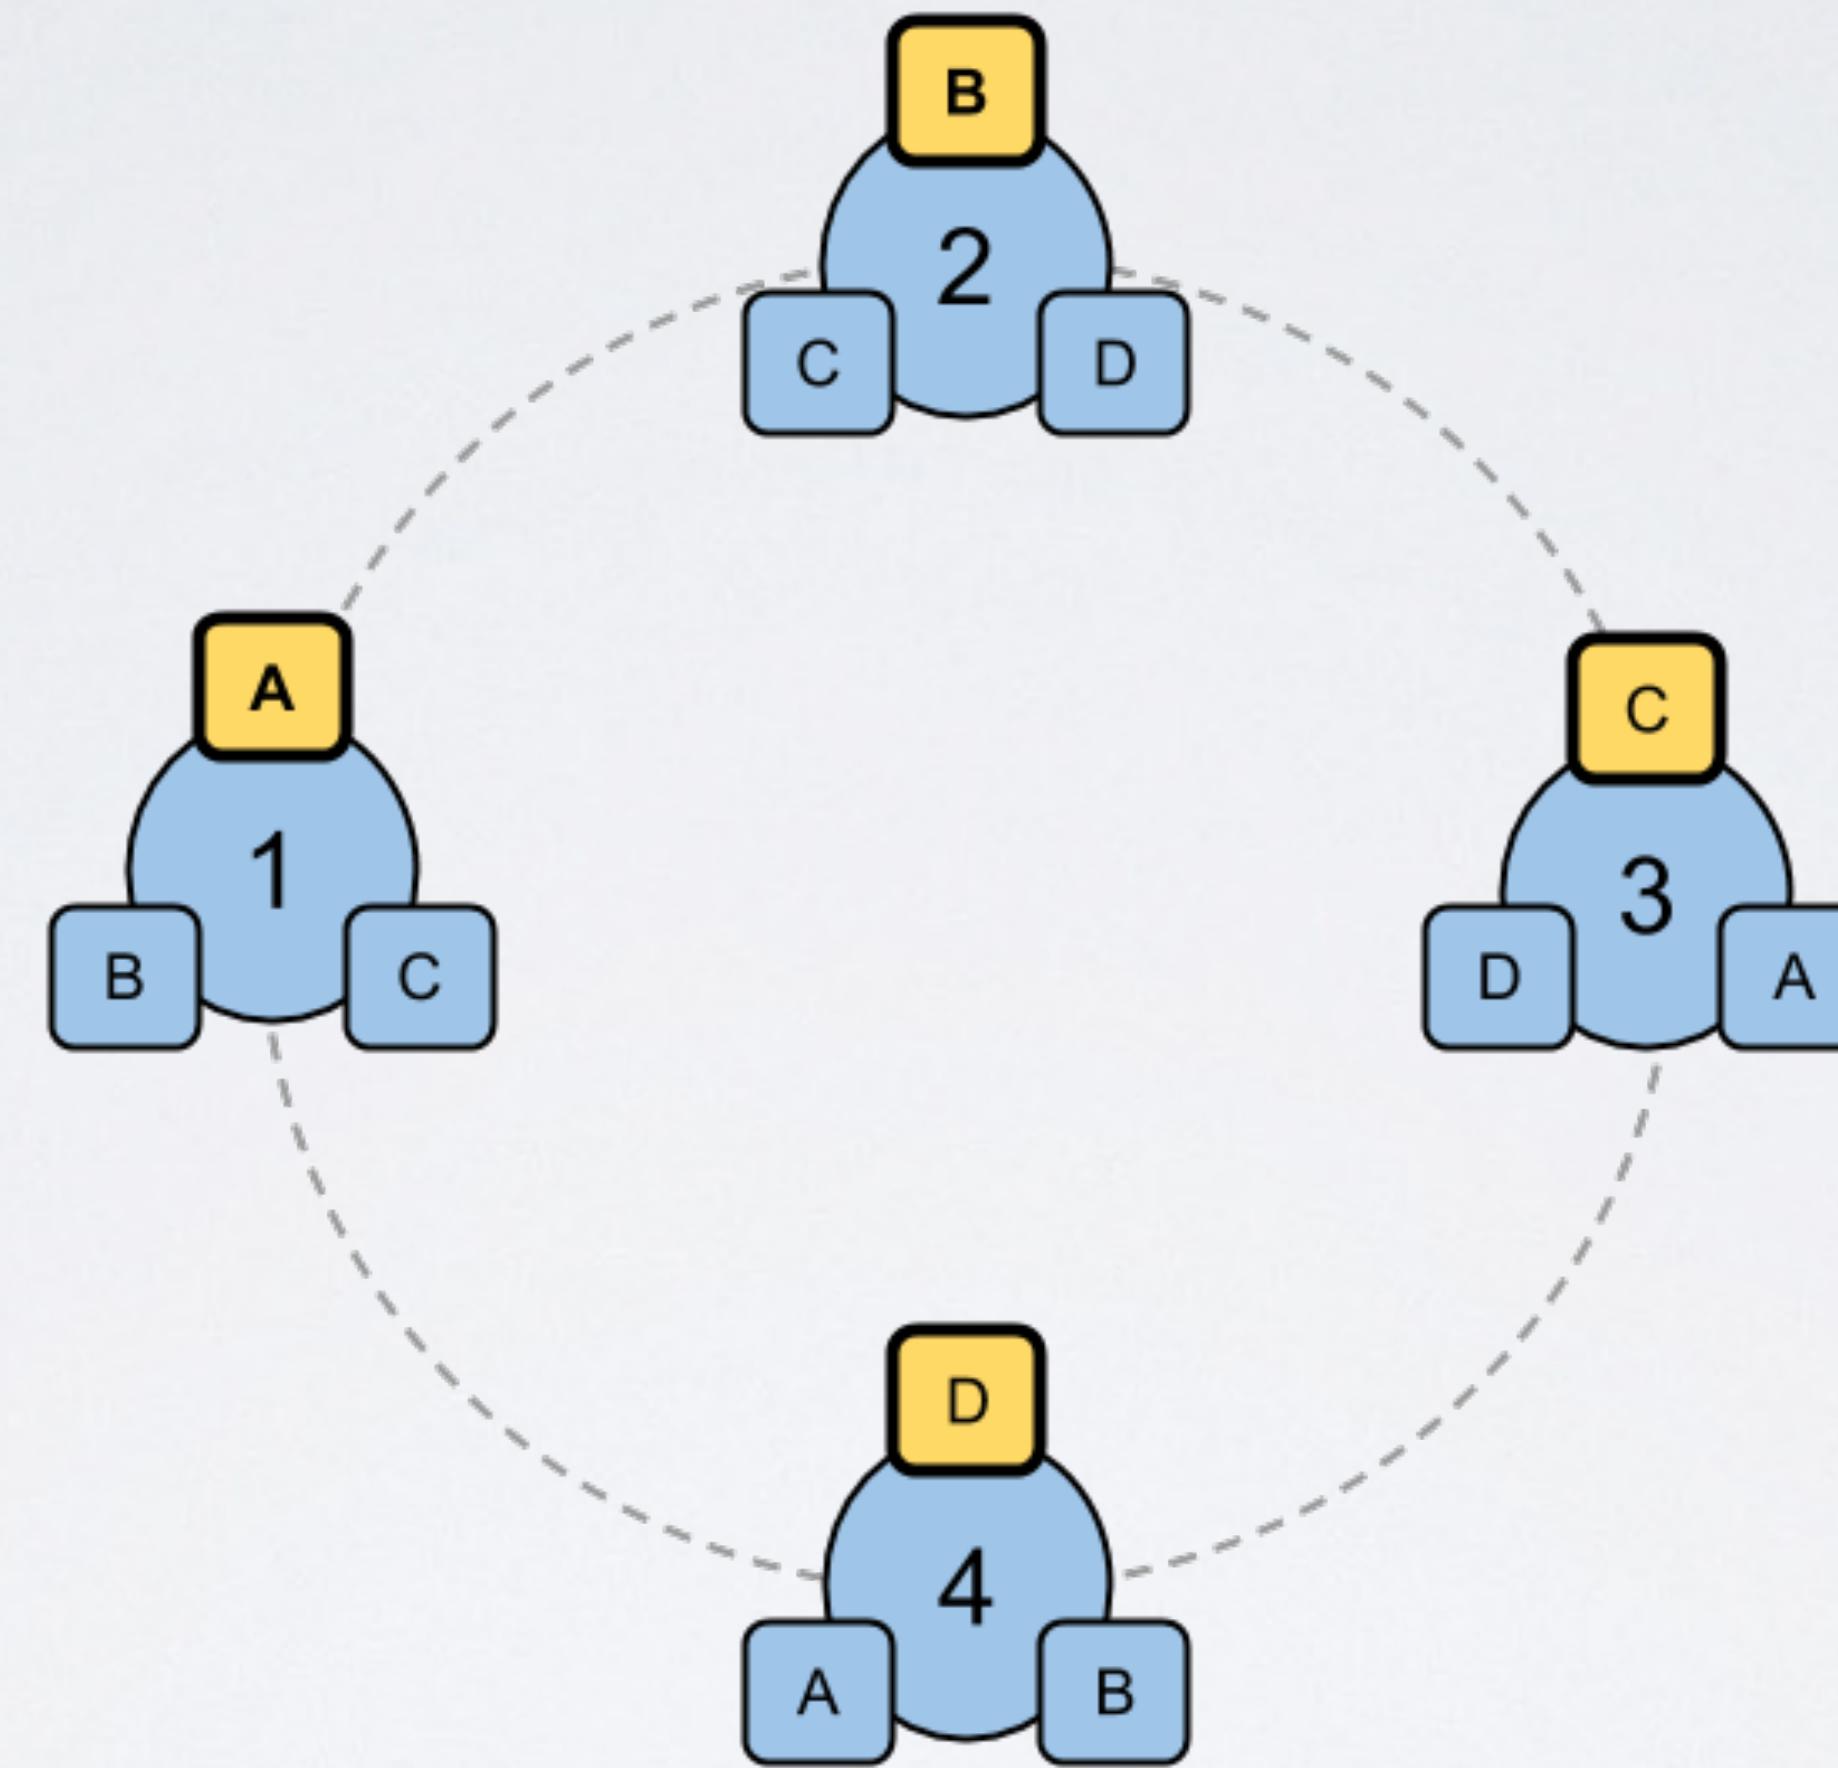
Improving repair



Improving repair



Improving repair



Improving repair

repairing each range **once** is enough

Improving repair

nodetool repair -pr

Improving repair

`nodetool repair -pr`
not suitable for node recovery

Repair too slow ?

Sequential repair is the default
since C* 2.0

Repair too slow ?

nodetool repair -par

The problem with dense nodes

Overstreaming

Leaves of the Merkle tree contain several partitions.

The solutions with dense nodes

cassandra_range_repair (Matt Stump & Brian Gallew)

Breaks the repair sessions in n steps

The solutions with dense nodes

vnodes : one repair session per vnode

Drawback : if you have many vnodes, repair takes longer

Repair in...



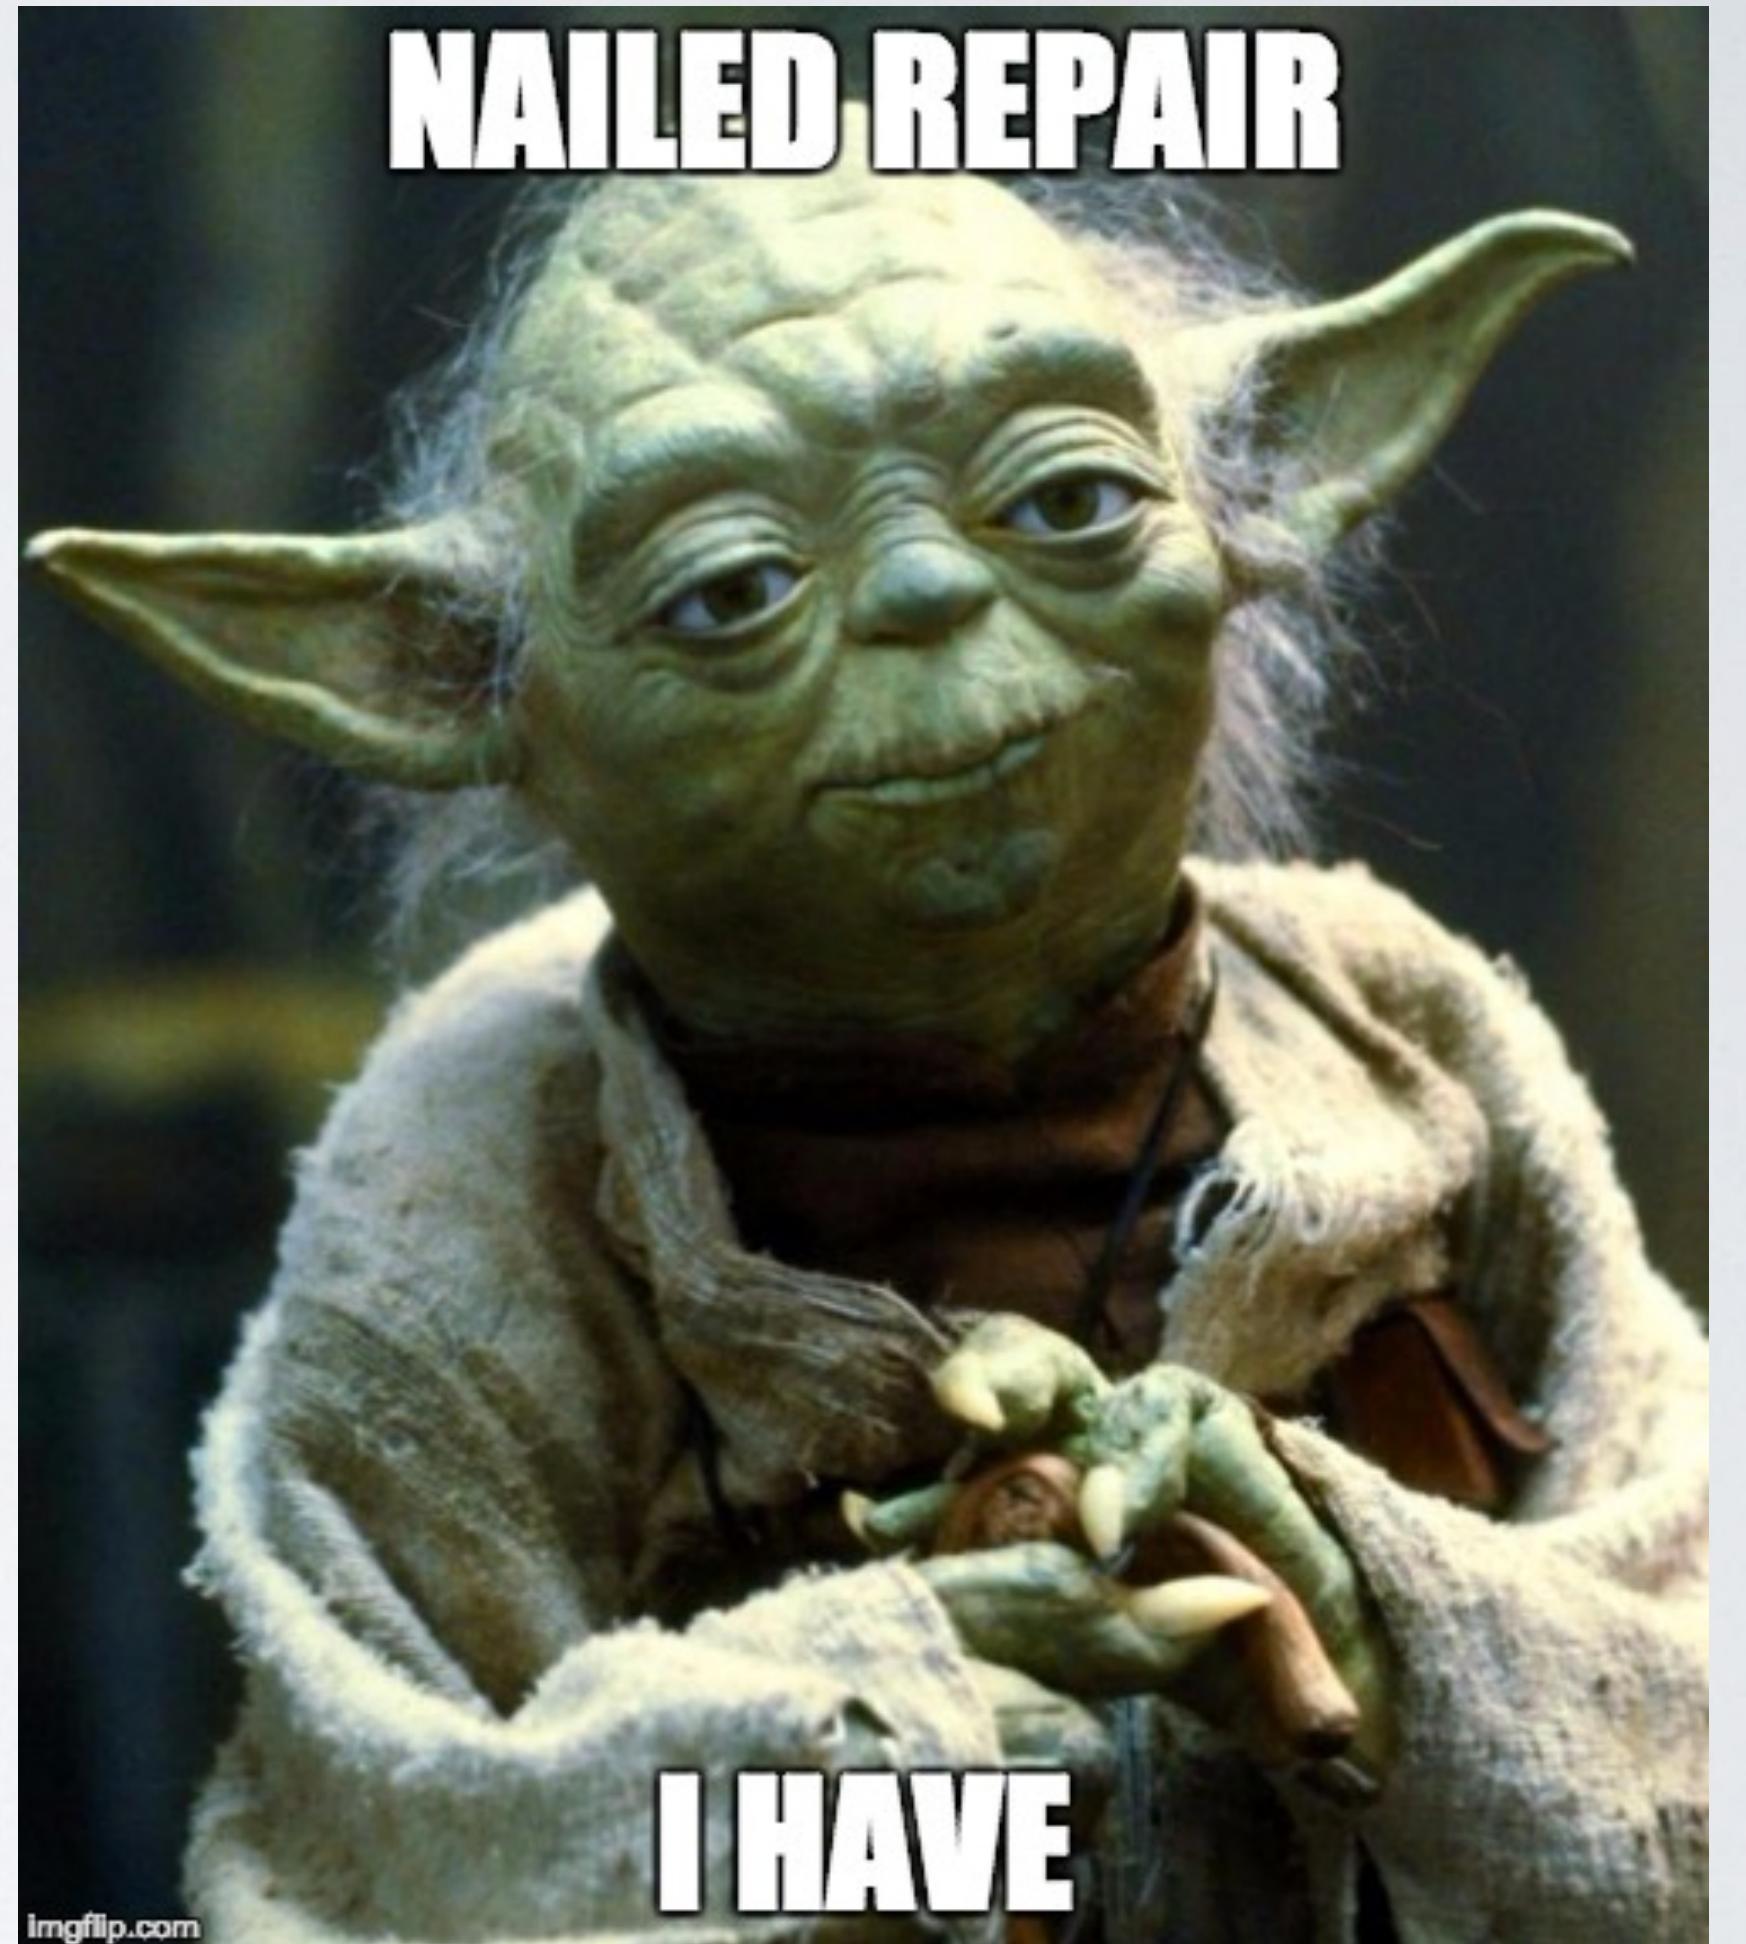
The early days of your cluster

Node density is low,
repair works just fine
however you run it.



The early days of your cluster

So maybe like I did,
you run « nodetool repair »
on all nodes... at the same
time



The (not so) early days of your cluster

As nodes gets higher
in density, repair takes
longer... and longer...



The (not so) early days of your cluster

... and latencies rise
as repair is a CPU and
I/O intensive operation



Your cluster is a grown up now

... until it breaks your
cluster



How can it break ?

Load gets too high

How can it break ?

Load gets too high

You don't meet your latency SLA anymore

How can it break ?

Load gets too high

How can it break ?

Load gets too high
Streams get stuck

How can it break ?

Load gets too high
Streams get stuck
and out of nowhere, all nodes start to eat
all your CPU doing nothing

The fun part ?

You need to run repair
to recover from
the repair outage !



The cluster keeps growing

And you realize **orchestration** is needed
to stop blowing up your cluster



Orchestrating repair

**Repair must not
run on all nodes
at the same time**



Tools to orchestrate repairs

OpsCenter repair service (DSE users)
Spotify reaper

Spotify reaper

<https://github.com/spotify/cassandra-reaper>

Spotify reaper

Performs subrange repair

Spotify reaper

Performs subrange repair
Limits repair pressure

Spotify reaper

Performs subrange repair
Limits repair pressure
Retries failed sessions

Spotify reaper

- Performs subrange repair
- Limits repair pressure
- Retries failed sessions
- Schedules cyclic repairs

Spotify reaper

- Performs subrange repair
- Limits repair pressure
- Retries failed sessions
- Schedules cyclic repairs
- Optimizes cluster load

Spotify reaper - with UI (thx Stefan Podkowinski)

The screenshot shows the Spotify reaper UI with the "Repair" tab selected. The interface is divided into two main sections: "Done" and "Repair".

Done:

ID	State	Cluster	Keyspace	CFs	Incremental	Repaired	Actions
1	NOT_STARTED	twcs308	test		false	0/201	<button>Activate</button> <button>Delete</button>

Repair:

Form fields for repairing a keyspace:

- Cluster*: twcs308
- Keyspace*: test
- Tables: table1, table2, table3
- Owner*: alex
- Segment count: amount of segments to create for repair
- Parallelism: Parallel
- Repair intensity: repair intensity
- Cause: reason repair was started
- Incremental: false

A large orange "Repair" button is located at the bottom of the form.

What and why ?

Full repair

Incremental repair

How to make it work

Automated repairs

What if we stopped repairing repaired data ?



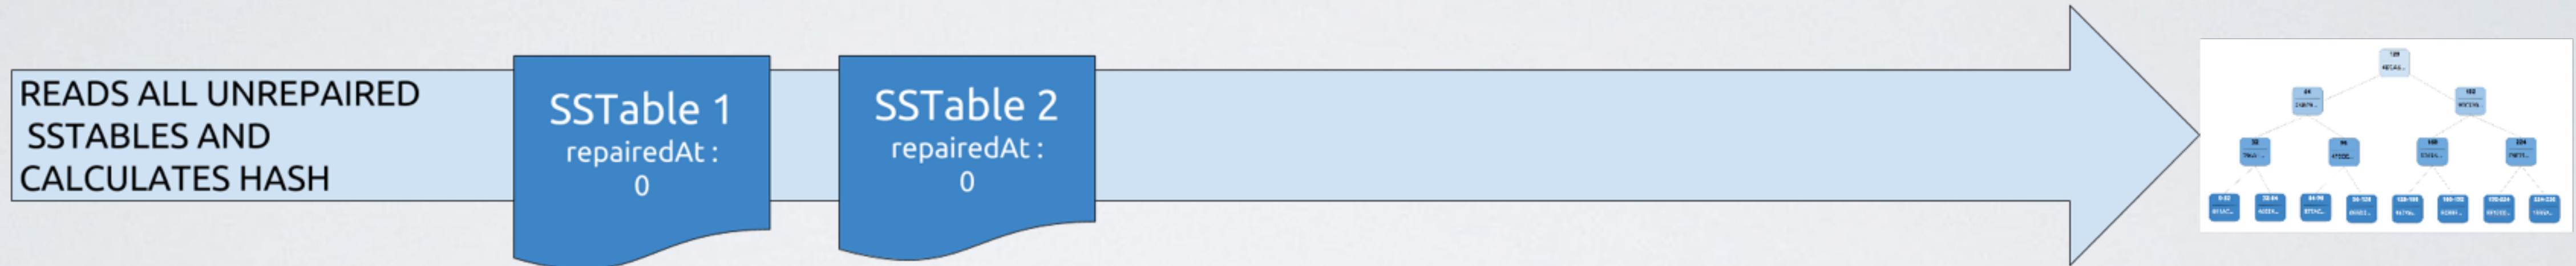
Here comes the savior !

C* 2.1 introduces incremental repair

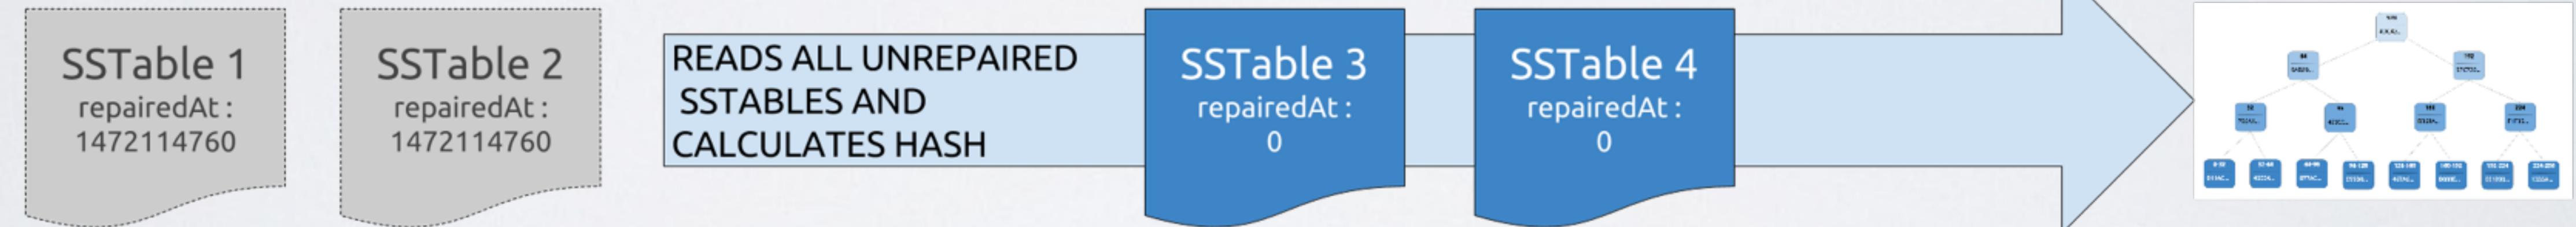
Default repair mode since C* 2.2

How does incremental repair work ?

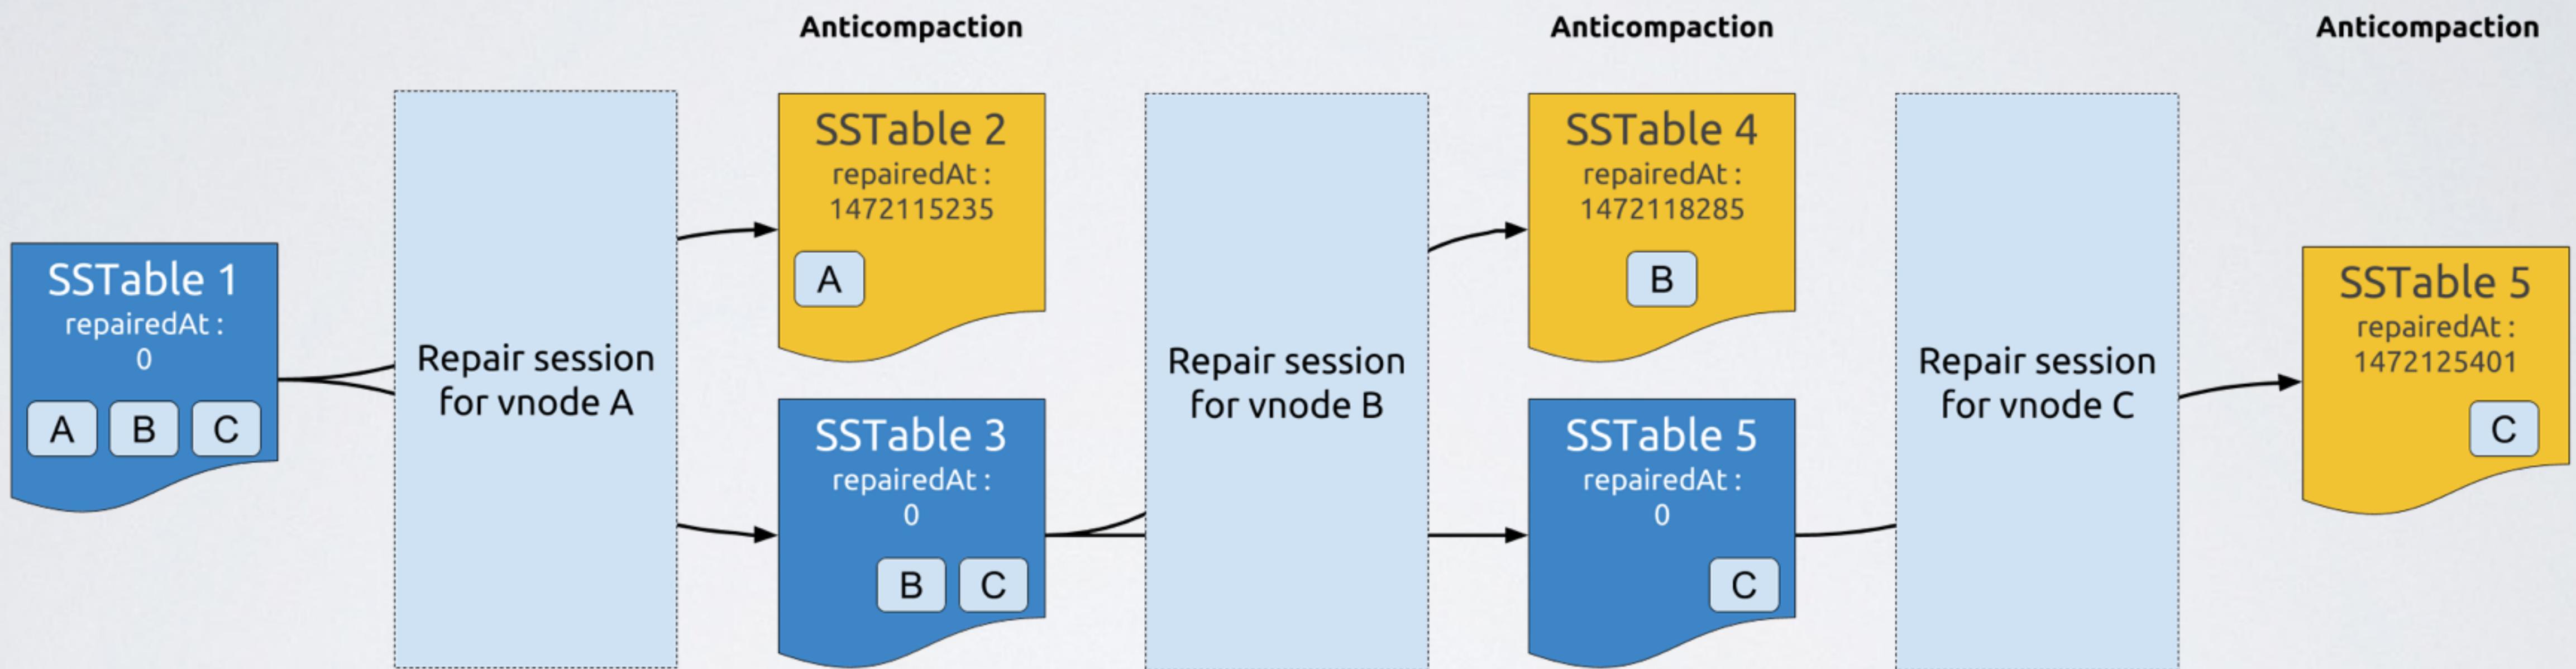
First repair



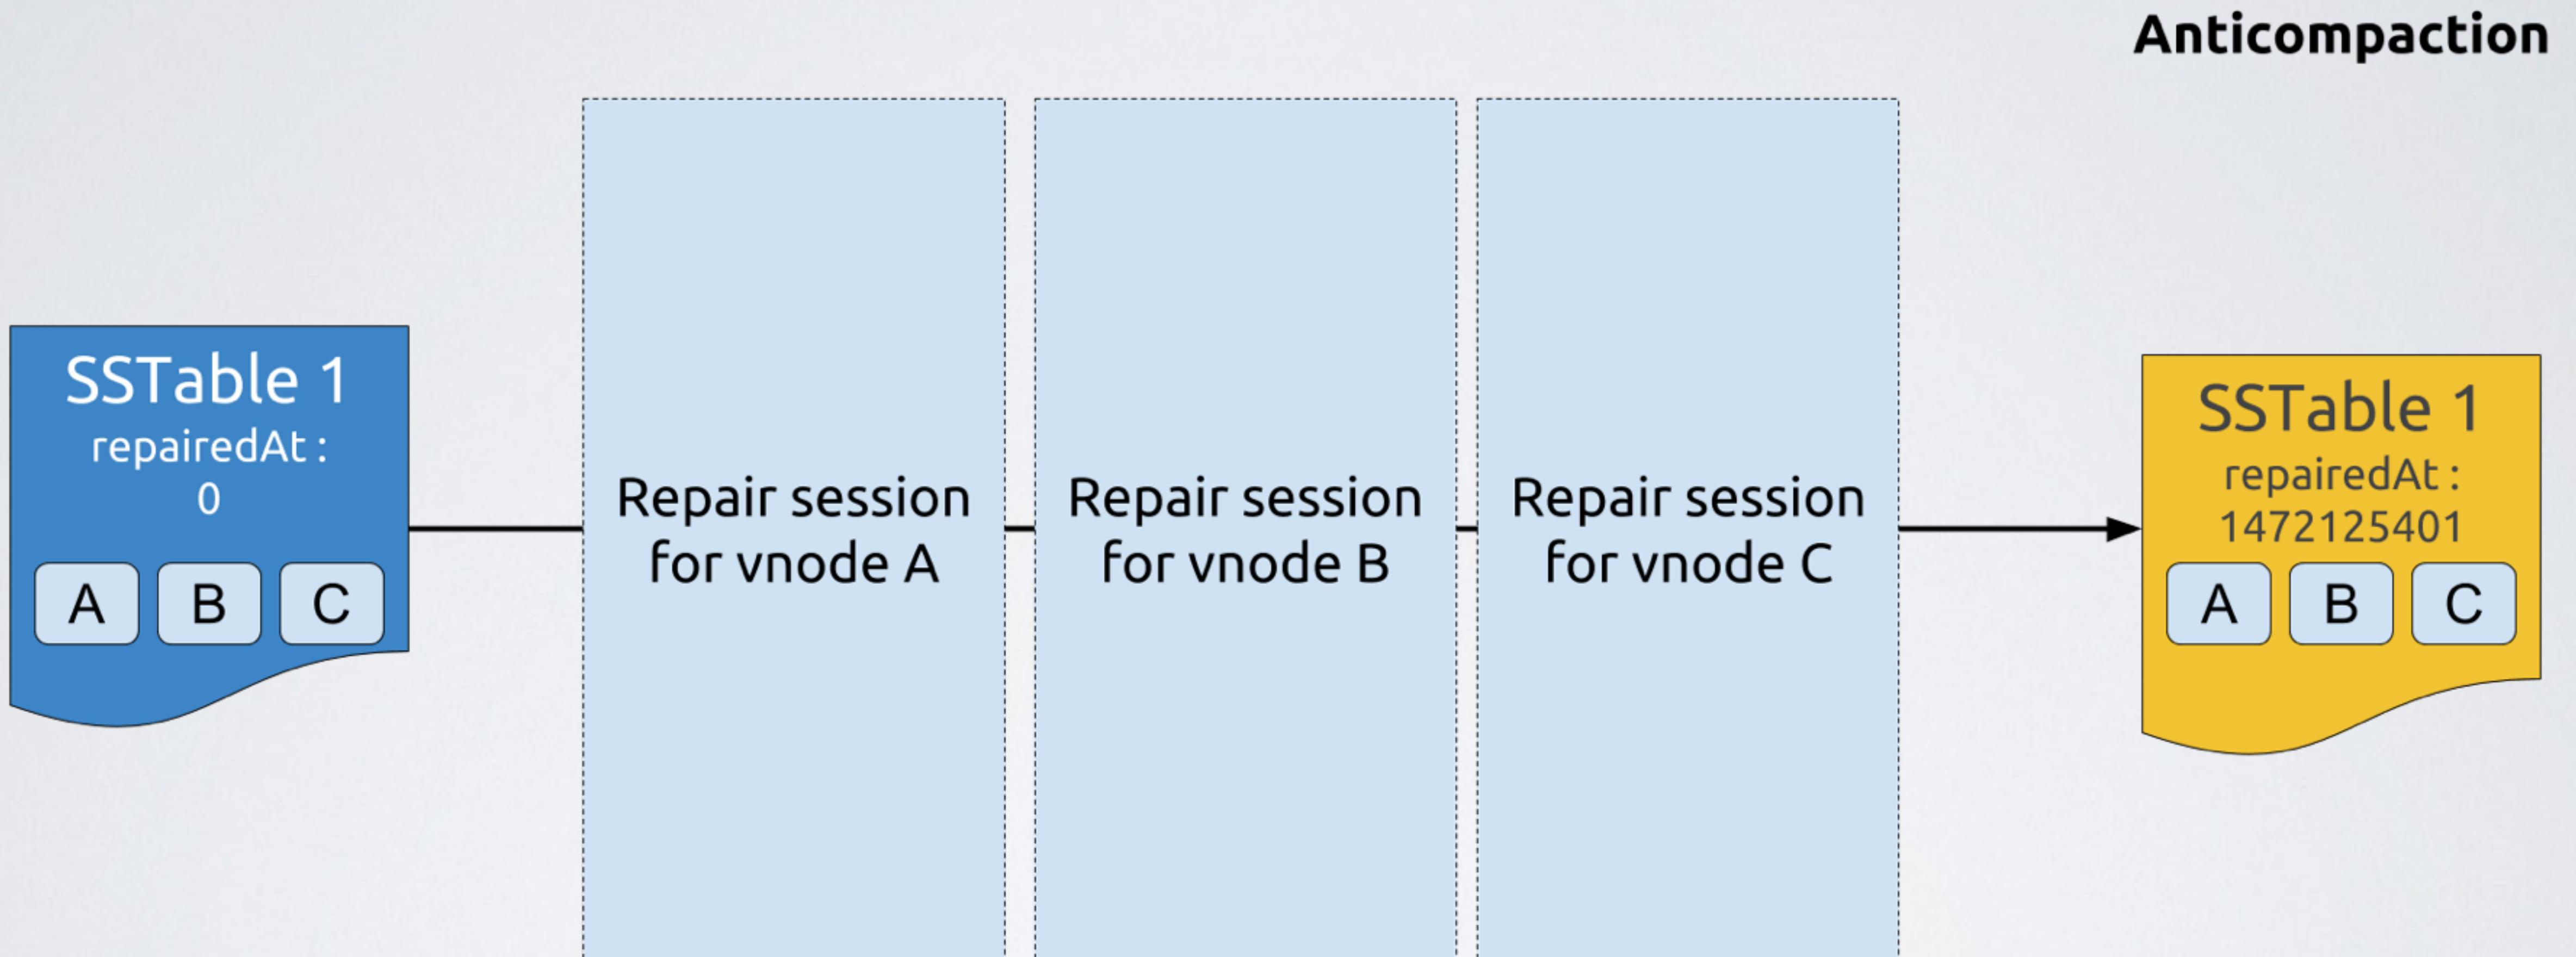
Second repair



Anticompaction



Anticompaction (repair on all ranges on local node)



Incremental repair looks awesome...
...but has flaws and drawbacks



Incremental repair caveats

**Carefully prepare your switch to
incremental repair**

Incremental repair caveats

Carefully prepare your switch to
incremental repair

i.e. do not run « nodetool repair -inc »
straight away...

Incremental repair caveats

It doesn't handle missing/corrupted data
that was already repaired



Incremental repair caveats

It splits SSTables in 2 sets
that cannot be compacted together
(think tombstone purge)



Incremental repair caveats

It is incompatible with subrange repair
(anticompaction)



Incremental repair caveats

It doesn't like concurrency very much
even in C* 3.x



Incremental repair caveats

Validator.java:261 -
Failed creating a merkle tree for [repair #e4c782d0-11fc-11e6-
b616-51a3849870bb on table_v2/table_attributes,
[(883546083348233317,883877311566358575],
(-7300486781514672850,-7298192396576668423],
(-959298474675167225,-959177964106074209]]], /10.10.10.33
(see log for details)

Incremental repair caveats

CompactionManager.java:1320 - **Cannot start multiple repair sessions over the same sstables**

Incremental repair caveats

CASSANDRA-8316

A running anticomposition prevents validation compaction

Incremental repair caveats

Do not use -pr with incremental repair

Incremental repair caveats

Do not use -pr with incremental repair

Useless : data is repaired once only

Incremental repair caveats

Do not use -pr with incremental repair

Useless : data is repaired once only

Expensive : anticompaсtion overhead

Incremental repair will not...

Fix a poor repair strategy

Incremental repair will not...

Prevent you from having to run full repair

Reaper does not support incremental repair

But this fork does :

<https://github.com/adejanovski/cassandra-reaper/tree/inc-repair-that-works>

Reaper does not support incremental repair

And this one embeds the modded UI :

<https://github.com/adejanovski/cassandra-reaper/tree/inc-repair-support-with-ui>

Reaper does not support incremental repair

Not enough time to write those urls ?

github.com/adejanovski

Reaper inc repair fork

No subrange repair

Reaper inc repair fork

No subrange repair
Single repair thread => no concurrency

What and why ?

Full repair

Incremental repair

How to make it work

Automated repairs

Repair best practices

Put your repair strategy in place on day 1

Repair best practices

Use appropriate tooling or build your own



Repair best practices

Spread repair over
a `gc_grace_seconds` cycle

Repair best practices

Adjust repair pressure
on your cluster
(Reaper does that)



Repair best practices

Don't repair everything !

Pick tables with deletes and those with
critical data

Repair best practices

If every data is critical, then none is ;)



Hmm Interesting...

Repair best practices

Be tight on your schedule with inc repair

Tombstones and anticompa^cction



Repair best practices

Avoid concurrency with inc repair
One node at a time



What and why ?

Full repair

Incremental repair

How to make it work

Automated repairs

The bright future ?

Not having to think about cyclic
maintenance repairs



CASSANDRA-10070

Automatic repair scheduling

CASSANDRA-8911

Mutation-based Repairs

Thanks!

@alexanderdeja

THE LAST PICKLE